


Formal Language Theory

Integrating Experimentation and Proof

Alley Stoughton

Draft of Fall 2018



Copyright © 2018 Alley Stoughton

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts.

The \LaTeX source of this book is part of the Forlan distribution, which is available on the Web at <http://alleystoughton.us/forlan>. A copy of the GNU Free Documentation License is included in the Forlan distribution.

Contents

Preface	ix
1 Mathematical Background	1
1.1 Basic Set Theory	1
1.1.1 Describing Sets by Listing Their Elements	1
1.1.2 Sets of Numbers	1
1.1.3 Relationships between Sets	2
1.1.4 Set Formation	3
1.1.5 Operations on Sets	4
1.1.6 Relations and Functions	6
1.1.7 Set Cardinality	10
1.1.8 Data Structures	14
1.1.9 Notes	16
1.2 Induction	16
1.2.1 Mathematical Induction	16
1.2.2 Strong Induction	17
1.2.3 Well-founded Induction	20
1.2.4 Notes	23
1.3 Inductive Definitions and Recursion	23
1.3.1 Inductive Definition of Trees	24
1.3.2 Recursion	28
1.3.3 Paths in Trees	31
1.3.4 Notes	33
2 Formal Languages	35
2.1 Symbols, Strings, Alphabets and (Formal) Languages	35
2.1.1 Symbols	35
2.1.2 Strings	36
2.1.3 Alphabets	39
2.1.4 Languages	40
2.1.5 Notes	41
2.2 Using Induction to Prove Language Equalities	41
2.2.1 String Induction Principles	41

2.2.2	Proving Language Equalities	44
2.2.3	Notes	50
2.3	Introduction to Forlan	50
2.3.1	Invoking Forlan	51
2.3.2	The SML Core of Forlan	51
2.3.3	Symbols	59
2.3.4	Sets	60
2.3.5	Sets of Symbols	61
2.3.6	Strings	62
2.3.7	Sets of Strings	64
2.3.8	Relations on Symbols	65
2.3.9	Notes	68
3	Regular Languages	69
3.1	Regular Expressions and Languages	69
3.1.1	Operations on Languages	69
3.1.2	Regular Expressions	73
3.1.3	Processing Regular Expressions in Forlan	79
3.1.4	JForlan	82
3.1.5	Notes	83
3.2	Equivalence and Correctness of Regular Expressions	83
3.2.1	Equivalence of Regular Expressions	83
3.2.2	Proving the Correctness of Regular Expressions	88
3.2.3	Notes	96
3.3	Simplification of Regular Expressions	96
3.3.1	Regular Expression Complexity	96
3.3.2	Weak Simplification	104
3.3.3	Local and Global Simplification	111
3.3.4	Notes	125
3.4	Finite Automata and Labeled Paths	125
3.4.1	Finite Automata	125
3.4.2	Labeled Paths and FA Meaning	129
3.4.3	Design of Finite Automata	133
3.4.4	Notes	134
3.5	Isomorphism of Finite Automata	135
3.5.1	Definition and Algorithm	135
3.5.2	Isomorphism Finding/Checking in Forlan	140
3.5.3	Notes	141
3.6	Checking Acceptance and Finding Accepting Paths	141
3.6.1	Processing a String from a Set of States	142
3.6.2	Checking String Acceptance and Finding Accepting Paths	144
3.6.3	Notes	146
3.7	Simplification of Finite Automata	147

3.7.1	Notes	152
3.8	Proving the Correctness of Finite Automata	152
3.8.1	Definition of Λ	152
3.8.2	Proving that Enough is Accepted	154
3.8.3	Proving that Everything Accepted is Wanted	156
3.8.4	Notes	158
3.9	Empty-string Finite Automata	158
3.9.1	Definition of EFAs	159
3.9.2	Converting FAs to EFAs	159
3.9.3	Processing EFAs in Forlan	161
3.9.4	Notes	163
3.10	Nondeterministic Finite Automata	163
3.10.1	Definition of NFAs	163
3.10.2	Converting EFAs to NFAs	163
3.10.3	Converting EFAs to NFAs, and Processing NFAs in Forlan	167
3.10.4	Notes	169
3.11	Deterministic Finite Automata	169
3.11.1	Definition of DFAs	169
3.11.2	Proving the Correctness of DFAs	172
3.11.3	Simplification of DFAs	174
3.11.4	Converting NFAs to DFAs	177
3.11.5	Processing DFAs in Forlan	181
3.11.6	Notes	184
Bibliography		185
Index		189

List of Figures

1.1 Example Diagonalization Table for Cardinality Proof 13

Preface

Background

Since the 1930s, the subject of formal language theory, also known as automata theory, has been developed by computer scientists, linguists and mathematicians. Formal languages (or simply languages) are sets of strings over finite sets of symbols, called alphabets, and various ways of describing such languages have been developed and studied, including regular expressions (which “generate” languages), finite automata (which “accept” languages), grammars (which “generate” languages) and Turing machines (which “accept” languages). For example, the set of identifiers of a given programming language is a formal language—one that can be described by a regular expression or a finite automaton. And, the set of all strings of tokens that are generated by a programming language’s grammar is another example of a formal language.

Because of its applications to computer science, most computer science programs offer both undergraduate and graduate courses in this subject. Perhaps the best known applications are to compiler construction. For example, regular expressions and finite automata are used when specifying and implementing lexical analyzers, and grammars are used to specify and implement parsers. Finite automata are used when designing hardware and network protocols. And Turing machines—or other machines/programs of equivalent power—are used to formalize the notion of algorithm, which in turn makes possible the study of what is, and is not, computable.

Formal language theory is largely concerned with algorithms, both ones that are explicitly presented, and ones implicit in theorems that are proved constructively. In typical courses on formal language theory, students apply these algorithms to toy examples by hand, and learn how they are used in applications. Although much can be achieved by a paper-and-pencil approach to the subject, students would obtain a deeper understanding of the subject if they could experiment with the algorithms of formal language theory using computer tools.

Consider, e.g., a typical exercise of a formal language theory class in which students are asked to synthesize a deterministic finite automaton that accepts some language, L . With the paper-and-pencil approach, the student is obliged

to build the machine by hand, and then (hopefully) prove it correct. But, given the right computer tools, another approach would be possible. First, the student could try to express L in terms of simpler languages, making use of various language operations (e.g., union, intersection, difference, concatenation, closure). The student could then synthesize automata accepting the simpler languages, enter these machines into the system, and then combine these machines using operations corresponding to the language operations used to express L . Finally, the resulting machine could be minimized. With some such exercises, a student could solve the exercise in both ways, and could compare the results. Other exercises of this type could only be solved with machine support.

Integrating Experimentation and Proof

To support experimentation with formal languages, I designed and implemented a computer toolset called Forlan [Sto05, Sto08]. Forlan is implemented in the functional programming language Standard ML (SML) [MTHM97, Pau96], a language whose notation and concepts are similar to those of mathematics. Forlan is a library on top of the Standard ML of New Jersey (SML/NJ) implementation of SML [AM91]. It's used interactively, and users are able to extend Forlan by defining SML functions.

In Forlan, the usual objects of formal language theory—finite automata, regular expressions, grammars, labeled paths, parse trees, etc.—are defined as abstract types, and have concrete syntax. Instead of Turing machines, Forlan implements a simple functional programming language of equivalent power, but which has the advantage of being much easier to program in than Turing machines. Programs are also abstract types, and have concrete syntax. Although mainly *not* graphical in nature, Forlan includes the Java program JForlan, a graphical editor for finite automata and regular expression, parse and program trees. It can be invoked directly, or via Forlan.

Numerous algorithms of formal language theory are implemented in Forlan, including conversions between regular expressions and different kinds of automata, the usual operations (e.g., union) on regular expressions, automata and grammars, equivalence testing and minimization of deterministic finite automata, a general parser for grammars, etc. Forlan provides support for regular expression simplification, although the algorithms used are works in progress. It also implements the functional programming language used as a substitute for Turing machines.

This undergraduate-level textbook and Forlan were designed and developed together. I have attempted to keep the conceptual and notational distance between the textbook and toolset as small as possible. The book treats most concepts and algorithms both theoretically, especially using proof, and through experimentation, using Forlan.

In contrast to some books on formal language theory, the book emphasizes

the concrete over the abstract, providing numerous, fully worked-out examples of how regular expressions, finite automata, grammars and programs can be designed and proved correct. In my view, students are most able to learn how to write proofs—and to see the benefit of doing so—if their proofs are about things that they have designed.

I have also attempted to simplify the foundations of the subject, using alternative definitions when needed. E.g., finite automata are defined in such a way that they form a set (as opposed to a proper class), and so that more restricted forms of automata (e.g., deterministic finite automata (DFAs)) are proper subsets of that set. And finite automata are given semantics using labeled paths, from which the traditional δ notation is derived, in the case of DFAs. Furthermore, the book treats the set theoretic foundations of the subject more rigorously than is typical.

Readers of this book are assumed to have significant experience reading and writing informal mathematical proofs, of the kind one finds in mathematics books. This experience could have been gained, e.g., in courses on discrete mathematics, logic or set theory. The book assumes no previous knowledge of Standard ML. In order to understand and extend the implementation of Forlan, though, one must have a good working knowledge of Standard ML, as could be obtained by working through [Pau96] or [Ull98].

Drafts of this book were successfully used at Kansas State University in a semester long, undergraduate course on formal language theory.

Outline of the Book

The book consists of five chapters. Chapter 1, *Mathematical Background*, consists of the material on set theory, induction and recursion, and trees and inductive definitions that is required in the remaining chapters.

In Chapter 2, *Formal Languages*, we say what symbols, strings, alphabets and (formal) languages are, show how to use various induction principles to prove language equalities, and give an introduction to the Forlan toolset. The remaining three chapters introduce and study more restricted sets of languages.

In Chapter 3, *Regular Languages*, we study regular expressions and languages, five kinds of finite automata, algorithms for processing and converting between regular expressions and finite automata, properties of regular languages, and applications of regular expressions and finite automata to searching in text files, lexical analysis, and the design of finite state systems.

In Chapter ??, *Context-free Languages*, we study context-free grammars and languages, algorithms for processing grammars and for converting regular expressions and finite automata to grammars, top-down (recursive descent) parsing, and properties of context-free languages. We prove that the set of context-free languages is a proper superset of the set of regular languages.

Finally, in Chapter ??, *Recursive and Recursively Enumerable Languages*, we study the functional programming language that we use instead of Turing machines to define the recursive and recursively enumerable languages. We study algorithms for processing programs and for converting grammars to programs, and properties of recursive and recursively enumerable languages. We prove that the set of context-free languages is a proper subset of the set of recursive languages, that the set of recursive languages is a proper subset of the set of recursively enumerable languages, and that there are languages that are not recursively enumerable. Furthermore, we show that there are problems, like the halting problem (the problem of determining whether a program halts when run on a given input), that can't be solved by programs.

Further Reading and Related Work

This book covers most of the material that is typically presented in an undergraduate course on formal language theory. On the other hand, typical textbooks on formal language theory cover much more of the subject than we do. Readers who are interested in learning more, or who would like to be exposed to alternative presentations of some of the material in this book, should consult one of the many fine books on formal language theory, such as [HMU01, Koz97, LP98, Mar91, Lin01].

Neil Jones [Jon97] pioneered the use of a programming language with structured data as an alternative to Turing machines for studying the limits of what is computable. In Chapter ??, we have followed Jones's approach in some ways. On the other hand, our programming language is functional, not imperative (assignment-oriented), and it has explicit support for the symbols and strings of formal language theory.

The existing formal languages toolsets fit into two categories. In the first category are tools, like JFLAP [BLP⁺97, HR00, Rod06], Pâté [BLP⁺97, HR00], the Java Computability Toolkit [RHND99], and Turing's World [BE93], that are graphically oriented and help students work out relatively small examples. The books [Rod06] (on JFLAP) and [Lin01] (an introduction to formal language theory) are intended to be used in conjunction with each other. The second category consists of toolsets that, like Forlan, are embedded in programming languages, and so support sophisticated experimentation with formal languages. Toolsets in this category include Automata [Sut92], Grail+ [RW94, RWYC17], HaLeX [Sar02], Leiß's Automata Library [Lei10] and Vaucanson [LRGS04]. I am not aware of other textbook/toolset packages whose toolsets are members of this second category.

Notes, Exercises and Website

In the “notes” subsections that conclude most sections of the book, I have restricted myself to describing how the book’s approach differs from standard practice. Readers interested in the history of the subject can consult [HMU01, Koz97, LP98].

The book contains numerous fully worked-out examples, many of which consist of designing and proving the correctness of regular expressions, finite automata, grammars and programs. Similar exercises, as well as other kinds of exercises, are scattered throughout the book.

The Forlan website

`http://alleystoughton.us/forlan`

contains:

- instructions for downloading and installing the Forlan toolset, and JForlan;
- the Forlan manual;
- instructions for reporting errors or making suggestions; and
- the Forlan distribution, including the source for Forlan and JForlan, as well as the \LaTeX source for this book.

Acknowledgments

Leonard Lee and Jessica Sherrill designed and implemented graphical editors for Forlan finite automata (JFA), and regular expression and parse trees (JTR), respectively. Their work was unified and enhanced (of particular note was the addition of support for program trees) by Srinivasa Aditya Uppu, resulting in an initial version of JForlan. Subsequently, Kenton Born carried out a major redevelopment of JForlan, resulting in JForlan Version 1.0. A further revision, by the author, led to JForlan Version 2.0.

It is a pleasure to acknowledge helpful discussions relating to the Forlan project with Eli Fox-Epstein, Brian Howard, Rod Howell, John Hughes, Nathan James, Patrik Jansson, Jace Kohlmeier, Dexter Kozen, Matthew Miller, Aarne Ranta, Ryan Stejskal, Colin Stirling, Lucio Torrico and Lyn Turbak. Much of this work was done while I was employed by Kansas State University, and some of the work was done while I was on sabbatical at the Department of Computing Science of Chalmers University of Technology.

Chapter 1

Mathematical Background

This chapter consists of the material on set theory, induction, trees, inductive definitions and recursion that is required in the remaining chapters.

1.1 Basic Set Theory

In this section, we will cover the material on sets, relations, functions and data structures that will be needed in what follows. Much of this material should be at least partly familiar.

1.1.1 Describing Sets by Listing Their Elements

We write \emptyset for the *empty set*—the set with no elements. Finite sets can be described by listing their elements inside set braces: $\{x_1, \dots, x_n\}$. E.g., $\{3\}$ is the *singleton set* whose only element is 3, and $\{1, 3, 5, 7\}$ is the set consisting of the first four odd numbers.

1.1.2 Sets of Numbers

We write:

- \mathbb{N} for the set $\{0, 1, \dots\}$ of all natural numbers;
- \mathbb{Z} for the set $\{\dots, -1, 0, 1, \dots\}$ of all integers;
- \mathbb{R} for the set of all real numbers.

Note that, for us, 0 *is* a natural number. This has many pleasant consequences, e.g., that the size of a finite set or the length of a list will always be a natural number.

1.1.3 Relationships between Sets

As usual, we write $x \in Y$ to mean that x is one of the elements (members) of the set Y . Sets A and B are equal ($A = B$) iff (if and only if) they have the same elements, i.e., for all x , $x \in A$ iff $x \in B$.

Suppose A and B are sets. We say that:

- A is a *subset* of B ($A \subseteq B$) iff, for all $x \in A$, $x \in B$;
- A is a *proper subset* of B ($A \subsetneq B$) iff $A \subseteq B$ but $A \neq B$.

In other words: A is a subset of B iff every everything in A is also in B , and A is a proper subset of B iff everything in A is in B , but there is at least one element of B that is not in A .

For example, $\emptyset \subsetneq \mathbb{N}$, $\mathbb{N} \subseteq \mathbb{N}$ and $\mathbb{N} \subsetneq \mathbb{Z}$. The definition of \subseteq gives us the most common way of showing that $A \subseteq B$: we suppose that $x \in A$, and show (with no additional assumptions about x) that $x \in B$. Similarly, if we want to show that $A = B$, it will suffice to show that $A \subseteq B$ and $B \subseteq A$, i.e., that everything in A is in B , and everything in B is in A . Of course, we can also use the usual properties of equality to show set equalities. E.g., if $A = B$ and $B = C$, we have that $A = C$.

Note that, for all sets A , B and C :

- if $A \subseteq B \subseteq C$, then $A \subseteq C$;
- if $A \subseteq B \subsetneq C$, then $A \subsetneq C$;
- if $A \subsetneq B \subseteq C$, then $A \subsetneq C$;
- if $A \subsetneq B \subsetneq C$, then $A \subsetneq C$.

Given sets A and B , we say that:

- A is a *superset* of B ($A \supseteq B$) iff, for all $x \in B$, $x \in A$;
- A is a *proper superset* of B ($A \supsetneq B$) iff $A \supseteq B$ but $A \neq B$.

Of course, for all sets A and B , we have that: $A = B$ iff $A \supseteq B \supseteq A$; and $A \subseteq B$ iff $B \supseteq A$. Furthermore, for all sets A , B and C :

- if $A \supseteq B \supseteq C$, then $A \supseteq C$;
- if $A \supseteq B \supsetneq C$, then $A \supsetneq C$;
- if $A \supsetneq B \supseteq C$, then $A \supsetneq C$;
- if $A \supsetneq B \supsetneq C$, then $A \supsetneq C$.

1.1.4 Set Formation

We will make extensive use of the $\{\dots \mid \dots\}$ notation for forming sets. Let's consider two representative examples of its use.

For the first example, let

$$A = \{n \mid n \in \mathbb{N} \text{ and } n^2 \geq 20\} = \{n \in \mathbb{N} \mid n^2 \geq 20\}.$$

(where the third of these expressions abbreviates the second one). Here, n is a bound variable and is universally quantified—changing it uniformly to m , for instance, wouldn't change the meaning of A . By the definition of A , we have that, for all n ,

$$n \in A \quad \text{iff} \quad n \in \mathbb{N} \text{ and } n^2 \geq 20.$$

Thus, e.g.,

$$5 \in A \quad \text{iff} \quad 5 \in \mathbb{N} \text{ and } 5^2 \geq 20.$$

Since $5 \in \mathbb{N}$ and $5^2 = 25 \geq 20$, it follows that $5 \in A$. On the other hand, $5.5 \notin A$, since $5.5 \notin \mathbb{N}$, and $4 \notin A$, since $4^2 \not\geq 20$.

For the second example, let

$$B = \{n^3 + m^2 \mid n, m \in \mathbb{N} \text{ and } n, m \geq 1\}.$$

Note that $n^3 + m^2$ is a term (expression), rather than a variable. The variables n and m are existentially quantified, rather than universally quantified, so that, for all l ,

$$\begin{aligned} l \in B & \quad \text{iff} \quad l = n^3 + m^2, \text{ for some } n, m \text{ such that } n, m \in \mathbb{N} \text{ and } n, m \geq 1 \\ & \quad \text{iff} \quad l = n^3 + m^2, \text{ for some } n, m \in \mathbb{N} \text{ such that } n, m \geq 1. \end{aligned}$$

Thus, to show that $9 \in B$, we would have to show that

$$9 = n^3 + m^2 \text{ and } n, m \in \mathbb{N} \text{ and } n, m \geq 1,$$

for some values of n, m . And, this holds, since $9 = 2^3 + 1^2$ and $2, 1 \in \mathbb{N}$ and $2, 1 \geq 1$.

We use set formation in the following definition. Given $n, m \in \mathbb{Z}$, we write $[n : m]$ for $\{l \in \mathbb{Z} \mid l \geq n \text{ and } l \leq m\}$. Thus $[n : m]$ is all of the integers that are at least n and no more than m . For example, $[-2 : 1]$ is $\{-2, -1, 0, 1\}$ and $[3 : 2]$ is \emptyset .

Some uses of the $\{\dots \mid \dots\}$ notation are too “big” to be sets, and instead are *proper classes*. A *class* is a collection of universe elements, and a class is *proper* iff it is not a set. E.g., $\{A \mid A \text{ is a set and } A \notin A\}$ is a proper class: assuming that it is a set is inconsistent—makes it possible to prove anything. This is the so-called Russell's Paradox. To know that a set formation is valid, it suffices to find a set that includes all the elements in the class being defined. E.g., the above set formations are valid, because the classes being defined are subsets of \mathbb{N} .

1.1.5 Operations on Sets

Next, we consider some standard operations on sets. Recall the following operations on sets A and B :

$$\begin{aligned} A \cup B &= \{x \mid x \in A \text{ or } x \in B\} && \text{(union)} \\ A \cap B &= \{x \mid x \in A \text{ and } x \in B\} && \text{(intersection)} \\ A - B &= \{x \in A \mid x \notin B\} && \text{(difference)} \\ A \times B &= \{(x, y) \mid x \in A \text{ and } y \in B\} && \text{(product)} \\ \mathcal{P}A &= \{X \mid X \subseteq A\} && \text{(power set).} \end{aligned}$$

The axioms of set theory assert that all of these set formations are valid (intersection and difference are obviously valid).

Of course, union and intersection are both commutative and associative ($A \cup B = B \cup A$, $(A \cup B) \cup C = A \cup (B \cup C)$, $A \cap B = B \cap A$ and $(A \cap B) \cap C = A \cap (B \cap C)$, for all sets A, B, C). Furthermore, we have that union is idempotent ($A \cup A = A$, for all sets A), and that \emptyset is the identity for union ($\emptyset \cup A = A = A \cup \emptyset$, for all sets A). Also, intersection is idempotent ($A \cap A = A$, for all sets A), and \emptyset is the zero for intersection ($\emptyset \cap A = \emptyset = A \cap \emptyset$, for all sets A).

It is easy to see that, for all sets X and Y , $X \subseteq X \cup Y$ and $Y \subseteq X \cup Y$. We say that sets X and Y are *disjoint* iff $X \cap Y = \emptyset$, i.e., iff X and Y have nothing in common.

$A - B$ is formed by removing the elements of B from A , if necessary. For example, $\{0, 1, 2\} - \{1, 4\} = \{0, 2\}$. $A \times B$ consists of all ordered pairs (x, y) , where x comes from A and y comes from B . For example, $\{0, 1\} \times \{1, 2\} = \{(0, 1), (0, 2), (1, 1), (1, 2)\}$. Remember that an ordered pair (x, y) is different from $\{x, y\}$, the set containing just x and y . In particular, we have that, for all x, x', y, y' , $(x, y) = (x', y')$ iff $x = x'$ and $y = y'$, whereas $\{1, 2\} = \{2, 1\}$. If A and B have n and m elements, respectively, for $n, m \in \mathbb{N}$, then $A \times B$ will have nm elements. Finally, $\mathcal{P}A$ consists of all of the subsets of A . For example, $\mathcal{P}\{0, 1\} = \{\emptyset, \{0\}, \{1\}, \{0, 1\}\}$. If A has n elements, for $n \in \mathbb{N}$, then $\mathcal{P}A$ will have 2^n elements.

We let \times associate to the right, so that, e.g., $A \times B \times C = A \times (B \times C)$. And, we abbreviate $(x_1, (x_2, \dots (x_{n-1}, x_n) \dots))$ to $(x_1, x_2, \dots, x_{n-1}, x_n)$, thinking of it as an *ordered n -tuple*. For example $(x, (y, z))$ is abbreviated to (x, y, z) , and we think of it as an *ordered triple*.

As an example of a proof involving sets, let's prove the following simple proposition, which says that intersections may be distributed over unions:

Proposition 1.1.1

Suppose A, B and C are sets.

- (1) $A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$.
- (2) $(A \cup B) \cap C = (A \cap C) \cup (B \cap C)$.

Proof. We show (1), the proof of (2) being similar. It will suffice to show that $A \cap (B \cup C) \subseteq (A \cap B) \cup (A \cap C) \subseteq A \cap (B \cup C)$.

$(A \cap (B \cup C) \subseteq (A \cap B) \cup (A \cap C))$ Suppose $x \in A \cap (B \cup C)$. We must show that $x \in (A \cap B) \cup (A \cap C)$. By our assumption, we have that $x \in A$ and $x \in B \cup C$. Since $x \in B \cup C$, there are two cases to consider.

- Suppose $x \in B$. Then $x \in A \cap B \subseteq (A \cap B) \cup (A \cap C)$, so that $x \in (A \cap B) \cup (A \cap C)$.
- Suppose $x \in C$. Then $x \in A \cap C \subseteq (A \cap B) \cup (A \cap C)$, so that $x \in (A \cap B) \cup (A \cap C)$.

$((A \cap B) \cup (A \cap C) \subseteq A \cap (B \cup C))$ Suppose $x \in (A \cap B) \cup (A \cap C)$. We must show that $x \in A \cap (B \cup C)$. There are two cases to consider.

- Suppose $x \in A \cap B$. Then $x \in A$ and $x \in B \subseteq B \cup C$, so that $x \in A \cap (B \cup C)$.
- Suppose $x \in A \cap C$. Then $x \in A$ and $x \in C \subseteq B \cup C$, so that $x \in A \cap (B \cup C)$.

□

Exercise 1.1.2

Suppose A , B and C are sets. Prove that union distributes over intersection, i.e., for all sets A , B and C :

$$(1) A \cup (B \cap C) = (A \cup B) \cap (A \cup C).$$

$$(2) (A \cap B) \cup C = (A \cup C) \cap (B \cup C).$$

Next, we consider generalized versions of union and intersection that work on sets of sets. If X is a set of sets, then the *generalized union* of X ($\bigcup X$) is

$$\{a \mid a \in A, \text{ for some } A \in X\}.$$

Thus, to show that $a \in \bigcup X$, we must show that a is in at least one element A of X . For example

$$\begin{aligned} \bigcup \{\{0, 1\}, \{1, 2\}, \{2, 3\}\} &= \{0, 1, 2, 3\} = \{0, 1\} \cup \{1, 2\} \cup \{2, 3\}, \\ \bigcup \emptyset &= \emptyset. \end{aligned}$$

(Again, we rely on set theory's axioms to know this set formation is valid.)

If X is a *nonempty* set of sets, then the *generalized intersection* of X ($\bigcap X$) is

$$\{a \mid a \in A, \text{ for all } A \in X\}.$$

Thus, to show that $a \in \bigcap X$, we must show that a is in every element A of X . For example

$$\bigcap \{\{0, 1\}, \{1, 2\}, \{2, 3\}\} = \emptyset = \{0, 1\} \cap \{1, 2\} \cap \{2, 3\}.$$

If we allowed $\bigcap \emptyset$, then it would contain all elements a of our universe that are in all of the nonexistent elements of \emptyset , i.e., it would contain all elements of our universe. But this collection is a proper class, not a set. Let's consider the above reasoning in more detail. Suppose a is a universe element. To prove that $a \in A$, for all $A \in \emptyset$, suppose $A \in \emptyset$. But this is a logical contradiction, from which we may prove anything, in particular our desired conclusion, $a \in A$.

1.1.6 Relations and Functions

Next, we consider relations and functions. A *relation* R is a set of ordered pairs. The *domain* of a relation R (**domain** R) is $\{x \mid (x, y) \in R, \text{ for some } y\}$, and the *range* of R (**range** R) is $\{y \mid (x, y) \in R, \text{ for some } x\}$. We say that R is a *relation from* a set X *to* a set Y iff **domain** $R \subseteq X$ and **range** $R \subseteq Y$, and that R is a *relation on* a set A iff **domain** $R \cup \text{range } R \subseteq A$. We often write $x R y$ for $(x, y) \in R$.

Consider the relation

$$R = \{(0, 1), (1, 2), (0, 2)\}.$$

Then, **domain** $R = \{0, 1\}$, **range** $R = \{1, 2\}$, R is a relation from $\{0, 1\}$ to $\{1, 2\}$, and R is a relation on $\{0, 1, 2\}$. Of course, R is also, e.g., a relation between \mathbb{N} and \mathbb{R} , as well as relation on \mathbb{R} .

We often form relations using set formation. For example, suppose $\phi(x, y)$ is a formula involving variables x and y . Then we can let $R = \{(x, y) \mid \phi(x, y)\}$, and it is easy to show that, for all x, y , $(x, y) \in R$ iff $\phi(x, y)$.

Given a set A , the *identity relation* on A (**id** _{A}) is $\{(x, x) \mid x \in A\}$. For example, **id** _{$\{1, 3, 5\}$} is $\{(1, 1), (3, 3), (5, 5)\}$. Given relations R and S , the *composition of* S and R ($S \circ R$) is $\{(x, z) \mid (x, y) \in R \text{ and } (y, z) \in S, \text{ for some } y\}$. Intuitively, it's the relation formed by starting with a pair $(x, y) \in R$, following it with a pair $(y, z) \in S$ (one that begins where the pair from R left off), and suppressing the intermediate value y , leaving us with (x, z) . For example, if $R = \{(1, 1), (1, 2), (2, 3)\}$ and $S = \{(2, 3), (2, 4), (3, 4)\}$, then from $(1, 2) \in R$ and $(2, 3) \in S$, we can conclude $(1, 3) \in S \circ R$. There are two more pairs in $S \circ R$, giving us $S \circ R = \{(1, 3), (1, 4), (2, 4)\}$. For all sets A , B and C , and relations R and S , if R is a relation from A to B , and S is a relation from B to C , then $S \circ R$ is a relation from A to C .

It is easy to show that \circ is associative and has the identity relations as its identities:

- (1) For all sets A and B , and relations R from A to B , **id** _{B} $\circ R = R = R \circ \text{id}_A$.

- (2) For all sets A, B, C and D , and relations R from A to B , S from B to C , and T from C to D , $(T \circ S) \circ R = T \circ (S \circ R)$.

Because of (2), we can write $T \circ S \circ R$, without worrying about how it is parenthesized.

The *inverse* of a relation R (R^{-1}) is the relation $\{(y, x) \mid (x, y) \in R\}$, i.e., it is the relation obtained by reversing each of the pairs in R . For example, if $R = \{(0, 1), (1, 2), (1, 3)\}$, then the inverse of R is $\{(1, 0), (2, 1), (3, 1)\}$. So for all sets A and B , and relations R , R is a relation from A to B iff R^{-1} is a relation from B to A . We also have that, for all sets A and B and relations R from A to B , $(R^{-1})^{-1} = R$. Furthermore, for all sets A, B and C , relations R from A to B , and relations S from B to C , $(S \circ R)^{-1} = R^{-1} \circ S^{-1}$.

A relation R is:

- *reflexive* on a set A iff, for all $x \in A$, $(x, x) \in R$;
- *transitive* iff, for all x, y, z , if $(x, y) \in R$ and $(y, z) \in R$, then $(x, z) \in R$;
- *symmetric* iff, for all x, y , if $(x, y) \in R$, then $(y, x) \in R$;
- *antisymmetric* iff, for all x, y , if $(x, y) \in R$ and $(y, x) \in R$, then $x = y$;
- *total* on a set A iff, for all $x, y \in A$, either $(x, y) \in R$ or $(y, x) \in R$;
- a *function* iff, for all x, y, z , if $(x, y) \in R$ and $(x, z) \in R$, then $y = z$.

Note that being antisymmetric is *not* the same as not being symmetric.

Suppose, e.g., that $R = \{(0, 1), (1, 2), (0, 2)\}$. Then:

- R is not reflexive on $\{0, 1, 2\}$, since $(0, 0) \notin R$.
- R is transitive, since whenever (x, y) and (y, z) are in R , it follows that $(x, z) \in R$. Since $(0, 1)$ and $(1, 2)$ are in R , we must have that $(0, 2)$ is in R , which is indeed true.
- R is not symmetric, since $(0, 1) \in R$, but $(1, 0) \notin R$.
- R is antisymmetric, since there are no x, y such that (x, y) and (y, x) are both in R . (If we added $(1, 0)$ to R , then R would not be antisymmetric, since then R would contain $(0, 1)$ and $(1, 0)$, but $0 \neq 1$.)
- R is not total on $\{0, 1, 2\}$, since $(0, 0) \notin R$.
- R is not a function, since $(0, 1) \in R$ and $(0, 2) \in R$. Intuitively, given an input of 0, it's not clear whether R 's output is 1 or 2.

We say that R is a *total ordering* on a set A iff R is a transitive, antisymmetric and total relation on A . It is easy to see that such an R will also be reflexive on A . Furthermore, if R is a total ordering on A , then R^{-1} is also a total ordering on A .

We often use a symbol like \leq to stand for a total ordering on a set A , which lets us use its mirror image, \geq , for its inverse, as well as to write $<$ for the relation on A defined by: for all $x, y \in A$, $x < y$ iff $x \leq y$ but $x \neq y$. $<$ is a *strict total ordering* on A , i.e., a transitive relation on A such that, for all $x, y \in A$, exactly one of $x < y$, $x = y$ and $y < x$ holds. We write $>$ for the inverse of $<$. We can also start with a strict total ordering $<$ on A , and then define a total ordering \leq on A by: $x \leq y$ iff $x < y$ or $x = y$. The relations \leq and $<$ on the natural numbers are examples of such relations.

The relation

$$f = \{(0, 1), (1, 2), (2, 0)\}$$

is a function. We think of it as sending the input 0 to the output 1, the input 1 to the output 2, and the input 2 to the output 0.

If f is a function and $x \in \mathbf{domain} f$, we write fx for the *application of f to x* , i.e., the unique y such that $(x, y) \in f$. We say that f is a *function from* a set X *to* a set Y iff f is a function, $\mathbf{domain} f = X$ and $\mathbf{range} f \subseteq Y$. Such an f is also a function from X to $\mathbf{range} f$. We write $X \rightarrow Y$ for the set of all functions from X to Y . If A has n elements and B has m elements, for $n, m \in \mathbb{N}$, then $A \rightarrow B$ will have m^n elements.

For the f defined above, we have that $f0 = 1$, $f1 = 2$, $f2 = 0$, f is a function from $\{0, 1, 2\}$ to $\{0, 1, 2\}$, and $f \in \{0, 1, 2\} \rightarrow \{0, 1, 2\}$. Of course, f is also in $\{0, 1, 2\} \rightarrow \mathbb{N}$, but it is not in $\mathbb{N} \rightarrow \{0, 1, 2\}$.

Exercise 1.1.3

Suppose X is a set, and $x \in X$. What are the elements of $\emptyset \rightarrow X$, $X \rightarrow \emptyset$, $\{x\} \rightarrow X$ and $X \rightarrow \{x\}$? Prove that your answers are correct.

We let \rightarrow associate to the right and have lower precedence than \times , so that, e.g., $A \times B \rightarrow C \times D \rightarrow E \times F$ means $(A \times B) \rightarrow ((C \times D) \rightarrow (E \times F))$. An element of this set takes in a pair (a, b) in $A \times B$, and returns a function that takes in a pair (c, d) in $C \times D$, and returns an element of $E \times F$.

Suppose $f, g \in A \rightarrow B$. It is easy to show that $f = g$ iff, for all $x \in A$, $fx = gx$. This is the most common way of showing the equality of functions.

Given a set A , it is easy to see that \mathbf{id}_A , the identity relation on A , is a function from A to A , and we call it the *identity function* on A . It is the function that returns its input. Given sets A , B and C , if f is a function from A to B , and g is a function from B to C , then the composition $g \circ f$ of (the relations) g and f is the function h from A to C such that $hx = g(fx)$, for all $x \in A$. In other words, $g \circ f$ is the function that runs f and then g , in sequence.

Because of how composition of relations works, we have that \circ is associative and has the identity functions as its identities:

- (1) For all sets A and B , and functions f from A to B , $\text{id}_B \circ f = f = f \circ \text{id}_A$.
- (2) For all sets A, B, C and D , and functions f from A to B , g from B to C , and h from C to D , $(h \circ g) \circ f = h \circ (g \circ f)$.

Because of (2), we can write $h \circ g \circ f$, without worrying about how it is parenthesized. It is the function that runs f , then g , then h , in sequence.

Given $f \in X \rightarrow Y$ and a subset A of X , we write $f(A)$ for the *image of A under f* , $\{f x \mid x \in A\}$. And given $f \in X \rightarrow Y$ and a subset B of Y , we write $f^{-1}(B)$ for the *inverse image of B under f* , $\{x \in X \mid f x \in B\}$. For example, if $f \in \mathbb{N} \rightarrow \mathbb{N}$ is the function that doubles its argument, then $f(\{3, 5, 7\}) = \{6, 10, 14\}$ and $f^{-1}(\{1, 2, 3, 4\}) = \{1, 2\}$.

Given a function f and a set $X \subseteq \text{domain } f$, we write $f|X$ for the *restriction of f to X* , $\{(x, y) \in f \mid x \in X\}$. Hence $\text{domain}(f|X) = X$. For example, if f is the function over \mathbb{Z} that increases its argument by 2, then $f|\mathbb{N}$ is the function over \mathbb{N} that increases its argument by 2. Given a function f and elements x, y of our universe, we write $f[x \mapsto y]$ for the *updating of f to send x to y* , $(f|(\text{domain } f - \{x\})) \cup \{(x, y)\}$. This function is the same as f , except that it sends x to y . For example, if $f = \{(0, 1), (1, 2)\}$, then $f[1 \mapsto 0] = \{(0, 1), (1, 0)\}$, and $f[2 \mapsto 3] = \{(0, 1), (1, 2), (2, 3)\}$.

We often define a function by saying how an element of its domain is transformed into an element of its range. E.g., we might say that $f \in \mathbb{N} \rightarrow \mathbb{Z}$ is the unique function such that, for all $n \in \mathbb{N}$,

$$f n = \begin{cases} -(n/2), & \text{if } n \text{ is even,} \\ (n+1)/2, & \text{if } n \text{ is odd.} \end{cases}$$

This is shorthand for saying that f is the set of all (n, m) such that

- $n \in \mathbb{N}$,
- if n is even, then $m = -(n/2)$, and
- if n is odd, then $m = (n+1)/2$.

Then $f 0 = 0$, $f 1 = 1$, $f 2 = -1$, $f 3 = 2$, $f 4 = -2$, etc.

Exercise 1.1.4

There are three things wrong with the following “definition”—what are they? Let $f \in \mathbb{N} \rightarrow \mathbb{N}$ be the unique function such that, for all $n \in \mathbb{N}$,

$$f n = \begin{cases} n - 2, & \text{if } n \geq 1 \text{ and } n \leq 10, \\ n + 2, & \text{if } n \geq 10. \end{cases}$$


If the domain of f is a product $X_1 \times \cdots \times X_n$, for $n \geq 1$ and sets X_1, \dots, X_n , we often abbreviate $f((x_1, \dots, x_n))$ to $f(x_1, \dots, x_n)$. We write $\#i_{X_1, \dots, X_n}$ (or just $\#i$, if n and the X_i are clear from the context) for the i -th *projection* function, which selects the i -th component of a tuple, i.e., transforms an input (x_1, \dots, x_n) to x_i .

1.1.7 Set Cardinality

Next, we see how we can use functions to compare the sizes (or cardinalities) of sets. A *bijection* f from a set X to a set Y is a function from X to Y such that, for all $y \in Y$, there is a unique $x \in X$ such that $(x, y) \in f$. For example,

$$f = \{(0, 5.1), (1, 2.6), (2, 0.5)\}$$

is a bijection from $\{0, 1, 2\}$ to $\{0.5, 2.6, 5.1\}$. We can visualize f as a one-to-one correspondence between these sets:



A function f is an *injection* (or is *injective*) iff, for all x, y and z , if $(x, z) \in f$ and $(y, z) \in f$, then $x = y$, i.e., for all $x, y \in \mathbf{domain} f$, if $f x = f y$, then $x = y$. In other words, a function is injective iff it never sends two different elements of its domain to the same element of its range. For example, the function

$$\{(0, 1), (1, 2), (2, 3), (3, 0)\}$$

is injective, but the function

$$\{(0, 1), (1, 2), (2, 1)\}$$

is not injective (both 0 and 2 are sent to 1). We say that f is an *injection from* a set X *to* a set Y , iff f is a function from X to Y and f is injective.

It is easy to see that:

- For all sets A , \mathbf{id}_A is injective.
- For all sets A, B and C , functions f from A to B , and functions g from B to C , if f and g are injective, then so is $g \circ f$.

Exercise 1.1.5

Suppose A and B are sets. Show that for all f , f is a bijection from A to B iff

- f is a function from A to B ;

- **range** $f = B$; and
- f is injective.

Consequently, if f is an injection from A to B , then f is a bijection from A to **range** $f \subseteq B$.

Exercise 1.1.6

Show that:

- (1) For all sets A , **id** $_A$ is a bijection from A to A .
- (2) For all sets A , B and C , bijections f from A to B , and bijections g from B to C , $g \circ f$ is a bijection from A to C .
- (3) For all sets A and B , and bijections f from A to B , f^{-1} is a bijection from B to A .

We say that a set X has the *same size* as a set Y ($X \cong Y$) iff there is a bijection from X to Y . By Exercise 1.1.6, we have that, for all sets A, B, C :

- (1) $A \cong A$;
- (2) If $A \cong B \cong C$, then $A \cong C$; and
- (3) If $A \cong B$, then $B \cong A$.

We say that a set X is:

- *finite* iff $X \cong [1 : n]$, for some $n \in \mathbb{N}$ (recall that $[1 : n]$ is all of the natural numbers that are at least 1 and no more than n , so that $[1 : 0] = \emptyset$);
- *infinite* iff it is not finite;
- *countably infinite* iff $X \cong \mathbb{N}$;
- *countable* iff X is either finite or countably infinite; and
- *uncountable* iff X is not countable.

Every set X has a *size* or *cardinality* ($|X|$) and we have that, for all sets X and Y , $|X| = |Y|$ iff $X \cong Y$. The sizes of finite sets are natural numbers.

We have that:

- The sets \emptyset and $\{0.5, 2.6, 5.1\}$ are finite, and are thus also countable;
- The sets \mathbb{N} , \mathbb{Z} , \mathbb{R} and $\mathcal{P}\mathbb{N}$ are infinite;
- The set \mathbb{N} is countably infinite, and is thus countable; and

- The set \mathbb{Z} is countably infinite, and is thus countable, because of the existence of the following bijection:

...	-2	-1	0	1	2	...
...	↓	↓	↓	↓	↓	...
...	4	2	0	1	3	...

- The sets \mathbb{R} and $\mathcal{P}\mathbb{N}$ are uncountable.

To prove that \mathbb{R} and $\mathcal{P}\mathbb{N}$ are uncountable, we can use an important technique called “diagonalization”, which we will see again in Chapter ???. Let’s consider the proof that $\mathcal{P}\mathbb{N}$ is uncountable.

We proceed using proof by contradiction. Suppose $\mathcal{P}\mathbb{N}$ is countable. If we can obtain a contradiction, it will follow that $\mathcal{P}\mathbb{N}$ is uncountable. Since $\mathcal{P}\mathbb{N}$ is not finite, it follows that there is a bijection f from \mathbb{N} to $\mathcal{P}\mathbb{N}$. Our plan is to define a subset X of \mathbb{N} such that $X \notin \text{range } f$, thus obtaining a contradiction, since this will show that f is not a bijection from \mathbb{N} to $\mathcal{P}\mathbb{N}$.

Consider the infinite table in which both the rows and the columns are indexed by the elements of \mathbb{N} , listed in ascending order, and where a cell (m, n) (m is the row, n is the column) contains 1 iff $m \in f n$, and contains 0 iff $m \notin f n$. Thus the n th column of this table represents the set $f n$ of natural numbers.

Figure 1.1 shows how part of this table might look, where i, j and k are sample elements of \mathbb{N} . Because of the table’s data, we have, e.g., that $i \in f i$ and $j \notin f i$.

To define our $X \subseteq \mathbb{N}$, we work our way down the diagonal of the table, putting n into our set just when cell (n, n) of the table is 0, i.e., when $n \notin f n$. This will ensure that, for all $n \in \mathbb{N}$, $f n \neq X$. With our example table:

- since $i \in f i$, but $i \notin X$, we have that $f i \neq X$;
- since $j \notin f j$, but $j \in X$, we have that $f j \neq X$; and
- since $k \in f k$, but $k \notin X$, we have that $f k \neq X$.

Next, we turn the above ideas into a shorter, but more opaque, proof that:

Proposition 1.1.7

$\mathcal{P}\mathbb{N}$ is uncountable.

Proof. Suppose, toward a contradiction, that $\mathcal{P}\mathbb{N}$ is countable. Because $\mathcal{P}\mathbb{N}$ is not finite, there is a bijection f from \mathbb{N} to $\mathcal{P}\mathbb{N}$. Define $X \in \{n \in \mathbb{N} \mid n \notin f n\}$, so that $X \in \mathcal{P}\mathbb{N}$. By the definition of f , it follows that $X = f n$, for some $n \in \mathbb{N}$. There are two cases to consider.

- Suppose $n \in X$. By the definition of X , it follows that $n \notin f n = X$ —contradiction.

	...	i	...	j	...	k	...
\vdots							
i		1		1		0	
\vdots							
j		0		0		1	
\vdots							
k		0		1		1	
\vdots							

Figure 1.1: Example Diagonalization Table for Cardinality Proof

- Suppose $n \notin X$. Because $X = fn$, we have that $n \notin fn$. Thus, since $n \in \mathbb{N}$ and $n \notin fn$, it follows that $n \in X$ —contradiction.

Since we obtained a contradiction in both cases, we have an overall contradiction. Thus $\mathcal{P}\mathbb{N}$ is uncountable. \square

We have seen how bijections may be used to determine whether sets have the same size. But how can we compare the relative sizes of sets, i.e., say whether one set is smaller or larger than another? The answer is to make use of injective functions. We say that a set X is *no bigger* than a set Y ($X \preceq Y$) iff there is an injection from X to Y , i.e., an injective function whose domain is X and whose range is a subset of Y . Because identity functions are injective, we have that $X \subseteq Y$ implies $X \preceq Y$; e.g., $\mathbb{N} \preceq \mathbb{R}$. This definition makes sense, because X is no bigger than Y iff X has the same size as a subset of Y . We say that X is *strictly smaller* than Y iff $X \preceq Y$ and $X \not\cong Y$.

Using our observations about injections, we have that, for all sets A , B and C :

- (1) $A \preceq A$;
- (2) If $A \preceq B \preceq C$, then $A \preceq C$.

We can also characterize \preceq using “surjectivity”. We say that f is a *surjection* from X to Y iff f is a function from X to Y and **range** $f = Y$. A consequence

of set theory's Axiom of Choice is that, for all sets X and Y , $X \preceq Y$ iff $X = \emptyset$ or there is a surjection from Y to X .

Clearly, for all sets A and B , if $A \cong B$, then $A \preceq B \preceq A$. And, a famous result of set theory, the Schröder-Bernstein Theorem, says that the converse holds, i.e., for all sets A and B , if $A \preceq B \preceq A$, then $A \cong B$. This gives us a powerful method for proving that two sets have the same size.

Exercise 1.1.8

Use the Schröder-Bernstein Theorem to show that $\mathbb{N} \cong \mathbb{N} \times \mathbb{N}$. Hint: use the following consequence of the Fundamental Theorem of Arithmetic: if two finite, ascending (each element is \leq the next) sequences of prime numbers (natural numbers that are at least 2 and have no divisors other than 1 and themselves) have the same product (the product of the empty sequence is 1), then they are equal.

One of the forms of the Axiom of Choice says that, for all sets A and B , either $A \preceq B$ or $B \preceq A$, i.e., either A is no bigger than B , or B is no bigger than A . Furthermore, the sizes of sets are ordered in such a way that, for all sets A and B , $|A| \leq |B|$ iff $A \preceq B$, which tells us that, given sets A and B , either $|A| \leq |B|$ or $|B| \leq |A|$. Given the above machinery, one can strengthen Proposition 1.1.7 into: for all sets X , X is strictly smaller than $\mathcal{P}X$, i.e., $|X| < |\mathcal{P}X|$.

1.1.8 Data Structures

We conclude this section by introducing some data structures that are built out of sets. We write **Bool** for the set of booleans, $\{\mathbf{true}, \mathbf{false}\}$. (We can actually let $\mathbf{true} = 1$ and $\mathbf{false} = 0$, although we'll never make use of these equalities.) We define the negation function $\mathbf{not} \in \mathbf{Bool} \rightarrow \mathbf{Bool}$ by:

$$\mathbf{not\ true} = \mathbf{false}, \quad \mathbf{not\ false} = \mathbf{true}.$$

And the conjunction and disjunction operations on the booleans are defined by:

$$\begin{aligned} \mathbf{true\ and\ true} &= \mathbf{true}, \\ \mathbf{true\ and\ false} &= \mathbf{false\ and\ true} = \mathbf{false\ and\ false} = \mathbf{false}, \end{aligned}$$

and

$$\begin{aligned} \mathbf{true\ or\ true} &= \mathbf{true\ or\ false} = \mathbf{false\ or\ true} = \mathbf{true}, \\ \mathbf{false\ or\ false} &= \mathbf{false}. \end{aligned}$$

Given a set X , we write **Option** X for $\{\mathbf{none}\} \cup \{\mathbf{some\ } x \mid x \in X\}$, where we define $\mathbf{none} = (0, 0)$ and $\mathbf{some\ } x = (1, x)$, which guarantees that $\mathbf{none} = \mathbf{some\ } x$ can't hold, and that $\mathbf{some\ } x = \mathbf{some\ } y$ only holds when $x = y$. (We won't make use of the particular way we've defined \mathbf{none} and $\mathbf{some\ } x$.) For example, $\mathbf{Option\ Bool} = \{\mathbf{none}, \mathbf{some\ true}, \mathbf{some\ false}\}$.

The idea is that an element of **Option** X is an optional element of X . E.g., when a function needs to either return an element of X or indicate that an error has occurred, it could return an element of **Option** X , using **none** to indicate an error, and returning **some** x to indicate success with value x . E.g., we could define a function $f \in \mathbb{N} \times \mathbb{N} \rightarrow \mathbf{Option\ Bool}$ by:

$$f(n, m) = \begin{cases} \mathbf{none}, & \text{if } m = 0, \\ \mathbf{some\ true} & \text{if } m \neq 0 \text{ and } n = ml \text{ for some } l \in \mathbb{N}, \\ \mathbf{some\ false} & \text{if } m \neq 0 \text{ and } n \neq ml \text{ for all } l \in \mathbb{N}. \end{cases}$$

Finally, we consider lists. A *list* is a function with domain $[1 : n]$, for some $n \in \mathbb{N}$. (Recall that $[1 : n]$ is all of the natural numbers that are at least 1 and no more than n .) For example, \emptyset is a list, as it is a function with domain $\emptyset = [1 : 0]$. And $\{(1, 3), (2, 5), (3, 7)\}$ is a list, as it is a function with domain $[1 : 3]$. Note that, if x is a list, then $|x|$, the size of the set x , doubles as the *length* of x .

We abbreviate a list $\{(1, x_1), (2, x_2), \dots, (n, x_n)\}$ to $[x_1, x_2, \dots, x_n]$. Thus \emptyset and $\{(1, 3), (2, 5), (3, 7)\}$ are abbreviated to $[]$ and $[3, 5, 7]$, respectively. If $[x_1, x_2, \dots, x_n] = [y_1, y_2, \dots, y_m]$, it is easy to see that $n = m$ and $x_i = y_i$, for all $i \in [1 : n]$.

Given lists f and g , the *concatenation* of f and g ($f @ g$) is the list

$$f \cup \{ (m + |f|, y) \mid (m, y) \in g \}.$$

For example,

$$\begin{aligned} [2, 3] @ [4, 5, 6] &= \{(1, 2), (2, 3)\} @ \{(1, 4), (2, 5), (3, 6)\} \\ &= \{(1, 2), (2, 3)\} \cup \{ (m + 2, y) \mid (m, y) \in \{(1, 4), (2, 5), (3, 6)\} \} \\ &= \{(1, 2), (2, 3)\} \cup \{(1 + 2, 4), (2 + 2, 5), (3 + 2, 6)\} \\ &= \{(1, 2), (2, 3), (3, 4), (4, 5), (5, 6)\} \\ &= [2, 3, 4, 5, 6]. \end{aligned}$$

Given lists f and g , it is easy to see that $|f @ g| = |f| + |g|$. And, given $n \in [1 : |f @ g|]$, we have that

$$(f @ g) n = \begin{cases} f n, & \text{if } n \in [1 : |f|], \\ g(n - |f|), & \text{if } n > |f|. \end{cases}$$

Using this fact, it is easy to prove that:

- $[]$ is the identity for concatenation: for all lists f ,

$$[] @ f = f = f @ [].$$

- Concatenation is associative: for all lists f, g, h ,

$$(f @ g) @ h = f @ (g @ h).$$

Because concatenation is associative, we can write $f @ g @ h$ without worrying where the parentheses go.

Given a set X , we write **List** X for the set of all X -lists, i.e., lists whose ranges are subsets of X , i.e., all of whose elements come from X . E.g., $[]$ and $[3, 5, 7]$ are elements of **List** \mathbb{N} , the set of all lists of natural numbers. It is easy to see that $[] \in \mathbf{List} X$, for all sets X , and that, for all sets X and $f, g \in \mathbf{List} X$, $f @ g \in \mathbf{List} X$.

1.1.9 Notes

In a traditional treatment of set theory, e.g., [End77], the natural numbers, integers, real numbers and ordered pairs (x, y) are encoded as sets. But for our purposes, it is clearer to suppress this detail.

Furthermore, in the traditional approach, \mathbb{N} is not a subset of \mathbb{Z} , and \mathbb{Z} is not a subset of \mathbb{R} . On the other hand, there are proper subsets of \mathbb{R} corresponding to \mathbb{N} and \mathbb{Z} , and these are what we take \mathbb{N} and \mathbb{Z} to be, so that $\mathbb{N} \subsetneq \mathbb{Z} \subsetneq \mathbb{R}$.

1.2 Induction

In the section, we consider several induction principles, i.e., methods for proving that every element x of some set A has a property $P(x)$.

1.2.1 Mathematical Induction

We begin with the familiar principle of mathematical induction, which is a basic result of set theory.

Theorem 1.2.1 (Principle of Mathematical Induction)

Suppose $P(n)$ is a property of a natural number n . If

(basis step)

$$P(0) \text{ and}$$

(inductive step)

$$\text{for all } n \in \mathbb{N}, \text{ if } (\dagger) P(n), \text{ then } P(n+1),$$

then,

$$\text{for all } n \in \mathbb{N}, P(n).$$

We refer to the formula (\dagger) as the *inductive hypothesis*. To use the principle to show that every natural number has property P , we carry out two steps. In the basis step, we show that 0 has property P . In the inductive step, we

assume that n is a natural number with property P . We then show that $n + 1$ has property P , without making any more assumptions about n .

Next we give an example of a mathematical induction.

Proposition 1.2.2

For all $n \in \mathbb{N}$, $3n^2 + 3n + 6$ is divisible by 6.

Proof. We proceed by mathematical induction.

(Basis Step) We have that $3 \cdot 0^2 + 3 \cdot 0 + 6 = 0 + 0 + 6 = 6 = 6 \cdot 1$. Thus $3 \cdot 0^2 + 3 \cdot 0 + 6$ is divisible by 6.

(Inductive Step) Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $3n^2 + 3n + 6$ is divisible by 6. Hence $3n^2 + 3n + 6 = 6m$, for some $m \in \mathbb{N}$. Thus, we have that

$$\begin{aligned} 3(n+1)^2 + 3(n+1) + 6 &= 3(n^2 + 2n + 1) + 3n + 3 + 6 \\ &= 3n^2 + 6n + 3 + 3n + 3 + 6 \\ &= (3n^2 + 3n + 6) + (6n + 6) \\ &= 6m + 6(n+1) \\ &= 6(m + (n+1)), \end{aligned}$$

showing that $3(n+1)^2 + 3(n+1) + 6$ is divisible by 6.

□

Exercise 1.2.3

Use Proposition 1.2.2 to prove, by mathematical induction, that, for all $n \in \mathbb{N}$, $n(n^2 + 5)$ is divisible by 6.

1.2.2 Strong Induction

Next, we consider the principle of strong induction.

Theorem 1.2.4 (Principle of Strong Induction)

Suppose $P(n)$ is a property of a natural number n . If

for all $n \in \mathbb{N}$,
if (\dagger) for all $m \in \mathbb{N}$, if $m < n$, then $P(m)$,
then $P(n)$,

then

for all $n \in \mathbb{N}$, $P(n)$.

We refer to the formula (\dagger) as the *inductive hypothesis*. To use the principle to show that every natural number has property P , we assume that n is a natural number, and that every natural number that is strictly smaller than n has property P . We then show that n has property P , without making any more assumptions about n .

It turns out that we can use mathematical induction to prove the validity of the principle of strong induction, by using a property $Q(n)$ derived from $P(n)$.

Proof. Suppose $P(n)$ is a property, and assume

$$\begin{aligned} (\dagger) \text{ for all } n \in \mathbb{N}, \\ \text{if for all } m \in \mathbb{N}, \text{ if } m < n, \text{ then } P(m), \\ \text{then } P(n). \end{aligned}$$

Let the property $Q(n)$ be

$$\text{for all } m \in \mathbb{N}, \text{ if } m < n, \text{ then } P(m).$$

First, we use mathematical induction to show that, for all $n \in \mathbb{N}$, $Q(n)$.

(Basis Step) We must show $Q(0)$. Suppose $m \in \mathbb{N}$ and $m < 0$. We must show that $P(m)$. Since $m < 0$ is a contradiction, we are allowed to conclude anything. So, we conclude $P(m)$.

(Inductive Step) Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $Q(n)$. We must show that $Q(n+1)$. Suppose $m \in \mathbb{N}$ and $m < n+1$. We must show that $P(m)$. Since $m \leq n$, there are two cases to consider.

- Suppose $m < n$. Because $Q(n)$, we have that $P(m)$.
- Suppose $m = n$. We must show that $P(n)$. By Property (\dagger) , it will suffice to show that

$$\text{for all } m \in \mathbb{N}, \text{ if } m < n, \text{ then } P(m).$$

But this formula is exactly $Q(n)$, and so we are done.

Now, we use the result of our mathematical induction to show that, for all $n \in \mathbb{N}$, $P(n)$. Suppose $n \in \mathbb{N}$. By our mathematical induction, we have $Q(n)$. By Property (\dagger) , it will suffice to show that

$$\text{for all } m \in \mathbb{N}, \text{ if } m < n, \text{ then } P(m).$$

But this formula is exactly $Q(n)$, and so we are done. \square

As an example use of the principle of strong induction, we will prove a proposition that we would normally take for granted:

Proposition 1.2.5

Every nonempty set of natural numbers has a least element.

Proof. Let X be a nonempty set of natural numbers.

We begin by using strong induction to show that, for all $n \in \mathbb{N}$,

if $n \in X$, then X has a least element.

Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: for all $m \in \mathbb{N}$, if $m < n$, then

if $m \in X$, then X has a least element.

We must show that

if $n \in X$, then X has a least element.

Suppose $n \in X$. It remains to show that X has a least element. If n is less-than-or-equal-to every element of X , then we are done. Otherwise, there is an $m \in X$ such that $m < n$. By the inductive hypothesis, we have that

if $m \in X$, then X has a least element.

But $m \in X$, and thus X has a least element. This completes our strong induction.

Now we use the result of our strong induction to prove that X has a least element. Since X is a nonempty subset of \mathbb{N} , there is an $n \in \mathbb{N}$ such that $n \in X$. By the result of our induction, we can conclude that

if $n \in X$, then X has a least element.

But $n \in X$, and thus X has a least element. \square

We conclude this subsection with one more proof using strong induction. Recall that a natural number is prime iff it is at least 2 and has no divisors other than 1 and itself.

Proposition 1.2.6

For all $n \in \mathbb{N}$,

if $n \geq 2$, then there are $m, l \in \mathbb{N}$ such that $n = ml$ and m is prime.

Proof. We proceed by strong induction. Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: for all $i \in \mathbb{N}$, if $i < n$, then

if $i \geq 2$, then there are $m, l \in \mathbb{N}$ such that $i = ml$ and m is prime.

We must show that

if $n \geq 2$, then there are $m, l \in \mathbb{N}$ such that $n = ml$ and m is prime.

Suppose $n \geq 2$. We must show that

there are $m, l \in \mathbb{N}$ such that $n = ml$ and m is prime.

There are two cases to consider.

- Suppose n is prime. Then $n, 1 \in \mathbb{N}$, $n = n1$ and n is prime.
- Suppose n is not prime. Since $n \geq 2$, it follows that $n = ij$ for some $i, j \in \mathbb{N}$ such that $1 < i < n$. Thus, by the inductive hypothesis, we have that

if $i \geq 2$, then there are $m, l \in \mathbb{N}$ such that $i = ml$ and m is prime.

But $i \geq 2$, and thus there are $m, l \in \mathbb{N}$ such that $i = ml$ and m is prime. Thus $m, lj \in \mathbb{N}$, $n = ij = (ml)j = m(lj)$ and m is prime.

□

Exercise 1.2.7

Use strong induction to prove that, for all $n \in \mathbb{N}$, if $n \geq 1$, then there are $i, j \in \mathbb{N}$ such that $n = 2^i(2j + 1)$.

1.2.3 Well-founded Induction

We can also do induction over a well-founded relation. A relation R on a set A is *well-founded* iff every nonempty subset X of A has an R -minimal element, where an element $x \in X$ is *R -minimal in X* iff there is no $y \in X$ such that $y R x$.

Given a relation R on a set A , and $x, y \in A$, we say that y is a *predecessor of x in R* iff $y R x$. Thus $x \in X$ is R -minimal in X iff none of x 's predecessors in R (there may be none) are in X .

For example, in Proposition 1.2.5, we proved that the strict total ordering $<$ on \mathbb{N} is well-founded. On the other hand, the strict total ordering $<$ on \mathbb{Z} is *not* well-founded, as \mathbb{Z} itself lacks a $<$ -minimal element.

Here's another negative example, showing that even if the underlying set is finite, the relation need not be well-founded. Let $A = \{0, 1\}$, and $R = \{(0, 1), (1, 0)\}$. Then 0 is the only predecessor of 1 in R , and 1 is the only predecessor of 0 in R . Of the nonempty subsets of A , we have that $\{0\}$ and $\{1\}$ have R -minimal elements. But consider A itself. Then 0 is not R -minimal in A , because $1 \in A$ and $1 R 0$. And 1 is not R -minimal in A , because $0 \in A$ and $0 R 1$. Hence R is not well-founded.

Theorem 1.2.8 (Principle of Well-founded Induction)

Suppose A is a set, R is a well-founded relation on A , and $P(x)$ is a property of an element $x \in A$. If

for all $x \in A$,
if (\dagger) for all $y \in A$, if $y R x$, then $P(y)$,
then $P(x)$,

then

for all $x \in A$, $P(x)$.

We refer to the formula (\dagger) as the *inductive hypothesis*. When $A = \mathbb{N}$ and $R = <$, this is the same as the principle of strong induction. But it's much more generally applicable than strong induction. Furthermore, the proof of this theorem isn't by induction.

Proof. Suppose A is a set, R is a well-founded relation on A , $P(x)$ is a property of an element $x \in A$, and

$$\begin{aligned} (\dagger) \text{ for all } x \in A, \\ \text{if for all } y \in A, \text{ if } y R x, \text{ then } P(y), \\ \text{then } P(x). \end{aligned}$$

We must show that, for all $x \in A$, $P(x)$.

Suppose, toward a contradiction, that it is not the case that, for all $x \in A$, $P(x)$. Hence there is an $x \in A$ such that $P(x)$ is false. Let $X = \{x \in A \mid P(x) \text{ is false}\}$. Thus $x \in X$, showing that X is non-empty. Because R is well-founded on A , it follows that there is a $z \in X$ that is R -minimal in X , i.e., such that there is no $y \in X$ such that $y R z$.

By (\dagger) and since $z \in X \subseteq A$, we have that

$$\begin{aligned} \text{if for all } y \in A, \text{ if } y R z, \text{ then } P(y), \\ \text{then } P(z). \end{aligned}$$

Because $z \in X$, we have that $P(z)$ is false. Thus, to obtain a contradiction, it will suffice to show that

$$\text{for all } y \in A, \text{ if } y R z, \text{ then } P(y).$$

Suppose $y \in A$, and $y R z$. We must show that $P(y)$. Because z is R -minimal in X , it follows that $y \notin X$. Thus $P(y)$.

Because we obtained our contradiction, we have that, for all $x \in A$, $P(x)$, as required. \square

We conclude this subsection by considering three ways of building well-founded relations. The first is by taking a subset of a well-founded relation:

Proposition 1.2.9

Suppose R is a well-founded relation on a set A . If $S \subseteq R$, then S is also a well-founded relation on A .

Proof. Suppose R is a well-founded relation on A , and $S \subseteq R$. Let X be a nonempty subset of A . Let $x \in X$ be R -minimal in X . Suppose, toward a contradiction, that x is not S -minimal in X . Thus there is a $y \in X$ such that $y S x$. But $S \subseteq R$, and thus $y R x$ —contradiction. Thus x is S -minimal in X , as required. \square

Let the *predecessor* relation $\mathbf{pred}_{\mathbb{N}}$ on \mathbb{N} be $\{(n, n+1) \mid n \in \mathbb{N}\}$. Thus, for all $n, m \in \mathbb{N}$, $m \mathbf{pred}_{\mathbb{N}} n$ iff m is exactly one less than n . Because $\mathbf{pred}_{\mathbb{N}} \subseteq <$, and $<$ is well-founded on \mathbb{N} , Proposition 1.2.9 tells us that $\mathbf{pred}_{\mathbb{N}}$ is well-founded on \mathbb{N} . 0 has no predecessors in $\mathbf{pred}_{\mathbb{N}}$, and, for all $n \in \mathbb{N}$, n is the only predecessor of $n+1$ in $\mathbf{pred}_{\mathbb{N}}$. Consequently, if a zero/non-zero case analysis is used, a proof by well-founded induction on $\mathbf{pred}_{\mathbb{N}}$ will look like a proof by mathematical induction.

Suppose A and B are sets, S is a relation on B , and $f \in A \rightarrow B$. Then the *inverse image of the relation S under f* , $f^{-1}(S)$, is the relation R on A defined by: for all $x, y \in A$, $x R y$ iff $f x S f y$.

Proposition 1.2.10

Suppose A and B are sets, S is a well-founded relation on B , and $f \in A \rightarrow B$. Then $f^{-1}(S)$ is a well-founded relation on A .

Proof. Let $R = f^{-1}(S)$. To see that R is well-founded on A , suppose X is a nonempty subset of A . We must show that there is an R -minimal element of X . Let $Y = f(X) = \{f x \mid x \in X\}$. Thus Y is a nonempty subset of B . Because S is well-founded on B , it follows that there is an S -minimal element y of Y . Hence $y = f x$ for some $x \in X$. Suppose, toward a contradiction, that x is not R -minimal in X . Thus there is an $x' \in X$ such that $x' R x$. Hence $f x' \in Y$ and $f x' S f x = y$, contradicting the S -minimality of y in Y . Thus x is R -minimal in X . \square

For example, let R be the relation on \mathbb{Z} such that, for all $n, m \in \mathbb{Z}$, $n R m$ iff $|n| < |m|$ (where we're writing $|\cdot|$ for the absolute value of an integer). Since $<$ is well-founded on \mathbb{N} , Proposition 1.2.10 tells us that R is well-founded on \mathbb{Z} . If we do a well-founded induction on R , when proving $P(n)$, for $n \in \mathbb{Z}$, we can make use of $P(m)$ for any $m \in \mathbb{Z}$ whose absolute value is strictly less than the absolute value of n . E.g., when proving $P(-10)$, we could make use of $P(5)$ or $P(-9)$.

If R and S are relations on sets A and B , respectively, then the *lexicographic relation of R and then S* , $R \triangleright S$, is the relation on $A \times B$ defined by: $(x, y) R \triangleright S (x', y')$ iff

- $x R x'$, or
- $x = x'$ and $y S y'$.

Proposition 1.2.11

Suppose R and S are well-founded relations on sets A and B , respectively. Then $R \triangleright S$ is a well-founded relation on $A \times B$.

Proof. Suppose T is a nonempty subset of $A \times B$. We must show that there is an $R \triangleright S$ -minimal element of T . By our assumption, T is a relation from A to

B . Because T is nonempty, it follows that $\mathbf{domain} T$ is a nonempty subset of A . Since R is a well-founded relation on A , it follows that there is an R -minimal $x \in \mathbf{domain} T$. Let $Y = \{y \in B \mid (x, y) \in T\}$. Because Y is a nonempty subset of B , and S is well-founded on B , there exists an S -minimal $y \in Y$. Thus $(x, y) \in T$.

Suppose, toward a contradiction, that (x, y) is not $R \triangleright S$ -minimal in T . Thus there are $x' \in A$ and $y' \in B$ such that $(x', y') \in T$ and $(x', y') R \triangleright S (x, y)$. Hence, there are two cases to consider.

- Suppose $x' R x$. Because $x' \in \mathbf{domain} T$, this contradicts the R -minimality of x in $\mathbf{domain} T$.
- Suppose $x' = x$ and $y' S y$. Thus $(x, y') = (x', y') \in T$, so that $y' \in Y$. But this contradicts the S -minimality of y in Y .

Because we obtained a contradiction in both cases, we have an overall contradiction. Thus (x, y) is $R \triangleright S$ -minimal in T . \square

For example, if we consider the strict total ordering $<$ on \mathbb{N} , then $< \triangleright <$ is a well-founded relation on $\mathbb{N} \times \mathbb{N}$. If we do a well-founded induction on $< \triangleright <$, when proving that $P((x, y))$ holds, we can use $P((x', y'))$, whenever $x' < x$ or $x = x'$ but $y' < y$.

Just as we abbreviate $A \times (B \times C)$ to $A \times B \times C$, and abbreviate $(x, (y, z))$ to (x, y, z) , we abbreviate $R \triangleright (S \triangleright T)$ to $R \triangleright S \triangleright T$. If R , S and T are well-founded relations on sets A , B and C , respectively, then $R \triangleright S \triangleright T$ is the well-founded relation on $A \times B \times C$ such that, for all $x \in A$, $y \in B$ and $z \in C$: $(x, y, z) R \triangleright S \triangleright T (x', y', z')$ iff

- $x R x'$, or
- $x = x'$ and $y S y'$, or
- $x = x'$, $y = y'$ and $z T z'$.

And we can do the analogous thing with four or more well-founded relations.

1.2.4 Notes

A typical book on formal language theory doesn't introduce well-founded relations and induction. But this material, and our treatment of well-founded recursion in the next section, will prove to be useful in subsequent chapters.

1.3 Inductive Definitions and Recursion

In this section, we will introduce and study ordered trees of arbitrary (finite) arity, whose nodes are labeled by elements of some set. In later chapters, we

will define regular expressions (in Chapter 3), parse trees (in Chapter ??) and programs (in Chapter ??) as restrictions of the trees we consider here.

The definition of the set of all trees over a set of labels is our first example of an inductive definition—a definition in which we collect together all of the values that can be constructed using a set of rules. We will see many examples of inductive definitions in the book. In this section, we will also see how to define functions by recursion.


1.3.1 Inductive Definition of Trees

Suppose X is a set. The set **Tree** X of X -trees is the least set such that, for all $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$, $(x, trs) \in \mathbf{Tree} X$. Recall that saying $trs \in \mathbf{List}(\mathbf{Tree} X)$ simply means that trs is a list all of whose elements come from **Tree** X .


Ignoring the adjective “least” for the moment, some example elements of **Tree** \mathbb{N} (the case when the set X of tree labels is \mathbb{N}) are:

- Since $3 \in \mathbb{N}$ and $[] \in \mathbf{List}(\mathbf{Tree} \mathbb{N})$, we have that $(3, []) \in \mathbf{Tree} \mathbb{N}$. For similar reasons, $(1, [])$, $(6, [])$ and all pairs of the form $(n, [])$, for $n \in \mathbb{N}$, are in **Tree** \mathbb{N} .
- Because $4 \in \mathbb{N}$, and $[(3, []), (1, []), (6, [])] \in \mathbf{List}(\mathbf{Tree} \mathbb{N})$, we have that $(4, [(3, []), (1, []), (6, [])]) \in \mathbf{Tree} \mathbb{N}$.
- And we can continue like this forever.


Trees are often easier to comprehend if they are drawn. We draw the X -tree $(x, [tr_1, \dots, tr_n])$ as




For example,



is the drawing of the \mathbb{N} -tree $(4, [(3, []), (1, []), (6, [])])$. And



is the \mathbb{N} -tree $(2, [(4, [(3, []), (1, []), (6, [])]), (9, [])])$. Consider the tree



again. We say that the *root label* of this tree is x , and tr_1 is the tree's *first child*, etc. We write **rootLabel** tr for the root label of tr . We often write a tree $(x, [tr_1, \dots, tr_n])$ in a more compact, linear syntax:

- $x(tr_1, \dots, tr_n)$, when $n \geq 1$, and
- x , when $n = 0$.

Thus $(2, [(4, [(3, []), (1, []), (6, [])]), (9, [])])$ can be written as $2(4(3, 1, 6), 9)$.

Consider the definition of **Tree** X again: the set **Tree** X of X -trees is the least set such that, (\dagger) for all $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$, $(x, trs) \in \mathbf{Tree} X$. Let's call a set U X -closed iff it satisfies property (\dagger) , where we have replaced **Tree** X by U : for all $x \in X$ and $trs \in \mathbf{List} U$, $(x, trs) \in U$. Thus the definition says that **Tree** X is the least X -closed set, where we've yet to say just what "least" means.

First we address the concern that there might not be any X -closed sets. A easy result of set theory says that:

Lemma 1.3.1

For all sets X , there is a set U such that $X \subseteq U$, $U \times U \subseteq U$ and $\mathbf{List} U \subseteq U$.

In other words, the lemma says that, given any set X , there exists a superset U of X such that every pair of elements of U is already an element of U , and every list of elements of U is already an element of U . Now, suppose X is a set, and let U be as in the lemma. To see that U is X -closed, suppose $x \in X$ and $trs \in \mathbf{List} U$. Thus $x \in U$ and $trs \in U$, so that $(x, trs) \in U$, as required.

E.g., we know that there is an \mathbb{Z} -closed set U . But $\mathbb{N} \subseteq \mathbb{Z}$, and thus U is also \mathbb{N} -closed. But if **Tree** \mathbb{N} turned out to be U , it would have elements like $(-5, [])$, which are not wanted.

To keep **Tree** X from having junk, we say that **Tree** X is the set U such that:

- U is X -closed, and
- for all X -closed sets V , $U \subseteq V$.

This is what we mean by saying that **Tree** X is the *least* X -closed set. It is our first example of an *inductive definition*, the least (relative to \subseteq) set satisfying a given set of rules saying that if some elements are already in the set, then some other elements are also in the set.

To see that there is a unique least X -closed set, we first prove the following lemma.

Lemma 1.3.2

Suppose X is a set and \mathcal{W} is a nonempty set of X -closed sets. Then $\bigcap \mathcal{W}$ is an X -closed set.

Proof. Suppose X is a set and \mathcal{W} is a nonempty set of X -closed sets. Because \mathcal{W} is nonempty, $\bigcap \mathcal{W}$ is well-defined. To see that $\bigcap \mathcal{W}$ is X -closed, suppose $x \in X$ and $trs \in \mathbf{List} \bigcap \mathcal{W}$. We must show that $(x, trs) \in \bigcap \mathcal{W}$, i.e., that $(x, trs) \in W$, for all $W \in \mathcal{W}$. Suppose $W \in \mathcal{W}$. We must show that $(x, trs) \in W$. Because $W \in \mathcal{W}$, we have that $\bigcap \mathcal{W} \subseteq W$, so that $\mathbf{List} \bigcap \mathcal{W} \subseteq \mathbf{List} W$. Thus, since $trs \in \mathbf{List} \bigcap \mathcal{W}$, it follows that $trs \in \mathbf{List} W$. But W is X -closed, and thus $(x, trs) \in W$, as required. \square

As explained above, we have that there is an X -closed set, V . Let \mathcal{W} be the set of all subsets of V that are X -closed. Thus \mathcal{W} is a nonempty set of X -closed sets, since $V \in \mathcal{W}$. Let $U = \bigcap \mathcal{W}$. By Lemma 1.3.2, we have that U is X -closed. To see that $U \subseteq T$ for all X -closed sets T , suppose T is X -closed. By Lemma 1.3.2, we have that $V \cap T = \bigcap \{V, T\}$ is X -closed. And $V \cap T \subseteq V$, so that $V \cap T \in \mathcal{W}$. Hence $U = \bigcap \mathcal{W} \subseteq V \cap T \subseteq T$, as required. Finally, suppose that U' is also an X -closed set such that, for all X -closed sets T , $U' \subseteq T$. Then $U \subseteq U' \subseteq U$, showing that $U' = U$. Thus U is the least X -closed set.

Suppose X is a set, $x, x' \in X$ and $trs, trs' \in \mathbf{List}(\mathbf{Tree} X)$. It is easy to see that $(x, trs) = (x', trs')$ iff $x = x'$, $|trs| = |trs'|$ and, for all $i \in [1 : |trs|]$, $trs i = trs' i$.

Because trees are defined via an inductive definition, we get an induction principle for trees almost for free:

Theorem 1.3.3 (Principle of Induction on Trees)

Suppose X is a set and $P(tr)$ is a property of an element $tr \in \mathbf{Tree} X$. If

for all $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$,
if (\dagger) for all $i \in [1 : |trs|]$, $P(trs i)$,
then $P((x, trs))$,

then

for all $tr \in \mathbf{Tree} X$, $P(tr)$.

We refer to (\dagger) as the inductive hypothesis.

Proof. Suppose X is a set, $P(tr)$ is a property of an element $tr \in \mathbf{Tree} X$, and

(\ddagger) for all $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$,
if for all $i \in [1 : |trs|]$, $P(trs i)$,
then $P((x, trs))$.

We must show that

$$\text{for all } tr \in \mathbf{Tree} X, P(tr).$$

Let $U = \{ tr \in \mathbf{Tree} X \mid P(tr) \}$. We will show that U is X -closed. Suppose $x \in X$ and $trs \in \mathbf{List} U$. We must show that $(x, trs) \in U$. It will suffice to show that $P((x, trs))$. By (\ddagger) , it will suffice to show that, for all $i \in [1 : |trs|]$, $P(trs\ i)$. Suppose $i \in [1 : |trs|]$. We must show that $P(trs\ i)$. Because $trs \in \mathbf{List} U$, we have that $trs\ i \in U$. Hence $P(trs\ i)$, as required.

Because U is X -closed, we have that $\mathbf{Tree} X \subseteq U$, as $\mathbf{Tree} X$ is the least X -closed set. Hence, for all $tr \in \mathbf{Tree} X$, $tr \in U$, so that, for all $tr \in \mathbf{Tree} X$, $P(tr)$. \square

Using our induction principle, we can now prove that every X -tree can be “deconstructed” into an element of X paired with a list of X -trees:

Proposition 1.3.4

Suppose X is a set. For all $tr \in \mathbf{Tree} X$, there are $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$ such that $tr = (x, trs)$.

Proof. Suppose X is a set. We use induction on trees to prove that, for all $tr \in \mathbf{Tree} X$, there are $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$ such that $tr = (x, trs)$. Suppose $x \in X$, $trs \in \mathbf{List}(\mathbf{Tree} X)$, and assume the inductive hypothesis: for all $i \in [1 : |trs|]$, there are $x' \in X$ and $trs' \in \mathbf{List}(\mathbf{Tree} X)$ such that $trs\ i = (x', trs')$. We must show that there are $x' \in X$ and $trs' \in \mathbf{List}(\mathbf{Tree} X)$ such that $(x, trs) = (x', trs')$. And this holds, since $x \in X$, $trs \in \mathbf{List}(\mathbf{Tree} X)$ and $(x, trs) = (x, trs)$. \square

Note that the preceding induction makes no use of its inductive hypothesis, and yet the induction is still necessary.

Suppose X is a set. Let the predecessor relation $\mathbf{pred}_{\mathbf{Tree} X}$ on $\mathbf{Tree} X$ be the set of all pairs of X -trees (tr, tr') such that there are $x \in X$ and $trs' \in \mathbf{List}(\mathbf{Tree} X)$ such that $tr' = (x, trs')$ and $trs'\ i = tr$ for some $i \in [1 : |trs'|]$, i.e., such that tr is one of the children of tr' . Thus the predecessors of a tree $(x, [tr_1, \dots, tr_n])$ are its children tr_1, \dots, tr_n .

Proposition 1.3.5

If X is a set, then $\mathbf{pred}_{\mathbf{Tree} X}$ is a well-founded relation on $\mathbf{Tree} X$.

Proof. Suppose X is a set and Y is a nonempty subset of $\mathbf{Tree} X$. Mimicking Proposition 1.2.5, we can use the principle of induction on trees to prove that, for all $tr \in \mathbf{Tree} X$, if $tr \in Y$, then Y has a $\mathbf{pred}_{\mathbf{Tree} X}$ -minimal element. Because Y is nonempty, we can conclude that Y has a $\mathbf{pred}_{\mathbf{Tree} X}$ -minimal element. \square

Exercise 1.3.6

Do the induction on trees used by the preceding proof.

1.3.2 Recursion

Suppose R is a well-founded relation on a set A . We can define a function f from A to a set B by *well-founded recursion on R* . The idea is simple: when f is called with an element $x \in A$, it may call itself recursively on any of the predecessors of x in R . Typically, such a definition will be concrete enough that we can regard it as defining an algorithm as well as a function.

If R is a well-founded relation on a set A , and B is a set, then we write $\mathbf{Rec}_{A,R,B}$ for

$$\{ (x, f) \mid x \in A \text{ and } f \in \{ y \in A \mid y R x \} \rightarrow B \} \rightarrow B.$$

An element F of $\mathbf{Rec}_{A,R,B}$ may only be called with a pair (x, f) such that $x \in A$ and f is a function from the predecessors of x in R to B . Intuitively, F 's job is to transform x into an element of B , using f to carry out recursive calls.

Theorem 1.3.7 (Well-Founded Recursion)

Suppose R is a well-founded relation on a set A , B is a set, and $F \in \mathbf{Rec}_{A,R,B}$. There is a unique $f \in A \rightarrow B$ such that, for all $x \in A$,

$$f x = F(x, f|_{\{y \in A \mid y R x\}}).$$

The second argument to F in the definition of f is the restriction of f to the predecessors of x in R , i.e., it's the subset of f whose domain is $\{y \in A \mid y R x\}$.

If we can understand F as an algorithm, then we can understand the definition of f as an algorithm. If we call f with an $x \in A$, then F may return an element of B without consulting its second argument. Alternatively, it may call this second argument with a predecessor of x in R . This is a recursive call of f , which must complete before F continues. Once it does complete, F may make more recursive calls, but must eventually return an element of B .

Proof. We start with an inductive definition: let f be the least subset of $A \times B$ such that, for all $x \in A$ and $g \in \{y \in A \mid y R x\} \rightarrow B$,

$$\text{if } g \subseteq f, \text{ then } (x, F(x, g)) \in f.$$

We say that a subset U of $A \times B$ is *closed* iff, for all $x \in A$ and $g \in \{y \in A \mid y R x\} \rightarrow B$,

$$\text{if } g \subseteq U, \text{ then } (x, F(x, g)) \in U.$$

Thus, we are saying that f is the least closed subset of $A \times B$. This definition is well-defined because $A \times B$ is closed, and if \mathcal{W} is a nonempty set of closed subsets of $A \times B$, then $\bigcap \mathcal{W}$ is also closed. Thus we can let f be the intersection of all closed subsets of $A \times B$.

Thus f is a relation from A to B . An easy well-founded induction on R suffices to show that, for all $x \in A$, $x \in \mathbf{domain} f$. Suppose $x \in A$, and assume

the inductive hypothesis: for all $y \in A$, if $y R x$, then $y \in \mathbf{domain} f$. We must show that $x \in \mathbf{domain} f$. By the inductive hypothesis, we have that for all $y \in \{y \in A \mid y R x\}$, there is a $z \in B$ such that $(y, z) \in f$. Thus there is a subset g of f such that $g \in \{y \in A \mid y R x\} \rightarrow B$. (Since we don't know at this point that f is a function, we are using the Axiom of Choice in this last step.) Hence $(x, F(x, g)) \in f$, showing that $x \in \mathbf{domain} f$. Thus $\mathbf{domain} f = A$.

Next we show that f is a function. Let h be

$$\{(x, y) \in f \mid \text{for all } y' \in B, \text{ if } (x, y') \in f, \text{ then } y = y'\}.$$

If we can show that h is closed, then we will have that $f \subseteq h$, because f is the least closed set, and thus we'll be able to conclude that f is a function. To show that h is closed, suppose $x \in A$, $g \in \{y \in A \mid y R x\} \rightarrow B$ and $g \subseteq h$. We must show that $(x, F(x, g)) \in h$. It will suffice to show that, for all $y' \in B$, if $(x, y') \in f$, then $F(x, g) = y'$. Suppose $y' \in B$ and $(x, y') \in f$. We must show that $F(x, g) = y'$. Because $(x, y') \in f$, and f is the least closed subset of $A \times B$, there must be a $g' \in \{y \in A \mid y R x\} \rightarrow B$ such that $g' \subseteq f$ and $y' = F(x, g')$. Thus it will suffice to show that $F(x, g) = F(x, g')$, which will follow from showing that $g = g'$, i.e., for all $z \in \{y \in A \mid y R x\}$, $gz = g'z$. Suppose $z \in \{y \in A \mid y R x\}$. We must show that $gz = g'z$. Since $z \in \{y \in A \mid y R x\}$, we have that $z \in A$ and $z R x$. Because $(z, gz) \in g \subseteq h$, we have that $(z, gz) \in h$. Since $(z, g'z) \in g' \subseteq f$, we have that $(z, g'z) \in f$. Hence, by the definition of h , we have that $gz = g'z$, as required.

Summarizing, we know that $f \in A \rightarrow B$. Next, we must show that, for all $x \in A$,

$$fx = F(x, f|\{y \in A \mid y R x\}).$$

Suppose $x \in A$. Because $f|\{y \in A \mid y R x\} \in \{y \in A \mid y R x\} \rightarrow B$, we have that $(x, F(x, f|\{y \in A \mid y R x\})) \in f$, so that $fx = F(x, f|\{y \in A \mid y R x\})$.

Finally, suppose that $f' \in A \rightarrow B$ and for all $x \in A$,

$$f'x = F(x, f'|\{y \in A \mid y R x\}).$$

To see that $f = f'$, it will suffice to show that, for all $x \in A$, $fx = f'x$. We proceed by well-founded induction on R . Suppose $x \in A$ and assume the inductive hypothesis: for all $y \in A$, if $y R x$, then $fy = f'y$. We must show that $fx = f'x$. By the inductive hypothesis, we have that $f|\{y \in A \mid y R x\} = f'|\{y \in A \mid y R x\}$. Thus

$$\begin{aligned} fx &= F(x, f|\{y \in A \mid y R x\}) \\ &= F(x, f'|\{y \in A \mid y R x\}) \\ &= f'x, \end{aligned}$$

as required. \square

Here are some examples of well-founded recursion:


- If we define $f \in \mathbb{N} \rightarrow B$ by well-founded recursion on $<$, then, when f is called with $n \in \mathbb{N}$, it may call itself recursively on any strictly smaller natural numbers. In the case $n = 0$, it can't make any recursive calls.
- If we define $f \in \mathbb{N} \rightarrow B$ by well-founded recursion on the predecessor relation $\mathbf{pred}_{\mathbb{N}}$, then when f is called with $n \in \mathbb{N}$, it may call itself recursively on $n - 1$, in the case when $n \geq 1$, and may make no recursive calls, when $n = 0$.

Thus, if such a definition case-splits according to whether its input is 0 or not, it can be split into two parts:

- $f\,0 = \dots$;
- for all $n \in \mathbb{N}$, $f(n+1) = \dots f\,n \dots$.

We say that such a definition is by *recursion on \mathbb{N}* .


- If we define $f \in \mathbf{Tree}\,X \rightarrow B$ by well-founded recursion on the predecessor relation $\mathbf{pred}_{\mathbf{Tree}\,X}$, then when f is called on an X -tree $(x, [tr_1, \dots, tr_n])$, it may call itself recursively on any of tr_1, \dots, tr_n . When $n = 0$, it may make no recursive calls. We say that such a definition is by *structural recursion*.
- We may define the *size* of an X -tree $(x, [tr_1, \dots, tr_n])$ by summing the recursively computed sizes of tr_1, \dots, tr_n , and then adding 1. (When $n = 0$, the sum of no sizes is 0, and so we get 1.) Then, e.g., the size of



is 6. This defines a function **size** $\in \mathbf{Tree}\,X \rightarrow \mathbb{N}$.

- We may define the *number of leaves* of an X -tree $(x, [tr_1, \dots, tr_n])$ as
 - 1, when $n = 0$, and
 - the sum of the recursively computed numbers of leaves of tr_1, \dots, tr_n , when $n \geq 1$.


Then, e.g., the number of leaves of



is 4. This defines a function $\mathbf{numLeaves} \in \mathbf{Tree} X \rightarrow \mathbb{N}$.

- We may define the *height* of an X -tree $(x, [tr_1, \dots, tr_n])$ as
 - 0, when $n = 0$, and
 - 1 plus the maximum of the recursively computed heights of tr_1, \dots, tr_n , when $n \geq 1$.

E.g., the height of



is 2. This defines a function $\mathbf{height} \in \mathbf{Tree} X \rightarrow \mathbb{N}$.

- Given a set X , we can define a well-founded relation $\mathbf{size}_{\mathbf{Tree} X}$ on $\mathbf{Tree} X$ by: for all $tr, tr' \in \mathbf{Tree} X$, $tr \mathbf{size}_{\mathbf{Tree} X} tr'$ iff $\mathbf{size} tr < \mathbf{size} tr'$. (This is an application of Proposition 1.2.10.)

If we define a function $f \in \mathbf{Tree} X \rightarrow B$ by well-founded recursion on $\mathbf{size}_{\mathbf{Tree} X}$, when f is called with an X -tree tr , it may call itself recursively on any X -trees with strictly smaller sizes.

- Given a set X , we can define a well-founded relation $\mathbf{height}_{\mathbf{Tree} X}$ on $\mathbf{Tree} X$ by: for all $tr, tr' \in \mathbf{Tree} X$, $tr \mathbf{height}_{\mathbf{Tree} X} tr'$ iff $\mathbf{height} tr < \mathbf{height} tr'$.

If we define a function $f \in \mathbf{Tree} X \rightarrow B$ by well-founded recursion on $\mathbf{height}_{\mathbf{Tree} X}$, when f is called with an X -tree tr , it may call itself recursively on any X -trees with strictly smaller heights.


- Given a set X , we can define a well-founded relation $\mathbf{length}_{\mathbf{List} X}$ on $\mathbf{List} X$ by: for all $xs, ys \in \mathbf{List} X$, $xs \mathbf{length}_{\mathbf{List} X} ys$ iff $|xs| < |ys|$.

If we define a function $f \in \mathbf{List} X \rightarrow B$ by well-founded recursion on $\mathbf{length}_{\mathbf{List} X}$, when f is called with an X -list xs , it may call itself recursively on any X -lists with strictly smaller lengths.

1.3.3 Paths in Trees

We can think of an $\mathbb{N} - \{0\}$ -list $[n_1, n_2, \dots, n_m]$ as a *path* through an X -tree tr : one starts with tr itself, goes to the n_1 -th child of tr , selects the n_2 -th child of that tree, etc., stopping when the list is exhausted.

Consider the \mathbb{N} -tree



Then:

- $[]$ takes us to the whole tree.
- $[1]$ takes us to the tree $4(3, 1, 6)$.
- $[1, 3]$ takes us to the tree 6.
- $[1, 4]$ takes us to no tree.

We define the valid paths of an X -tree via structural recursion. For a set X , we define $\mathbf{validPaths}_X \in \mathbf{Tree} X \rightarrow \mathbf{List}(\mathbb{N} - \{0\})$ (we often drop the subscript X , when it's clear from the context) by: for all $x \in X$ and $trs \in \mathbf{List}(\mathbf{Tree} X)$, $\mathbf{validPaths}(x, trs)$ is

$$\{[]\} \cup \{[i] @ xs \mid i \in [1 : |trs|] \text{ and } xs \in \mathbf{validPaths}(trs\ i)\}.$$

We say that $xs \in \mathbf{List}(\mathbb{N} - \{0\})$ is a *valid path* for an X -tree tr iff $xs \in \mathbf{validPaths}\ tr$. For example, $\mathbf{validPaths}(3(4, 1(7), 6)) = \{[], [1], [2], [2, 1], [3]\}$. Thus, e.g., $[2, 1]$ is a valid path for $3(4, 1(7), 6)$, whereas $[2, 2]$ and $[4]$ are not valid paths for this tree.

Now, we define a function **select** that takes in an X -tree tr and a valid path xs for tr , and returns the subtree of tr pointed to by xs . Let Y_X be

$$\{(tr, xs) \in \mathbf{Tree} X \times \mathbf{List}(\mathbb{N} - \{0\}) \mid xs \text{ is a valid path for } tr\}.$$

Let the relation R on Y_X be

$$\{((tr, xs), (tr', xs')) \in Y_X \times Y_X \mid |xs| < |xs'|\}.$$

By Proposition 1.2.10, we have that R is well-founded, and so we can use well-founded recursion on R to define a function \mathbf{select}_X (we often drop the subscript X) from Y_X to $\mathbf{Tree} X$. Suppose, we are given an input $((x, trs), xs) \in Y$. If xs is $[]$, then we return (x, trs) . Otherwise, $xs = [i] @ xs'$, for some $i \in \mathbb{N} - \{0\}$ and $xs' \in \mathbf{List}(\mathbb{N} - \{0\})$. Because $((x, trs), xs) \in Y$, it follows that $i \in [1 : |trs|]$ and xs' is a valid path for $trs\ i$. Thus $(trs\ i, xs')$ is in Y , and is a predecessor of $((x, trs), xs)$ in R , so that we can call ourselves recursively on $(trs\ i, xs')$ and return the resulting tree. For example $\mathbf{select}(4(3(2, 1(7)), 6), []) = 4(3(2, 1(7)), 6)$ and $\mathbf{select}(4(3(2, 1(7)), 6), [1, 2]) = 1(7)$.

We say that an X -tree tr' is a *subtree* of an X -tree tr iff there is a valid path xs for tr such that $tr' = \mathbf{select}(tr, xs)$. And tr' is a *leaf* of tr iff tr' is a subtree of tr and tr' has no children.

Finally, we can define a function that takes in a X -tree tr , a valid path xs for tr , and an X -tree tr' , and returns the result of replacing the subtree of tr pointed to by xs with tr' . Let Y_X be

$$\{ (tr, xs, tr') \in \mathbf{Tree} X \times \mathbf{List}(\mathbb{N} - \{0\}) \times \mathbf{Tree} X \mid xs \text{ is a valid path for } tr \}.$$

We use well-founded recursion on the size of the second components (the paths) of the elements of Y_X to define a function **update** _{X} (we often drop the subscript) from Y_X to $\mathbf{Tree} X$. Suppose, we are given an input $((x, trs), xs, tr') \in Y$. If xs is $[]$, then we return tr' . Otherwise, $xs = [i] @ xs'$, for some $i \in \mathbb{N} - \{0\}$ and $xs' \in \mathbf{List}(\mathbb{N} - \{0\})$. Because $((x, trs), xs, tr') \in Y$, it follows that $i \in [1 : |trs|]$ and xs' is a valid path for $trs i$. Thus $(trs i, xs', tr')$ is in Y , and $|xs'| < |xs|$. Hence, we can let tr'' be the result of calling ourselves recursively on $(trs i, xs', tr')$. Finally, we can return (x, trs') , where $trs' = trs[i \mapsto tr'']$ (which is the same as trs , excepts that its i th element is tr''). For example **update** $(4(3(2, 1(7)), 6), [], 3(7, 8)) = 3(7, 8)$ and **update** $(4(3(2, 1(7)), 6), [1, 2], 3(7, 8)) = 4(3(2, 3(7, 8)), 6)$.

1.3.4 Notes

Our treatment of trees, inductive definition, and well-founded recursion is more formal than what one finds in typical books on formal language theory. But those with a background in set theory will find nothing surprising in this section, and our foundational approach will serve the reader well in later chapters.

Chapter 2

Formal Languages

In this chapter we say what symbols, strings, alphabets and (formal) languages are, show how to use various induction principles to prove language equalities, and give an introduction to the Forlan toolset. In subsequent chapters, we will study four more restricted kinds of languages: the regular (Chapter 3), context-free (Chapter ??), recursive and recursively enumerable (Chapter ??) languages.

2.1 Symbols, Strings, Alphabets and (Formal) Languages

In this section, we define the basic notions of the subject: symbols, strings, alphabets and (formal) languages. In most presentations of formal language theory, the “symbols” that make up strings are allowed to be arbitrary elements of the mathematical universe. This is convenient in some ways, but it means that, e.g., the collection of all strings is too “big” to be a set. Furthermore, if we were to adopt this convention, we wouldn’t be able to have notation in Forlan for all strings and symbols. These considerations lead us to the following definition.

2.1.1 Symbols

The set **Char** of *symbol characters* consists of the following 65 elements:

- the comma (“,”);
- the *digits* 0–9;
- the *letters* a–z and A–Z; and
- the angle brackets (“<” and “>”).

We order **Char** as follows:

$$, < 0 < \cdots < 9 < a < \cdots < z < A < \cdots < Z < \langle < \rangle.$$

The set **Sym** of *symbols* is the least subset of **List Char** such that:

- for all digits and letters c , $[c] \in \mathbf{Sym}$; and
- for all $n \in \mathbb{N}$ and $x_1, \dots, x_n \in \{[,]\} \cup \mathbf{Sym}$,

$$[\langle \rangle @ x_1 @ \dots @ x_n @ \rangle] \in \mathbf{Sym}.$$

This is an inductive definition (see Section 1.3). **Sym** consists of just those lists of symbol characters that can be built using the above, two rules. For example, $[9]$, $[\langle, \rangle]$, $[\langle, i, d, \rangle]$ and $[\langle, \langle, a, ,, \rangle, b, \rangle]$ are symbols. On the other hand, $[\langle, \rangle, \rangle]$ is not a symbol.

We can prove by induction that, for all $z \in \mathbf{Sym}$, for all $x, y \in \mathbf{List Char}$, if $z = x @ y \in \mathbf{Sym}$, then:

- if $x \in \mathbf{Sym}$, then $y = []$;
- if $y \in \mathbf{Sym}$, then $x = []$.

Thus a symbol never starts or ends with another symbol.

We normally abbreviate a symbol $[c_1, \dots, c_n]$ to $c_1 \dots c_n$, so that 9 , $\langle \rangle$, $\langle id \rangle$ and $\langle \langle a, \rangle b \rangle$ are symbols. And if x and y are elements of **List Char**, we typically abbreviate $x @ y$ to xy .

Whenever possible, we will use the mathematical variables a , b and c to name symbols. We order **Sym** first by length, and then lexicographically (in dictionary order). So, we have that

$$0 < \dots < 9 < A < \dots < Z < a < \dots < z,$$

and, e.g.,

$$z < \langle be \rangle < \langle by \rangle < \langle on \rangle < \langle can \rangle < \langle con \rangle.$$

Obviously, **Sym** is infinite, but is it countably infinite? The answer is “yes”, because we can enumerate the symbols in order.

2.1.2 Strings

A *string* is a list of symbols. Whenever possible, we will use the mathematical variables u , v , w , x , y and z to name strings.

We typically abbreviate the empty string $[]$ to $\%$, and abbreviate $[a_1, \dots, a_n]$ to $a_1 \dots a_n$, when $n \geq 1$. For example $[0, \langle 0 \rangle, 1, \langle \langle, \rangle \rangle]$ is abbreviated to $0\langle 0 \rangle 1\langle \langle, \rangle \rangle$. We name the empty string by $\%$, instead of following convention and using ϵ , since this symbol can also be used in Forlan.

We write **Str** for **List Sym**, the set of all strings. We order **Str** first by length and then lexicographically, using our order on **Sym**. Thus, e.g.,

$$\% < ab < a\langle be \rangle < a\langle by \rangle < \langle can \rangle \langle be \rangle < abc.$$

Since every string can be written as a finite sequence of ASCII characters, it follows that **Str** is countably infinite.

Because strings are lists, we have that $|x|$ is the *length* of a string x , and that $x @ y$ is the *concatenation* of strings x and y . We typically abbreviate $x @ y$ to xy . For example:

- $|\%| = |[[]]| = 0$;
- $|0\langle 0 \rangle 1\langle \langle , \rangle \rangle| = |[0, \langle 0 \rangle, 1, \langle \langle , \rangle \rangle]| = 4$; and
- $(01)(00) = [0, 1] @ [0, 0] = [0, 1, 0, 0] = 0100$.

From our study of lists, we know that:

- Concatenation is associative: for all $x, y, z \in \mathbf{Str}$,

$$(xy)z = x(yz).$$

- $\%$ is the identity for concatenation: for all $x \in \mathbf{Str}$,

$$\%x = x = x\%.$$

It is easy to see that, for all $x, y, x', y' \in \mathbf{Str}$:

- $xy = xy'$ iff $y = y'$; and
- $xy = x'y$ iff $x = x'$.

We define the string x^n *formed by raising a string x to the power $n \in \mathbb{N}$* by recursion on n :

$$\begin{aligned} x^0 &= \%, \text{ for all } x \in \mathbf{Str}; \text{ and} \\ x^{n+1} &= xx^n, \text{ for all } x \in \mathbf{Str} \text{ and } n \in \mathbb{N}. \end{aligned}$$

We assign this operation higher precedence than concatenation, so that xx^n means $x(x^n)$ in the above definition. For example, we have that

$$(ab)^2 = (ab)(ab)^1 = (ab)(ab)(ab)^0 = (ab)(ab)\% = abab.$$

Proposition 2.1.1

For all $x \in \mathbf{Str}$ and $n, m \in \mathbb{N}$, $x^{n+m} = x^n x^m$.

Proof. Suppose $x \in \mathbf{Str}$ and $m \in \mathbb{N}$. We use mathematical induction to show that, for all $n \in \mathbb{N}$, $x^{n+m} = x^n x^m$.

(Basis Step) We have that $x^{0+m} = x^m = \%x^m = x^0 x^m$.

(Inductive Step) Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $x^{n+m} = x^n x^m$. We must show that $x^{(n+1)+m} = x^{n+1} x^m$. We have that

$$\begin{aligned}
 x^{(n+1)+m} &= x^{(n+m)+1} \\
 &= x x^{n+m} && \text{(definition of } x^{(n+m)+1}\text{)} \\
 &= x(x^n x^m) && \text{(inductive hypothesis)} \\
 &= (x x^n) x^m \\
 &= x^{n+1} x^m && \text{(definition of } x^{n+1}\text{)}.
 \end{aligned}$$

□

Thus, if $x \in \mathbf{Str}$ and $n \in \mathbb{N}$, then

$$x^{n+1} = x x^n \quad \text{(definition),}$$

and

$$x^{n+1} = x^n x^1 = x^n x \quad \text{(Proposition 2.1.1).}$$

Next, we consider the prefix, suffix and substring relations on strings. Suppose x and y are strings. We say that:

- x is a *prefix* of y iff $y = xv$ for some $v \in \mathbf{Str}$;
- x is a *suffix* of y iff $y = ux$ for some $u \in \mathbf{Str}$; and
- x is a *substring* of y iff $y = uxv$ for some $u, v \in \mathbf{Str}$.

In other words, x is a prefix of y iff x is an initial part of y , x is a suffix of y iff x is a trailing part of y , and x is a substring of y iff x appears in the middle of y . But note that the strings u and v can be empty in these definitions. Thus, e.g., a string x is always a prefix of itself, since $x = x\%$. A prefix, suffix or substring of a string other than the string itself is called *proper*.

For example:

- $\%$ is a proper prefix, suffix and substring of \mathbf{ab} ;
- \mathbf{a} is a proper prefix and substring of \mathbf{ab} ;
- \mathbf{b} is a proper suffix and substring of \mathbf{ab} ; and
- \mathbf{ab} is a (non-proper) prefix, suffix and substring of \mathbf{ab} .

Proposition 2.1.2

For all $x, y, x', y' \in \mathbf{Str}$, $xy = x'y'$ iff either

- $xu = x'$ and $y = uy'$, for some $u \in \mathbf{Str}$, or

- $x'u = x$ and $y' = uy$, for some $u \in \mathbf{Str}$.

Proof. Straightforward. \square

As a consequence of this proposition, we have that:

- For all $x, x', y' \in \mathbf{Str}$, x is a prefix of $x'y'$ iff either
 - x is a prefix of x' , or
 - $x = x'u$, for some prefix u of y' .
- For all $x, x', y' \in \mathbf{Str}$, x is a suffix of $x'y'$ iff either
 - x is a suffix of y' , or
 - $x = uy'$, for some suffix u of x' .
- For all $x, x', y' \in \mathbf{Str}$, x is a substring of $x'y'$ iff either
 - x is a substring of x' , or
 - $x = uv$, for some $u, v \in \mathbf{Str}$ such that u is a suffix of x' and v is a prefix of y' , or
 - x is a substring of y' .

2.1.3 Alphabets

Having said what symbols and strings are, we now come to alphabets. An *alphabet* is a finite subset of \mathbf{Sym} . We use Σ (upper case Greek letter sigma) to name alphabets. For example, \emptyset , $\{0\}$ and $\{0, 1\}$ are alphabets. We write \mathbf{Alp} for the set of all alphabets. \mathbf{Alp} is countably infinite.

We define $\mathbf{alphabet} \in \mathbf{Str} \rightarrow \mathbf{Alp}$ by recursion on (the length of) strings:

$$\mathbf{alphabet} \% = \emptyset,$$

$$\mathbf{alphabet}(ax) = \{a\} \cup \mathbf{alphabet} x, \text{ for all } a \in \mathbf{Sym} \text{ and } x \in \mathbf{Str}.$$

I.e., $\mathbf{alphabet} w$ consists of all of the symbols occurring in the string w . E.g., $\mathbf{alphabet}(01101) = \{0, 1\}$. Because the string x appears on the right side of ax in the rule $\mathbf{alphabet}(ax) = \{a\} \cup \mathbf{alphabet} x$, we call this *right* recursion. (Since \cup is associative and commutative, it would have been equivalent to use *left* recursion, $\mathbf{alphabet}(xa) = \{a\} \cup \mathbf{alphabet} x$.) We say that $\mathbf{alphabet} x$ is the *alphabet of* x .

If Σ is an alphabet, then we write Σ^* for $\mathbf{List} \Sigma$. I.e., Σ^* consists of all of the strings that can be built using the symbols of Σ . For example, the elements of $\{0, 1\}^* = \mathbf{List} \{0, 1\}$ are:

$$\%, 0, 1, 00, 01, 10, 11, 000, \dots$$

2.1.4 Languages

We say that L is a *formal language* (or just *language*) iff $L \subseteq \Sigma^*$, for some $\Sigma \in \mathbf{Alp}$. In other words, a language is a set of strings over some alphabet. If $\Sigma \in \mathbf{Alp}$, then we say that L is a Σ -*language* iff $L \subseteq \Sigma^*$.

Here are some example languages (all are $\{0, 1\}$ -languages):

- \emptyset ;
- $\{0, 1\}^*$;
- $\{010, 1001, 1101\}$;
- $\{0^n 1^n \mid n \in \mathbb{N}\} = \{0^0 1^0, 0^1 1^1, 0^2 1^2, \dots\} = \{\epsilon, 01, 0011, \dots\}$; and
- $\{w \in \{0, 1\}^* \mid w \text{ is a palindrome}\}$.

(A *palindrome* is a string that reads the same backwards and forwards, i.e., that is equal to its own reversal.) On the other hand, the set of strings $X = \{\langle \rangle, \langle 0 \rangle, \langle 00 \rangle, \dots\}$, is not a language, since it involves infinitely many symbols, i.e., since there is no alphabet Σ such that $X \subseteq \Sigma^*$.

Since \mathbf{Str} is countably infinite and every language is a subset of \mathbf{Str} , it follows that every language is countable. Furthermore, Σ^* is countably infinite, as long as the alphabet Σ is nonempty ($\emptyset^* = \{\epsilon\}$).

We write \mathbf{Lan} for the set of all languages. It turns out that \mathbf{Lan} is uncountable. In fact even $\mathcal{P}(\{0\}^*)$, the set of all $\{0\}$ -languages, has the same size as $\mathcal{P}(\mathbb{N})$, and is thus uncountable.

Exercise 2.1.3

Show that $\mathcal{P}(\mathbb{N})$ has the same size as $\mathcal{P}(\{0\}^*)$.

Given a language L , we write $\mathbf{alphabet} L$ for the *alphabet*

$$\bigcup \{ \mathbf{alphabet} w \mid w \in L \}.$$

of L . I.e., $\mathbf{alphabet} L$ consists of all of the symbols occurring in the strings of L . For example,

$$\begin{aligned} \mathbf{alphabet} \{011, 112\} &= \bigcup \{ \mathbf{alphabet}(011), \mathbf{alphabet}(112) \} \\ &= \bigcup \{ \{0, 1\}, \{1, 2\} \} = \{0, 1, 2\}. \end{aligned}$$

Note that, for all languages L , $L \subseteq (\mathbf{alphabet} L)^*$.

If A is an infinite subset of \mathbf{Sym} (and so is not an alphabet), we allow ourselves to write A^* for $\mathbf{List} A$. I.e., A^* consists of all of the strings that can be built using the symbols of A . For example, $\mathbf{Sym}^* = \mathbf{Str}$.

2.1.5 Notes

In a traditional approach to the subject, symbols may be anything, real numbers, sets, etc. But such a choice would mean that not all symbols could be expressed in Forlan's syntax, and would needlessly complicate the set theoretic foundations of the subject. By working with a fixed, countably infinite set of symbols, all symbols can be expressed in Forlan, and we have that that strings, regular expressions, etc., are sets, not set-indexed families of sets.

Representing strings as lists of symbols, which in turn are represented as functions, is nontraditional, but should seem a natural approach to those with a background in set theory or functional programming.

2.2 Using Induction to Prove Language Equalities

In this section, we introduce three string induction principles, ways of showing that every $w \in A^*$ has property $P(w)$, where A is some set of symbols. Typically, A will be an alphabet, i.e., a finite set of symbols. But when we want to prove that all strings have some property, we can let $A = \mathbf{Sym}$, so that $A^* = \mathbf{Str}$. Each of these principles corresponds to an instance of well-founded induction. We also look at how different kinds of induction can be used to show that two languages are equal.

2.2.1 String Induction Principles

Suppose A is a set of symbols. We define well-founded relations \mathbf{right}_A , \mathbf{left}_A and \mathbf{strong}_A on A^* by:

- $x \mathbf{right}_A y$ iff $y = ax$ for some $a \in A$ (it's called *right* because the string x is on the right side of ax);
- $x \mathbf{left}_A y$ iff $y = xa$ for some $a \in A$ (it's called *left* because the string x is on the left side of xa);
- $x \mathbf{strong}_A y$ iff x is a proper substring of y .

Thus, for all $a \in A$ and $x \in A^*$, the only predecessor of ax in \mathbf{right}_A is x , and the only predecessor of xa in \mathbf{left}_A is x . And, for all $y \in A^*$, the predecessors of y in \mathbf{strong}_A are the proper substrings of y . The empty string, ϵ , has no predecessors in any of these relations.

The well-foundedness of \mathbf{right}_A , \mathbf{left}_A and \mathbf{strong}_A follows by Proposition 1.2.9, since each of these relations is a subset of $\mathbf{length}_{\mathbf{List } A}$, which is a well-founded relation on $\mathbf{List } A$.

We can do well-founded induction and recursion on these relations. In fact, what we called right and left recursion on strings in Section 2.1 correspond to recursion on $\mathbf{right}_{\mathbf{Sym}}$ and $\mathbf{left}_{\mathbf{Sym}}$.

We now introduce string induction principles corresponding to well-founded induction on each of the above relations.

Theorem 2.2.1 (Principle of Right String Induction)

Suppose $A \subseteq \mathbf{Sym}$ and $P(w)$ is a property of a string w . If

(basis step)

$$P(\epsilon) \text{ and}$$

(inductive step)

$$\text{for all } a \in A \text{ and } w \in A^*, \text{ if } (\dagger) P(w), \text{ then } P(aw),$$

then,

$$\text{for all } w \in A^*, P(w).$$

We refer to the formula (\dagger) as the *inductive hypothesis*. According to the induction principle, to show that every $w \in A^*$ has property P , we show that the empty string has property P , and then assume that $a \in A$, $w \in A^*$ and that (the inductive hypothesis) w has property P , and show that aw has property P .

Proof. Equivalent to well-founded induction on \mathbf{right}_A . \square

By switching aw to wa in the inductive step, we get the principle of left string induction.

Theorem 2.2.2 (Principle of Left String Induction)

Suppose $A \subseteq \mathbf{Sym}$ and $P(w)$ is a property of a string w . If

(basis step)

$$P(\epsilon) \text{ and}$$

(inductive step)

$$\text{for all } a \in A \text{ and } w \in A^*, \text{ if } (\dagger) P(w), \text{ then } P(wa),$$

then,

$$\text{for all } w \in A^*, P(w).$$

We refer to the formula (\dagger) as the *inductive hypothesis*.

Proof. Equivalent to well-founded induction on \mathbf{left}_A . \square

Theorem 2.2.3 (Principle of Strong String Induction)

Suppose $A \subseteq \mathbf{Sym}$ and $P(w)$ is a property of a string w . If

for all $w \in A^*$,
 if (†) for all $x \in A^*$, if x is a proper substring of w , then $P(x)$,
 then $P(w)$,

then,

for all $w \in A^*$, $P(w)$.

We refer to (†) as the inductive hypothesis. It says that all the proper substrings of w have property P . According to the induction principle, to show that every $w \in A^*$ has property P , we let $w \in A^*$, and assume (the inductive hypothesis) that every proper substring of w has property P . Then we must show that w has property P .

Proof. Equivalent to well-founded induction on \mathbf{strong}_A . \square

The next subsection, on proving language equalities, contains two examples of proofs by strong string induction. Before moving on to that subsection, we give an example proof by right string induction.

We define the *reversal* $x^R \in \mathbf{Str}$ of a string x by right recursion on strings:

$$\begin{aligned} \%^R &= \% \\ (ax)^R &= x^R a, \text{ for all } a \in \mathbf{Sym} \text{ and } x \in \mathbf{Str}. \end{aligned}$$

E.g., we have that $(021)^R = 120$. And, an easy calculation shows that, for all $a \in \mathbf{Sym}$, $a^R = a$. We let the reversal operation have higher precedence than string concatenation, so that, e.g., $xx^R = x(x^R)$.

Proposition 2.2.4

For all $x, y \in \mathbf{Str}$, $(xy)^R = y^R x^R$.

As usual, we must start by figuring out which of x and y to do induction on, as well as what sort of induction to use. Because we defined string reversal using right recursion, it turns out that we should do right string induction on x .

Proof. Suppose $y \in \mathbf{Str}$. Since $\mathbf{Sym}^* = \mathbf{Str}$, it will suffice to show that, for all $x \in \mathbf{Sym}^*$, $(xy)^R = y^R x^R$. We proceed by right string induction.

(Basis Step) We have that $(\%y)^R = y^R = y^R \% = y^R \%^R$.

(Inductive Step) Suppose $a \in \mathbf{Sym}$ and $x \in \mathbf{Sym}^*$. Assume the inductive hypothesis: $(xy)^R = y^R x^R$. Then,

$$\begin{aligned}
 ((ax)y)^R &= (a(xy))^R \\
 &= (xy)^R a && \text{(definition of } (a(xy))^R \text{)} \\
 &= (y^R x^R) a && \text{(inductive hypothesis)} \\
 &= y^R (x^R a) \\
 &= y^R (ax)^R && \text{(definition of } (ax)^R \text{)}.
 \end{aligned}$$

□

Exercise 2.2.5

Use right string induction and Proposition 2.2.4 to prove that, for all $x \in \mathbf{Str}$, $(x^R)^R = x$.

Exercise 2.2.6

In Section 2.1, we used right recursion to define the function $\mathbf{alphabet} \in \mathbf{Str} \rightarrow \mathbf{Alp}$. Use right string induction to show that, for all $x, y \in \mathbf{Str}$, $\mathbf{alphabet}(xy) = \mathbf{alphabet} x \cup \mathbf{alphabet} y$.

2.2.2 Proving Language Equalities

In this subsection, we show two examples of how strong string induction and induction over inductively defined languages can be used to show that two languages are equal.

For the first example, let X be the least subset of $\{0, 1\}^*$ such that:

- (1) $\% \in X$; and
- (2) for all $a \in \{0, 1\}$ and $x \in X$, $axa \in X$.

This is another example of an inductive definition: X consists of just those strings of 0's and 1's that can be constructed using (1) and (2). For example, by (1) and (2), we have that $00 = 0\%0 \in X$. Thus, by (2), we have that $1001 = 1(00)1 \in X$. In general, we have that X contains the elements:

$$\%, 00, 11, 0000, 0110, 1001, 1111, \dots$$

We will show that $X = Y$, where $Y = \{w \in \{0, 1\}^* \mid w \text{ is a palindrome and } |w| \text{ is even}\}$.

Lemma 2.2.7

$Y \subseteq X$.

Proof. Since $Y \subseteq \{0,1\}^*$, it will suffice to show that, for all $w \in \{0,1\}^*$,

if $w \in Y$, then $w \in X$.

We proceed by strong string induction.

Suppose $w \in \{0,1\}^*$, and assume the inductive hypothesis: for all $x \in \{0,1\}^*$, if x is a proper substring of w , then

if $x \in Y$, then $x \in X$.

We must show that

if $w \in Y$, then $w \in X$.

Suppose $w \in Y$, so that $w \in \{0,1\}^*$, w is a palindrome and $|w|$ is even. It remains to show that $w \in X$. If $w = \%$, then $w = \% \in X$, by Part (1) of the definition of X . So, suppose $w \neq \%$. Since $|w|$ is even, we have that $|w| \geq 2$, and thus that $w = axb$ for some $a, b \in \{0,1\}$ and $x \in \{0,1\}^*$. Because $|w|$ is even, it follows that $|x|$ is even. Furthermore, because w is a palindrome, it follows that $a = b$ and x is a palindrome. Thus $w = axa$ and $x \in Y$. Since x is a proper substring of w , the inductive hypothesis tells us that

if $x \in Y$, then $x \in X$.

But $x \in Y$, and thus $x \in X$. Thus, by Part (2) of the definition of X , we have that $w = axa \in X$. \square

We could also prove $X \subseteq Y$ by strong string induction. But an alternative approach is more elegant and generally applicable: we use the induction principle that comes from the inductive definition of X .

Proposition 2.2.8 (Principle of Induction on X)

Suppose $P(w)$ is a property of a string w . If

(1)

$P(\%)$, and

(2)

for all $a \in \{0,1\}$ and $x \in X$, if $(\dagger) P(x)$, then $P(axa)$,

then,

for all $w \in X$, $P(w)$.

We refer to (\dagger) as the *inductive hypothesis* of Part (2). By Part (1) of the definition of X , $\% \in X$. Thus Part (1) of the induction principle requires us to show $P(\%)$. By Part (2) of the definition of X , if $a \in \{0,1\}$ and $x \in X$, then $axa \in X$. Thus in Part (2) of the induction principle, when proving that the “new” element axa has property P , we’re allowed to assume that the “old” element has property P .

Lemma 2.2.9

$X \subseteq Y$.

Proof. We use induction on X to show that, for all $w \in X$, $w \in Y$.

There are two steps to show.

- (1) Since $\% \in \{0,1\}^*$, $\%$ is a palindrome and $|\%| = 0$ is even, we have that $\% \in Y$.
- (2) Let $a \in \{0,1\}$ and $x \in X$. Assume the inductive hypothesis: $x \in Y$. We must show that $axa \in Y$. Since $x \in Y$, we have that $x \in \{0,1\}^*$, x is a palindrome and $|x|$ is even. Because $a \in \{0,1\}$ and $x \in \{0,1\}^*$, it follows that $axa \in \{0,1\}^*$. Since x is a palindrome, we have that axa is also a palindrome. And, because $|axa| = |x| + 2$ and $|x|$ is even, it follows that $|axa|$ is even. Thus $axa \in Y$, as required.

□

Proposition 2.2.10

$X = Y$.

Proof. Follows immediately from Lemmas 2.2.7 and 2.2.9. □

We end this subsection by proving a more complex language equality. One of the languages is defined using a “difference” function on strings, which we will use a number of times in later chapters. Define $\mathbf{diff} \in \{0,1\}^* \rightarrow \mathbb{Z}$ by: for all $w \in \{0,1\}^*$,

$\mathbf{diff} w = \text{the number of 1's in } w - \text{the number of 0's in } w$.

Then:

- $\mathbf{diff} \% = 0$;
- $\mathbf{diff} 1 = 1$;
- $\mathbf{diff} 0 = -1$; and
- for all $x, y \in \{0,1\}^*$, $\mathbf{diff}(xy) = \mathbf{diff} x + \mathbf{diff} y$.

Note that, for all $w \in \{0,1\}^*$, $\mathbf{diff} w = 0$ iff w has an equal number of 0's and 1's. If we think of a 1 as representing the production of one unit of some resource, and of a 0 as representing the consumption of one unit of that resource, then a string will have a diff of 0 iff it is balanced in terms of production and consumption. Note that such a string may have prefixes with negative diff's, i.e., it may temporarily go "into the red".

Let X (forget the previous definition of X) be the least subset of $\{0,1\}^*$ such that:

- (1) $\% \in X$;
- (2) for all $x, y \in X$, $xy \in X$;
- (3) for all $x \in X$, $0x1 \in X$; and
- (4) for all $x \in X$, $1x0 \in X$.

Let $Y = \{w \in \{0,1\}^* \mid \mathbf{diff} w = 0\}$.

For example, since $\% \in X$, it follows, by (3) and (4) that $01 = 0\%1 \in X$ and $10 = 1\%0 \in X$. Thus, by (2), we have that $0110 = (01)(10) \in X$. And, Y consists of all strings of 0's and 1's with an equal number of 0's and 1's.

Our goal is to prove that $X = Y$, i.e., that: (the easy direction) every string that can be constructed using X 's rules has an equal number of 0's and 1's; and (the hard direction) that every string of 0's and 1's with an equal number of 0's and 1's can be constructed using X 's rules.

Because X was defined inductively, it gives rise to an induction principle, which we will use to prove the following lemma. (Because of Part (2) of the definition of X , we wouldn't be able to prove this lemma using strong string induction.)

Lemma 2.2.11

$X \subseteq Y$.

Proof. We use induction on X to show that, for all $w \in X$, $w \in Y$. There are four steps to show, corresponding to the four rules of X 's definition.

- (1) We must show $\% \in Y$. Since $\% \in \{0,1\}^*$ and $\mathbf{diff} \% = 0$, we have that $\% \in Y$.
- (2) Suppose $x, y \in X$, and assume the inductive hypothesis: $x, y \in Y$. We must show that $xy \in Y$. Since $x, y \in Y$, we have that $xy \in \{0,1\}^*$ and $\mathbf{diff}(xy) = \mathbf{diff} x + \mathbf{diff} y = 0 + 0 = 0$. Thus $xy \in Y$.
- (3) Suppose $x \in X$, and assume the inductive hypothesis: $x \in Y$. We must show that $0x1 \in Y$. Since $x \in Y$, we have that $0x1 \in \{0,1\}^*$ and $\mathbf{diff}(0x1) = \mathbf{diff} 0 + \mathbf{diff} x + \mathbf{diff} 1 = -1 + 0 + 1 = 0$. Thus $0x1 \in Y$.

- (4) Suppose $x \in X$, and assume the inductive hypothesis: $x \in Y$. We must show that $1x0 \in Y$. Since $x \in Y$, we have that $1x0 \in \{0,1\}^*$ and $\mathbf{diff}(1x0) = \mathbf{diff} 1 + \mathbf{diff} x + \mathbf{diff} 0 = 1 + 0 + -1 = 0$. Thus $1x0 \in Y$.

□

Lemma 2.2.12

$Y \subseteq X$.

Proof. Since $Y \subseteq \{0,1\}^*$, it will suffice to show that, for all $w \in \{0,1\}^*$,

if $w \in Y$, then $w \in X$.

We proceed by strong string induction. Suppose $w \in \{0,1\}^*$, and assume the inductive hypothesis: for all $x \in \{0,1\}^*$, if x is a proper substring of w , then

if $x \in Y$, then $x \in X$.

We must show that

if $w \in Y$, then $w \in X$.

Suppose $w \in Y$. We must show that $w \in X$. There are three cases to consider.

- Suppose $w = \%$. Then $w = \% \in X$, by Part (1) of the definition of X .
- Suppose $w = 0t$ for some $t \in \{0,1\}^*$. Since $w \in Y$, we have that $-1 + \mathbf{diff} t = \mathbf{diff} 0 + \mathbf{diff} t = \mathbf{diff}(0t) = \mathbf{diff} w = 0$, and thus that $\mathbf{diff} t = 1$.

Let u be the shortest prefix of t such that $\mathbf{diff} u \geq 1$. (Since t is a prefix of itself and $\mathbf{diff} t = 1 \geq 1$, it follows that u is well-defined.) Let $z \in \{0,1\}^*$ be such that $t = uz$. Clearly, $u \neq \%$, and thus $u = yb$ for some $y \in \{0,1\}^*$ and $b \in \{0,1\}$. Hence $t = uz = ybz$. Since y is a shorter prefix of t than u , we have that $\mathbf{diff} y \leq 0$.

Suppose, toward a contradiction, that $b = 0$. Then $\mathbf{diff} y + -1 = \mathbf{diff} y + \mathbf{diff} 0 = \mathbf{diff} y + \mathbf{diff} b = \mathbf{diff}(yb) = \mathbf{diff} u \geq 1$, so that $\mathbf{diff} y \geq 2$. But $\mathbf{diff} y \leq 0$ —contradiction. Hence $b = 1$.

Summarizing, we have that $u = yb = y1$, $t = uz = y1z$ and $w = 0t = 0y1z$. Since $\mathbf{diff} y + 1 = \mathbf{diff} y + \mathbf{diff} 1 = \mathbf{diff}(y1) = \mathbf{diff} u \geq 1$, it follows that $\mathbf{diff} y \geq 0$. But $\mathbf{diff} y \leq 0$, and thus $\mathbf{diff} y = 0$. Thus $y \in Y$. Since $1 + \mathbf{diff} z = 0 + 1 + \mathbf{diff} z = \mathbf{diff} y + \mathbf{diff} 1 + \mathbf{diff} z = \mathbf{diff}(y1z) = \mathbf{diff} t = 1$, it follows that $\mathbf{diff} z = 0$. Thus $z \in Y$.

Because y and z are proper substrings of w , and $y, z \in Y$, the inductive hypothesis tells us that $y, z \in X$. Thus, by Part (3) of the definition of X , we have that $0y1 \in X$. Hence, Part (2) of the definition of X tells us that $w = 0y1z = (0y1)z \in X$.

- Suppose $w = 1t$ for some $t \in \{0, 1\}^*$. Since $w \in Y$, we have that $1 + \mathbf{diff} t = \mathbf{diff} 1 + \mathbf{diff} t = \mathbf{diff}(1t) = \mathbf{diff} w = 0$, and thus that $\mathbf{diff} t = -1$.

Let u be the shortest prefix of t such that $\mathbf{diff} u \leq -1$. (Since t is a prefix of itself and $\mathbf{diff} t = -1 \leq -1$, it follows that u is well-defined.) Let $z \in \{0, 1\}^*$ be such that $t = uz$. Clearly, $u \neq \%$, and thus $u = yb$ for some $y \in \{0, 1\}^*$ and $b \in \{0, 1\}$. Hence $t = uz = ybz$. Since y is a shorter prefix of t than u , we have that $\mathbf{diff} y \geq 0$.

Suppose, toward a contradiction, that $b = 1$. Then $\mathbf{diff} y + 1 = \mathbf{diff} y + \mathbf{diff} 1 = \mathbf{diff} y + \mathbf{diff} b = \mathbf{diff}(yb) = \mathbf{diff} u \leq -1$, so that $\mathbf{diff} y \leq -2$. But $\mathbf{diff} y \geq 0$ —contradiction. Hence $b = 0$.

Summarizing, we have that $u = yb = y0$, $t = uz = y0z$ and $w = 1t = 1y0z$. Since $\mathbf{diff} y + -1 = \mathbf{diff} y + \mathbf{diff} 0 = \mathbf{diff}(y0) = \mathbf{diff} u \leq -1$, it follows that $\mathbf{diff} y \leq 0$. But $\mathbf{diff} y \geq 0$, and thus $\mathbf{diff} y = 0$. Thus $y \in Y$. Since $-1 + \mathbf{diff} z = 0 + -1 + \mathbf{diff} z = \mathbf{diff} y + \mathbf{diff} 0 + \mathbf{diff} z = \mathbf{diff}(y0z) = \mathbf{diff} t = -1$, it follows that $\mathbf{diff} z = 0$. Thus $z \in Y$.

Because y and z are proper substrings of w , and $y, z \in Y$, the inductive hypothesis tells us that $y, z \in X$. Thus, by Part (4) of the definition of X , we have that $1y0 \in X$. Hence, Part (2) of the definition of X tells us that $w = 1y0z = (1y0)z \in X$.

□

In the proof of the preceding lemma we made use of all four rules of X 's definition. If this had not been the case, we would have known that the unused rules were redundant (or that we had made a mistake in our proof!).

Proposition 2.2.13

$X = Y$.

Proof. Follows immediately from Lemmas 2.2.11 and 2.2.12. □

Exercise 2.2.14

Define the function $\mathbf{diff} \in \{0, 1\}^* \rightarrow \mathbb{Z}$ as in the above example. Let X be the least subset of $\{0, 1\}^*$ such that:

- (1) $\% \in X$;
- (2) for all $x \in X$, $0x1 \in X$; and
- (3) for all $x, y \in X$, $xy \in X$.

Let $Y = \{w \in \{0, 1\}^* \mid \mathbf{diff} w = 0 \text{ and, for all prefixes } v \text{ of } w, \mathbf{diff} v \leq 0\}$. Prove that $X = Y$.

Exercise 2.2.15

Define a function $\mathbf{diff} \in \{0, 1\}^* \rightarrow \mathbb{Z}$ by: for all $w \in \{0, 1\}^*$,

$$\mathbf{diff} w = \text{the number of 1's in } w - 2(\text{the number of 0's in } w).$$

Thus $\mathbf{diff} \% = 0$, $\mathbf{diff} 0 = -2$, $\mathbf{diff} 1 = 1$, and for all $x, y \in \{0, 1\}^*$, $\mathbf{diff}(xy) = \mathbf{diff} x + \mathbf{diff} y$. Furthermore, for all $w \in \{0, 1\}^*$, $\mathbf{diff} w = 0$ iff w has twice as many 1's as 0's. Let X be the least subset of $\{0, 1\}^*$ such that:

- (1) $\% \in X$;
- (2) $1 \in X$;
- (3) for all $x, y \in X$, $1x1y0 \in X$; and
- (4) for all $x, y \in X$, $xy \in X$.

Let $Y = \{w \in \{0, 1\}^* \mid \text{for all prefixes } v \text{ of } w, \mathbf{diff} v \geq 0\}$. Prove that $X = Y$.

2.2.3 Notes

A novel feature of this book is the introduction and use of explicit string induction principles, as an alternative to doing proofs by induction (mathematical or strong) on the length of strings. Also novel is our focus on languages defined using “difference” functions.

2.3 Introduction to Forlan

The Forlan toolset is an extension of the Standard ML of New Jersey (SML/NJ) implementation of Standard ML (SML). It is implemented as a set of SML modules. It is used interactively, and users can extend Forlan by defining SML functions.

Instructions for installing and running Forlan on machines running Linux, macOS and Windows can be found on the Forlan website:

<http://alleystoughton.us/forlan>.

A manual for Forlan is available on the Forlan website, and describes Forlan's modules in considerably more detail than does this book. See the manual for instructions for setting system parameters controlling such things as the search path used for loading files, the line length used by Forlan's pretty printer, and the number of elements of a list that the Forlan top-level displays.

In the concrete syntax for describing Forlan objects—automata, grammars, etc.—*comments* begin with a “#”, and run through the end of the line. Comments and whitespace may be arbitrarily inserted into the descriptions of Forlan objects without changing how the objects will be lexically analyzed and parsed. For instance,

```
ab cd # this is a comment
efg h
```

describes the Forlan string `abcdefgh`. Forlan’s input functions prompt with “`@`” when reading from the standard input, in which case the user signifies end-of-file by typing a line consisting of a single dot (“`.`”).

We begin this section by showing how to invoke Forlan, and giving a quick introduction to the SML core of Forlan. We then show how symbols, strings, finite sets of symbols and strings, and finite relations on symbols can be manipulated using Forlan.

2.3.1 Invoking Forlan

To invoke Forlan, type the command `forlan` to your shell (command processor):

```
% forlan
Forlan Version m (based on Standard ML of New Jersey Version n)
val it = () : unit
-
```

(m and n will be the Forlan and SML/NJ versions, respectively.) The identifier `it` is normally bound to the value of the most recently evaluated expression. Initially, though, its value is the empty tuple `()`, the single element of the type `unit`. The value `()` is used in circumstances when a value is required, but it makes no difference what that value is.

Forlan’s primary prompt is “`-`”. To exit Forlan, type `CTRL-d` under Linux and macOS, and `CTRL-z` under Windows. To interrupt back to the Forlan top-level, type `CTRL-c`.

On Windows, you may find it more convenient to invoke Forlan by double-clicking on the Forlan icon. On all platforms, a much more flexible and satisfying way of running Forlan is as a subprocess of the Emacs text editor. See the Forlan website for information about how to do this.

2.3.2 The SML Core of Forlan

This subsection gives a quick introduction to the SML core of Forlan. Let’s begin by using Forlan as a calculator:

```
- 4 + 5;
val it = 9 : int
- it * it;
val it = 81 : int
- it - 1;
val it = 80 : int
- 5 div 2;
val it = 2 : int
- 5 mod 2;
```



```

val it = 1 : int
- ~4 + 2;
val it = ~2 : int

```

Forlan responds to each expression by printing its value and type (`int` is the type of integers), and noting that the expression's value has been bound to the identifier `it`. Expressions must be terminated with semicolons. The operators `div` and `mod` compute integer division and remainder, respectively, and negative numbers begin with `~`.

In addition to the type `int` of integers, SML has types `string` and `bool`, product types $t_1 * \dots * t_n$, and list types t `list`.

```

- "hello" ^ " " ^ "there";
val it = "hello there" : string
- true andalso (false orelse true);
val it = true : bool
- if 5 < 7 then "hello" else "bye";
val it = "hello" : string
- (3 + 1, 4 = 4, "a" ^ "b");
val it = (4,true,"ab") : int * bool * string
- #2 it;
val it = true : bool
- [1, 3, 5] @ [7, 9, 11];
val it = [1,3,5,7,9,11] : int list
- rev it;
val it = [11,9,7,5,3,1] : int list
- length it;
val it = 6 : int
- null[];
val it = true : bool
- null[1, 2];
val it = false : bool
- hd[1, 2, 3];
val it = 1 : int
- tl[1, 2, 3];
val it = [2,3] : int list

```

The operator `^` is string concatenation. The conjunction `andalso` evaluates its left-hand side first, and yields `false` without evaluating its right-hand side, if the value of the left-hand side is `false`. Similarly, the disjunction `orelse` evaluates its left-hand side first, and yields `true` without evaluating its right-hand side, if the value of the left-hand side is `true`. A conditional (`if-then-else`) is evaluated by first evaluating its boolean expression, and then evaluating its `then`-part, if the boolean expression's value is `true`, and evaluating its `else`-part, if its value is `false`. Tuples are evaluated from left to right, and the function `#n` selects the n th (starting from 1) element of a tuple. The operator `@` appends lists. The function `rev` reverses a list, the function `length` computes the length of a list, and the function `null` tests whether a list is empty. Finally, the functions `hd`

and `tl` return the head (first element) and tail (all but the first element) of a list.

`nil` and `::` (pronounced “cons”, for “constructor”), which have types `'a list` and `'a * 'a list -> 'a list`, respectively, are the constructors for type `'a list`. These constructors are *polymorphic*, having all of the types that can be formed by instantiating the type variable `'a` with a type. E.g., `nil` has type `int list`, `bool list`, `(int * bool)list`, etc. `::` is an infix operator, i.e., one writes `x :: xs` for the list whose first element is `x` and remaining elements are those in the list `xs`.

```
- nil;
val it = [] : 'a list
- 1 :: nil;
val it = [1] : int list
- 1 :: 2 :: nil;
val it = [1,2] : int list
- 3 :: [5, 7, 9];
val it = [3,5,7,9] : int list
```

Lists are implemented as linked-lists, so that doing a cons involves the creation of a single list node.

SML also has option types `t option`, whose values are built using the type’s two constructors: `NONE` of type `'a option`, and `SOME` of type `'a -> 'a option`. This is a predefined datatype, declared by

```
datatype 'a option = NONE | SOME of 'a
```

E.g., `NONE`, `SOME 1` and `SOME ~6` are three of the values of type `int option`, and `NONE`, `SOME true` and `SOME false` are the only values of type `bool option`.

```
- NONE;
val it = NONE : 'a option
- SOME 3;
val it = SOME 3 : int option
- SOME true;
val it = SOME true : bool option
```

In addition to the usual operators `<`, `<=`, `>` and `>=` for comparing integers, SML offers a function `Int.compare` of type `int * int -> order`, where the `order` type contains three elements: `LESS`, `EQUAL` and `GREATER`.

```
- Int.compare(3, 4);
val it = LESS : order
- Int.compare(4, 4);
val it = EQUAL : order
- Int.compare(4, 3);
val it = GREATER : order
```

It is possible to bind the value of an expression to an identifier using a value declaration:

```

- val x = 3 + 4;
val x = 7 : int
- val y = x + 1;
val y = 8 : int
- val x = 5 * x;
val x = 35 : int
- y;
val it = 8 : int

```

In the first declaration of `x`, its right-hand side is first evaluated, resulting in 7, and then `x` is bound to this value. Note that the redeclaration of `x` doesn't change the value of the previous declaration of `x`, it just makes that declaration inaccessible.

One can use a value declaration to give names to the components of a tuple, or give a name to the data of a non-`NONE` optional value:

```

- val (x, y, z) = (3 + 1, 4 = 4, "a" ^ "b");
val x = 4 : int
val y = true : bool
val z = "ab" : string
- val SOME n = SOME(4 * 25);
stdIn:39.5-39.26 Warning: binding not exhaustive
      SOME n = ...
val n = 100 : int

```

This last declaration uses pattern matching: `SOME(4 * 25)` is evaluated to `SOME 100`, and is then matched against the pattern `SOME n`. Because the constructors match, the pattern matching succeeds, and `n` becomes bound to 100. The warning is because the SML typechecker doesn't know the expression won't evaluate to `NONE`.

One can use a `let` expression to carry out some declarations in a local environment, evaluate an expression in that environment, and yield the result of that evaluation:

```

- val x = 3;
val x = 3 : int
- val z = 10;
val z = 10 : int
- let val x = 4 * 5
=      val y = x * z
= in (x, y, x + y) end;
val it = (20,200,220) : int * int * int
- x;
val it = 3 : int

```

When a declaration or expression spans more than one line, Forlan uses its secondary prompt, `=`, on all of the lines except for the first one. Forlan doesn't process a declaration or expression until it is terminated with a semicolon.

One can declare functions, and apply those functions to arguments:

```

- fun f n = n * 2;
val f = fn : int -> int
- f 3;
val it = 6 : int
- f(4 + 5);
val it = 18 : int
- f 4 + 5;
val it = 13 : int
- fun g(x, y) = (x ^ y, y ^ x);
val g = fn : string * string -> string * string
- val (u, v) = g("a", "b");
val u = "ab" : string
val v = "ba" : string

```

The function `f` doubles its argument. All function values are printed as `fn`. A type $t_1 \rightarrow t_2$ is the type of all functions taking arguments of type t_1 and producing results (if they terminate without raising exceptions) of type t_2 . SML infers the types of functions. The function application `f(4 + 5)` is evaluated as follows. First, the argument `4 + 5` is evaluated, resulting in 9. Then a local environment is created in which `n` is bound to 9, and `f`'s body is evaluated in that environment, producing 18. Function application has higher precedence than operators like `+`.

Technically, the function `g` matches its single argument, which must be a pair, against the pair pattern `(x, y)`, binding `x` and `y` to the left and right sides of this argument, and then evaluates its body. But we can think such a function as having multiple arguments. The type operator `*` has higher precedence than the operator `->`.

Except for basic entities like integers and booleans, all values in SML are represented by pointers, so that passing such a value to a function, or putting it in a datastructure, only involves copying a pointer.

Given functions f and g of types $t_1 \rightarrow t_2$ and $t_2 \rightarrow t_3$, respectively, $g \circ f$ is the composition of g and f , the function of type $t_1 \rightarrow t_3$ that, when given an argument x of type t_1 , evaluates the expression $g(f\ x)$. For example, we have that:

```

- fun f x = x >= 1 andalso x <= 10;
val f = fn : int -> bool
- fun g x = if x then "inside" else "outside";
val g = fn : bool -> string
- val h = g o f;
val h = fn : int -> string
- h ~5;
val it = "outside" : string
- h 6;
val it = "inside" : string
- h 14;
val it = "outside" : string

```

SML also has anonymous functions, which may also be given names using value declarations:

```
- (fn x => x + 1)(3 + 4);
val it = 8 : int
- val f = fn x => x + 1;
val f = fn : int -> int
- f(3 + 4);
val it = 8 : int
```

The anonymous function `fn x => x + 1` has type `int -> int` and adds one to its argument.

Functions are data: they may be passed to functions, returned from functions (a function that returns a function is called *curried*), be components of tuples or lists, etc. For example,

```
val map : ('a -> 'b) -> 'a list -> 'b list
```

is a polymorphic, curried function. The type operator `->` associates to the right, so that `map`'s type is

```
val map : ('a -> 'b) -> ('a list -> 'b list)
```

`map` takes in a function `f` of type `'a -> 'b`, and returns a function that when called with a list of elements of type `'a`, transforms each element using `f`, forming a list of elements of type `'b`.

```
- val f = map(fn x => x + 1);
val f = fn : int list -> int list
- f[2, 4, 6];
val it = [3,5,7] : int list
- f[~2, ~1, 0];
val it = [~1,0,1] : int list
- map (fn x => x mod 2 = 1) [3, 4, 5, 6, 7];
val it = [true,false,true,false,true] : bool list
```

In the last use of `map`, we are using the fact that function application associates to the left, so that `fxy` means `(fx)y`, i.e., apply `f` to `x`, and then apply the resulting function to `y`.

The following example shows that local environments are kept alive as long as there are accessible function values referring to them:

```
- val f =
=       let val x = 2 * 10
=       in fn y => y * x end;
val f = fn : int -> int
- f 4;
val it = 80 : int
- f 7;
```

```
val it = 140 : int
```

If the local environment containing the binding of `x` was discarded, then calling `f` would fail.

It's also possible to declare recursive functions, like the factorial function:

```
- fun fact n =  
=       if n = 0  
=       then 1  
=       else n * fact(n - 1);  
val fact = fn : int -> int  
- fact 4;  
val it = 24 : int
```

One can load the contents of a file into Forlan using the function

```
val use : string -> unit
```

For example, if the file `fact.sml` contains the declaration of the factorial function, then this declaration can be loaded into the system as follows:

```
- use "fact.sml";  
[opening fact.sml]  
val fact = fn : int -> int  
val it = () : unit  
- fact 4;  
val it = 24 : int
```

The factorial function can also be defined using pattern matching, either by using a case expression in the body of the function, or by using multiple clauses in the function's definition:

```
- fun fact n =  
=       case n of  
=       0 => 1  
=       | n => n * fact(n - 1);  
val fact = fn : int -> int  
- fact 3;  
val it = 6 : int  
- fun fact 0 = 1  
=   | fact n = n * fact(n - 1);  
val fact = fn : int -> int  
- fact 4;  
val it = 24 : int
```

The order of the clauses of a case expression or function definition is significant. If the clauses of either the case expression or the function definition were reversed, the function being defined would never return.

Pattern matching is especially useful when doing list processing. E.g., we could (inefficiently) define the list reversal function like this:

```

- fun rev nil          = nil
=   | rev (x :: xs) = rev xs @ [x];
val rev = fn : 'a list -> 'a list

```

Calling `rev` with the empty list will result in the empty list being returned. And calling it with a nonempty list will temporarily bind `x` to the list's head, bind `xs` to its tail, and then evaluate the expression `rev xs @ [x]`, recursively calling `rev`, and then returning the result of appending the result of this recursive call and `[x]`. Unfortunately, this definition of `rev` is slow for long lists, as `@` rebuilds its left argument. Instead, the official definition of `rev` is much more efficient:

```

- fun rev xs =
=   let fun rv(nil,    vs) = vs
=       | rv(u :: us, vs) = rv(us, u :: vs)
=   in rv(xs, nil) end;
val rev = fn : 'a list -> 'a list
- rev [1, 2, 3, 4];
val it = [4,3,2,1] : int list

```

When `rev` is called with `[1, 2, 3, 4]`, it calls the auxiliary function `rv` with the pair `([1, 2, 3, 4], [])`. The recursive calls that `rv` makes to itself are a special kind of recursion called *tail recursion*. The SML/NJ compiler generates code for such calls that simply jumps back to the beginning of `rv`, only changing its arguments. We have the following sequence of calls: `rv([1, 2, 3, 4], [])` calls `rv([2, 3, 4], [1])` calls `rv([3, 4], [2, 1])` calls `rv([4], [3, 2, 1])` calls `rv([], [4, 3, 2, 1])`, which returns `[4, 2, 2, 1]`, which is returned as the result of `rev`.

Finally, we can define recursive datatypes, and define functions by structural recursion on recursive datatypes. E.g., here's how we can define the datatype of labeled binary trees (both leaves and nodes (non-leaves) can have labels):

```

- datatype ('a, 'b) tree =
=   Leaf of 'b
=   | Node of 'a * ('a, 'b) tree * ('a, 'b) tree;
datatype ('a, 'b) tree
= Leaf of 'b | Node of 'a * ('a, 'b) tree * ('a, 'b) tree
- Leaf;
val it = fn : 'a -> ('b, 'a) tree
- Node;
val it = fn : 'a * ('a, 'b) tree * ('a, 'b) tree -> ('a, 'b) tree
- val tr = Node(true, Node(false, Leaf 7, Leaf ~1), Leaf 8);
val tr = Node (true, Node (false, Leaf 7, Leaf ~1), Leaf 8)
: (bool, int) tree

```

Then we can define a function for reversing a tree, and apply it to `tr`:

```

- fun revTree (Leaf n)          = Leaf n

```

```

= | revTree (Node(m, tr1, tr2)) =
=   Node(m, revTree tr2, revTree tr1);
val revTree = fn : ('a,'b) tree -> ('a,'b) tree
- revTree tr;
val it = Node (true,Leaf 8,Node (false,Leaf ~1,Leaf 7))
: (bool,int) tree
- revTree it;
val it = Node (true,Node (false,Leaf 7,Leaf ~1),Leaf 8)
: (bool,int) tree

```

2.3.3 Symbols

The Forlan module `Sym` defines the abstract type `sym` of Forlan symbols, as well as some functions for processing symbols, including:

```

val input    : string -> sym
val output   : string * sym -> unit
val compare  : sym * sym -> order
val equal    : string * string -> bool

```

Symbols are expressed in Forlan's syntax as sequences of symbol characters, i.e., as `a` or `<id>`, rather than `[a]` or `[< i, d, >]`. The above functions behave as follows:

- `input fil` reads a symbol from file `fil`; if `fil = ""`, then the symbol is read from the standard input;
- `output(fil, a)` writes the symbol `a` to the file `fil`; if `fil = ""`, then the string is written to the standard output;
- `compare` implements our total ordering on symbols; and
- `equal` tests whether two symbols are equal.

All of Forlan's input functions read from the standard input when called with `""` instead of a file, and all of Forlan's output functions write to the standard output when given `""` instead of a file.

The type `sym` is bound in the top-level environment. On the other hand, one must write `Sym.f` to select the function `f` of module `Sym`. As described above, interactive input is terminated by a line consisting of a single “.” (dot), and Forlan's input prompt is “@”.

The module `Sym` also provides the functions

```

val fromString : string -> sym
val toString   : sym -> string

```

where `fromString` is like `input`, except that it takes its input from a string, and `toString` is like `output`, except that it writes its output to a string. These functions are especially useful when defining functions. In the sequel, whenever

a module/type has `input` and `output` functions, you may assume that it also has `fromString` and `toString` functions.

Here are some example uses of the functions of `Sym`:

```
- val a = Sym.input "";
@ <i
@ d>
@ .
val a = - : sym
- val b = Sym.fromString "<num>";
val b = - : sym
- Sym.output("", a);
<id>
val it = () : unit
- Sym.compare(a, b);
val it = LESS : order
- Sym.equal(a, b);
val it = false : bool
- Sym.equal(a, Sym.fromString "<id>");
val it = true : bool
```

Values of abstract types (like `sym`) are printed as “-”.

2.3.4 Sets

The module `Set` defines the abstract type

```
type 'a set
```

of finite sets of elements of type `'a`. It is bound in the top-level environment. E.g., `sym set` is the type of sets of symbols.

Each set has an associated total ordering, and some of the functions of `Set` take total orderings as arguments. See the Forlan manual for the details. In the book, we won't have to work with such functions explicitly.

`Set` provides various constants and functions for processing sets, but we will only make direct use of a few of them:

```
val toList : 'a set -> 'a list
val size   : 'a set -> int
val empty  : 'a set
val sing   : 'a -> 'a set
val filter : ('a -> bool) -> 'a set -> 'a set
```

These values are polymorphic: `'a` can be `int`, `sym`, etc. The function `toList` returns the elements of a set, listing them in ascending order, according to the set's total ordering. The function `size` returns the size of a set. The value `empty` is the empty set. The function `sing` makes a value x into the singleton set $\{x\}$. And `filter` goes through the elements of a set, keeping those elements on which the supplied predicate function returns `true`.

2.3.5 Sets of Symbols

The module `SymSet` defines various functions for processing finite sets of symbols (elements of type `sym set`; alphabets), including:

```

val input      : string -> sym set
val output     : string * sym set -> unit
val fromList   : sym list -> sym set
val memb       : sym * sym set -> bool
val subset     : sym set * sym set -> bool
val equal      : sym set * sym set -> bool
val union      : sym set * sym set -> sym set
val inter      : sym set * sym set -> sym set
val minus      : sym set * sym set -> sym set
val genUnion   : sym set list -> sym set
val genInter   : sym set list -> sym set

```

The total ordering associated with sets of symbols is our total ordering on symbols. Sets of symbols are expressed in Forlan's syntax as sequences of symbols, separated by commas.

The function `fromList` returns a set with the same elements of the list of symbols it is called with. The function `memb` tests whether a symbol is a member (element) of a set of symbols, `subset` tests whether a first set of symbols is a subset of a second one, and `equal` tests whether two sets of symbols are equal. The functions `union`, `inter` and `minus` compute the union, intersection and difference of two sets of symbols. The function `genUnion` computes the generalized intersection of a list of sets of symbols *xss*, returning the set of all symbols appearing in at least one element of *xss*. And, the function `genInter` computes the generalized intersection of a nonempty list of sets of symbols *xss*, returning the set of all symbols appearing in all elements of *xss*.

Here are some example uses of the functions of `SymSet`:

```

- val bs = SymSet.input "";
@ a, <id>, 0, <num>
@ .
val bs = - : sym set
- SymSet.output("", bs);
0, a, <id>, <num>
val it = () : unit
- val cs = SymSet.input "";
@ a, <char>
@ .
val cs = - : sym set
- SymSet.subset(cs, bs);
val it = false : bool
- SymSet.output("", SymSet.union(bs, cs));
0, a, <id>, <num>, <char>
val it = () : unit

```

```

- SymSet.output("", SymSet.inter(bs, cs));
a
val it = () : unit
- SymSet.output("", SymSet.minus(bs, cs));
0, <id>, <num>
val it = () : unit
- val ds = SymSet.fromString "<char>, <>";
val ds = - : sym set
- SymSet.output("", SymSet.genUnion[bs, cs, ds]);
0, a, <>, <id>, <num>, <char>
val it = () : unit
- SymSet.output("", SymSet.genInter[bs, cs, ds]);

val it = () : unit

```

2.3.6 Strings

We will be working with two kinds of strings:

- SML strings, i.e., elements of type `string`;
- The strings of formal language theory, which we call “formal language strings”, when necessary.

The module `Str` defines the type `str` of formal language strings, which is bound in the top-level environment, and is equal to `sym list`, the type of lists of symbols. Because strings are lists, we can use SML’s list processing functions on them. Strings are expressed in Forlan’s syntax as either a single `%` or a nonempty sequence of symbols.

The module `Str` also defines some functions for processing strings, including:

```

val input      : string -> str
val output     : string * str -> unit
val alphabet   : str -> sym set
val compare    : str * str -> order
val equal      : str * str -> bool
val prefix     : str * str -> bool
val suffix     : str * str -> bool
val substr     : str * str -> bool
val power      : str * int -> str
val last       : str -> sym
val allButLast : str -> str

```

The function `alphabet` returns the alphabet of a string, and `compare` implements our total ordering on strings. `prefix(x , y)` tests whether x is a prefix of y , and `suffix` and `substring` work similarly. `power(x , n)` raises x to the power n . And `last` and `allButLast` return the last symbol and all but the last symbol of a string, respectively.

Here are some example uses of the functions of `Str`:

```

- val x = Str.input "";
@ hello<there>
@ .
val x = [-,-,-,-,-] : str
- length x;
val it = 6 : int
- Str.output("", x);
hello<there>
val it = () : unit
- SymSet.output("", Str.alphabet x);
e, h, l, o, <there>
val it = () : unit
- Str.output("", Str.power(x, 3));
hello<there>hello<there>hello<there>
val it = () : unit
- val y = Str.fromString "ello";
val y = [-,-,-,-] : str
- Str.compare(y, x);
val it = LESS : order
- Str.equal(y, x);
val it = false : bool
- Str.prefix(y, x);
val it = false : bool
- Str.substr(y, x);
val it = true : bool
- val z = Str.fromString "h" @ y;
val z = [-,-,-,-,-] : sym list
- Str.prefix(z, x);
val it = true : bool
- val x = Str.fromString "hellothere";
val x = [-,-,-,-,-,-,-,-,-] : str
- null x;
val it = false : bool
- Sym.output("", hd x);
h
val it = () : unit
- Str.output("", tl x);
ellothere
val it = () : unit
- Sym.output("", Str.last x);
e
val it = () : unit
- Str.output("", Str.allButLast x);
hellother
val it = () : unit

```

2.3.7 Sets of Strings

The module `StrSet` defines various functions for processing finite sets of strings (elements of type `str set`; finite languages), including:

```
val input      : string -> str set
val output     : string * str set -> unit
val fromList   : str list -> str set
val memb       : str * str set -> bool
val subset     : str set * str set -> bool
val equal      : str set * str set -> bool
val union      : str set * str set -> str set
val inter      : str set * str set -> str set
val minus      : str set * str set -> str set
val genUnion   : str set list -> str set
val genInter   : str set list -> str set
val alphabet   : str set -> sym set
```

The total ordering associated with sets of strings is our total ordering on strings. Sets of strings are expressed in Forlan's syntax as sequences of strings, separated by commas.

Here are some example uses of the functions of `StrSet`:

```
- val xs = StrSet.input "";
@ hello, <id><num>, %
@ .
val xs = - : str set
- val ys = StrSet.input "";
@ <id><num>, another
@ .
val ys = - : str set
- val zs = StrSet.union(xs, ys);
val zs = - : str set
- Set.size zs;
val it = 4 : int
- StrSet.output("", zs);
%, <id><num>, hello, another
val it = () : unit
- val us = Set.filter (fn x => length x mod 2 = 0) zs;
val us = - : sym list set
- StrSet.output("", us);
%, <id><num>
val it = () : unit
- SymSet.output("", StrSet.alphabet zs);
a, e, h, l, n, o, r, t, <id>, <num>
val it = () : unit
```

In this transcript, `us` was declared to be all the even-length elements of `zs`.

2.3.8 Relations on Symbols

The module `SymRel` defines the type `sym_rel` of finite relations on symbols. It is bound in the top-level environment, and is equal to `(sym * sym)set`, i.e., its elements are finite sets of pairs of symbols. The total ordering associated with relations on symbols orders pairs of symbols first according to their left-hand sides (using the total ordering on symbols), and then according to their right-hand sides. Relations on symbols are expressed in Forlan's syntax as sequences of ordered pairs (a, b) of symbols, separated by commas.

`SymRel` also defines various functions for processing finite relations on symbols, including:

```

val input      : string -> sym_rel
val output     : string * sym_rel -> unit
val fromList   : (sym * sym)list -> sym_rel
val memb       : (sym * sym) * sym_rel -> bool
val subset     : sym_rel * sym_rel -> bool
val equal      : sym_rel * sym_rel -> bool
val union      : sym_rel * sym_rel -> sym_rel
val inter      : sym_rel * sym_rel -> sym_rel
val minus      : sym_rel * sym_rel -> sym_rel
val genUnion   : sym_rel list -> sym_rel
val genInter   : sym_rel list -> sym_rel
val domain     : sym_rel -> sym set
val range      : sym_rel -> sym set
val relationFromTo : sym_rel * sym_set * sym_set -> bool
val reflexive  : sym_rel * sym_set -> bool
val symmetric  : sym_rel -> bool
val antisymmetric : sym_rel -> bool
val transitive : sym_rel -> bool
val total      : sym_rel -> bool
val inverse    : sym_rel -> sym_rel
val compose    : sym_rel * sym_rel -> sym_rel
val function    : sym_rel -> bool
val functionFromTo : sym_rel * sym_set * sym_set -> bool
val injection  : sym_rel -> bool
val bijectionFromTo : sym_rel * sym_set * sym_set -> bool
val applyFunction : sym_rel -> sym -> sym
val restrictFunction : sym_rel * sym_set -> sym_rel
val updateFunction : sym_rel * sym * sym -> sym_rel

```

The functions `domain` and `range` return the domain and range, respectively, of a relation. `relationFromTo(rel, bs, cs)` tests whether *rel* is a relation from *bs* to *cs*.

`reflexive(rel, bs)` tests whether *rel* is reflexive on *bs*. The functions `symmetric`, `antisymmetric` and `transitive` test whether a relation is symmetric, antisymmetric or transitive, respectively. `total(rel, bs)` tests whether *rel* is total on *bs*.

The function `inverse` computes the inverse of a relation, and `compose` composes two relations.

The function `function` tests whether a relation is a function. The function `applyFunction` is curried. Given a relation `rel`, `applyFunction` checks that `rel` is a function, issuing an error message, and raising an exception, otherwise. If it is a function, it returns a function of type `sym -> sym` that, when called with a symbol `a`, will apply the function `rel` to `a`, issuing an error message if `a` is not in the domain of `rel`. `functionFromTo(rel, bs, cs)` tests whether `rel` is a function from `bs` to `cs`. The function `injection` tests whether a relation is an injective function. `bijectionFromTo(rel, bs, cs)` tests whether `rel` is a bijection from `bs` to `cs`.

`restrictFunction(rel, bs)` restricts the function `rel` to `bs`; it issues an error message if `rel` is not a function, or `bs` is not a subset of the domain of `rel`. And, `updateFunction(rel, a, b)` returns the updating of the function `rel` to send `a` to `b`; it issues an error message if `rel` isn't a function.

Here is how we can work with total orderings using functions from `SymRel`:

```
- val rel = SymRel.input "";
@ (0, 1), (1, 2), (0, 2), (0, 0), (1, 1), (2, 2)
@ .
val rel = - : sym_rel
- SymRel.output("", rel);
(0, 0), (0, 1), (0, 2), (1, 1), (1, 2), (2, 2)
val it = () : unit
- SymSet.output("", SymRel.domain rel);
0, 1, 2
val it = () : unit
- SymSet.output("", SymRel.range rel);
0, 1, 2
val it = () : unit
- SymRel.relationFromTo
= (rel, SymSet.fromString "0, 1, 2", SymSet.fromString "0, 1, 2");
val it = true : bool
- SymRel.relationOn(rel, SymSet.fromString "0, 1, 2");
val it = true : bool
- SymRel.reflexive(rel, SymSet.fromString "0, 1, 2");
val it = true : bool
- SymRel.symmetric rel;
val it = false : bool
- SymRel.antisymmetric rel;
val it = true : bool
- SymRel.transitive rel;
val it = true : bool
- SymRel.total(rel, SymSet.fromString "0, 1, 2");
val it = true : bool
- val rel' = SymRel.inverse rel;
```

```

val rel' = - : sym_rel
- SymRel.output("", rel');
(0, 0), (1, 0), (1, 1), (2, 0), (2, 1), (2, 2)
val it = () : unit
- val rel'' = SymRel.compose(rel', rel);
val rel'' = - : sym_rel
- SymRel.output("", rel'');
(0, 0), (0, 1), (0, 2), (1, 0), (1, 1), (1, 2), (2, 0), (2, 1),
(2, 2)
val it = () : unit

```

And here is how we can work with relations that are functions:

```

- val rel = SymRel.input "";
@ (1, 2), (2, 3), (3, 4)
@ .
val rel = - : sym_rel
- SymRel.output("", rel);
(1, 2), (2, 3), (3, 4)
val it = () : unit
- SymSet.output("", SymRel.domain rel);
1, 2, 3
val it = () : unit
- SymSet.output("", SymRel.range rel);
2, 3, 4
val it = () : unit
- SymRel.function rel;
val it = true : bool
- SymRel.functionFromTo
= (rel, SymSet.fromString "1, 2, 3", SymSet.fromString "2, 3, 4");
val it = true : bool
- SymRel.injection rel;
val it = true : bool
- SymRel.bijectionFromTo
= (rel, SymSet.fromString "1, 2, 3", SymSet.fromString "2, 3, 4");
val it = true : bool
- val f = SymRel.applyFunction rel;
val f = fn : sym -> sym
- Sym.output("", f(Sym.fromString "1"));
2
val it = () : unit
- Sym.output("", f(Sym.fromString "2"));
3
val it = () : unit
- Sym.output("", f(Sym.fromString "3"));
4
val it = () : unit
- Sym.output("", f(Sym.fromString "4"));
argument not in domain

```



```
uncaught exception Error
- val rel' = SymRel.input "";
@ (4, 3), (3, 2), (2, 1)
@ .
val rel' = - : sym_rel
- val rel'' = SymRel.compose(rel', rel);
val rel'' = - : sym_rel
- SymRel.functionFromTo
= (rel'', SymSet.fromString "1, 2, 3",
= SymSet.fromString "1, 2, 3");
val it = true : bool
- SymRel.output("", rel'');
(1, 1), (2, 2), (3, 3)
val it = () : unit
```

2.3.9 Notes

The book and toolset were designed and developed together, which made it possible to minimize the notational and conceptual distance between the two.

Chapter 3

Regular Languages

In this chapter, we study our most restrictive set of languages, the regular languages. We begin by introducing regular expressions, and saying that a language is regular iff it is generated by a regular expression. We study regular expression equivalence, look at how regular expressions can be synthesized and proved correct, and study several algorithms for regular expression simplification.

We go on to study five kinds of finite automata, culminating in finite automata whose transitions are labeled by regular expressions. We introduce methods for synthesizing and proving the correctness of finite automata, and study numerous algorithms for processing and converting between regular expressions and finite automata. Because of these conversions, the set of languages accepted by the finite automata is exactly the regular languages. The chapter concludes by considering the application of regular expressions and finite automata to searching in text files, lexical analysis, and the design of finite state systems.

3.1 Regular Expressions and Languages

In this section, we define several operations on languages, say what regular expressions are, what they mean, and what regular languages are, and begin to show how regular expressions can be processed by Forlan.

3.1.1 Operations on Languages

The union, intersection and set-difference operations on sets are also operations on languages, i.e., if $L_1, L_2 \in \mathbf{Lan}$, then $L_1 \cup L_2$, $L_1 \cap L_2$ and $L_1 - L_2$ are all languages. (Since $L_1, L_2 \in \mathbf{Lan}$, we have that $L_1 \subseteq \Sigma_1^*$ and $L_2 \subseteq \Sigma_2^*$, for alphabets Σ_1 and Σ_2 . Let $\Sigma = \Sigma_1 \cup \Sigma_2$, so that Σ is an alphabet, $L_1 \subseteq \Sigma^*$ and $L_2 \subseteq \Sigma^*$. Thus $L_1 \cup L_2$, $L_1 \cap L_2$ and $L_1 - L_2$ are all subsets of Σ^* , and so are all languages.)

The first new operation on languages is language concatenation. The *concatenation* of languages L_1 and L_2 ($L_1 @ L_2$) is the language

$$\{x_1 @ x_2 \mid x_1 \in L_1 \text{ and } x_2 \in L_2\}.$$

I.e., $L_1 @ L_2$ consists of all strings that can be formed by concatenating an element of L_1 with an element of L_2 . For example,

$$\begin{aligned} \{ab, abc\} @ \{cd, d\} &= \{(ab)(cd), (ab)(d), (abc)(cd), (abc)(d)\} \\ &= \{abcd, abd, abccd\}. \end{aligned}$$

Note that, if $L_1, L_2 \subseteq \Sigma^*$, for an alphabet Σ , then $L_1 @ L_2 \subseteq \Sigma^*$.

Concatenation of languages is associative: for all $L_1, L_2, L_3 \in \mathbf{Lan}$,

$$(L_1 @ L_2) @ L_3 = L_1 @ (L_2 @ L_3).$$

And, $\{\epsilon\}$ is the identity for concatenation: for all $L \in \mathbf{Lan}$,

$$\{\epsilon\} @ L = L @ \{\epsilon\} = L.$$

Furthermore, \emptyset is the zero for concatenation: for all $L \in \mathbf{Lan}$,

$$\emptyset @ L = L @ \emptyset = \emptyset.$$

We often abbreviate $L_1 @ L_2$ to $L_1 L_2$.

Now that we know what language concatenation is, we can say what it means to raise a language to a power. We define the *language L^n formed by raising a language L to the power $n \in \mathbb{N}$* by recursion on n :

$$\begin{aligned} L^0 &= \{\epsilon\}, \text{ for all } L \in \mathbf{Lan}; \text{ and} \\ L^{n+1} &= LL^n, \text{ for all } L \in \mathbf{Lan} \text{ and } n \in \mathbb{N}. \end{aligned}$$

We assign this exponentiation operation higher precedence than concatenation, so that LL^n means $L(L^n)$ in the above definition. Note that, if $L \subseteq \Sigma^*$, for an alphabet Σ , then $L^n \subseteq \Sigma^*$, for all $n \in \mathbb{N}$.

For example, we have that

$$\begin{aligned} \{a, b\}^2 &= \{a, b\}\{a, b\}^1 = \{a, b\}\{a, b\}\{a, b\}^0 \\ &= \{a, b\}\{a, b\}\{\epsilon\} = \{a, b\}\{a, b\} \\ &= \{aa, ab, ba, bb\}. \end{aligned}$$

Proposition 3.1.1

For all $L \in \mathbf{Lan}$ and $n, m \in \mathbb{N}$, $L^{n+m} = L^n L^m$.

Proof. An easy mathematical induction on n . The language L and the natural number m can be fixed at the beginning of the proof. \square

Thus, if $L \in \mathbf{Lan}$ and $n \in \mathbb{N}$, then

$$L^{n+1} = LL^n \quad (\text{definition}),$$

and

$$L^{n+1} = L^n L^1 = L^n L \quad (\text{Proposition 3.1.1}).$$

Another useful fact about language exponentiation is:

Proposition 3.1.2

For all $w \in \mathbf{Str}$ and $n \in \mathbb{N}$, $\{w\}^n = \{w^n\}$.

Proof. By mathematical induction on n . \square

For example, we have that $\{01\}^4 = \{(01)^4\} = \{01010101\}$.

Now we consider a language operation that is named after Stephen Cole Kleene, one of the founders of formal language theory. The *Kleene closure* (or just *closure*) of a language L (L^*) is the language

$$\bigcup \{L^n \mid n \in \mathbb{N}\}.$$

Thus, for all w ,

$$\begin{aligned} w \in L^* \quad &\text{iff} \quad w \in A, \text{ for some } A \in \{L^n \mid n \in \mathbb{N}\} \\ &\text{iff} \quad w \in L^n \text{ for some } n \in \mathbb{N}. \end{aligned}$$

Or, in other words:

- $L^* = L^0 \cup L^1 \cup L^2 \cup \dots$; and
- L^* consists of all strings that can be formed by concatenating together some number (maybe none) of elements of L (the same element of L can be used as many times as is desired).

For example,

$$\begin{aligned} \{a, ba\}^* &= \{a, ba\}^0 \cup \{a, ba\}^1 \cup \{a, ba\}^2 \cup \dots \\ &= \{\epsilon\} \cup \{a, ba\} \cup \{aa, aba, baa, baba\} \cup \dots \end{aligned}$$

If L is a language, then $L \subseteq \Sigma^*$ for some alphabet Σ , and thus L^* is also a subset of Σ^* —showing that L^* is a language, not just a set of strings.

Suppose $w \in \mathbf{Str}$. By Proposition 3.1.2, we have that, for all x ,

$$\begin{aligned} x \in \{w\}^* \quad &\text{iff} \quad x \in \{w\}^n, \text{ for some } n \in \mathbb{N}, \\ &\text{iff} \quad x \in \{w^n\}, \text{ for some } n \in \mathbb{N}, \\ &\text{iff} \quad x = w^n, \text{ for some } n \in \mathbb{N}. \end{aligned}$$

If we write $\{0, 1\}^*$, then this could mean:

- all strings over the alphabet $\{0,1\}$ (Section 2.1); or
- the closure of the language $\{0,1\}$.

Fortunately, these languages are equal (both are all strings of 0's and 1's), and this kind of ambiguity is harmless.

We assign our operations on languages relative precedences as follows:

Highest: closure $((\cdot)^*)$ and raising to a power $((\cdot)^n)$;

Intermediate: concatenation ($@$, or just juxtapositioning); and

Lowest: union (\cup), intersection (\cap) and difference ($-$).

For example, if $n \in \mathbb{N}$ and $A, B, C \in \mathbf{Lan}$, then $A^*BC^n \cup B$ abbreviates $((A^*)B(C^n)) \cup B$. The language $((A \cup B)C)^*$ can't be abbreviated, since removing either pair of parentheses will change its meaning. If we removed the outer pair, then we would have $(A \cup B)(C^*)$, and removing the inner pair would yield $(A \cup (BC))^*$.

Suppose L , L_1 and L_2 are languages, and $n \in \mathbb{N}$. It is easy to see that $\mathbf{alphabet}(L_1 \cup L_2) = \mathbf{alphabet}(L_1) \cup \mathbf{alphabet}(L_2)$. And, if L_1 and L_2 are both nonempty, then $\mathbf{alphabet}(L_1 L_2) = \mathbf{alphabet}(L_1) \cup \mathbf{alphabet}(L_2)$, and otherwise, $\mathbf{alphabet}(L_1 L_2) = \emptyset$. Furthermore, if $n \geq 1$, then $\mathbf{alphabet}(L^n) = \mathbf{alphabet}(L)$; otherwise, $\mathbf{alphabet}(L^n) = \emptyset$. Finally, we have that $\mathbf{alphabet}(L^*) = \mathbf{alphabet}(L)$.

In Section 2.3, we introduced the Forlan module `StrSet`, which defines various functions for processing finite sets of strings, i.e., finite languages. This module also defines the functions

```
val concat : str set * str set -> str set
val power  : str set * int -> str set
```

which implement our concatenation and exponentiation operations on finite languages. Here are some examples of how these functions can be used:

```
- val xs = StrSet.fromString "ab, cd";
val xs = - : str set
- val ys = StrSet.fromString "uv, wx";
val ys = - : str set
- StrSet.output("", StrSet.concat(xs, ys));
abuv, abwx, cduv, cdwx
val it = () : unit
- StrSet.output("", StrSet.power(xs, 0));
%
val it = () : unit
- StrSet.output("", StrSet.power(xs, 1));
ab, cd
val it = () : unit
```

```

- StrSet.output("", StrSet.power(xs, 2));
abab, abcd, cdab, cdc d
val it = () : unit
- StrSet.output("", StrSet.power(xs, 3));
ababab, ababcd, abcdab, abcdcd, cdabab, cdabcd, cdc dcd
val it = () : unit

```

3.1.2 Regular Expressions

Next, we define the set of all regular expressions. Let the set **RegLab** of *regular expression labels* be

$$\mathbf{Sym} \cup \{\%, \$, *, @, +\}.$$

Let the set **Reg** of *regular expressions* be the least subset of **TreeRegLab** such that:

- (empty string) $\% \in \mathbf{Reg}$;
- (empty set) $\$ \in \mathbf{Reg}$;
- (symbol) for all $a \in \mathbf{Sym}$, $a \in \mathbf{Reg}$;
- (closure) for all $\alpha \in \mathbf{Reg}$, $*(\alpha) \in \mathbf{Reg}$;
- (concatenation) for all $\alpha, \beta \in \mathbf{Reg}$, $@(\alpha, \beta) \in \mathbf{Reg}$; and
- (union) for all $\alpha, \beta \in \mathbf{Reg}$, $+(\alpha, \beta) \in \mathbf{Reg}$.


This is yet another example of an inductive definition. The elements of **Reg** are precisely those **RegLab**-trees (trees (See Section 1.3) whose labels come from **RegLab**) that can be built using these six rules.

Whenever possible, we will use the mathematical variables α , β and γ to name regular expressions. Since regular expressions are **RegLab**-trees, we may talk of their sizes and heights.

For example,

$$+(\@(*0), \@1, *(0)),$$

i.e.,



is a regular expression. On the other hand, the **RegLab**-tree $*(*,*)$ is *not* a regular expression, since it can't be built using our six rules.

We order the elements of **RegLab** as follows:

$$\% < \$ < \text{symbols in order} < * < @ < +.$$

It is important that $+$ be the greatest element of **RegLab**; if this were not so, then the definition of weakly simplified regular expressions (see Section 3.3) would have to be altered.

We order regular expressions first by their root labels, and then, recursively, by their children, working from left to right. For example, we have that

$$\% < *(%) < *(@(\$, *(\$))) < *(@(\mathbf{a}, \%)) < @(\%, \$).$$

Because **Reg** is defined inductively, it gives rise to an induction principle.

Theorem 3.1.3 (Principle of Induction on Regular Expressions)

Suppose $P(\alpha)$ is a property of a regular expression α . If

- $P(\%)$,
- $P(\$)$,
- for all $a \in \mathbf{Sym}$, $P(a)$,
- for all $\alpha \in \mathbf{Reg}$, if $P(\alpha)$, then $P(*(\alpha))$,
- for all $\alpha, \beta \in \mathbf{Reg}$, if $P(\alpha)$ and $P(\beta)$, then $P(@(\alpha, \beta))$, and
- for all $\alpha, \beta \in \mathbf{Reg}$, if $P(\alpha)$ and $P(\beta)$, then $P(+(\alpha, \beta))$,

then

$$\text{for all } \alpha \in \mathbf{Reg}, P(\alpha).$$

To increase readability, we use infix and postfix notation, abbreviating:

- $*(\alpha)$ to α^* or $\alpha*$;
- $@(\alpha, \beta)$ to $\alpha @ \beta$; and
- $+(\alpha, \beta)$ to $\alpha + \beta$.

We assign the operators $(\cdot)^*$, $@$ and $+$ the following precedences and associativities:

Highest: $(\cdot)^*$;

Intermediate: $@$ (right associative); and

Lowest: $+$ (right associative).

We parenthesize regular expressions when we need to override the default precedences and associativities, and for reasons of clarity. Furthermore, we often abbreviate $\alpha @ \beta$ to $\alpha\beta$.

For example, we can abbreviate the regular expression

$$+(\@(*(\mathbf{0}),\@(\mathbf{1},*(\mathbf{0}))),\%)$$

to $0^* @ 1 @ 0^* + \%$ or $0^*10^* + \%$. On the other hand, the regular expression $((0 + 1)2)^*$ can't be further abbreviated, since removing either pair of parentheses would result in a different regular expression. Removing the outer pair would result in $(0 + 1)(2^*) = (0 + 1)2^*$, and removing the inner pair would yield $(0 + (12))^* = (0 + 12)^*$.

Now we can say what regular expressions mean, using some of our language operations. The *language generated by* a regular expression α ($L(\alpha)$) is defined by recursion:

$$\begin{aligned} L(\%) &= \{\%\}; \\ L(\$) &= \emptyset; \\ L(a) &= \{[a]\} = \{a\}, \text{ for all } a \in \mathbf{Sym}; \\ L(*(\alpha)) &= L(\alpha)^*, \text{ for all } \alpha \in \mathbf{Reg}; \\ L(@(\alpha, \beta)) &= L(\alpha) @ L(\beta), \text{ for all } \alpha, \beta \in \mathbf{Reg}; \text{ and} \\ L(+(\alpha, \beta)) &= L(\alpha) \cup L(\beta), \text{ for all } \alpha, \beta \in \mathbf{Reg}. \end{aligned}$$

This is a good definition since, if L is a language, then so is L^* , and, if L_1 and L_2 are languages, then so are L_1L_2 and $L_1 \cup L_2$. We say that w is *generated by* α iff $w \in L(\alpha)$.

For example,

$$\begin{aligned} L(0^*10^* + \%) &= L(+(@(*(\mathbf{0}),\@(\mathbf{1},*(\mathbf{0}))),\%)) \\ &= L(@(*(\mathbf{0}),\@(\mathbf{1},*(\mathbf{0})))) \cup L(\%) \\ &= L(*(\mathbf{0}))L(@(\mathbf{1},*(\mathbf{0}))) \cup \{\%\} \\ &= L(\mathbf{0})^*L(\mathbf{1})L(*(\mathbf{0})) \cup \{\%\} \\ &= \{0\}^*\{1\}L(\mathbf{0})^* \cup \{\%\} \\ &= \{0\}^*\{1\}\{0\}^* \cup \{\%\} \\ &= \{0^n10^m \mid n, m \in \mathbb{N}\} \cup \{\%\}. \end{aligned}$$

E.g., 0001000, 10, 001 and % are generated by $0^*10^* + \%$.

We define functions $\mathbf{symToReg} \in \mathbf{Sym} \rightarrow \mathbf{Reg}$ and $\mathbf{strToReg} \in \mathbf{Str} \rightarrow \mathbf{Reg}$, as follows. Given a symbol $a \in \mathbf{Sym}$, $\mathbf{symToReg} a = a$. And, given a string x , $\mathbf{strToReg} x$ is the *canonical regular expression for* x : %, if $x = \%$, and $@(a_1, @(a_2, \dots a_n \dots)) = a_1a_2 \dots a_n$, if $x = a_1a_2 \dots a_n$, for symbols a_1, a_2, \dots, a_n

and $n \geq 1$. It is easy to see that, for all $a \in \mathbf{Sym}$, $L(\mathbf{symToReg} a) = \{a\}$, and, for all $x \in \mathbf{Str}$, $L(\mathbf{strToReg} x) = \{x\}$.

We define the *regular expression α^n formed by raising a regular expression α to the power $n \in \mathbb{N}$* by recursion on n :

$$\begin{aligned}\alpha^0 &= \%, \text{ for all } \alpha \in \mathbf{Reg}; \\ \alpha^1 &= \alpha, \text{ for all } \alpha \in \mathbf{Reg}; \text{ and} \\ \alpha^{n+1} &= \alpha\alpha^n, \text{ for all } \alpha \in \mathbf{Reg} \text{ and } n \in \mathbb{N} - \{0\}.\end{aligned}$$

We assign this operation the same precedence as closure, so that $\alpha\alpha^n$ means $\alpha(\alpha^n)$ in the above definition. Note that, in contrast to the definitions of x^n and L^n , we have split the case $n+1$ into two subcases, depending upon whether $n = 0$ or $n \geq 1$. Thus α^1 is α , not $\alpha\%$. For example, $(0+1)^3 = (0+1)(0+1)(0+1)$.

Proposition 3.1.4

For all $\alpha \in \mathbf{Reg}$ and $n \in \mathbb{N}$, $L(\alpha^n) = L(\alpha)^n$.

Proof. An easy mathematical induction on n . α may be fixed at the beginning of the proof. \square

An example consequence of the proposition is that $L((0+1)^3) = L(0+1)^3 = \{0,1\}^3$, the set of all strings of 0's and 1's of length 3.

We define $\mathbf{alphabet} \in \mathbf{Reg} \rightarrow \mathbf{Alp}$ by recursion:

$$\begin{aligned}\mathbf{alphabet} \% &= \emptyset; \\ \mathbf{alphabet} \$ &= \emptyset; \\ \mathbf{alphabet} a &= \{a\}, \text{ for all } a \in \mathbf{Sym}; \\ \mathbf{alphabet} (*(\alpha)) &= \mathbf{alphabet} \alpha, \text{ for all } \alpha \in \mathbf{Reg}; \\ \mathbf{alphabet} (@(\alpha, \beta)) &= \mathbf{alphabet} \alpha \cup \mathbf{alphabet} \beta, \text{ for all } \alpha, \beta \in \mathbf{Reg}; \text{ and} \\ \mathbf{alphabet} (+(\alpha, \beta)) &= \mathbf{alphabet} \alpha \cup \mathbf{alphabet} \beta, \text{ for all } \alpha, \beta \in \mathbf{Reg}.\end{aligned}$$

This is a good definition, since the union of two alphabets is an alphabet. For example, $\mathbf{alphabet}(0^*10^* + \%) = \{0,1\}$. We say that $\mathbf{alphabet} \alpha$ is *the alphabet of a regular expression α* .

Proposition 3.1.5

For all $\alpha \in \mathbf{Reg}$, $\mathbf{alphabet}(L(\alpha)) \subseteq \mathbf{alphabet} \alpha$.

In other words, the proposition says that every symbol of every string in $L(\alpha)$ comes from $\mathbf{alphabet} \alpha$.

Proof. An easy induction on regular expressions. \square

For example, since $L(1\$) = \{1\}\emptyset = \emptyset$, we have that

$$\begin{aligned} \text{alphabet}(L(0^* + 1\$)) &= \text{alphabet}(\{0\}^*) \\ &= \{0\} \\ &\subseteq \{0, 1\} \\ &= \text{alphabet}(0^* + 1\$). \end{aligned}$$

Next, we define some useful auxiliary functions on regular expressions. The *generalized concatenation* function $\text{genConcat} \in \mathbf{List\ Reg} \rightarrow \mathbf{Reg}$ is defined by right recursion:

$$\begin{aligned} \text{genConcat } [] &= \%, \\ \text{genConcat } [\alpha] &= \alpha, \text{ and} \\ \text{genConcat } ([\alpha] @ \bar{\alpha}) &= @(\alpha, \text{genConcat } \bar{\alpha}), \text{ if } \bar{\alpha} \neq []. \end{aligned}$$

And the *generalized union* function $\text{genUnion} \in \mathbf{List\ Reg} \rightarrow \mathbf{Reg}$ is defined by right recursion:

$$\begin{aligned} \text{genUnion } [] &= \%, \\ \text{genUnion } [\alpha] &= \alpha, \text{ and} \\ \text{genUnion } ([\alpha] @ \bar{\alpha}) &= +(\alpha, \text{genUnion } \bar{\alpha}), \text{ if } \bar{\alpha} \neq []. \end{aligned}$$

E.g., $\text{genConcat}[1, 0, 12, 3 + 4] = 10(12)(3 + 4)$ and $\text{genUnion}[1, 0, 12, 3 + 4] = 1 + 0 + (12) + 3 + 4$.

$\text{rightConcat} \in \mathbf{Reg} \times \mathbf{Reg} \rightarrow \mathbf{Reg}$ is defined by structural recursion on its first argument:

$$\begin{aligned} \text{rightConcat}(@(\alpha_1, \alpha_2), \beta) &= @(\alpha_1, \text{rightConcat}(\alpha_2, \beta)), \text{ and} \\ \text{rightConcat}(\alpha, \beta) &= @(\alpha, \beta), \text{ if } \alpha \text{ is not a concatenation.} \end{aligned}$$

And $\text{rightUnion} \in \mathbf{Reg} \times \mathbf{Reg} \rightarrow \mathbf{Reg}$ is defined by structural recursion on its first argument:

$$\begin{aligned} \text{rightUnion}(+(\alpha_1, \alpha_2), \beta) &= +(\alpha_1, \text{rightUnion}(\alpha_2, \beta)), \text{ and} \\ \text{rightUnion}(\alpha, \beta) &= +(\alpha, \beta), \text{ if } \alpha \text{ is not a union.} \end{aligned}$$

E.g., $\text{rightConcat}(012, 345) = 012345$ and $\text{rightUnion}(0 + 1 + 2, 1 + 2 + 3) = 0 + 1 + 2 + 1 + 2 + 3$.

$\text{concatsToList} \in \mathbf{Reg} \rightarrow \mathbf{List\ Reg}$ is defined by structural recursion:

$$\begin{aligned} \text{concatsToList}(@(\alpha, \beta)) &= [\alpha] @ \text{concatsToList } \beta, \text{ and} \\ \text{concatsToList } \alpha &= \alpha, \text{ if } \alpha \text{ is not a concatenation.} \end{aligned}$$

And **unionsToList** $\in \mathbf{Reg} \rightarrow \mathbf{List\ Reg}$ is defined by structural recursion:

$$\begin{aligned} \mathbf{unionsToList}(+(\alpha, \beta)) &= [\alpha] @ \mathbf{unionsToList} \beta, \text{ and} \\ \mathbf{unionsToList} \alpha &= \alpha, \text{ if } \alpha \text{ is not a union.} \end{aligned}$$

E.g., **concatToList** $((12)34) = [12, 3, 4]$ and **unionsToList** $((0 + 1) + 2 + 3) = [0 + 1, 2, 3]$.

Finally, **sortUnions** $\in \mathbf{Reg} \rightarrow \mathbf{Reg}$ is defined by:

$$\mathbf{sortUnions} \alpha = \mathbf{genUnion} \bar{\beta},$$

where $\bar{\beta}$ is the result of sorting the elements of **unionsToList** α into strictly ascending order (without duplicates), according to our total ordering on regular expressions. E.g., **sortUnions** $(1 + 0 + 23 + 1) = 0 + 1 + 23$.

We define functions **allSym** $\in \mathbf{Alp} \rightarrow \mathbf{Reg}$ and **allStr** $\in \mathbf{Alp} \rightarrow \mathbf{Reg}$ as follows. Given an alphabet Σ , **allSym** Σ is the *all symbols regular expression* for Σ : $a_1 + \dots + a_n$, where a_1, \dots, a_n are the elements of Σ , listed in order and without repetition (when $n = 0$, we use $\$$, and when $n = 1$, we use a_1). And, given an alphabet Σ , **allStr** Σ is the *all strings regular expression* for Σ : $(\mathbf{allSym} \Sigma)^*$. For example,

$$\begin{aligned} \mathbf{allSym} \{0, 1, 2\} &= 0 + 1 + 2, \text{ and} \\ \mathbf{allStr} \{0, 1, 2\} &= (0 + 1 + 2)^*. \end{aligned}$$

Thus, for all $\Sigma \in \mathbf{Alp}$,

$$\begin{aligned} L(\mathbf{allSym} \Sigma) &= \{ [a] \mid a \in \Sigma \} = \{ a \mid a \in \Sigma \}, \text{ and} \\ L(\mathbf{allStr} \Sigma) &= \Sigma^*. \end{aligned}$$

Now we are able to say what it means for a language to be regular: a language L is *regular* iff $L = L(\alpha)$ for some $\alpha \in \mathbf{Reg}$. We define

$$\begin{aligned} \mathbf{RegLan} &= \{ L(\alpha) \mid \alpha \in \mathbf{Reg} \} \\ &= \{ L \in \mathbf{Lan} \mid L \text{ is regular} \}. \end{aligned}$$

Since every regular expression can be described, e.g., in fully parenthesized form, by a finite sequence of ASCII characters, we can enumerate the regular expressions, and consequently we have that **Reg** is countably infinite. Since $\{0^0\}$, $\{0^1\}$, $\{0^2\}$, \dots , are all regular languages, we have that **RegLan** is infinite. Furthermore, we can establish an injection h from **RegLan** to **Reg**: $h L$ is the first (in our enumeration of regular expressions) α such that $L(\alpha) = L$. Because **Reg** is countably infinite, it follows that there is an injection from **Reg** to \mathbb{N} . Composing these injections, gives us an injection from **RegLan** to \mathbb{N} . And when observing that **RegLan** is infinite, we implicitly gave an injection from \mathbb{N}

to **RegLan**. Thus, by the Schröder-Bernstein Theorem, we have that **RegLan** and \mathbb{N} have the same size, so that **RegLan** is countably infinite.

Since **RegLan** is countably infinite but **Lan** is uncountable, it follows that **RegLan** \subsetneq **Lan**, i.e., there are non-regular languages. In Section ??, we will see a concrete example of a non-regular language.

3.1.3 Processing Regular Expressions in Forlan

Now, we turn to the Forlan implementation of regular expressions. The Forlan module **Reg** defines the abstract type **reg** (in the top-level environment) of regular expressions, as well as various functions and constants for processing regular expressions, including:

```

val input      : string -> reg
val output     : string * reg -> unit
val size       : reg -> int
val numLeaves  : reg -> int
val height     : reg -> int
val emptyStr   : reg
val emptySet   : reg
val fromSym    : sym -> reg
val closure    : reg -> reg
val concat     : reg * reg -> reg
val union      : reg * reg -> reg
val compare    : reg * reg -> order
val equal      : reg * reg -> bool
val fromStr    : str -> reg
val power      : reg * int -> reg
val alphabet   : reg -> sym set
val genConcat  : reg list -> reg
val genUnion   : reg list -> reg
val rightConcat : reg * reg -> reg
val rightUnion : reg * reg -> reg
val concatsToList : reg -> reg list
val unionsToList : reg -> reg list
val sortUnions : reg -> reg
val allSym     : sym set -> reg
val allStr     : sym set -> reg
val fromStrSet : str set -> reg

```

The Forlan syntax for regular expressions is the infix/postfix one introduced in the preceding subsection, where $\alpha @ \beta$ is always written as $\alpha\beta$, and we use parentheses to override default precedences/associativities, or simply for clarity. For example, $0^*10^* + \%$ and $(0^*(1(0^*))) + \%$ are the same regular expression. And, $((0^*)1)0^* + \%$ is a different regular expression, but one with the same meaning. Furthermore, $0^*1(0^* + \%)$ is not only different from the two preceding regular expressions, but it has a different meaning (it fails to generate $\%$.) When regular expressions are outputted, as few parentheses as possible are used.

The functions `size`, `numLeaves` and `height` return the size, number of leaves and height, respectively, of a regular expression. The values `emptyStr` and `emptySet` are `%` and `$`, respectively. The function `fromSym` takes in a symbol a and returns the regular expression a . It is available in the top-level environment as `symToReg`. The function `closure` takes in a regular expression α and returns $\ast(\alpha)$. The function `concat` takes a pair (α, β) of regular expressions and returns $\alpha\beta$. The function `union` takes a pair (α, β) of regular expressions and returns $\alpha + \beta$. The function `compare` implements our total ordering on regular expressions, and `equal` tests whether two regular expressions are equal. The function `fromStr` implements the function `strToReg`, and is also available in the top-level environment as `strToReg`. The function `power` raises a regular expression to a power, and the function `alphabet` returns the alphabet of a regular expression. Finally, the functions `genConcat`, `genUnion`, `rightConcat`, `rightUnion`, `concatToList`, `unionToList`, `sortUnions`, `allSym` and `allStr` implement the functions with the same names. The function `fromStrSet` returns `$`, if called with the empty set. Otherwise, it returns `fromStr $x_1 + \dots + \text{fromStr } x_n$` , where x_1, \dots, x_n are the elements of its argument, listed in strictly ascending order.

Here are some example uses of the functions of `Reg`:

```
- val reg = Reg.input "";
@ 0*10* + %
@ .
val reg = - : reg
- Reg.size reg;
val it = 9 : int
- Reg.numLeaves reg;
val it = 4 : int
- val reg' = Reg.fromStr(Str.power(Str.input "", 3));
@ 01
@ .
val reg' = - : reg
- Reg.output("", reg');
010101
val it = () : unit
- Reg.size reg';
val it = 11 : int
- Reg.numLeaves reg';
val it = 6 : int
- Reg.compare(reg, reg');
val it = GREATER : order
- val reg'' = Reg.concat(Reg.closure reg, reg');
val reg'' = - : reg
- Reg.output("", reg'');
(0*10* + %)*010101
val it = () : unit
- SymSet.output("", Reg.alphabet reg'');
```

```

0, 1
val it = () : unit
- val reg''' = Reg.power(reg, 3);
val reg''' = - : reg
- Reg.output("", reg''');
(0*10* + %)(0*10* + %)(0*10* + %)
val it = () : unit
- Reg.size reg''';
val it = 29 : int
- Reg.numLeaves reg''';
val it = 12 : int
- Reg.output("", Reg.fromString "(0*(1(0*))) + %");
0*10* + %
val it = () : unit
- Reg.output("", Reg.fromString "(0*1)0* + %");
(0*1)0* + %
val it = () : unit
- Reg.output("", Reg.fromString "0*1(0* + %)");
0*1(0* + %)
val it = () : unit
- Reg.equal
= (Reg.fromString "0*10* + %",
  = Reg.fromString "0*1(0* + %)");
val it = false : bool

```

We can use the functions `genConcat`, `genUnion`, `rightConcat`, `rightUnion`, `concatstoList`, `unionsToList` and `sortUnions` as follows:

```

- Reg.output("", Reg.genConcat nil);
%
val it = () : unit
- Reg.output("", Reg.genUnion nil);
$
val it = () : unit
- val regs =
= [Reg.fromString "01", Reg.fromString "01 + 12",
  = Reg.fromString "(1 + 2)*", Reg.fromString "3 + 4"];
val regs = [-,-,-,-] : reg list
- Reg.output("", Reg.genConcat regs);
(01)(01 + 12)(1 + 2)*(3 + 4)
val it = () : unit
- Reg.output("", Reg.genUnion regs);
01 + (01 + 12) + (1 + 2)* + 3 + 4
val it = () : unit
- Reg.output
= ("",
  = Reg.rightConcat
  = (Reg.fromString "0123", Reg.fromString "4567"));
01234567

```

```

val it = () : unit
- Reg.output
= ("",
= Reg.rightUnion
= (Reg.fromString "0 + 1 + 2", Reg.fromString "1 + 2 + 3"));
0 + 1 + 2 + 1 + 2 + 3
val it = () : unit
- map
= Reg.toString
= (Reg.concatsToList(Reg.fromString "0(12)34"));
val it = ["0","12","3","4"] : string list
- map
= Reg.toString
= (Reg.unionsToList(Reg.fromString "0 + (1 + 2) + 3 + 0"));
val it = ["0","1 + 2","3","0"] : string list
- Reg.output
= ("", Reg.sortUnions(Reg.fromString "12 + 0 + 3 + 0"));
0 + 3 + 12
val it = () : unit

```

We can use the functions `allSym`, `allStr` and `fromStrSet` like this:

```

- Reg.output("", Reg.allSym(SymSet.fromString ""));
$
val it = () : unit
- Reg.output("", Reg.allSym(SymSet.fromString "0"));
0
val it = () : unit
- Reg.output("", Reg.allSym(SymSet.fromString "0, 1, 2"));
0 + 1 + 2
val it = () : unit
- Reg.output("", Reg.allStr(SymSet.fromString ""));
$*
val it = () : unit
- Reg.output("", Reg.allStr(SymSet.fromString "0"));
0*
val it = () : unit
- Reg.output("", Reg.allStr(SymSet.fromString "2, 1, 0"));
(0 + 1 + 2)*
val it = () : unit
- Reg.output
= ("",
= Reg.fromStrSet(StrSet.fromString "one, two, three, four"));
one + two + four + three
val it = () : unit

```

3.1.4 JForlan

The Java program JForlan can be used to view and edit regular expression trees. It can be invoked directly, or run via Forlan. See the Forlan website for more

information.

3.1.5 Notes

A novel feature of this book is that regular expressions are trees, so that our linear syntax for regular expressions is derived rather than primary. Thus regular expression equality is just tree equality, and it's easy to explain when parentheses are necessary in a linear description of a regular expression. Furthermore, tree-oriented concepts, notation and operations automatically apply to regular expressions, letting us, e.g., give definitions by structural recursion.

3.2 Equivalence and Correctness of Regular Expressions

In this section, we say what it means for regular expressions to be equivalent, show a series of results about regular expression equivalence, and consider how regular expressions may be designed and proved correct.

3.2.1 Equivalence of Regular Expressions

Regular expressions α and β are *equivalent* iff $L(\alpha) = L(\beta)$. In other words, α and β are equivalent iff α and β generate the same language. We define a relation \approx on **Reg** by: $\alpha \approx \beta$ iff α and β are equivalent. For example, $L((00)^* + \%) = L((00)^*)$, and thus $(00)^* + \% \approx (00)^*$.

One approach to showing that $\alpha \approx \beta$ is to show that $L(\alpha) \subseteq L(\beta)$ and $L(\beta) \subseteq L(\alpha)$. The following proposition is useful for showing language inclusions, not just ones involving regular languages.

Proposition 3.2.1

- (1) For all $A_1, A_2, B_1, B_2 \in \mathbf{Lan}$, if $A_1 \subseteq B_1$ and $A_2 \subseteq B_2$, then $A_1 \cup A_2 \subseteq B_1 \cup B_2$.
- (2) For all $A_1, A_2, B_1, B_2 \in \mathbf{Lan}$, if $A_1 \subseteq B_1$ and $A_2 \subseteq B_2$, then $A_1 \cap A_2 \subseteq B_1 \cap B_2$.
- (3) For all $A_1, A_2, B_1, B_2 \in \mathbf{Lan}$, if $A_1 \subseteq B_1$ and $B_2 \subseteq A_2$, then $A_1 - A_2 \subseteq B_1 - B_2$.
- (4) For all $A_1, A_2, B_1, B_2 \in \mathbf{Lan}$, if $A_1 \subseteq B_1$ and $A_2 \subseteq B_2$, then $A_1 A_2 \subseteq B_1 B_2$.
- (5) For all $A, B \in \mathbf{Lan}$ and $n \in \mathbb{N}$, if $A \subseteq B$, then $A^n \subseteq B^n$.
- (6) For all $A, B \in \mathbf{Lan}$, if $A \subseteq B$, then $A^* \subseteq B^*$.

In Part (3), note that the second part of the sufficient condition for knowing $A_1 - A_2 \subseteq B_1 - B_2$ is $B_2 \subseteq A_2$, not $A_2 \subseteq B_2$.

Proof. (1) and (2) are straightforward. We show (3) as an example, below. (4) is easy. (5) is proved by mathematical induction, using (4). (6) is proved using (5).

For (3), suppose that $A_1, A_2, B_1, B_2 \in \mathbf{Lan}$, $A_1 \subseteq B_1$ and $B_2 \subseteq A_2$. To show that $A_1 - A_2 \subseteq B_1 - B_2$, suppose $w \in A_1 - A_2$. We must show that $w \in B_1 - B_2$. It will suffice to show that $w \in B_1$ and $w \notin B_2$.

Since $w \in A_1 - A_2$, we have that $w \in A_1$ and $w \notin A_2$. Since $A_1 \subseteq B_1$, it follows that $w \in B_1$. Thus, it remains to show that $w \notin B_2$.

Suppose, toward a contradiction, that $w \in B_2$. Since $B_2 \subseteq A_2$, it follows that $w \in A_2$ —contradiction. Thus we have that $w \notin B_2$. \square

Next we show that our relation \approx has some of the familiar properties of equality.

Proposition 3.2.2

(1) \approx is reflexive on \mathbf{Reg} , symmetric and transitive.

(2) For all $\alpha, \beta \in \mathbf{Reg}$, if $\alpha \approx \beta$, then $\alpha^* \approx \beta^*$.

(3) For all $\alpha_1, \alpha_2, \beta_1, \beta_2 \in \mathbf{Reg}$, if $\alpha_1 \approx \beta_1$ and $\alpha_2 \approx \beta_2$, then $\alpha_1 \alpha_2 \approx \beta_1 \beta_2$.

(4) For all $\alpha_1, \alpha_2, \beta_1, \beta_2 \in \mathbf{Reg}$, if $\alpha_1 \approx \beta_1$ and $\alpha_2 \approx \beta_2$, then $\alpha_1 + \alpha_2 \approx \beta_1 + \beta_2$.

Proof. Follows from the properties of $=$. As an example, we show Part (4).

Suppose $\alpha_1, \alpha_2, \beta_1, \beta_2 \in \mathbf{Reg}$, and assume that $\alpha_1 \approx \beta_1$ and $\alpha_2 \approx \beta_2$. Then $L(\alpha_1) = L(\beta_1)$ and $L(\alpha_2) = L(\beta_2)$, so that

$$\begin{aligned} L(\alpha_1 + \alpha_2) &= L(\alpha_1) \cup L(\alpha_2) = L(\beta_1) \cup L(\beta_2) \\ &= L(\beta_1 + \beta_2). \end{aligned}$$

Thus $\alpha_1 + \alpha_2 \approx \beta_1 + \beta_2$. \square

A consequence of Proposition 3.2.2 is the following proposition, which says that, if we replace a subtree of a regular expression α by an equivalent regular expression, that the resulting regular expression is equivalent to α .

Proposition 3.2.3

Suppose $\alpha, \beta, \beta' \in \mathbf{Reg}$, $\beta \approx \beta'$, $pat \in \mathbf{Path}$ is valid for α , and β is the subtree of α at position pat . Let α' be the result of replacing the subtree at position pat in α by β' . Then $\alpha \approx \alpha'$.

Proof. By induction on α . \square

Next, we state and prove some equivalences involving union.

Proposition 3.2.4

(1) For all $\alpha, \beta \in \mathbf{Reg}$, $\alpha + \beta \approx \beta + \alpha$.

- (2) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $(\alpha + \beta) + \gamma \approx \alpha + (\beta + \gamma)$.
- (3) For all $\alpha \in \mathbf{Reg}$, $\$ + \alpha \approx \alpha$.
- (4) For all $\alpha \in \mathbf{Reg}$, $\alpha + \alpha \approx \alpha$.
- (5) If $L(\alpha) \subseteq L(\beta)$, then $\alpha + \beta \approx \beta$.

Proof.

- (1) Follows from the commutativity of \cup .
- (2) Follows from the associativity of \cup .
- (3) Follows since \emptyset is the identity for \cup .
- (4) Follows since \cup is idempotent: $A \cup A = A$, for all sets A .
- (5) Follows since, if $L_1 \subseteq L_2$, then $L_1 \cup L_2 = L_2$.

□

Next, we consider equivalences for concatenation.

Proposition 3.2.5

- (1) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $(\alpha\beta)\gamma \approx \alpha(\beta\gamma)$.
- (2) For all $\alpha \in \mathbf{Reg}$, $\% \alpha \approx \alpha \approx \alpha \%$.
- (3) For all $\alpha \in \mathbf{Reg}$, $\$ \alpha \approx \$ \approx \alpha \$$.

Proof.

- (1) Follows from the associativity of language concatenation.
- (2) Follows since $\{\%\}$ is the identity for language concatenation.
- (3) Follows since \emptyset is the zero for language concatenation.

□

Next we consider the distributivity of concatenation over union. First, we prove a proposition concerning languages. Then, we use this proposition to show the corresponding proposition for regular expressions.

Proposition 3.2.6

- (1) For all $L_1, L_2, L_3 \in \mathbf{Lan}$, $L_1(L_2 \cup L_3) = L_1L_2 \cup L_1L_3$.
- (2) For all $L_1, L_2, L_3 \in \mathbf{Lan}$, $(L_1 \cup L_2)L_3 = L_1L_3 \cup L_2L_3$.

Proof. We show the proof of Part (1); the proof of the other part is similar. Suppose $L_1, L_2, L_3 \in \mathbf{Lan}$. It will suffice to show that

$$L_1(L_2 \cup L_3) \subseteq L_1L_2 \cup L_1L_3 \subseteq L_1(L_2 \cup L_3).$$

To see that $L_1(L_2 \cup L_3) \subseteq L_1L_2 \cup L_1L_3$, suppose $w \in L_1(L_2 \cup L_3)$. We must show that $w \in L_1L_2 \cup L_1L_3$. By our assumption, $w = xy$ for some $x \in L_1$ and $y \in L_2 \cup L_3$. There are two cases to consider.

- Suppose $y \in L_2$. Then $w = xy \in L_1L_2 \subseteq L_1L_2 \cup L_1L_3$.
- Suppose $y \in L_3$. Then $w = xy \in L_1L_3 \subseteq L_1L_2 \cup L_1L_3$.

To see that $L_1L_2 \cup L_1L_3 \subseteq L_1(L_2 \cup L_3)$, suppose $w \in L_1L_2 \cup L_1L_3$. We must show that $w \in L_1(L_2 \cup L_3)$. There are two cases to consider.

- Suppose $w \in L_1L_2$. Then $w = xy$ for some $x \in L_1$ and $y \in L_2$. Thus $y \in L_2 \cup L_3$, so that $w = xy \in L_1(L_2 \cup L_3)$.
- Suppose $w \in L_1L_3$. Then $w = xy$ for some $x \in L_1$ and $y \in L_3$. Thus $y \in L_2 \cup L_3$, so that $w = xy \in L_1(L_2 \cup L_3)$.

□

Proposition 3.2.7

(1) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\alpha(\beta + \gamma) \approx \alpha\beta + \alpha\gamma$.

(2) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $(\alpha + \beta)\gamma \approx \alpha\gamma + \beta\gamma$.

Proof. Follows from Proposition 3.2.6. Consider, e.g., the proof of Part (1). By Proposition 3.2.6(1), we have that

$$\begin{aligned} L(\alpha(\beta + \gamma)) &= L(\alpha)L(\beta + \gamma) \\ &= L(\alpha)(L(\beta) \cup L(\gamma)) \\ &= L(\alpha)L(\beta) \cup L(\alpha)L(\gamma) \\ &= L(\alpha\beta) \cup L(\alpha\gamma) \\ &= L(\alpha\beta + \alpha\gamma) \end{aligned}$$

Thus $\alpha(\beta + \gamma) \approx \alpha\beta + \alpha\gamma$. □

Finally, we turn our attention to equivalences for Kleene closure, first stating and proving some results for languages, and then stating and proving the corresponding results for regular expressions.

Proposition 3.2.8

- For all $L \in \mathbf{Lan}$, $LL^* \subseteq L^*$.
- For all $L \in \mathbf{Lan}$, $L^*L \subseteq L^*$.

Proof. E.g., to see that $LL^* \subseteq L^*$, suppose $w \in LL^*$. Then $w = xy$ for some $x \in L$ and $y \in L^*$. Hence $y \in L^n$ for some $n \in \mathbb{N}$. Thus $w = xy \in LL^n = L^{n+1} \subseteq L^*$. \square

Proposition 3.2.9

- (1) $\emptyset^* = \{\epsilon\}$.
- (2) $\{\epsilon\}^* = \{\epsilon\}$.
- (3) For all $L \in \mathbf{Lan}$, $L^*L = LL^*$.
- (4) For all $L \in \mathbf{Lan}$, $L^*L^* = L^*$.
- (5) For all $L \in \mathbf{Lan}$, $(L^*)^* = L^*$.
- (6) For all $L_1L_2 \in \mathbf{Lan}$, $(L_1L_2)^*L_1 = L_1(L_2L_1)^*$.

Proof. The six parts can be proven in order using Proposition 3.2.1. All parts but (2), (5) and (6) can be proved without using induction.

As an example, we show the proof of (5). To show that $(L^*)^* = L^*$, it will suffice to show that $(L^*)^* \subseteq L^* \subseteq (L^*)^*$.

To see that $(L^*)^* \subseteq L^*$, we use mathematical induction to show that, for all $n \in \mathbb{N}$, $(L^*)^n \subseteq L^*$.

(Basis Step) We have that $(L^*)^0 = \{\epsilon\} = L^0 \subseteq L^*$.

(Inductive Step) Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $(L^*)^n \subseteq L^*$. We must show that $(L^*)^{n+1} \subseteq L^*$. By the inductive hypothesis, Proposition 3.2.1(4) and Part (4), we have that $(L^*)^{n+1} = L^*(L^*)^n \subseteq L^*L^* = L^*$.

Now, we use the result of the induction to prove that $(L^*)^* \subseteq L^*$. Suppose $w \in (L^*)^*$. We must show that $w \in L^*$. Since $w \in (L^*)^*$, we have that $w \in (L^*)^n$ for some $n \in \mathbb{N}$. Thus, by the result of the induction, $w \in (L^*)^n \subseteq L^*$.

Finally, for the other inclusion, we have that $L^* = (L^*)^1 \subseteq (L^*)^*$. \square

Exercise 3.2.10

Prove Proposition 3.2.9(6), i.e., for all $L_1L_2 \in \mathbf{Lan}$, $(L_1L_2)^*L_1 = L_1(L_2L_1)^*$.

Proposition 3.2.11

- (1) $\$^* \approx \%$.
- (2) $\%^* \approx \%$.
- (3) For all $\alpha \in \mathbf{Reg}$, $\alpha^*\alpha \approx \alpha\alpha^*$.
- (4) For all $\alpha \in \mathbf{Reg}$, $\alpha^*\alpha^* \approx \alpha^*$.
- (5) For all $\alpha \in \mathbf{Reg}$, $(\alpha^*)^* \approx \alpha^*$.

(6) For all $\alpha, \beta \in \mathbf{Reg}$, $(\alpha\beta)^*\alpha \approx \alpha(\beta\alpha)^*$.

Proof. Follows from Proposition 3.2.9. Consider, e.g., the proof of Part (5). By Proposition 3.2.9(5), we have that

$$L((\alpha^*)^*) = L(\alpha^*)^* = (L(\alpha)^*)^* = L(\alpha)^* = L(\alpha^*).$$

Thus $(\alpha^*)^* \approx \alpha^*$. \square

3.2.2 Proving the Correctness of Regular Expressions

In this subsection, we use two examples to show how regular expressions can be designed and proved correct.

For our first example, define a function **zeros** $\in \{0, 1\}^* \rightarrow \mathbb{N}$ by recursion:

$$\begin{aligned} \mathbf{zeros} \% &= 0, \\ \mathbf{zeros}(0w) &= \mathbf{zeros} w + 1, \text{ for all } w \in \{0, 1\}^*, \text{ and} \\ \mathbf{zeros}(1w) &= \mathbf{zeros} w, \text{ for all } w \in \{0, 1\}^*. \end{aligned}$$

Thus **zeros** w is the number of occurrences of 0 in w . It is easy to show that:

- **zeros** 0 = 1;
- **zeros** 1 = 0;
- for all $x, y \in \{0, 1\}^*$, **zeros**(xy) = **zeros** x + **zeros** y ;
- for all $n \in \mathbb{N}$, **zeros**(0^n) = n ; and
- for all $n \in \mathbb{N}$, **zeros**(1^n) = 0.

Let

$$X = \{w \in \{0, 1\}^* \mid \mathbf{zeros} w \text{ is even}\},$$

so that X is all strings of 0's and 1's with an even number of 0's. Clearly, $\% \in X$ and $\{1\}^* \subseteq X$.

Let's consider the problem of finding a regular expression that generates X . A string with this property would begin with some number of 1's (possibly none). After this, the string would have some number of parts (possibly none), each consisting of a 0, followed by some number of 1's, followed by a 0, followed by some number of 1's. The above considerations lead us to the regular expression

$$\alpha = 1^*(01^*01^*)^*.$$

To prove $L(\alpha) = X$, it's helpful to give names to the meanings of two parts of α . Let

$$Y = \{0\}\{1\}^*\{0\}\{1\}^* \quad \text{and} \quad Z = \{1\}^*Y^*,$$

so that $L(01^*01^*) = Y$ and $L(\alpha) = Z$. Thus it will suffice to prove that $Z = X$, and we do this by showing $Z \subseteq X \subseteq Z$. We begin by showing $Z \subseteq X$.

Lemma 3.2.12

- (1) $Y \subseteq X$.
- (2) $XX \subseteq X$.

Proof.

- (1) Suppose $w \in Y$, so that $w = 0x0y$ for some $x, y \in \{1\}^*$. Thus $\mathbf{zeros} w = \mathbf{zeros}(0x0y) = \mathbf{zeros} 0 + \mathbf{zeros} x + \mathbf{zeros} 0 + \mathbf{zeros} y = 1 + 0 + 1 + 0 = 2$ is even, so that $w \in X$.
- (2) Suppose $w \in XX$, so that $w = xy$ for some $x, y \in X$. Then $\mathbf{zeros} x$ and $\mathbf{zeros} y$ are even, so that $\mathbf{zeros} w = \mathbf{zeros} x + \mathbf{zeros} y$ is even. Hence $w \in X$.

□

Lemma 3.2.13 $Y^* \subseteq X$.**Proof.** It will suffice to show that, for all $n \in \mathbb{N}$, $Y^n \subseteq X$, and we show this using mathematical induction.**(Basis Step)** We have that $Y^0 = \{\epsilon\} \subseteq X$.**(Inductive Step)** Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $Y^n \subseteq X$. Then $Y^{n+1} = YY^n \subseteq XX \subseteq X$, by Lemma 3.2.12

□

Lemma 3.2.14 $Z \subseteq X$.**Proof.** By Lemmas 3.2.12 and 3.2.13, we have that $Z = \{1\}^*Y^* \subseteq XX \subseteq X$.

□

To prove $X \subseteq Z$, it's helpful to define another language:

$$U = \{w \in X \mid 1 \text{ is not a prefix of } w\}.$$

Lemma 3.2.15 $U \subseteq Y^*$.**Proof.** Because $U \subseteq \{0, 1\}^*$, it will suffice to show that, for all $w \in \{0, 1\}^*$,

$$\text{if } w \in U, \text{ then } w \in Y^*.$$

We proceed by strong string induction. Suppose $w \in \{0, 1\}^*$, and assume the inductive hypothesis: for all $x \in \{0, 1\}^*$, if x is a proper substring of w , then

$$\text{if } x \in U, \text{ then } x \in Y^*.$$

We must show that

$$\text{if } w \in U, \text{ then } w \in Y^*.$$

Suppose $w \in U$, so that $\mathbf{zeros} w$ is even and 1 is not a prefix of w . We must show that $w \in Y^*$. If $w = \epsilon$, then $w \in Y^*$. Otherwise, $w = 0x$ for some $x \in \{0, 1\}^*$. Since $1 + \mathbf{zeros} x = \mathbf{zeros} 0 + \mathbf{zeros} x = \mathbf{zeros}(0x) = \mathbf{zeros} w$ is even, we have that $\mathbf{zeros} x$ is odd. Let n be the largest element of \mathbb{N} such that 1^n is a prefix of x . (n is well-defined, since $1^0 = \epsilon$ is a prefix of x .) Thus $x = 1^n y$ for some $y \in \{0, 1\}^*$. Since $\mathbf{zeros} y = 0 + \mathbf{zeros} y = \mathbf{zeros} 1^n + \mathbf{zeros} y = \mathbf{zeros}(1^n y) = \mathbf{zeros} x$ is odd, we have that $y \neq \epsilon$. And, by the definition of n , 1 is not a prefix of y . Hence $y = 0z$ for some $z \in \{0, 1\}^*$. Since $1 + \mathbf{zeros} z = \mathbf{zeros} 0 + \mathbf{zeros} z = \mathbf{zeros}(0z) = \mathbf{zeros} y$ is odd, we have that $\mathbf{zeros} z$ is even. Let m be the largest element of \mathbb{N} such that 1^m is a prefix of z . Thus $z = 1^m u$ for some $u \in \{0, 1\}^*$, and 1 is not a prefix of u . Since $\mathbf{zeros} u = 0 + \mathbf{zeros} u = \mathbf{zeros} 1^m + \mathbf{zeros} u = \mathbf{zeros}(1^m u) = \mathbf{zeros} z$ is even, it follows that $u \in U$. Summarizing, we have that $w = 0x = 01^n y = 01^n 0z = 01^n 01^m u$ and $u \in U$. Since u is a proper substring of w , the inductive hypothesis tells us that $u \in Y^*$. Hence $w = 01^n 01^m u = (01^n 01^m)u \in YY^* \subseteq Y^*$. \square

Lemma 3.2.16

$X \subseteq Z$.

Proof. Suppose $w \in X$. Let n be the largest element of \mathbb{N} such that 1^n is a prefix of w . Thus $w = 1^n x$ for some $x \in \{0, 1\}^*$. Since $w \in X$, we have that $\mathbf{zeros} x = 0 + \mathbf{zeros} x = \mathbf{zeros} 1^n + \mathbf{zeros} x = \mathbf{zeros} w$ is even, so that $x \in X$. By the definition of n , we have that 1 is not a prefix of x , and thus $x \in U$. Hence $w = 1^n x \in \{1\}^* U \subseteq \{1\}^* Y^* = Z$, by Lemma 3.2.15. \square

By Lemmas 3.2.14 and 3.2.16, we have that $Z \subseteq X \subseteq Z$, completing our proof that α is correct.

Our second example of regular expression design and proof of correction involves the languages

$$A = \{001, 011, 101, 111\}, \text{ and}$$

$$B = \{w \in \{0, 1\}^* \mid \text{for all } x, y \in \{0, 1\}^*, \text{ if } w = x0y, \text{ then there is a } z \in A \text{ such that } z \text{ is a prefix of } y\}.$$

The elements of A can be thought of as the odd numbers between 1 and 7, expressed in binary, and B consists of those strings of 0's and 1's in which every occurrence of 0 is immediately followed by an element of A .

E.g., $\epsilon \in B$, since the empty string has no occurrences of 0, and 00111 is in B , since its first 0 is followed by 011 and its second 0 is followed by 111. But 0000111 is not in B , since its first 0 is followed by 000, which is not in A . And 011 is not in B , since $|11| < 3$.

Note that, for all $x, y \in B$, $xy \in B$, i.e., $BB \subseteq B$. This holds, since: each occurrence of 0 in x is followed by an element of A in x , and is thus followed by the same element of A in xy ; and each occurrence of 0 in y is followed by an element of A in y , and is thus followed by the same element of A in xy .

Furthermore, for all strings x, y , if $xy \in B$, then y is in B , i.e., every suffix of an element of B is also in B . This holds since if there was an occurrence of 0 in y that wasn't followed by an element of A , then this same occurrence of 0 in the suffix y of xy would also not be followed by an element of A , contradicting $xy \in B$.

How should we go about finding a regular expression α such that $L(\alpha) = B$? Because $\epsilon \in B$, for all $x, y \in B$, $xy \in B$, and for all strings x, y , if $xy \in B$ then $y \in B$, our regular expression can have the form β^* , where β generates all the strings that are *basic* in the sense that they are nonempty elements of B with no non-empty proper prefixes that are in B .

Let's try to understand what the basic strings look like. Clearly, 1 is basic, so there will be no more basic strings that begin with 1. But what about the basic strings beginning with 0? No sequence of 0's is basic, and any string that begins with four or more 0's will not be basic. It is easy to see that 000111 is basic. In fact, it is the only basic string of the form 000u. (The first 0 forces u to begin with 1, the second 0 forces u to continue with 1, and the third forces u to continue with 1. And, if $|u| > 3$, then the overall string would have a nonempty, proper prefix in B , and so wouldn't be basic.) Similarly, 00111 is the only basic string beginning with 001. But what about the basic strings beginning with 01? It's not hard to see that there are infinitely many such strings: 0111, 010111, 01010111, 0101010111, etc. Fortunately, there is a simple pattern here: we have all strings of the form $0(10)^n111$ for $n \in \mathbb{N}$.

By the above considerations, it seems that we can let our regular expression be

$$(1 + 0(10)^*111 + 00111 + 000111)^*.$$

But, using some of the equivalences we learned about above, we can turn this regular expression into

$$(1 + 0(0 + 00 + (10)^*)111)^*,$$

which we take as our α . Now, we prove that $L(\alpha) = B$.

Let

$$X = \{0\} \cup \{00\} \cup \{10\}^* \quad \text{and} \quad Y = \{1\} \cup \{0\}X\{111\}.$$

Then, we have that

$$\begin{aligned} X &= L(0 + 00 + (10)^*), \\ Y &= L(1 + 0(0 + 00 + (10)^*)111), \text{ and} \\ Y^* &= L((1 + 0(0 + 00 + (10)^*)111)^*) = L(\alpha). \end{aligned}$$

Thus, it will suffice to show that $Y^* = B$. We will show that $Y^* \subseteq B \subseteq Y^*$.

To begin with, we would like to use mathematical induction to prove that, for all $n \in \mathbb{N}$, $\{0\}\{10\}^n\{111\} \subseteq B$. But in order for the inductive step to succeed, we must prove something stronger. Let

$$C = \{w \in B \mid 01 \text{ is a prefix of } w\}.$$

Lemma 3.2.17

For all $n \in \mathbb{N}$, $\{0\}\{10\}^n\{111\} \subseteq C$.

Proof. We proceed by mathematical induction.

(Basis Step) Since 01 is a prefix of 0111, and $0111 \in B$, we have that $0111 \in C$. Hence $\{0\}\{10\}^0\{111\} = \{0\}\{1\}\{111\} = \{0\}\{111\} = \{0111\} \subseteq C$.

(Inductive Step) Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $\{0\}\{10\}^n\{111\} \subseteq C$. We must show that $\{0\}\{10\}^{n+1}\{111\} \subseteq C$. Since

$$\begin{aligned} \{0\}\{10\}^{n+1}\{111\} &= \{0\}\{10\}\{10\}^n\{111\} \\ &= \{01\}\{0\}\{10\}^n\{111\} \\ &\subseteq \{01\}C \quad (\text{inductive hypothesis}), \end{aligned}$$

it will suffice to show that $\{01\}C \subseteq C$. Suppose $w \in \{01\}C$. We must show that $w \in C$. We have that $w = 01x$ for some $x \in C$. Thus w begins with 01. It remains to show that $w \in B$. Since $x \in C$, we have that x begins with 01. Thus the first occurrence of 0 in $w = 01x$ is followed by 101 $\in A$. Furthermore, any other occurrence of 0 in $w = 01x$ is within x , and so is followed by an element of A because $x \in C \subseteq B$. Thus $w \in B$.

□

Lemma 3.2.18

$Y \subseteq B$.

Proof. Suppose $w \in Y$. We must show that $w \in B$. If $w = 1$, then $w \in B$. Otherwise, we have that $w = 0x111$ for some $x \in X$. There are three cases to consider.

- Suppose $x = 0$. Then $w = 00111$ is in B .
- Suppose $x = 00$. Then $w = 000111$ is in B .
- Suppose $x \in \{10\}^*$. Then $x \in \{10\}^n$ for some $n \in \mathbb{N}$. By Lemma 3.2.17, we have that $w = 0x111 \in \{0\}\{10\}^n\{111\} \subseteq C \subseteq B$.

□

Lemma 3.2.19

$Y^* \subseteq B$.

Proof. It will suffice to show that, for all $n \in \mathbb{N}$, $Y^n \subseteq B$, and we proceed by mathematical induction.

(Basis Step) Since $\% \in B$, we have that $Y^0 = \{\%\} \subseteq B$.

(Inductive Step) Suppose $n \in \mathbb{N}$, and assume the inductive hypothesis: $Y^n \subseteq B$. Then $Y^{n+1} = YY^n \subseteq BB \subseteq B$, by Lemma 3.2.18 and the inductive hypothesis.

□

Lemma 3.2.20

$B \subseteq Y^*$.

Proof. Since $B \subseteq \{0,1\}^*$, it will suffice to show that, for all $w \in \{0,1\}^*$,

if $w \in B$, then $w \in Y^*$.

We proceed by strong string induction. Suppose $w \in \{0,1\}^*$, and assume the inductive hypothesis: for all $x \in \{0,1\}^*$, if x is a proper substring of w , then

if $x \in B$, then $x \in Y^*$.

We must show that

if $w \in B$, then $w \in Y^*$.

Suppose $w \in B$. We must show that $w \in Y^*$. There are three main cases to consider.

- Suppose $w = \%$. Then $w \in \{\%\} = Y^0 \subseteq Y^*$.
- Suppose $w = 0x$ for some $x \in \{0,1\}^*$. Since $w \in B$, the first 0 of w must be followed by an element of A . Hence $x \neq \%$, so that there are two cases to consider.
 - Suppose $x = 0y$ for some $y \in \{0,1\}^*$. Thus $w = 0x = 00y$. Since $00y = w \in B$, we have that $y \neq \%$. Thus, there are two cases to consider.
 - * Suppose $y = 0z$ for some $z \in \{0,1\}^*$. Thus $w = 00y = 000z$. Since the first 0 in $000z = w$ is followed by an element of A , and the only element of A that begins with 00 is 001, we have that $z = 1u$ for some $u \in \{0,1\}^*$. Thus $w = 0001u$. Since the second 0 in $0001u = w$ is followed by an element of A , and 011 is the only element of A that begins with 01, we have that $u = 1v$ for some $v \in \{0,1\}^*$. Thus $w = 00011v$. Since the third 0 in

$00011v = w$ is followed by an element of A , and 111 is the only element of A that begins with 11 , we have that $v = 1t$ for some $t \in \{0,1\}^*$. Thus $w = 000111t$. Since $00 \in X$, we have that $000111 = (0)(00)(111) \in \{0\}X\{111\} \subseteq Y$. Because t is a suffix of w , it follows that $t \in B$. Thus the inductive hypothesis tells us that $t \in Y^*$. Hence $w = (000111)t \in YY^* \subseteq Y^*$.

- * Suppose $y = 1z$ for some $z \in \{0,1\}^*$. Thus $w = 00y = 001z$. Since the first 0 in $001z = w$ is followed by an element of A , and the only element of A that begins with 01 is 011 , we have that $z = 1u$ for some $u \in \{0,1\}^*$. Thus $w = 0011u$. Since the second 0 in $0011u = w$ is followed by an element of A , and 111 is the only element of A that begins with 11 , we have that $u = 1v$ for some $v \in \{0,1\}^*$. Thus $w = 00111v$. Since $0 \in X$, we have that $00111 = (0)(0)(111) \in \{0\}X\{111\} \subseteq Y$. Because v is a suffix of w , it follows that $v \in B$. Thus the inductive hypothesis tells us that $v \in Y^*$. Hence $w = (00111)v \in YY^* \subseteq Y^*$.
- Suppose $x = 1y$ for some $y \in \{0,1\}^*$. Thus $w = 0x = 01y$. Since $w \in B$, we have that $y \neq \%$. Thus, there are two cases to consider.

- * Suppose $y = 0z$ for some $z \in \{0,1\}^*$. Thus $w = 010z$. Let u be the longest prefix of z that is in $\{10\}^*$. (Since $\%$ is a prefix of z and is in $\{10\}^*$, it follows that u is well-defined.) Let $v \in \{0,1\}^*$ be such that $z = uv$. Thus $w = 010z = 010uv$.

Suppose, toward a contradiction, that v begins with 10 . Then $u10$ is a prefix of $z = uv$ that is longer than u . Furthermore $u10 \in \{10\}^*\{10\} \subseteq \{10\}^*$, contradicting the definition of u . Thus we have that v does not begin with 10 .

Next, we show that $010u$ ends with 010 . Since $u \in \{10\}^*$, we have that $u \in \{10\}^n$ for some $n \in \mathbb{N}$. There are three cases to consider.

- Suppose $n = 0$. Since $u \in \{10\}^0 = \{\%\}$, we have that $u = \%$. Thus $010u = 010$ ends with 010 .
- Suppose $n = 1$. Since $u \in \{10\}^1 = \{10\}$, we have that $u = 10$. Hence $010u = 01010$ ends with 010 .
- Suppose $n \geq 2$. Then $n - 2 \geq 0$, so that $u \in \{10\}^{(n-2)+2} = \{10\}^{n-2}\{10\}^2$. Hence u ends with 1010 , showing that $010u$ ends with 010 .

Summarizing, we have that $w = 010uv$, $u \in \{10\}^*$, $010u$ ends with 010 , and v does not begin with 10 . Since the second-to-last 0 in $010u$ is followed in w by an element of A , and 101 is the only element of A that begins with 10 , we have that $v = 1s$ for some $s \in \{0,1\}^*$. Thus $w = 010u1s$, and $010u1$ ends with 0101 . Since the second-to-last symbol of $010u1$ is a 0 , we have that

$s \neq \%$. Furthermore, s does not begin with 0, since, if it did, then $v = 1s$ would begin with 10. Thus we have that $s = 1t$ for some $t \in \{0,1\}^*$. Hence $w = 010u11t$. Since $010u11$ ends with 011, it follows that the last 0 in $010u11$ must be followed in w by an element of A . Because 111 is the only element of A that begins with 11, we have that $t = 1r$ for some $r \in \{0,1\}^*$. Thus $w = 010u111r$. Since $(10)u \in \{10\}\{10\}^* \subseteq \{10\}^* \subseteq X$, we have that $010u111 = (0)((10)u)111 \in \{0\}X\{111\} \subseteq Y$. Since r is a suffix of w , it follows that $r \in B$. Thus, the inductive hypothesis tells us that $r \in Y^*$. Hence $w = (010u111)r \in YY^* \subseteq Y^*$.

* Suppose $y = 1z$ for some $z \in \{0,1\}^*$. Thus $w = 011z$. Since the first 0 of w is followed by an element of A , and 111 is the only element of A that begins with 11, we have that $z = 1u$ for some $u \in \{0,1\}^*$. Thus $w = 0111u$. Since $\% \in \{10\}^* \subseteq X$, we have that $0111 = (0)(\%)(111) \in \{0\}X\{111\} \subseteq Y$. Because u is a suffix of w , it follows that $u \in B$. Thus, since u is a proper substring of w , the inductive hypothesis tells us that $u \in Y^*$. Hence $w = (0111)u \in YY^* \subseteq Y^*$.

- Suppose $w = 1x$ for some $x \in \{0,1\}^*$. Since x is a suffix of w , we have that $x \in B$. Because x is a proper substring of w , the inductive hypothesis tells us that $x \in Y^*$. Thus $w = 1x \in YY^* \subseteq Y^*$.

□

By Lemmas 3.2.19 and 3.2.20, we have that $Y^* \subseteq B \subseteq Y^*$, so that $Y^* = B$. This completes our regular expression design and proof of correctness example.

Exercise 3.2.21

Let $X = \{w \in \{0,1\}^* \mid 010 \text{ is not a substring of } w\}$. Find a regular expression α such that $L(\alpha) = X$, and prove that your answer is correct.

Exercise 3.2.22

Define $\mathbf{diff} \in \{0,1\}^* \rightarrow \mathbb{Z}$ as in Section 2.2, so that, for all $w \in \{0,1\}^*$,

$\mathbf{diff} w = \text{the number of 1's in } w - \text{the number of 0's in } w$.

Thus $\mathbf{diff} \% = 0$, $\mathbf{diff} 0 = -1$, $\mathbf{diff} 1 = 1$, and for all $x, y \in \{0,1\}^*$, $\mathbf{diff}(xy) = \mathbf{diff} x + \mathbf{diff} y$. Let $X = \{w \in \{0,1\}^* \mid \mathbf{diff} w = 3m, \text{ for some } m \in \mathbb{Z}\}$. Find a regular expression α such that $L(\alpha) = X$, and prove that your answer is correct.

Exercise 3.2.23

Define a function $\mathbf{diff} \in \{0,1\}^* \rightarrow \mathbb{Z}$ by: for all $w \in \{0,1\}^*$,

$\mathbf{diff} w = \text{the number of 0's in } w - 2(\text{the number of 1's in } w)$.

Thus $\mathbf{diff} w = 0$ iff w has twice as many 1's as 0's. Furthermore $\mathbf{diff} \% = 0$, $\mathbf{diff} 0 = 1$, $\mathbf{diff} 1 = -2$, and, for all $x, y \in \{0, 1\}^*$, $\mathbf{diff}(xy) = \mathbf{diff} x + \mathbf{diff} y$. Let $X = \{w \in \{0, 1\}^* \mid \mathbf{diff} w = 0 \text{ and, for all prefixes } v \text{ of } w, 0 \leq \mathbf{diff} v \leq 3\}$. Find a regular expression α such that $L(\alpha) = X$, and prove that your answer is correct.

3.2.3 Notes

Our approach in this section is somewhat more formal than is common, but is otherwise standard.

3.3 Simplification of Regular Expressions

In this section, we give three algorithms—of increasing power, but decreasing efficiency—for regular expression simplification. The first algorithm—weak simplification—is defined via a straightforward structural recursion, and is sufficient for many purposes. The remaining two algorithms—local simplification and global simplification—are based on a set of simplification rules that is still incomplete and evolving.

3.3.1 Regular Expression Complexity

To begin with, let's consider how we might measure the complexity/simplicity of regular expressions. The most obvious criterion is size (remember that regular expressions are trees). But consider this pair of equivalent regular expressions:

$$\begin{aligned}\alpha &= (00^*11^*)^*, \text{ and} \\ \beta &= \% + 0(0 + 11^*0)^*11^*.\end{aligned}$$

Although the size of β (18) is strictly greater than the size of α (10), β has only one closure inside another closure, whereas α has two closures inside its outer closure, and thus there is a sense in which β is easier to understand than α .

The standard measure of the closure-related complexity of a regular expression is its *star-height*: the maximum number $n \in \mathbb{N}$ such that there is a path from the root of the regular expression to one of its leaves that passes through n closures. But α and β both have star-heights of 2. Furthermore, star-height isn't respected by the ways of forming regular expressions. E.g., if γ_1 has strictly smaller star-height than γ_2 , we can't conclude that $\gamma_1\gamma'$ has strictly smaller star-height than $\gamma_2\gamma'$, as the star height of γ' may be greater than the star-height of γ_2 .

So, we need a better measure of the closure-related complexity of regular expressions than star-height. Toward that end, let's define a *closure complexity* to be a nonempty list ns of natural numbers that is (not-necessarily strictly)

descending: for all $i \in [1 : |ns| - 1]$, $ns\ i \geq ns(i+1)$. This is a way of representing nonempty multisets of natural numbers that makes it easy to define the usual ordering on multisets. We write \mathbf{CC} for the set of all closure complexities. E.g., $[3, 2, 2, 1]$ is a closure complexity, but $[3, 2, 3]$ and $[]$ are not. For all $n \in \mathbb{N}$, $[n]$ is a *singleton* closure complexity. The *union* of closure complexities ns and ms ($ns \cup ms$) is the closure complexity that results from putting ns @ ms in descending order, keeping any duplicate elements. (Here we are overloading the term union and the operation \cup , but the set-theoretic union isn't an operation on closure complexities, and so no confusion should result.) E.g., $[3, 2, 2, 1] \cup [4, 2, 1, 0] = [4, 3, 2, 2, 2, 1, 1, 0]$. The *successor* \overline{ns} of a closure complexity ns is the closure complexity formed by adding one to each element of ns , maintaining the order of the elements. E.g., $\overline{[3, 2, 2, 1]} = [4, 3, 3, 2]$.

It is easy to see that \cup is commutative and associative on \mathbf{CC} , and that the successor operation on \mathbf{CC} preserves union:

Proposition 3.3.1

- (1) For all $ns, ms \in \mathbf{CC}$, $ns \cup ms = ms \cup ns$.
- (2) For all $ns, ms, ls \in \mathbf{CC}$, $(ns \cup ms) \cup ls = ns \cup (ms \cup ls)$.
- (3) For all $ns, ms \in \mathbf{CC}$, $\overline{ns \cup ms} = \overline{ns} \cup \overline{ms}$.

Proposition 3.3.2

- (1) For all $ns, ms \in \mathbf{CC}$, $\overline{ns} = \overline{ms}$ iff $ns = ms$.
- (2) For all $ns, ms, ls \in \mathbf{CC}$, $ns \cup ls = ms \cup ls$ iff $ns = ms$.

We define a relation $<_{cc}$ on \mathbf{CC} by: for all $ns, ms \in \mathbf{CC}$, $ns <_{cc} ms$ iff either:

- $ms = ns @ ls$ for some $ls \in \mathbf{CC}$; or
- there is an $i \in \mathbb{N} - \{0\}$ such that
 - $i \leq |ns|$ and $i \leq |ms|$,
 - for all $j \in [1 : i - 1]$, $ns\ j = ms\ j$, and
 - $ns\ i < ms\ i$.

In other words, $ns <_{cc} ms$ iff either ms consists of the result of appending a nonempty list at the end of ns , or ns and ms agree up to some point, at which ns 's value is strictly smaller than ms 's value. E.g., $[2, 2] \leq_{cc} [2, 2, 1]$ and $[2, 1, 1, 0, 0] <_{cc} [2, 2, 1]$.

Proposition 3.3.3

- (1) For all $ns, ms \in \mathbf{CC}$, $\overline{ns} <_{cc} \overline{ms}$ iff $ns <_{cc} ms$.
- (2) For all $ns, ms, ls \in \mathbf{CC}$, $ns \cup ls <_{cc} ms \cup ls$ iff $ns <_{cc} ms$.

(3) For all $ns, ms \in \mathbf{CC}$, $ns <_{cc} ns \cup ms$.

Proposition 3.3.4

$<_{cc}$ is a strict total ordering on \mathbf{CC} .

Proposition 3.3.5

$<_{cc}$ is a well-founded relation on \mathbf{CC} .

Now we can define the closure complexity of a regular expression. Define the function $\mathbf{cc} \in \mathbf{Reg} \rightarrow \mathbf{CC}$ by structural recursion:

$$\begin{aligned} \mathbf{cc} \% &= [0]; \\ \mathbf{cc} \$ &= [0]; \\ \mathbf{cc} a &= [0], \text{ for all } a \in \mathbf{Sym}; \\ \mathbf{cc}(*(\alpha)) &= \overline{\mathbf{cc} \alpha}, \text{ for all } \alpha \in \mathbf{Reg}; \\ \mathbf{cc}(@(\alpha, \beta)) &= \mathbf{cc} \alpha \cup \mathbf{cc} \beta, \text{ for all } \alpha, \beta \in \mathbf{Reg}; \text{ and} \\ \mathbf{cc}+(\alpha, \beta) &= \mathbf{cc} \alpha \cup \mathbf{cc} \beta, \text{ for all } \alpha, \beta \in \mathbf{Reg}. \end{aligned}$$

We say that $\mathbf{cc} \alpha$ is the closure complexity of α . E.g.,

$$\begin{aligned} \mathbf{cc}((12^*)^*) &= \overline{\mathbf{cc}(12^*)} = \overline{\mathbf{cc} 1 \cup \mathbf{cc}(2^*)} = \overline{[0] \cup \overline{\mathbf{cc} 2}} \\ &= \overline{[0] \cup \overline{[0]}} = \overline{[0] \cup [1]} = \overline{[1, 0]} = [2, 1]. \end{aligned}$$

In other words, the $\mathbf{cc} \alpha$ can be computed by first collecting together all the paths through α that terminate in leafs, then counting the numbers of closures visited when following each of these paths, and finally putting those sums in descending order.

Returning to our initial examples, we have that $\mathbf{cc}((00^*11^*)^*) = [2, 2, 1, 1]$ and $\mathbf{cc}(\% + 0(0 + 11^*0)^*11^*) = [2, 1, 1, 1, 1, 0, 0, 0]$. Since $[2, 1, 1, 1, 1, 0, 0, 0] <_{cc} [2, 2, 1, 1]$, the closure complexity of $\% + 0(0 + 11^*0)^*11^*$ is strictly smaller than the closure complexity of $(00^*11^*)^*$.

Proposition 3.3.6

For all $\alpha \in \mathbf{Reg}$, $|\mathbf{cc} \alpha| = \mathbf{numLeaves} \alpha$.

Proof. An easy induction on regular expressions. \square

Exercise 3.3.7

Find regular expressions α and β such that $\mathbf{cc} \alpha = \mathbf{cc} \beta$ but $\mathbf{size} \alpha \neq \mathbf{size} \beta$.

In contrast to star-height, closure complexity is compatible with the ways of forming regular expressions. In fact, we can prove even stronger results.

Proposition 3.3.8

(1) For all $\alpha \in \mathbf{Reg}$, $\mathbf{cc} \alpha = \mathbf{cc} \beta$ iff $\mathbf{cc}(\alpha^*) = \mathbf{cc}(\beta^*)$.

- (2) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha = \mathbf{cc} \beta$ iff $\mathbf{cc}(\alpha\gamma) = \mathbf{cc}(\beta\gamma)$.
- (3) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha = \mathbf{cc} \beta$ iff $\mathbf{cc}(\gamma\alpha) = \mathbf{cc}(\gamma\beta)$.
- (4) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha = \mathbf{cc} \beta$ iff $\mathbf{cc}(\alpha + \gamma) = \mathbf{cc}(\beta + \gamma)$.
- (5) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha = \mathbf{cc} \beta$ iff $\mathbf{cc}(\gamma + \alpha) = \mathbf{cc}(\gamma + \beta)$.

Proof. Follows by Proposition 3.3.2. \square

The following proposition says that if we replace a subtree of a regular expression by a regular expression with the same closure complexity, then the closure complexity of the resulting, whole regular expression will be unchanged.

Proposition 3.3.9

Suppose $\alpha, \beta, \beta' \in \mathbf{Reg}$, $\mathbf{cc} \beta = \mathbf{cc} \beta'$, $pat \in \mathbf{Path}$ is valid for α , and β is the subtree of α at position pat . Let α' be the result of replacing the subtree at position pat in α by β' . Then $\mathbf{cc} \alpha = \mathbf{cc} \alpha'$.

Proof. By induction on α using Proposition 3.3.8. \square

Proposition 3.3.10

- (1) For all $\alpha \in \mathbf{Reg}$, $\mathbf{cc} \alpha <_{cc} \mathbf{cc} \beta$ iff $\mathbf{cc}(\alpha^*) <_{cc} \mathbf{cc}(\beta^*)$.
- (2) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha <_{cc} \mathbf{cc} \beta$ iff $\mathbf{cc}(\alpha\gamma) <_{cc} \mathbf{cc}(\beta\gamma)$.
- (3) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha <_{cc} \mathbf{cc} \beta$ iff $\mathbf{cc}(\gamma\alpha) <_{cc} \mathbf{cc}(\gamma\beta)$.
- (4) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha <_{cc} \mathbf{cc} \beta$ iff $\mathbf{cc}(\alpha + \gamma) <_{cc} \mathbf{cc}(\beta + \gamma)$.
- (5) For all $\alpha, \beta, \gamma \in \mathbf{Reg}$, $\mathbf{cc} \alpha <_{cc} \mathbf{cc} \beta$ iff $\mathbf{cc}(\gamma + \alpha) <_{cc} \mathbf{cc}(\gamma + \beta)$.

Proof. Follows by Proposition 3.3.3. \square

The following proposition says that if we replace a subtree of a regular expression by a regular expression with strictly smaller closure complexity, that the resulting, whole regular expression will have strictly smaller closure complexity than the original regular expression.

Proposition 3.3.11

Suppose $\alpha, \beta, \beta' \in \mathbf{Reg}$, $\mathbf{cc} \beta' <_{cc} \mathbf{cc} \beta$, $pat \in \mathbf{Path}$ is valid for α , and β is the subtree of α at position pat . Let α' be the result of replacing the subtree at position pat in α by β' . Then $\mathbf{cc} \alpha' <_{cc} \mathbf{cc} \alpha$.

Proof. By induction on α , using Proposition 3.3.10. \square

When judging the relative complexity of regular expressions α and β , we will first look at how their closure complexities are related. And, when their closure complexities are equal, we will look at how their sizes are related. To finish explaining how we will judge the relative complexity of regular expressions, we need three definitions.

The function

$$\mathbf{numConcats} \in \mathbf{Reg} \rightarrow \mathbb{N}$$

is defined by recursion:

$$\begin{aligned} \mathbf{numConcats} \% &= 0; \\ \mathbf{numConcats} \$ &= 0; \\ \mathbf{numConcats} a &= 0, \text{ for all } a \in \mathbf{Sym}; \\ \mathbf{numConcats}(\alpha^*) &= \mathbf{numConcats} \alpha, \text{ for all } \alpha \in \mathbf{Reg}; \\ \mathbf{numConcats}(\alpha\beta) &= 1 + \mathbf{numConcats} \alpha + \mathbf{numConcats} \beta; \text{ and} \\ \mathbf{numConcats}(\alpha + \beta) &= \mathbf{numConcats} \alpha + \mathbf{numConcats} \beta. \end{aligned}$$

Thus $\mathbf{numConcats} \alpha$ is the number of concatenations in α , i.e., the number of subtrees of α that are concatenations, where a given concatenation may occur (and will be counted) multiple times. E.g., $\mathbf{numConcats}(((01)^*(01))^*) = 3$. The function

$$\mathbf{numSyms} \in \mathbf{Reg} \rightarrow \mathbb{N}$$

is defined by structural recursion:

$$\begin{aligned} \mathbf{numSyms} \% &= 0; \\ \mathbf{numSyms} \$ &= 0; \\ \mathbf{numSyms} a &= 1, \text{ for all } a \in \mathbf{Sym}; \\ \mathbf{numSyms}(\alpha^*) &= \mathbf{numSyms} \alpha, \text{ for all } \alpha \in \mathbf{Reg}; \\ \mathbf{numSyms}(\alpha\beta) &= \mathbf{numSyms} \alpha + \mathbf{numSyms} \beta; \text{ and} \\ \mathbf{numSyms}(\alpha + \beta) &= \mathbf{numSyms} \alpha + \mathbf{numSyms} \beta. \end{aligned}$$

Thus $\mathbf{numSyms} \alpha$ is the number of occurrences of symbols in α , where a given symbol may occur (and will be counted) more than once. E.g., $\mathbf{numSyms}((0^*1) + 0) = 3$.

Finally, we say that a regular expression α is *standardized* iff none of α 's subtrees have any of the following forms:

- $(\beta_1 + \beta_2) + \beta_3$ (we can avoid needing parentheses, and make a regular expression easier to understand/process from left-to-right, by grouping unions to the right);

- $\beta_1 + \beta_2$, where $\beta_1 > \beta_2$, or $\beta_1 + (\beta_2 + \beta_3)$, where $\beta_1 > \beta_2$ (it's pleasing if the regular expressions appear in order (recall that unions are greater than all other kinds of regular expressions));
- $(\beta_1\beta_2)\beta_3$ (we can avoid needing parentheses, and make a regular expression easier to understand/process from left-to-right, by grouping concatenations to the right); and
- $\beta^*\beta$, $\beta^*(\beta\gamma)$, $(\beta_1\beta_2)^*\beta_1$ or $(\beta_1\beta_2)^*\beta_1\gamma$ (moving closures to the right makes a regular expression easier to understand/process from left-to-right).

Thus every subtree of a standardized regular expression will be standardized.

Returning to our assessment of regular expression complexity, suppose that α and β are regular expressions generating $\%$. Then $(\alpha\beta)^*$ and $(\alpha + \beta)^*$ are equivalent, but we will prefer the latter over the former, because unions are generally more amenable to understanding and processing than concatenations. Consequently, when two regular expressions have the same closure complexity and size, we will judge their relative complexity according to their numbers of concatenations.

Next, consider the regular expressions $0 + 01$ and $0(\% + 1)$. These regular expressions have the same closure complexity $[0, 0, 0]$, size (5) and number of concatenations (1). We would like to consider the latter to be simpler than the former, since in general we would like to prefer $\alpha(\% + \beta)$ over $\alpha + \alpha\beta$. And we can base this preference on the fact that the number of symbols of $0(\% + 1)$ (2) is one less than the number of symbols of $0 + 01$. When regular expressions have the same closure complexity, size and number of concatenations, the one with fewer symbols is likely to be easier to understand and process. Thus, when regular expressions have identical closure complexity, size and number of concatenations, we will use their relative numbers of symbols to judge their relative complexity.

Finally, when regular expressions have the same closure complexity, size, number of concatenations, and number of symbols, we will judge their relative complexity according to whether they are standardized, thinking that a standardized regular expression is simpler than one that is not standardized.

We define a relation $<_{\text{simp}}$ on **Reg** by, for all $\alpha, \beta \in \mathbf{Reg}$, $\alpha <_{\text{simp}} \beta$ iff:

- $\mathbf{cc} \alpha <_{\text{cc}} \mathbf{cc} \beta$; or
- $\mathbf{cc} \alpha = \mathbf{cc} \beta$ but $\mathbf{size} \alpha < \mathbf{size} \beta$; or
- $\mathbf{cc} \alpha = \mathbf{cc} \beta$ and $\mathbf{size} \alpha = \mathbf{size} \beta$, but $\mathbf{numConcat} \alpha < \mathbf{numConcat} \beta$;
or
- $\mathbf{cc} \alpha = \mathbf{cc} \beta$, $\mathbf{size} \alpha = \mathbf{size} \beta$ and $\mathbf{numConcat} \alpha = \mathbf{numConcat} \beta$, but $\mathbf{numSyms} \alpha < \mathbf{numSyms} \beta$; or

- $\mathbf{cc} \alpha = \mathbf{cc} \beta$, $\mathbf{size} \alpha = \mathbf{size} \beta$, $\mathbf{numConcat} \alpha = \mathbf{numConcat} \beta$ and $\mathbf{numSyms} \alpha = \mathbf{numSyms} \beta$, but α is standardized and β is not standardized.

We read $\alpha <_{\mathbf{simp}} \beta$ as α is *simpler* (less *complex*) than β . We define a relation $\equiv_{\mathbf{simp}}$ on **Reg** by, for all $\alpha, \beta \in \mathbf{Reg}$, $\alpha \equiv_{\mathbf{simp}} \beta$ iff α and β have the same closure complexity, size, numbers of concatenations, numbers of symbols, and status of being (or not being) standardized. We read $\alpha \equiv_{\mathbf{simp}} \beta$ as α and β have the *same complexity*. Finally, we define a relation $\leq_{\mathbf{simp}}$ on **Reg** by, for all $\alpha, \beta \in \mathbf{Reg}$, $\alpha \leq_{\mathbf{simp}} \beta$ iff $\alpha <_{\mathbf{simp}} \beta$ or $\alpha \equiv_{\mathbf{simp}} \beta$. We read $\alpha \leq_{\mathbf{simp}} \beta$ as α is at least as simpler as (no more complex) than β .

For example, the following regular expressions are equivalent and have the same complexity:

$$1(01 + 10) + (\% + 01)1 \quad \text{and} \quad 011 + 1(\% + 01 + 10).$$

Proposition 3.3.12

- (1) $<_{\mathbf{simp}}$ is transitive.
- (2) $\equiv_{\mathbf{simp}}$ is reflexive on **Reg**, transitive and symmetric.
- (3) For all $\alpha, \beta \in \mathbf{Reg}$, exactly one of the following holds: $\alpha <_{\mathbf{simp}} \beta$, $\beta <_{\mathbf{simp}} \alpha$ or $\alpha \equiv_{\mathbf{simp}} \beta$.
- (4) $\leq_{\mathbf{simp}}$ is transitive, and, for all $\alpha, \beta \in \mathbf{Reg}$, $\alpha \equiv_{\mathbf{simp}} \beta$ iff $\alpha \leq \beta$ and $\beta \leq \alpha$.

The Forlan module **Reg** defines the abstract type **cc** of closure complexities, along with these functions:

```
val ccToList   : cc -> int list
val singCC     : int -> cc
val unionCC    : cc * cc -> cc
val succCC     : cc -> cc
val cc         : reg -> cc
val compareCC  : cc * cc -> order
```

The function **ccToList** is the identity function on closure complexities: all that changes is the type. **singCC** n returns the singleton closure complexity $[n]$, if n is nonnegative; otherwise it issues an error message. The functions **unionCC** and **succCC** implement the union and successor operations on closure complexities. The function **cc** corresponds to **cc**, and **compareCC** implements $<_{cc}$.

Here are some examples of how these functions can be used:

```
- val ns = Reg.succCC(Reg.unionCC(Reg.singCC 1, Reg.singCC 1));
val ns = - : Reg.cc
- Reg.ccToList ns;
val it = [2,2] : int list
```

```

- val ms = Reg.unionCC(ns, Reg.succCC ns);
val ms = - : Reg.cc
- Reg.ccToList ms;
val it = [3,3,2,2] : int list
- Reg.ccToList(Reg.cc(Reg.fromString "(00*11*)*"));
val it = [2,2,1,1] : int list
- Reg.ccToList(Reg.cc(Reg.fromString "% + 0(0 + 11*0)*11*"));
val it = [2,1,1,1,1,0,0,0] : int list
- Reg.compareCC
= (Reg.cc(Reg.fromString "(00*11*)*"),
  Reg.cc(Reg.fromString "% + 0(0 + 11*0)*11*"));
val it = GREATER : order
- Reg.compareCC
= (Reg.cc(Reg.fromString "(00*11*)*"),
  Reg.cc(Reg.fromString "(1*10*0)*"));
val it = EQUAL : order

```

The module `Reg` also includes these functions:

```

val numConcats      : reg -> int
val numSyms         : reg -> int
val standardized    : reg -> bool
val compareComplexity : reg * reg -> order
val compareComplexityTotal : reg * reg -> order

```

The first two functions implement the functions with the same names. The function `standardized` tests whether a regular expression is standardized, and the function `compareComplexity` implements $\leq_{\text{simp}}/\equiv_{\text{simp}}$. Finally, `compareComplexityTotal` is like `compareComplexity`, but falls back on `Reg.compare` (our total ordering on regular expressions) to order regular expressions with the same complexity. Thus `compareComplexityTotal` is a total ordering.

Here are some examples of how these functions can be used:

```

- Reg.numConcats(Reg.fromString "(01)*(10)*");
val it = 3 : int
- Reg.numSyms(Reg.fromString "(01)*(10)*");
val it = 4 : int
- Reg.standardized(Reg.fromString "00*1");
val it = true : bool
- Reg.standardized(Reg.fromString "00*0");
val it = false : bool
- Reg.compareComplexity
= (Reg.fromString "(00*11*)*",
  Reg.fromString "% + 0(0 + 11*0)*11*");
val it = GREATER : order
- Reg.compareComplexity
= (Reg.fromString "0**1**", Reg.fromString "(01)**");
val it = GREATER : order

```

```

- Reg.compareComplexity
= (Reg.fromString "(0*1*)*", Reg.fromString "(0**1*)*");
val it = GREATER : order
- Reg.compareComplexity
= (Reg.fromString "0+01", Reg.fromString "0(%+1)");
val it = GREATER : order
- Reg.compareComplexity
= (Reg.fromString "(01)2", Reg.fromString "012");
val it = GREATER : order
- Reg.compareComplexity
= (Reg.fromString "1(01+10)+(%+01)1",
  Reg.fromString "011+1(%+01+10)");
val it = EQUAL : order

```

3.3.2 Weak Simplification

In this subsection, we give our first simplification algorithm: weak simplification. We say that a regular expression α is *weakly simplified* iff α is standardized and none of α 's subtrees have any of the following forms:

- $\$ + \beta$ or $\beta + \$$ (the $\$$ is redundant);
- $\beta + \beta$ or $\beta + (\beta + \gamma)$ (the duplicate occurrence of β is redundant);
- $\%\beta$ or $\beta\%$ (the $\%$ is redundant);
- $\$\beta$ or $\beta\$$ (both are equivalent to $\$$); and
- $\%^*$ or $\* or $(\beta^*)^*$ (the first two can be replaced by $\%$, and the extra closure can be omitted in the third case).

Thus, if a regular expression α is weakly simplified, then each of its subtrees will also be weakly simplified.

Weakly simplified regular expressions have some pleasing properties.

Proposition 3.3.13

- (1) For all $\alpha \in \mathbf{Reg}$, if α is weakly simplified and $L(\alpha) = \emptyset$, then $\alpha = \$$.
- (2) For all $\alpha \in \mathbf{Reg}$, if α is weakly simplified and $L(\alpha) = \{\%\}$, then $\alpha = \%$.
- (3) For all $\alpha \in \mathbf{Reg}$, for all $a \in \mathbf{Sym}$, if α is weakly simplified and $L(\alpha) = \{a\}$, then $\alpha = a$.

E.g., part (2) of the proposition says that, if α is weakly simplified and $L(\alpha)$ is the language whose only string is $\%$, then α is $\%$.

Proof. The three parts are proved in order, using induction on regular expressions. We will show the concatenation case of part (3). Suppose $\alpha, \beta \in \mathbf{Reg}$ and assume the inductive hypothesis: for all $a \in \mathbf{Sym}$, if α is weakly simplified

and $L(\alpha) = \{a\}$, then $\alpha = a$, and for all $a \in \mathbf{Sym}$, if β is weakly simplified and $L(\beta) = \{a\}$, then $\beta = a$. Suppose $a \in \mathbf{Sym}$, and assume that $\alpha\beta$ is weakly simplified and $L(\alpha\beta) = \{a\}$. We must show that $\alpha\beta = a$. Because $\alpha\beta$ is weakly simplified, so are α and β .

Since $L(\alpha)L(\beta) = L(\alpha\beta) = \{a\}$, there are two cases to consider.

- Suppose $L(\alpha) = \{a\}$ and $L(\beta) = \{\%\}$. Since β is weakly simplified and $L(\beta) = \{\%\}$, part (2) tells us that $\beta = \%$. But this means that $\alpha\beta = \alpha\%$ is not weakly simplified after all—contradiction. Thus we can conclude that $\alpha\beta = a$.
- Suppose $L(\alpha) = \{\%\}$ and $L(\beta) = \{a\}$. The proof of this case is similar to that of the other one.

□

Proposition 3.3.14

For all $\alpha \in \mathbf{Reg}$, if α is weakly simplified, then $\mathbf{alphabet}(L(\alpha)) = \mathbf{alphabet} \alpha$.

Proof. By Proposition 3.1.5, it suffices to show that, for all $\alpha \in \mathbf{Reg}$, if α is weakly simplified, then $\mathbf{alphabet} \alpha \subseteq \mathbf{alphabet}(L(\alpha))$. And this follows by an easy induction on α , using Proposition 3.3.13(2). □

The next proposition says that $\$$ need only be used at the top-level of a regular expression.

Proposition 3.3.15

For all $\alpha \in \mathbf{Reg}$, if α is weakly simplified and α has one or more occurrences of $\$$, then $\alpha = \$$.

Proof. An easy induction on regular expressions. □

Finally, we have that weakly simplified regular expressions with closures generate infinite languages:

Proposition 3.3.16

For all $\alpha \in \mathbf{Reg}$, if α is weakly simplified and α has one or more closures, then $L(\alpha)$ is infinite.

Proof. An easy induction on regular expressions. □

Next, we see how we can test whether a regular expression is weakly simplified via a simple structural recursion. Define $\mathbf{weaklySimplified} \in \mathbf{Reg} \rightarrow \mathbf{Bool}$ by structural recursion, as follows. Given a regular expression α , it proceeds as follows:

- Suppose α is $\%$, $\$$ or a symbol. Then it returns **true**.

- Suppose α has the form β^* . Then it checks that:
 - β is weakly simplified (this is done using recursion); and
 - β is neither $\%$, nor $\$$, nor a closure.
- Suppose α has the form $\alpha_1 \alpha_2$. Then it checks that:
 - α_1 and α_2 are weakly simplified; and
 - α_1 is neither $\%$ nor $\$$ nor a concatenation; and
 - α_2 is neither $\%$ nor $\$$; and
 - α has none of the following forms: $\beta^*\beta$, $\beta^*(\beta\gamma)$, $(\beta_1\beta_2)^*\beta_1$ or $(\beta_1\beta_2)^*\beta_1\gamma$.
- Suppose α has the form $\alpha_1 + \alpha_2$. Then it checks that:
 - α_1 and α_2 are weakly simplified; and
 - α_1 is neither $\$$ nor a union; and
 - α_2 is not $\$$;
 - if α_2 has the form $\beta_1 + \beta_2$, then $\alpha_1 < \beta_1$; and
 - if α_2 is not a union, then $\alpha_1 < \alpha_2$.

Proposition 3.3.17

For all $\alpha \in \mathbf{Reg}$, α is weakly simplified iff **weaklySimplified** $\alpha = \mathbf{true}$.

Proof. By induction on regular expressions. \square

In preparation for giving our weak simplification algorithm, we need to define some auxiliary functions. We say that a regular expression α is *almost weakly simplified* iff either:

- $w \in \{\%, \$\}$; or
- all elements of **concatstoList** α are weakly simplified, and are not $\%$, $\$$ or concatenations.

For example, $0^*0(1+2)^*(1+2) = 0^*(0((1+2)^*(1+2)))$ is almost weakly simplified, even though it's not weakly simplified. On the other hand: $(\$+1)1$ isn't almost weakly simplified, because $\$+1$ isn't weakly simplified; 1% isn't weakly simplified, because of the location of $\%$; and $(01)1$ isn't almost weakly simplified, because of the location of the concatenation 01 .

Let

$$\begin{aligned} \mathbf{WS} &= \{ \alpha \in \mathbf{Reg} \mid \alpha \text{ is weakly simplified} \}, \text{ and} \\ \mathbf{AWS} &= \{ \alpha \in \mathbf{Reg} \mid \alpha \text{ is almost weakly simplified} \}. \end{aligned}$$

We define a function **shiftClosuresRight** $\in \mathbf{AWS} \rightarrow \mathbf{WS}$ by recursion. Given $\alpha \in \mathbf{AWS}$, **shiftClosuresRight** proceeds as follows. If α is not a concatenation, then it returns α . Otherwise, $\alpha = \alpha_1\alpha_2$ for some $\alpha_1, \alpha_2 \in \mathbf{Reg}$. Since α is almost weakly simplified, so is α_2 . So it lets $\alpha'_2 \in \mathbf{WS}$ be the result of calling **shiftClosuresRight** on α_2 .

- If $\alpha_1\alpha'_2$ has the form $\beta^*\beta$, for some $\beta \in \mathbf{Reg}$, then **shiftClosuresRight** returns

$$\mathbf{shiftClosuresRight}(\mathbf{rightConcat}(\beta, \beta^*)).$$

- Otherwise, if $\alpha_1\alpha'_2$ has the form $\beta^*\beta\gamma$, for some $\beta, \gamma \in \mathbf{Reg}$, then **shiftClosuresRight** returns

$$\mathbf{shiftClosuresRight}(\beta\beta^*\gamma).$$

- Otherwise, if $\alpha_1\alpha'_2$ has the form $(\beta_1\beta_2)^*\beta_1$, for some $\beta_1, \beta_2 \in \mathbf{Reg}$, then **shiftClosuresRight** returns

$$\mathbf{shiftClosuresRight}(\beta_1(\mathbf{rightConcat}(\beta_2, \beta_1))^*).$$

- Otherwise, if $\alpha_1\alpha'_2$ has the form $(\beta_1\beta_2)^*\beta_1\gamma$, for some $\beta_1, \beta_2, \gamma \in \mathbf{Reg}$, then **shiftClosuresRight** returns

$$\mathbf{shiftClosuresRight}(\beta_1(\mathbf{rightConcat}(\beta_2, \beta_1))^*\gamma).$$

- Otherwise, **shiftClosuresRight** returns $\alpha_1\alpha'_2$.

(The work needed to justify the kind of well-founded recursion used in **shiftClosuresRight**'s definition will be added in a subsequent revision.)

Proposition 3.3.18

For all $\alpha \in \mathbf{AWS}$, **shiftClosuresRight** α is equivalent to α and has the same closure complexity, size, number of concatenations and number of symbols as α .

Define a function **deepClosure** $\in \mathbf{WS} \rightarrow \mathbf{WS}$ as follows. For all $\alpha \in \mathbf{WS}$:

$$\begin{aligned} \mathbf{deepClosure} \% &= \%, \\ \mathbf{deepClosure} \$ &= \%, \\ \mathbf{deepClosure} (*(\alpha)) &= \alpha^*, \text{ and} \\ \mathbf{deepClosure} \alpha &= \alpha^*, \text{ if } \alpha \notin \{\%, \$\} \text{ and } \alpha \text{ is not a closure.} \end{aligned}$$

Lemma 3.3.19

For all $\alpha \in \mathbf{WS}$, **deepClosure** α is equivalent to α^* , has the same alphabet as α^* , has a closure complexity that is no bigger than that of α^* , has a size that is no bigger than that of α^* , has the same number of concatenations as α^* , and has the same number of symbols as α^* .

Define a function **deepConcat** $\in \mathbf{WS} \times \mathbf{WS} \rightarrow \mathbf{WS}$ as follows. For all $\alpha, \beta \in \mathbf{WS}$:

$$\begin{aligned} \mathbf{deepConcat}(\alpha, \$) &= \$, \\ \mathbf{deepConcat}(\$, \alpha) &= \$, \text{ if } \alpha \neq \$, \\ \mathbf{deepConcat}(\alpha, \%) &= \alpha, \text{ if } \alpha \neq \$, \\ \mathbf{deepConcat}(\%, \alpha) &= \alpha, \text{ if } \alpha \notin \{\$, \%\}, \text{ and} \\ \mathbf{deepConcat}(\alpha, \beta) &= \mathbf{shiftClosuresRight}(\mathbf{rightConcat}(\alpha, \beta)), \\ &\quad \text{if } \alpha, \beta \notin \{\$, \%\}. \end{aligned}$$

To see that the last clause of this definition is proper, suppose that $\alpha, \beta \in \mathbf{WS} - \{\%, \$\}$. Thus all the elements of **concatstoList** α and **concatstoList** β are weakly simplified, and are not %, \$ or concatenations. Hence

$$\mathbf{concatstoList}(\mathbf{rightConcat}(\alpha, \beta)) = \mathbf{concatstoList} \alpha @ \mathbf{concatstoList} \beta$$

also has this property, showing that **rightConcat**(α, β) is almost weakly simplified, which is what **shiftClosuresRight** needs to deliver a weakly simplified result.

Lemma 3.3.20

*For all $\alpha, \beta \in \mathbf{WS}$, **deepConcat**(α, β) is equivalent to $\alpha\beta$, has an alphabet that is a subset of the alphabet of $\alpha\beta$, has a closure complexity that is no bigger than that of $\alpha\beta$, has a size that is no bigger than that of $\alpha\beta$, has no more concatenations than $\alpha\beta$, and has no more symbols than $\alpha\beta$.*

Define a function **deepUnion** $\in \mathbf{WS} \times \mathbf{WS} \rightarrow \mathbf{WS}$ as follows. For all $\alpha, \beta \in \mathbf{WS}$:

$$\begin{aligned} \mathbf{deepUnion}(\alpha, \$) &= \alpha, \\ \mathbf{deepUnion}(\$, \alpha) &= \alpha, \text{ if } \alpha \neq \$, \text{ and} \\ \mathbf{deepUnion}(\alpha, \beta) &= \mathbf{sortUnions}(\mathbf{rightUnion}(\alpha, \beta)), \text{ if } \alpha \neq \$ \text{ and } \beta \neq \$. \end{aligned}$$

To see that the last clause of this definition is proper, suppose $\alpha, \beta \in \mathbf{WS} - \{\$\}$. Then all the elements of **unionsToList**(**rightUnion**(α, β)) will be weakly simplified, and won't be \$ or unions. Consequently, **sortUnions** will deliver a weakly simplified result.

Lemma 3.3.21

*For all $\alpha, \beta \in \mathbf{WS}$, **deepUnion**(α, β) is equivalent to $\alpha + \beta$, has an alphabet that is a subset of the alphabet of $\alpha + \beta$, has a closure complexity that is no bigger than that of $\alpha + \beta$, has a size that is no bigger than that of $\alpha + \beta$, has no more concatenations than $\alpha + \beta$, and has no more symbols than $\alpha + \beta$.*

Now, we can define our weak simplification function/algorithm. Define **weaklySimplify** $\in \mathbf{Reg} \rightarrow \mathbf{WS}$ by structural recursion:

- $\text{weaklySimplify } \% = \%$;
- $\text{weaklySimplify } \$ = \$$;
- $\text{weaklySimplify } a = a$, for all $a \in \text{Sym}$;
- $\text{weaklySimplify}(*(\alpha))$

$$= \text{deepClosure}(\text{weaklySimplify } \alpha),$$
for all $\alpha \in \text{Reg}$;
- $\text{weaklySimplify}(@(\alpha, \beta))$

$$= \text{deepConcat}(\text{weaklySimplify } \alpha, \text{weaklySimplify } \beta),$$
for all $\alpha, \beta \in \text{Reg}$; and
- $\text{weaklySimplify}+(\alpha, \beta)$

$$= \text{deepUnion}(\text{weaklySimplify } \alpha, \text{weaklySimplify } \beta),$$
for all $\alpha, \beta \in \text{Reg}$.

Proposition 3.3.22

For all $\alpha \in \text{Reg}$:

- (1) $\text{weaklySimplify } \alpha \approx \alpha$;
- (2) $\text{alphabet}(\text{weaklySimplify}(\alpha)) \subseteq \text{alphabet } \alpha$;
- (3) $\text{cc}(\text{weaklySimplify } \alpha) \leq_{cc} \text{cc } \alpha$;
- (4) $\text{size}(\text{weaklySimplify } \alpha) \leq \text{size } \alpha$;
- (5) $\text{numSyms}(\text{weaklySimplify } \alpha) \leq \text{numSyms } \alpha$; and
- (6) $\text{numConcats}(\text{weaklySimplify } \alpha) \leq \text{numConcats } \alpha$.

Proof. By induction on regular expressions. \square

Exercise 3.3.23

Prove Proposition 3.3.22.

Proposition 3.3.24

For all $\alpha \in \text{Reg}$, if α is weakly simplified, then $\text{weaklySimplify}(\alpha) = \alpha$.

Proof. By induction on regular expressions. \square

Using our weak simplification algorithm, we can define an algorithm for calculating the language generated by a regular expression, when this language is finite, and for announcing that this language is infinite, otherwise. First, we weakly simplify our regular expression, α , and call the resulting regular expression β . If β contains no closures, then we compute its meaning in the usual way. But, if β contains one or more closures, then its language will be infinite, and thus we can output a message saying that $L(\alpha)$ is infinite.

The Forlan module **Reg** defines the following functions relating to weak simplification:

```
val weaklySimplified : reg -> bool
val weaklySimplify   : reg -> reg
val toStrSet         : reg -> str set
```

The function **weaklySimplified** tests whether its argument is weakly simplified, and **weaklySimplify** implements **weaklySimplify**. Finally, the function **toStrSet** implements our algorithm for calculating the language generated by a regular expression, if that language is finite, and for announcing the this language is infinite, otherwise.

Here are some examples of how these functions can be used:

```
- val reg = Reg.input "";
@ (% + $0)(% + 00*0 + 0**)*
@ .
val reg = - : reg
- Reg.output("", Reg.weaklySimplify reg);
(% + 0* + 000)*
val it = () : unit
- Reg.toStrSet reg;
language is infinite

uncaught exception Error
- val reg' = Reg.input "";
@ (1 + %)(2 + $)(3 + %*)(4 + $*)
@ .
val reg' = - : reg
- StrSet.output("", Reg.toStrSet reg');
2, 12, 23, 24, 123, 124, 234, 1234
val it = () : unit
- Reg.output("", Reg.weaklySimplify reg');
(% + 1)2(% + 3)(% + 4)
val it = () : unit
- Reg.output
= ("",
= Reg.weaklySimplify(Reg.fromString "(00*11*)*"));
(00*11*)*
val it = () : unit
```

3.3.3 Local and Global Simplification

In preparation for the definition of our local and global simplification algorithms, we must define some auxiliary functions. First, we show how we can recursively test whether $\% \in L(\alpha)$, for a regular expression α . We define a function

$$\text{hasEmp} \in \mathbf{Reg} \rightarrow \mathbf{Bool}$$

by recursion:

$$\begin{aligned} \text{hasEmp } \% &= \mathbf{true}; \\ \text{hasEmp } \$ &= \mathbf{false}; \\ \text{hasEmp } a &= \mathbf{false}, \text{ for all } a \in \mathbf{Sym}; \\ \text{hasEmp } (\alpha^*) &= \mathbf{true}, \text{ for all } \alpha \in \mathbf{Reg}; \\ \text{hasEmp } (\alpha\beta) &= \text{hasEmp } \alpha \text{ and } \text{hasEmp } \beta, \text{ for all } \alpha, \beta \in \mathbf{Reg}; \text{ and} \\ \text{hasEmp } (\alpha + \beta) &= \text{hasEmp } \alpha \text{ or } \text{hasEmp } \beta, \text{ for all } \alpha, \beta \in \mathbf{Reg}. \end{aligned}$$

Proposition 3.3.25

For all $\alpha \in \mathbf{Reg}$, $\% \in L(\alpha)$ iff $\text{hasEmp } \alpha = \mathbf{true}$.

Proof. By induction on regular expressions. \square

Next, we show how we can recursively test whether $a \in L(\alpha)$, for a symbol a and a regular expression α . We define a function

$$\text{hasSym} \in \mathbf{Sym} \times \mathbf{Reg} \rightarrow \mathbf{Bool}$$

by recursion:

$$\begin{aligned} \text{hasSym}(a, \%) &= \mathbf{false}, \text{ for all } a \in \mathbf{Sym}; \\ \text{hasSym}(a, \$) &= \mathbf{false}, \text{ for all } a \in \mathbf{Sym}; \\ \text{hasSym}(a, b) &= a = b, \text{ for all } a, b \in \mathbf{Sym}; \\ \text{hasSym}(a, \alpha^*) &= \text{hasSym}(a, \alpha), \text{ for all } a \in \mathbf{Sym} \text{ and } \alpha \in \mathbf{Reg}; \\ \text{hasSym}(a, \alpha\beta) &= (\text{hasSym}(a, \alpha) \text{ and } \text{hasEmp}(\beta)) \text{ or} \\ &\quad (\text{hasEmp}(\alpha) \text{ and } \text{hasSym}(a, \beta)), \\ &\quad \text{for all } a \in \mathbf{Sym} \text{ and } \alpha, \beta \in \mathbf{Reg}; \text{ and} \\ \text{hasSym}(a, \alpha + \beta) &= \text{hasSym}(a, \alpha) \text{ or } \text{hasSym}(a, \beta), \\ &\quad \text{for all } a \in \mathbf{Sym} \text{ and } \alpha, \beta \in \mathbf{Reg}. \end{aligned}$$

Proposition 3.3.26

For all $a \in \mathbf{Sym}$ and $\alpha \in \mathbf{Reg}$, $a \in L(\alpha)$ iff $\text{hasSym}(a, \alpha) = \mathbf{true}$.

Proof. By induction on regular expressions, using Proposition 3.3.25. \square

Finally, we define a function/algorithm

$$\mathbf{obviousSubset} \in \mathbf{Reg} \times \mathbf{Reg} \rightarrow \{\mathbf{true}, \mathbf{false}\}$$

meeting the following specification: for all $\alpha, \beta \in \mathbf{Reg}$,

$$\text{if } \mathbf{obviousSubset}(\alpha, \beta) = \mathbf{true}, \text{ then } L(\alpha) \subseteq L(\beta).$$

I.e., this function is a *conservative approximation to subset testing*. The function that always returns **false** would meet this specification, but our function will do much better than this, and will be reasonably efficient. In Section ??, we will learn of a less efficient algorithm that will provide a complete test for $L(\alpha) \subseteq L(\beta)$.

Given $\alpha, \beta \in \mathbf{Reg}$, $\mathbf{obviousSubset}(\alpha, \beta)$ proceeds as follows. First, it lets $\alpha' = \mathbf{weaklySimplify} \alpha$ and $\beta' = \mathbf{weaklySimplify} \beta$. Then it returns $\mathbf{obviSub}(\alpha', \beta')$, where

$$\mathbf{obviSub} \in \mathbf{WS} \times \mathbf{WS} \rightarrow \mathbf{Bool}$$

is the function defined below.

obviSub is defined by well-founded recursion on the sum of the sizes of its arguments. If $\alpha = \beta$, then it returns **true**; otherwise, it considers the possible forms of α .

- Suppose $\alpha = \%$. It returns **hasEmp** β .
- Suppose $\alpha = \$$. It returns **true**.
- Suppose $\alpha = a$, for some $a \in \mathbf{Sym}$. It returns **hasSym**(a, β).
- Suppose $\alpha = \alpha_1^*$, for some $\alpha_1 \in \mathbf{Reg}$. Here it looks at the form of β .
 - Suppose $\beta = \%$. It returns **false**. (Because α will be weakly simplified, and so α won't generate $\{\%\}$.)
 - Suppose $\beta = \$$. It returns **false**.
 - Suppose $\beta = a$, for some $a \in \mathbf{Sym}$. It returns **false**.
 - Suppose β is a closure. It returns $\mathbf{obviSub}(\alpha_1, \beta)$.
 - Suppose $\beta = \beta_1\beta_2$, for some $\beta_1, \beta_2 \in \mathbf{Reg}$. If **hasEmp** $\beta_1 = \mathbf{true}$ and $\mathbf{obviSub}(\alpha, \beta_2)$, then it returns **true**. Otherwise, if **hasEmp** $\beta_2 = \mathbf{true}$ and $\mathbf{obviSub}(\alpha, \beta_1)$, then it returns **true**. Otherwise, it returns **false** (even though the answer sometimes should be **true**).
 - Suppose $\beta = \beta_1 + \beta_2$, for some $\beta_1, \beta_2 \in \mathbf{Reg}$. It returns

$$\mathbf{obviSub}(\alpha, \beta_1) \text{ or } \mathbf{obviSub}(\alpha, \beta_2)$$

(even though this is **false** too often).

- Suppose $\alpha = \alpha_1\alpha_2$, for some $\alpha_1, \alpha_2 \in \mathbf{Reg}$. Here it looks at the form of β .
 - Suppose $\beta = \%$. It returns **false**. (Because α is weakly simplified, α won't generate $\{\%\}$.)
 - Suppose $\beta = \$$. It returns **false**. (Because α is weakly simplified, α won't generate \emptyset .)
 - Suppose $\beta = a$, for some $a \in \mathbf{Sym}$. It returns **false**. (Because α is weakly simplified, α won't generate $\{a\}$.)
 - Suppose $\beta = \beta_1^*$, for some $\beta_1 \in \mathbf{Reg}$. It returns

$$\mathbf{obviSub}(\alpha, \beta_1)$$

or

$$(\mathbf{obviSub}(\alpha_1, \beta) \text{ and } \mathbf{obviSub}(\alpha_2, \beta))$$

(even though this returns **false** too often).

- Suppose $\beta = \beta_1\beta_2$, for some $\beta_1, \beta_2 \in \mathbf{Reg}$. If $\mathbf{obviSub}(\alpha_1, \beta_1) = \mathbf{true}$ and $\mathbf{obviSub}(\alpha_2, \beta_2) = \mathbf{true}$, then it returns **true**. Otherwise, if $\mathbf{hasEmp} \beta_1 = \mathbf{true}$ and $\mathbf{obviSub}(\alpha, \beta_2) = \mathbf{true}$, then it returns **true**. Otherwise, if $\mathbf{hasEmp} \beta_2 = \mathbf{true}$ and $\mathbf{obviSub}(\alpha, \beta_1) = \mathbf{true}$, then it returns **true**. Otherwise, if β_1 is a closure but β_2 is not a closure, then it returns

$$\mathbf{obviSub}(\alpha_1, \beta_1) \text{ and } \mathbf{obviSub}(\alpha_2, \beta)$$

(even though this returns **false** too often). Otherwise, if β_2 is a closure but β_1 is not a closure, then it returns

$$\mathbf{obviSub}(\alpha_1, \beta) \text{ and } \mathbf{obviSub}(\alpha_2, \beta_2)$$

(even though this returns **false** too often). Otherwise, if β_1 and β_2 are closures, then it returns

$$(\mathbf{obviSub}(\alpha_1, \beta_1) \text{ and } \mathbf{obviSub}(\alpha_2, \beta))$$

or

$$(\mathbf{obviSub}(\alpha_1, \beta) \text{ and } \mathbf{obviSub}(\alpha_2, \beta_2))$$

(even though this returns **false** too often). Otherwise, it returns **false**, even though sometimes we would like the answer to be **true**).

- Suppose $\beta = \beta_1 + \beta_2$, for some $\beta_1, \beta_2 \in \mathbf{Reg}$. It returns

$$\mathbf{obviSub}(\alpha, \beta_1) \text{ or } \mathbf{obviSub}(\alpha, \beta_2)$$

(even though this is **false** too often).

- Suppose $\alpha = \alpha_1 + \alpha_2$. It returns

obviSub(α_1, β) and **obviSub**(α_2, β).

We say that α is *obviously a subset of* β iff **obviousSubset**(α, β) = **true**. On the positive side, we have that, e.g., **obviousSubset**($0^*011^*1, 0^*1^*$) = **true**. On the other hand, **obviousSubset**((01) * , ($\% + 0$)(10) * ($\% + 1$)) = **false**, even though $L((01)^*) \subseteq L((\% + 0)(10)^*(\% + 1))$.

Proposition 3.3.27

For all $\alpha, \beta \in \mathbf{Reg}$, if **obviousSubset**(α, β) = **true**, then $L(\alpha) \subseteq L(\beta)$.

Proof. First, we use induction on the sum of the sizes of α and β to show that, for all $\alpha, \beta \in \mathbf{Reg}$, if **obviSub**(α, β) = **true**, then $L(\alpha) \subseteq L(\beta)$. The result then follows by Proposition 3.3.22. \square

The Forlan module **Reg** provides the following functions corresponding to the auxiliary functions **hasEmp**, **hasSym** and **obviousSubset**:

```
val hasEmp      : reg -> bool
val hasSym      : sym * reg -> bool
val obviousSubset : reg * reg -> bool
```

Here are some examples of how they can be used:

```
- Reg.hasEmp(Reg.fromString "0*1*");
val it = true : bool
- Reg.hasEmp(Reg.fromString "01*");
val it = false : bool
- Reg.hasSym(Sym.fromString "0", Reg.fromString "0*1*");
val it = true : bool
- Reg.hasSym(Sym.fromString "1", Reg.fromString "0*1*");
val it = true : bool
- Reg.hasSym(Sym.fromString "0", Reg.fromString "0*$");
val it = false : bool
- Reg.obviousSubset
= (Reg.fromString "(0 + 1)*",
  = Reg.fromString "0*(0 + 1)*1*");
val it = true : bool
- Reg.obviousSubset
= (Reg.fromString "0*(0 + 1)*1*",
  = Reg.fromString "(0 + 1)*");
val it = true : bool
- Reg.obviousSubset
= (Reg.fromString "0*011*1",
  = Reg.fromString "0*1*");
val it = true : bool
- Reg.obviousSubset
```

```

= (Reg.fromString "(01 + 011)1*",
  = Reg.fromString "01*");
val it = true : bool
- Reg.obviousSubset
= (Reg.fromString "(01)*",
  = Reg.fromString "(% + 0)(10)*(% + 1)");
val it = false : bool

```

Our local and global simplification algorithms make use of simplification rules, which may be applied to arbitrary subtrees of regular expressions. There are three kinds of rules: structural rules, distributive rules and reduction rules.

There are nine *structural rules*, which preserve the alphabet, closure complexity, size, number of concatenations and number of symbols of a regular expression:

- (1) $(\alpha + \beta) + \gamma \rightarrow \alpha + (\beta + \gamma)$.
- (2) $\alpha + (\beta + \gamma) \rightarrow (\alpha + \beta) + \gamma$.
- (3) $\alpha(\beta\gamma) \rightarrow (\alpha\beta)\gamma$.
- (4) $(\alpha\beta)\gamma \rightarrow \alpha(\beta\gamma)$.
- (5) $\alpha + \beta \rightarrow \beta + \alpha$.
- (6) $\alpha^*\alpha \rightarrow \alpha\alpha^*$.
- (7) $\alpha\alpha^* \rightarrow \alpha^*\alpha$.
- (8) $\alpha(\beta\alpha)^* \rightarrow (\alpha\beta)^*\alpha$.
- (9) $(\alpha\beta)^*\alpha \rightarrow \alpha(\beta\alpha)^*$.

For even small regular expressions, there may be a very large number of ways to reorganize them using the structural rules. E.g., suppose $\alpha_1, \dots, \alpha_n$ are distinct regular expressions, for $n \geq 1$. There are $n!$ ways of ordering the α_i . And there are $(2n)!/(n!)(n+1)!$ (these are the Catalan numbers) binary trees with exactly n leaves. Consequently, using structural rules (1), (2) and (5) (withing making changes inside the α_i), we can reorganize $\alpha_1 + \dots + \alpha_n$ into $(n!)(2n)!/(n!)(n+1)!$ regular expressions. For $n = 16$, this is about $7 * 10^{20}$.

There are two *distributive rules*, which preserve the alphabet of a regular expression:

- (1) $\alpha(\beta_1 + \beta_2) \rightarrow \alpha\beta_1 + \alpha\beta_2$.
- (2) $(\alpha_1 + \alpha_2)\beta \rightarrow \alpha_1\beta + \alpha_2\beta$.

Finally, there are 26 *reduction rules*, some of which make use of a conservative approximation *sub* to subset testing. When $\alpha \rightarrow \beta$ because of a reduction rule, we have that $\mathbf{alphabet} \beta \subseteq \mathbf{alphabet} \alpha$ and $\beta \mathbf{simp} \alpha$, where **simp** is the well-founded relation on **Reg** that is defined below.

We define the relation **simp** on **Reg** by: for all $\alpha, \beta \in \mathbf{Reg}$, $\alpha \mathbf{simp} \beta$ iff either:

- $\mathbf{cc} \alpha <_{cc} \mathbf{cc} \beta$; or
- $\mathbf{cc} \alpha = \mathbf{cc} \beta$, but $\mathbf{size} \alpha < \mathbf{size} \beta$; or
- $\mathbf{cc} \alpha = \mathbf{cc} \beta$ and $\mathbf{size} \alpha = \mathbf{size} \beta$, but $\mathbf{numConcat} \alpha < \mathbf{numConcat} \beta$;
or
- $\mathbf{cc} \alpha = \mathbf{cc} \beta$ and $\mathbf{size} \alpha = \mathbf{size} \beta$, and $\mathbf{numConcat} \alpha = \mathbf{numConcat} \beta$,
but $\mathbf{numSyms} \alpha < \mathbf{numSyms} \beta$.

Note that this is almost the same definition as that of $<_{\mathbf{simp}}$ —the difference being that **simp** doesn't have the final step involving standardization.

Proposition 3.3.28

simp is a well-founded relation on **Reg**.

Proof. Follows by Propositions 3.3.5, 1.2.11 and 1.2.10, plus the fact that $<$ is well-founded on \mathbb{N} (Proposition 1.2.5). \square

The following proposition says that if we replace a subtree of a regular expression by a regular expression that is a **simp**-predecessor of that subtree, that the resulting, whole regular expression will be a **simp**-predecessor of the original, whole regular expression.

Proposition 3.3.29

Suppose $\alpha, \beta, \beta' \in \mathbf{Reg}$, $\beta' \mathbf{simp} \beta$, $pat \in \mathbf{Path}$ is valid for α , and β is the subtree of α at position pat . Let α' be the result of replacing the subtree at position pat in α by β' . Then $\alpha' \mathbf{simp} \alpha$.

Our reduction rules follow. In the rules, we abbreviate $\mathbf{hasEmp} \alpha = \mathbf{true}$ and $\mathbf{sub}(\alpha, \beta) = \mathbf{true}$ to $\mathbf{hasEmp} \alpha$ and $\mathbf{sub}(\alpha, \beta)$, respectively. Most of the rules strictly decrease a regular expression's closure complexity and size. The exceptions are labeled “cc” (for when the closure complexity strictly decreases, but the size strictly increases), “concatenations” (for when the closure complexity and size are preserved, but the number of concatenations strictly decreases) or “symbols” (for when the closure complexity and size normally strictly decrease, but occasionally they and the number of concatenations stay the same, but the number of symbols strictly decreases).

- (1) If $\mathbf{sub}(\alpha, \beta)$, then $\alpha + \beta \rightarrow \beta$.

- (2) $\alpha\beta_1 + \alpha\beta_2 \rightarrow \alpha(\beta_1 + \beta_2)$.
- (3) $\alpha_1\beta + \alpha_2\beta \rightarrow (\alpha_1 + \alpha_2)\beta$.
- (4) If **hasEmp** α and $\text{sub}(\alpha, \beta^*)$, then $\alpha\beta^* \rightarrow \beta^*$.
- (5) If **hasEmp** β and $\text{sub}(\beta, \alpha^*)$, then $\alpha^*\beta \rightarrow \alpha^*$.
- (6) If $\text{sub}(\alpha, \beta^*)$, then $(\alpha + \beta)^* \rightarrow \beta^*$.
- (7) $(\alpha^* + \beta)^* \rightarrow (\alpha + \beta)^*$.
- (8) (concatenations) If **hasEmp** α and **hasEmp** β , then $(\alpha\beta)^* \rightarrow (\alpha + \beta)^*$.
- (9) (concatenations) If **hasEmp** α and **hasEmp** β , then $(\alpha\beta + \gamma)^* \rightarrow (\alpha + \beta + \gamma)^*$.
- (10) If **hasEmp** α and $\text{sub}(\alpha, \beta^*)$, then $(\alpha\beta)^* \rightarrow \beta^*$.
- (11) If **hasEmp** β and $\text{sub}(\beta, \alpha^*)$, then $(\alpha\beta)^* \rightarrow \alpha^*$.
- (12) If **hasEmp** α and $\text{sub}(\alpha, (\beta + \gamma)^*)$, then $(\alpha\beta + \gamma)^* \rightarrow (\beta + \gamma)^*$.
- (13) If **hasEmp** β and $\text{sub}(\beta, (\alpha + \gamma)^*)$, then $(\alpha\beta + \gamma)^* \rightarrow (\alpha + \gamma)^*$.
- (14) (cc) If **not**(**hasEmp** α) and $\text{cc } \alpha \cup \overline{\text{cc } \beta} <_{cc} \overline{\text{cc } \beta}$, then $(\alpha\beta^*)^* \rightarrow \% + \alpha(\alpha + \beta)^*$.
- (15) (cc) If **not**(**hasEmp** β) and $\overline{\text{cc } \alpha} \cup \text{cc } \beta <_{cc} \overline{\text{cc } \alpha}$, then $(\alpha^*\beta)^* \rightarrow \% + (\alpha + \beta)^*\beta$.
- (16) (cc) If **not**(**hasEmp** α) or **not**(**hasEmp** γ), and $\text{cc } \alpha \cup \overline{\text{cc } \beta} \cup \text{cc } \gamma <_{cc} \overline{\text{cc } \beta}$, then $(\alpha\beta^*\gamma)^* \rightarrow \% + \alpha(\beta + \gamma\alpha)^*\gamma$.
- (17) If $\text{sub}(\alpha\alpha^*, \beta)$, then $\alpha^* + \beta \rightarrow \% + \beta$.
- (18) If **hasEmp** β and $\text{sub}(\alpha\alpha\alpha^*, \beta)$, then $\alpha^* + \beta \rightarrow \alpha + \beta$.
- (19) (symbols) If $\alpha \notin \{\%, \$\}$ and $\text{sub}(\alpha^n, \beta)$, then $\alpha^{n+1}\alpha^* + \beta \rightarrow \alpha^n\alpha^* + \beta$.
- (20) If $n \geq 2$, $l \geq 0$ and $2n - 1 < m_1 < \dots < m_l$, then $(\alpha^n + \alpha^{n+1} + \dots + \alpha^{2n-1} + \alpha^{m_1} + \dots + \alpha^{m_l})^* \rightarrow \% + \alpha^n\alpha^*$.
- (21) (symbols) If $\alpha \notin \{\%, \$\}$, then $\alpha + \alpha\beta \rightarrow \alpha(\% + \beta)$.
- (22) (symbols) If $\alpha \notin \{\%, \$\}$, then $\alpha + \beta\alpha \rightarrow (\% + \beta)\alpha$.
- (23) $\alpha^*(\% + \beta(\alpha + \beta)^*) \rightarrow (\alpha + \beta)^*$.
- (24) $(\% + (\alpha + \beta)^*\alpha)\beta^* \rightarrow (\alpha + \beta)^*$.

(25) If $\text{sub}(\alpha, \beta^*)$ and $\text{sub}(\beta, \alpha)$, then $\% + \alpha\beta^* \rightarrow \beta^*$.

(26) If $\text{sub}(\beta, \alpha^*)$ and $\text{sub}(\alpha, \beta)$, then $\% + \alpha^*\beta \rightarrow \alpha^*$.

In rules (14)-(16), the preconditions involving **cc** are necessary and sufficient conditions for the right-hand side to have strictly smaller closure complexity than the left-hand side.

Consider, e.g., reduction rule (4). Suppose **hasEmp** $\alpha = \text{true}$ and $\text{sub}(\alpha, \beta^*) = \text{true}$, so that that $\% \in L(\alpha)$ and $L(\alpha) \subseteq L(\beta^*)$. We need that $\alpha\beta^* \approx \beta^*$, **alphabet** $(\beta^*) \subseteq \text{alphabet}(\alpha\beta^*)$ and $\beta^* \text{simp } \alpha\beta^*$. The alphabet of β^* is clearly a subset of that of $\alpha\beta^*$.

To obtain $\alpha\beta^* \approx \beta^*$, it will suffice to show that, for all $A, B \in \mathbf{Lan}$, if $\% \in A$ and $A \subseteq B^*$, then $AB^* = B^*$. Suppose $A, B \in \mathbf{Lan}$, $\% \in A$ and $A \subseteq B^*$. We show that $AB^* \subseteq B^* \subseteq AB^*$. Suppose $w \in AB^*$, so that $w = xy$, for some $x \in A$ and $y \in B^*$. Since $A \subseteq B^*$, it follows that $w = xy \in B^*B^* = B^*$. Suppose $w \in B^*$. Then $w = \%w \in AB^*$.

And, to see that $\beta^* <_{cc} \alpha\beta^*$, it will suffice to show that **cc** $(\beta^*) <_{cc} \text{cc}(\alpha\beta^*)$. And we have that

$$\text{cc}(\beta^*) = \overline{\text{cc}\beta} <_{cc} \text{cc}\alpha \cup \overline{\text{cc}\beta} = \text{cc}(\alpha\beta^*).$$

Because the structural rules preserve the size and alphabet of regular expressions, if we start with a regular expression α , there are only finitely many regular expressions that we can transform α into using structural rules (we can apply one of the rules to some subtree of α , giving us β_1 , apply a rule to one of the subtrees of β_1 , giving us β_2 , etc.).

Suppose sub is a conservative approximation to subset testing. We say that a regular expression α is *locally simplified with respect to sub*: iff

- α is weakly simplified, and
- α can't be transformed by our structural rules into a regular expression to which one of our reduction rules applies.

The *local simplification* of a regular expression α with respect to a conservative approximation to subset testing sub proceeds as follows. It calls its main function with the weak simplification, β of α . The closure complexity, size, number of concatenations, and number of symbols of β are no bigger than those of α , and **alphabet** $\beta \subseteq \text{alphabet } \alpha$.

The main function is defined by well-founded recursion **simp**. It works as follows, when called with a weakly simplified argument, α .

- It generates the set X of all regular expressions **weaklySimplify** γ , such that α can be reorganized using the structural rules into a regular expression β , which can be transformed by a single application of one of our reduction rules into γ .

- If X is empty, then it returns α .
- Otherwise, it calls itself recursively on the simplest element, γ of X (when X doesn't have a unique simplest element, the smallest of the simplest elements—in our total ordering on regular expressions—is selected). Because
 - the structural rules preserve closure complexity, size, number of concatenations, and number of symbols,
 - the reduction rules produce **simp**-predecessors, and
 - and weak simplification doesn't increase closure complexity, size, numbers of concatenations, or numbers of symbols,

we have that γ **simp** α , so that this recursive call is legal. Furthermore, weak simplification, and all of the rules, either preserve or decrease (via \subseteq) the alphabet of regular expressions. Thus **alphabet** $\gamma \subseteq$ **alphabet** α .

The algorithm is referred to as “local”, because at each recursive call of its main function, γ is chosen using the best local knowledge. This strategy is reasonably efficient, but there is no guarantee that another local choice wouldn't result in a simpler global answer.

We define a function/algorithm

$$\text{locallySimplify} \in (\mathbf{Reg} \times \mathbf{Reg} \rightarrow \mathbf{Bool}) \rightarrow \mathbf{Reg} \rightarrow \mathbf{Reg}$$

by: for all conservative approximations to subset testing sub , and $\alpha \in \mathbf{Reg}$, **locallySimplify** sub α is the result of running our local simplification algorithm on α , using sub as the conservative approximation to subset testing.

Theorem 3.3.30

For all conservative approximations to subset testing sub , and $\alpha \in \mathbf{Reg}$:

- **locallySimplify** sub α is locally simplified with respect to sub ;
- **locallySimplify** sub α is equivalent to α ;
- **alphabet**(**locallySimplify** sub α) \subseteq **alphabet** α ; and
- **locallySimplify** sub $\alpha \leq_{\text{simp}} \alpha$.

The Forlan module **Reg** provides the following functions relating to local simplification:

```
val locallySimplified    :
    (reg * reg -> bool) -> reg -> bool
val locallySimplify     :
    int option * (reg * reg -> bool) -> reg -> bool * reg
val locallySimplifyTrace :
    int option * (reg * reg -> bool) -> reg -> bool * reg
```

The function `locallySimplified` takes in a conservative approximation to subset testing *sub* and returns a function that tests whether a regular expression is *sub*-locally simplified. The function `locallySimplifyTrace` implements **locallySimplify**. It emits tracing messages explaining its operation, takes in an extra argument of type `int option`, and produces an extra result of type `bool`. If this extra argument is `NONE`, then it runs as does **locallySimplify**, and its boolean result is always `true`. But if it is `SOME n`, for $n \geq 1$, then at each recursive call of the algorithm's function, no more than n ways of reorganizing the function's argument will be considered, and the boolean part of the result will be `false` iff, in the final recursive call, n was not sufficient to explore all structural reorganizations, so that the regular expression returned may not be locally simplified with respect to *sub*. The function `locallySimplify` works identically, except it doesn't issue tracing messages.

Here are some examples of how these functions can be used.

```
- val locSimpd = Reg.locallySimplified Reg.obviousSubset;
val locSimpd = fn : reg -> bool
- locSimpd(Reg.fromString "(1 + 00*1)*00*");
val it = false : bool
- locSimpd(Reg.fromString "(0 + 1)*0");
val it = true : bool
- fun locSimp nOpt =
=      Reg.locallySimplify(nOpt, Reg.obviousSubset);
val locSimp = fn : int option -> reg -> bool * reg
- locSimp NONE (Reg.fromString "% + 0*0(0 + 1)* + 1*1(0 + 1)*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
(0 + 1)*
val it = () : unit
- locSimp NONE (Reg.fromString "% + 1*0(0 + 1)* + 0*1(0 + 1)*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
(0 + 1)*
val it = () : unit
- locSimp NONE (Reg.fromString "(1 + 00*1)*00*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
(0 + 1)*0
val it = () : unit
- Reg.locallySimplifyTrace
= (NONE, Reg.obviousSubset)
= (Reg.fromString "1*(01*01)*");
considered all 10 structural reorganizations of 1*(01*01)*
1*(01*01)* transformed by structural rule 4 at position [2, 1] to
1*((01*)01)* transformed by structural rule 4 at position [2, 1]
to 1*(((01*)0)1)* transformed by reduction rule 14 at position
[2] to 1*(% + ((01*)0)((01*)0 + 1)*) weakly simplifies to
1*(% + 01*0(1 + 01*0)*)
```

```

considered all 40 structural reorganizations of
1*(% + 01*0(1 + 01*0)*)
1*(% + 01*0(1 + 01*0)*) transformed by structural rule 4 at
position [2, 2, 2] to 1*(% + 0(1*0)(1 + 01*0)*) transformed by
structural rule 4 at position [2, 2] to 1*(% + (01*0)(1 + 01*0)*)
transformed by reduction rule 23 at position [] to (1 + 01*0)*
considered all 4 structural reorganizations of (1 + 01*0)*
(1 + 01*0)* is locally simplified
val it = (true,-) : bool * reg

```

For even fairly small regular expressions, running through all the structural reorganizations can take prohibitively long. So, one often has to bound the number of such reorganizations, as in:

```

- val reg = Reg.input "";
@ 1 + (% + 0 + 2)(% + 0 + 2)*1 +
@ (1 + (% + 0 + 2)(% + 0 + 2)*1)
@ (% + 0 + 2 + 1(% + 0 + 2)*1)
@ (% + 0 + 2 + 1(% + 0 + 2)*1)*
@ .
val reg = - : reg
- Reg.equal(Reg.weaklySimplify reg, reg);
val it = true : bool
- val (b', reg') = locSimp (SOME 10) reg;
val b' = false : bool
val reg' = - : reg
- Reg.output("", reg');
(0 + 2)*1(0 + 2 + 1(0 + 2)*1)*
val it = () : unit
- val (b'', reg'') = locSimp (SOME 1000) reg';
val b'' = true : bool
val reg'' = - : reg
- Reg.output("", reg'');
(0 + 2)*1(0 + 2 + 1(0 + 2)*1)*
val it = () : unit

```

Note that, in this transcript, `reg'` turns out to be locally simplified, despite the fact that `b'` is `false`.

Our global simplification algorithm comes in two variants, a non-distributive one, which doesn't use the distributive rules, and a distributive one, which does. Given a boolean b , a conservative approximation to subset testing sub , and a regular expression α , we say that α is *globally simplified with respect to b and sub* iff no strictly simpler regular expression can be found by an arbitrary number of applications of weak simplification, structural rules, reduction rules and—if $b = \mathbf{true}$ —distributive rules.

The *global simplification of a regular expression α with respect to a boolean b and conservative approximation to subset testing sub* consists of generating the set X of all regular expressions β that can be formed from α by an arbitrary

number of applications of weak simplification, the structural rules, reduction rules, and—in the case of the distributive variant—the distributive ones. The simplest element of X is then selected (when there isn't a unique simplest element, the smallest of the simplest elements—in our total ordering on regular expressions—is selected). (A proof that the generation of X terminates even in the distributive case is not yet complete.)

Of course, this algorithm is much less efficient than the local one, but by revisiting choices, it is capable of producing simpler answers.

We define a function/algorithm

$$\mathbf{globallySimplify} \in \mathbf{Bool} \times (\mathbf{Reg} \times \mathbf{Reg} \rightarrow \mathbf{Bool}) \rightarrow \mathbf{Reg} \rightarrow \mathbf{Reg}$$

by: for all $b \in \mathbf{Bool}$, conservative approximation to subset testing sub , and $\alpha \in \mathbf{Reg}$, $\mathbf{globallySimplify}(b, sub)\alpha$ is the result of running our global simplification algorithm on α , including the distributive rules iff $b = \mathbf{true}$, and using sub as our conservative approximation to subset testing.

Theorem 3.3.31

For all $b \in \mathbf{Bool}$, conservative approximations to subset testing sub , and $\alpha \in \mathbf{Reg}$:

- $\mathbf{globallySimplify}(b, sub)\alpha$ is globally simplified with respect to b and sub ;
- $\mathbf{globallySimplify}(b, sub)\alpha$ is equivalent to α ;
- $\mathbf{alphabet}(\mathbf{globallySimplify}(b, sub)\alpha) \subseteq \mathbf{alphabet}\alpha$; and
- $\mathbf{globallySimplify}(b, sub)\alpha \leq_{\mathbf{simp}} \alpha$.

The Forlan module `Reg` provides the following functions relating to global simplification:

```
val globallySimplified      :
    bool * (reg * reg -> bool) -> reg -> bool
val globallySimplifyTrace :
    int option * bool * (reg * reg -> bool) -> reg -> bool * reg
val globallySimplify       :
    int option * bool * (reg * reg -> bool) -> reg -> bool * reg
```

The function `globallySimplified` takes in a boolean b and a conservative approximation to subset testing sub , and returns a function that tests whether a regular expression is globally simplified with respect to b and sub . The function `globallySimplifyTrace` implements `globallySimplify`. It emits tracing messages explaining its operation, and takes in an extra argument of type `int option`, and produces an extra result of type `bool`. If this argument is `NONE`, then it runs as does `globallySimplify`, and the boolean result is always

true. But if it is **SOME** n , for $n \geq 1$, then at most n elements of the set X are generated, before picking the simplest one, and the boolean result is **false** if this n isn't enough to generate all of X . The function `globallySimplify` works identically, except it doesn't issue tracing messages.

For even quite small regular expressions, `globallySimplified` will fail to run to completion in an acceptable time-frame, and one will have to bound the size of the set X in order for `globallySimplify` and `globallySimplifyTrace` to run to completion in an acceptable time-frame.

Here are some examples of how these functions can be used.

```
- Reg.globallySimplifyTrace
= (NONE, false, Reg.obviousSubset)
= (Reg.fromString "(00*1)*");
considering candidates with explanations of length 0
simplest result now: (00*1)*
considering candidates with explanations of length 1
simplest result now: (00*1)* transformed by reduction rule 16 at
position [] to % + 0(0 + 10)*1
considering candidates with explanations of length 2
simplest result now: (00*1)* transformed by reduction rule 16 at
position [] to % + 0(0 + 10)*1 transformed by reduction rule 22 at
position [2, 2, 1, 1] to % + 0((% + 1)0)*1
considering candidates with explanations of length 3
considering candidates with explanations of length 4
considering candidates with explanations of length 5
considering candidates with explanations of length 6
considering candidates with explanations of length 7
search completed after considering 36 candidates with maximum size
12
(00*1)* transformed by reduction rule 16 at position [] to
% + 0(0 + 10)*1 transformed by reduction rule 22 at position
[2, 2, 1, 1] to % + 0((% + 1)0)*1 is globally simplified
val it = (true,-) : bool * reg
- locSimp NONE (Reg.fromString "(00*11*)*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
% + 00*1(% + (0 + 1)*1)
val it = () : unit
- fun globSimp(nOpt, b) =
=      Reg.globallySimplify(nOpt, b, Reg.obviousSubset);
val globSimp = fn : int option * bool -> reg -> bool * reg
- globSimp (NONE, false) (Reg.fromString "(00*11*)*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
% + 0(0 + 1)*1
val it = () : unit
```

Finally, here are two examples showing how using the distributive rules can make a difference:


```

- globSimp (NONE, false) (Reg.fromString "% + 0*(0 + 1)");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
% + 0*(0 + 1)
val it = () : unit
- Reg.globallySimplifyTrace
= (NONE, true, Reg.obviousSubset)
= (Reg.fromString "% + 0*(0 + 1)");
considering candidates with explanations of length 0
simplest result now: % + 0*(0 + 1)
considering candidates with explanations of length 1
considering candidates with explanations of length 2
considering candidates with explanations of length 3
considering candidates with explanations of length 4
simplest result now: % + 0*(0 + 1) transformed by distributive
rule 1 at position [2] to % + 0*0 + 0*1 transformed by structural
rule 2 at position [] to (% + 0*0) + 0*1 transformed by reduction
rule 26 at position [1] to 0* + 0*1 transformed by reduction rule
21 at position [] to 0*(% + 1)
considering candidates with explanations of length 5
considering candidates with explanations of length 6
considering candidates with explanations of length 7
considering candidates with explanations of length 8
considering candidates with explanations of length 9
considering candidates with explanations of length 10
considering candidates with explanations of length 11
search completed after considering 76 candidates with maximum size
11
% + 0*(0 + 1) transformed by distributive rule 1 at position [2]
to % + 0*0 + 0*1 transformed by structural rule 2 at position []
to (% + 0*0) + 0*1 transformed by reduction rule 26 at position
[1] to 0* + 0*1 transformed by reduction rule 21 at position [] to
0*(% + 1) is globally simplified
val it = (true,-) : bool * reg
- globSimp (NONE, false) (Reg.fromString "(0(0(0 + 1)))*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
% + 0(0(% + 0 + 1))*
val it = () : unit
- globSimp (NONE, true) (Reg.fromString "(0(0(0 + 1)))*");
val it = (true,-) : bool * reg
- Reg.output("", #2 it);
% + 0(0(% + 1))*
val it = () : unit

```

3.3.4 Notes

Although books on formal language theory usually study various regular expression equivalences, we have gone much further, giving three at least partly novel algorithms for regular expression simplification. Although many of the simplification and structural rules used in the simplification algorithms are well-known, some were invented, as was the concept of closure complexity.

3.4 Finite Automata and Labeled Paths

In this section, we: say what finite automata (FA) are, and show how they can be processed using Forlan; say what labeled paths are, and show how they can be processed using Forlan; and use the notion of labeled path to say what finite automata mean.

3.4.1 Finite Automata

A *finite automaton* (FA) M consists of:

- a finite set Q_M of symbols (we call the elements of Q_M the *states* of M);
- an element s_M of Q_M (we call s_M the *start state* of M);
- a subset A_M of Q_M (we call the elements of A_M the *accepting states* of M);
- a finite subset T_M of $\{(q, x, r) \mid q, r \in Q_M \text{ and } x \in \mathbf{Str}\}$ (we call the elements of T_M the *transitions* of M , and we often write (q, x, r) as

$$q \xrightarrow{x} r$$

or $q, x \rightarrow r$).

We order transitions first by their left-hand sides, then by their middles, and then by their right-hand sides, using our total orderings on symbols and strings. This gives us a total ordering on transitions.

We often abbreviate Q_M , s_M , A_M and T_M to Q , s , A and T , when it's clear which FA we are working with. Whenever possible, we will use the mathematical variables p , q and r to name states. We write **FA** for the set of all finite automata, which is a countably infinite set.

As an example, we can define an FA M as follows:

- $Q_M = \{A, B, C\}$;
- $s_M = A$;

- $A_M = \{A, C\}$; and
- $T_M = \{(A, 1, A), (B, 11, B), (C, 111, C), (A, 0, B), (A, 2, B), (A, 0, C), (A, 2, C), (B, 0, C), (B, 2, C)\}$.

Finite automata are *nondeterministic* machines that take strings as inputs. When a machine is run on a given input, it begins in its start state.

If, after some number of steps, the machine is in state p , the machine's remaining input begins with x , and one of the machine's transitions is $p, x \rightarrow q$, then the machine *may* read x from its input and switch to state q . If $p, y \rightarrow r$ is also a transition, and the remaining input begins with y , then consuming y and switching to state r will also be possible, etc. The case when $x = \epsilon$, i.e., when we have a ϵ -transition, is interesting: a state switch can happen without reading anything.

If *at least one* execution sequence consumes all of the machine's input and takes it to one of its accepting states, then we say that the input is *accepted* by the machine; otherwise, we say that the input is *rejected*. The meaning of a machine is the language consisting of all strings that it accepts.

Here is how our example FA M can be expressed in Forlan's syntax:

```
{states}
A, B, C
{start state}
A
{accepting states}
A, C
{transitions}
A, 1 -> A; B, 11 -> B; C, 111 -> C;
A, 0 -> B; A, 2 -> B;
A, 0 -> C; A, 2 -> C;
B, 0 -> C; B, 2 -> C
```

Since whitespace characters are ignored by Forlan's input routines, the preceding description of M could have been formatted in many other ways. States are separated by commas, and transitions are separated by semicolons. The order of states and transitions is irrelevant.

Transitions that only differ in their right-hand states can be merged into single transition families. E.g., we can merge

```
A, 0 -> B
```

and

```
A, 0 -> C
```

into the transition family

```
A, 0 -> B | C
```

The Forlan module `FA` defines an abstract type `fa` (in the top-level environment) of finite automata, as well as a large number of functions and constants for processing FAs, including:

```
val input  : string -> fa
val output : string * fa -> unit
```


Remember that it's possible to read input from a file, and to write output to a file. During printing, Forlan merges transitions into transition families whenever possible.

Suppose that our example FA is in the file `3.4-fa`. We can input this FA into Forlan, and then output it to the standard output, as follows:

```
- val fa = FA.input "3.4-fa";
val fa = - : fa
- FA.output("", fa);
{states} A, B, C {start state} A {accepting states} A, C
{transitions}
A, 0 -> B | C; A, 1 -> A; A, 2 -> B | C; B, 0 -> C; B, 2 -> C;
B, 11 -> B; C, 111 -> C
val it = () : unit
```

We also make use of graphical notation for finite automata. Each of the states of a machine is circled, and its accepting states are double-circled. The machine's start state is pointed to by an arrow coming from “Start”, and each transition $p, x \rightarrow q$ is drawn as an arrow from state p to state q that is labeled by the string x . Multiple labeled arrows from one state to another can be abbreviated to a single arrow, whose label consists of the comma-separated list of the labels of the original arrows.

Here is how our FA M can be described graphically:



The Java program `JForlan`, can be used to view and edit finite automata. It can be invoked directly, or run via Forlan. See the Forlan website for more information.

We define a function $\mathbf{alphabet} \in \mathbf{FA} \rightarrow \mathbf{Alp}$ by: for all $M \in \mathbf{FA}$, $\mathbf{alphabet} M$ is $\{a \in \mathbf{Sym} \mid \text{there are } q, x, r \text{ such that } q, x \rightarrow r \in T_M \text{ and } a \in \mathbf{alphabet} x\}$. I.e., $\mathbf{alphabet} M$ is all of the symbols appearing in the strings of M 's transitions. We say that $\mathbf{alphabet} M$ is *the alphabet of M* . For example, the alphabet of our example FA M is $\{0, 1, 2\}$.

We say that an FA M is a *sub-FA* of an FA N iff:

- $Q_M \subseteq Q_N$;
- $s_M = s_N$;
- $A_M \subseteq A_N$; and
- $T_M \subseteq T_N$.

Thus $M = N$ iff M is a sub-FA of N and N is a sub-FA of M .

The Forlan module FA contains the functions

```
val equal      : fa * fa -> bool
val numStates  : fa -> int
val numTransitions : fa -> int
val alphabet   : fa -> sym set
val sub        : fa * fa -> bool
```

The function `equal` tests whether two FAs are equal, i.e., whether they have the same states, start states, accepting states and transitions. The functions `numStates` and `numTransitions` return the numbers of states and transitions, respectively, of an FA. The function `alphabet` returns the alphabet of an FA. And the function `sub` tests whether a first FA is a sub-FA of a second FA.

For example, we can continue our Forlan session as follows:

```
- val fa' = FA.input "";
@ {states} A, B, C
@ {start state} A
@ {accepting states} C
@ {transitions}
@ A, 0 -> B; A, 2 -> B; A, 0 -> C; A, 2 -> C;
@ B, 0 -> C; B, 2 -> C
@ .
val fa' = - : fa
- FA.equal(fa', fa);
val it = false : bool
- FA.sub(fa', fa);
val it = true : bool
- FA.sub(fa, fa');
val it = false : bool
- FA.numStates fa;
val it = 3 : int
- FA.numTransitions fa;
val it = 9 : int
- SymSet.output("", FA.alphabet fa);
0, 1, 2
val it = () : unit
```

3.4.2 Labeled Paths and FA Meaning

We will formally explain when strings are accepted by finite automata using the notion of a labeled path. A *labeled path* consists of a pair (xs, q) , where $xs \in \mathbf{List}(\mathbf{Sym} \times \mathbf{Str})$ and $q \in \mathbf{Sym}$, and the set \mathbf{LP} of labeled paths is $\mathbf{List}(\mathbf{Sym} \times \mathbf{Str}) \times \mathbf{Sym}$. Clearly, \mathbf{LP} is countably infinite. We typically write $[(q_1, x_1), (q_2, x_2), \dots, (q_n, x_n)], q_{n+1} \in \mathbf{LP}$ as:

$$q_1 \xRightarrow{x_1} q_2 \xRightarrow{x_2} \cdots q_n \xRightarrow{x_n} q_{n+1}$$

or

$$q_1, x_1 \Rightarrow q_2, x_2 \Rightarrow \cdots q_n, x_n \Rightarrow q_{n+1}.$$

This path describes a way of getting from state q_1 to state q_{n+1} in some unspecified machine, by reading the strings x_1, \dots, x_n from the machine's input. We start out in state q_1 , make use of the transition $q_1, x_1 \rightarrow q_2$ to read x_1 from the input and switch to state q_2 , etc.

Let $lp = (xs, q) \in \mathbf{LP}$. We say that:

- the *start state* of lp (**startState** lp) is the left-hand side of the first element of xs , if xs is nonempty, and is q , if xs is empty;
- the *end state* of lp (**endState** lp) is q ;
- the *length* of lp ($|lp|$) is $|xs|$; and
- the *label* of lp (**label** lp) is the result of concatenating the right-hand sides of xs (%), if xs is empty).

This defines functions **startState** $\in \mathbf{LP} \rightarrow \mathbf{Sym}$, **endState** $\in \mathbf{LP} \rightarrow \mathbf{Sym}$ and **label** $\in \mathbf{LP} \rightarrow \mathbf{Str}$. For example $A = ([], A)$ is a labeled path whose start and end states are both A , whose length is 0, and whose label is %. And

$$A \xRightarrow{0} B \xRightarrow{11} B \xRightarrow{2} C$$

is a labeled path whose start state is A , end state is C , length is 3, and label is $(0)(11)(2) = 0112$.

We can join compatible paths together. Let $\mathbf{Join} = \{(lp_1, lp_2) \in \mathbf{LP} \times \mathbf{LP} \mid \mathbf{endState} \, lp_1 = \mathbf{startState} \, lp_2\}$, and define **join** $\in \mathbf{Join} \rightarrow \mathbf{LP}$ by: for all $xs, ys \in \mathbf{List}(\mathbf{Sym} \times \mathbf{Str})$ and $q, r \in \mathbf{Sym}$, if $((xs, q), (ys, r)) \in \mathbf{Join}$, then **join** $((xs, q), (ys, r)) = (xs @ ys, r)$. E.g., if lp_1 and lp_2 are defined by

$$lp_1 = A \xRightarrow{0} B \xRightarrow{11} B, \quad \text{and} \quad lp_2 = B \xRightarrow{11} B \xRightarrow{2} C,$$

then $\mathbf{join}(lp_1, lp_2)$ is


$$A \xRightarrow{0} B \xRightarrow{11} B \xRightarrow{11} B \xRightarrow{2} C.$$

A labeled path $(xs, q) \in \mathbf{LP}$ is *valid* for an FA M iff

- for all $i \in [1 : |xs| - 1]$, $\#1(xs\ i), \#2(xs\ i) \rightarrow \#1(xs\ (i + 1))$;
- if xs is nonempty, then $\#1(xs\ |xs|), \#2(xs\ |xs|) \rightarrow q$; and
- $q \in Q_M$.

(The last of these conditions is redundant whenever xs is nonempty.)

Recall our example FA M :



The labeled path $A = ([], A)$ is valid for M , since $A \in Q_M$. And

$$A \xRightarrow{0} B \xRightarrow{11} B \xRightarrow{2} C$$

is valid for M , since $A \xrightarrow{0} B$, $B \xrightarrow{11} B$ and $B \xrightarrow{2} C$ are in T_M (and $C \in Q_M$). But the labeled path

$$A \xRightarrow{\%} A$$

is not valid for M , since $A, \% \rightarrow A \notin T_M$.

A string w is *accepted* by a finite automaton M iff there is a labeled path lp such that

- lp is valid for M ;
- the label of lp is w ;
- the start state of lp is the start state of M ; and
- the end state of lp is an accepting state of M .

For example, 0112 is accepted by M because of the labeled path

$$A \xRightarrow{0} B \xRightarrow{11} B \xRightarrow{2} C,$$

since this labeled path is valid for M , is labeled by $0112 = (0)(11)(2)$, has a start state (A) that is M 's start state, and has an end state (C) that is one of M 's accepting states.

Clearly, if w is accepted by M , then $\mathbf{alphabet} \, w \subseteq \mathbf{alphabet} \, M$. Thus $\{w \in \mathbf{Str} \mid w \text{ is accepted by } M\} \subseteq (\mathbf{alphabet} \, M)^*$, so we may define the *language accepted by a finite automaton M* ($L(M)$) to be

$$\{w \in \mathbf{Str} \mid w \text{ is accepted by } M\}.$$

Furthermore:

Proposition 3.4.1

Suppose M is a finite automaton. Then $\mathbf{alphabet}(L(M)) \subseteq \mathbf{alphabet} \, M$.

In other words, every symbol of every string that is accepted by M comes from the alphabet of M , i.e., appears in the label of one of M 's transitions.

Going back to our example, we have that

$$\begin{aligned} L(M) = & \{1\}^* \cup \\ & \{1\}^* \{0, 2\} \{11\}^* \{0, 2\} \{111\}^* \cup \\ & \{1\}^* \{0, 2\} \{111\}^*. \end{aligned}$$

For example, %, 11, 110112111 and 2111111 are accepted by M . But 21112 and 2211 are not accepted by M .

Suppose that M is a sub-FA of N . Then any labeled path that is valid for M will also be valid for N . Furthermore, we have that:

Proposition 3.4.2

If M is a sub-FA of N , then $L(M) \subseteq L(N)$.

We say that finite automata M and N are *equivalent* iff $L(M) = L(N)$. In other words, M and N are equivalent iff M and N accept the same language. We define a relation \approx on \mathbf{FA} by: $M \approx N$ iff M and N are equivalent. It is easy to see that \approx is reflexive on \mathbf{FA} , symmetric and transitive.

The Forlan module `LP` defines an abstract type `lp` (in the top-level environment) of labeled paths, as well as various functions for processing labeled paths, including:

```
val input      : string -> lp
val output     : string * lp -> unit
val equal      : lp * lp -> bool
val startState : lp -> sym
```



```

val endState    : lp -> sym
val label       : lp -> str
val length      : lp -> int
val join        : lp * lp -> lp

```

The function `equal` tests whether two labeled paths are equal. The functions `startState`, `endState`, `label` and `length` return the start state, end state, label and length, respectively, of a labeled path. And the function `join` joins two compatible paths, and issues an error message when given paths that are incompatible.

The module `FA` also defines the functions

```

val checkLP : fa -> lp -> unit
val validLP  : fa -> lp -> bool

```

for checking whether a labeled path is valid in a finite automaton. These are curried functions—functions that return functions as their results. The function `checkLP` takes in an FA M and returns a function that checks whether a labeled path lp is valid for M . When lp is not valid for M , the function explains why it isn't; otherwise, it prints nothing. And, the function `validLP` takes in an FA M and returns a function that tests whether a labeled path lp is valid for M , silently returning `true`, if it is, and silently returning `false`, otherwise.

Here are some examples of labeled path and FA processing (`fa` is still our example FA):

```

- val lp = LP.input "";
@ A, 1 => A, 0 => B, 11 => B, 2 => C, 111 => C
@ .
val lp = - : lp
- Sym.output("", LP.startState lp);
A
val it = () : unit
- Sym.output("", LP.endState lp);
C
val it = () : unit
- LP.length lp;
val it = 5 : int
- Str.output("", LP.label lp);
10112111
val it = () : unit
- val checkLP = FA.checkLP fa;
val checkLP = fn : lp -> unit
- checkLP lp;
val it = () : unit
- val lp' = LP.fromString "A";
val lp' = - : lp
- LP.length lp';
val it = 0 : int

```

```

- Str.output("", LP.label lp');
%
val it = () : unit
- checkLP lp';
val it = () : unit
- val lp'' = LP.input "";
@ A, % => A, 1 => B
@ .
val lp'' = - : lp
- checkLP lp'';
invalid transition: "A, % -> A"


uncaught exception Error
- val lp''' = LP.input "";
@ B, 2 => C, 34 => D
@ .
val lp''' = - : lp
- LP.output("", LP.join(lp'', lp'''));
A, % => A, 1 => B, 2 => C, 34 => D
val it = () : unit
- LP.output("", LP.join(lp''', lp''));
incompatible labeled paths

uncaught exception Error


```

3.4.3 Design of Finite Automata


In this subsection, we give two examples of finite automata design. First, let's find a finite automaton that accepts the set of all strings of 0's and 1's with an even number of 0's. Thus, we should be looking for an FA whose alphabet is $\{0,1\}^*$. It seems reasonable that our machine have two states: an accepting state A corresponding to the strings of 0's and 1's with an even number of zeros, and a state B corresponding to the strings of 0's and 1's with an odd number of zeros. Processing a 1 in either state should cause us to stay in that state, but processing a 0 in one of the states should cause us to switch to the other state. The above considerations lead us to the FA



For the second example, let's find an FA that accepts the language $X = \{w \in \{0,1,2\}^* \mid \text{for all substrings } x \text{ of } w, \text{ if } |x| = 2, \text{ then } x \in \{01, 12, 20\}\}$. We have that 0, 01, 012, 0120, etc., are in X , and so are 1, 12, 120, 1201, etc., and 2, 20, 201, 2012, etc. On the other hand, no string containing 00, 02, 11, 10, 22 or 20 is in X . The above observations suggest that part of our machine should look like:



But how should the machine get started? The simplest approach is to make use of ϵ -transitions from the start state, giving us the FA



Exercise 3.4.3

Let $X = \{w \in \{0,1\}^* \mid 010 \text{ is not a substring of } w\}$. Find a finite automaton M such that $L(M) = X$.

Exercise 3.4.4

Let

$A = \{001, 011, 101, 111\}$, and

$B = \{w \in \{0,1\}^* \mid \text{for all } x, y \in \{0,1\}^*, \text{ if } w = x0y, \text{ then there is a } z \in A \text{ such that } z \text{ is a prefix of } y\}$.

Find a finite automaton M such that $L(M) = B$. Hint: see the second example of Subsection 3.2.2.

3.4.4 Notes

Finite automata are normally defined via transition functions, δ , which is simple to do for deterministic finite automata, but increasingly complicated as one adds degrees of nontermininism. Furthermore, this approach means that a deterministic finite automaton (DFA) is not a nondeterministic finite automaton (NFA), and that an NFA is not an ϵ -NFA, because the transition function of a DFA is not one for an NFA, and the transition function for an NFA is not one for an ϵ -NFA. And formalizing our FAs, whose transition labels can be strings of length greater than one, is very messy if done via transition functions, which probably accounts for why such machines are not normally considered. Furthermore, in the standard approach, to say when strings are accepted by finite

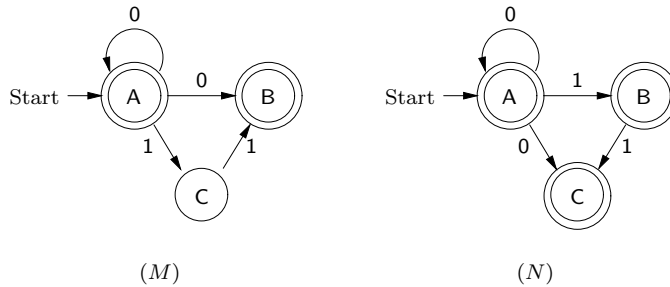
automata, one must first extend the different kinds of transition functions to work on strings.

In contrast, our approach is very simple. Instead of transition functions, we work with finite sets of transitions, enabling us to define the deterministic finite automata, nondeterministic finite automata, and nondeterministic finite automata with ϵ -moves (which we call empty-string finite automata) as restrictions on finite automata. Furthermore, using labeled paths to say when—and how—strings are accepted by finite automata is simple, natural and diagrammatic. It's the analogue of using parse trees to say when—and how—strings are generated by grammars.

3.5 Isomorphism of Finite Automata

3.5.1 Definition and Algorithm

Let M and N be the finite automata



How are M and N related? Although they are not equal, they do have the same “structure”, in that M can be turned into N by replacing A, B and C by A, C and B, respectively. When FAs have the same structure, we will say they are “isomorphic”.

In order to say more formally what it means for two FAs to be isomorphic, we define the notion of an isomorphism from one FA to another. An *isomorphism* h from an FA M to an FA N is a bijection from Q_M to Q_N such that:

- $h s_M = s_N$;
- $\{ h q \mid q \in A_M \} = A_N$; and
- $\{ (h q), x \rightarrow (h r) \mid q, x \rightarrow r \in T_M \} = T_N$.

We define a relation **iso** on **FA** by: $M \text{ iso } N$ iff there is an isomorphism from M to N . We say that M and N are *isomorphic* iff $M \text{ iso } N$.

Consider our example FAs M and N , and let h be the function

$$\{(A, A), (B, C), (C, B)\}.$$

Then it is easy to check that h is an isomorphism from M to N . Hence $M \text{ iso } N$.

Proposition 3.5.1

The relation **iso** is reflexive on **FA**, symmetric and transitive.

Proof. If M is an FA, then \mathbf{id}_M is an isomorphism from M to M .

If M, N are FAs, and h is an isomorphism from M to N , then h^{-1} is an isomorphism from N to M .

If M_1, M_2, M_3 are FAs, f is an isomorphism from M_1 to M_2 , and g is an isomorphism from M_2 to M_3 , then $g \circ f$ is an isomorphism from M_1 to M_3 . \square

Next, we see that, if M and N are isomorphic, then every string accepted by M is also accepted by N .

Proposition 3.5.2

Suppose M and N are isomorphic FAs. Then $L(M) \subseteq L(N)$.

Proof. Let h be an isomorphism from M to N . Suppose $w \in L(M)$. Then, there is a labeled path

$$lp = q_1 \xRightarrow{x_1} q_2 \xRightarrow{x_2} \cdots q_n \xRightarrow{x_n} q_{n+1},$$

such that $w = x_1 x_2 \cdots x_n$, lp is valid for M , $q_1 = s_M$ and $q_{n+1} \in A_M$. Let

$$lp' = h q_1 \xRightarrow{x_1} h q_2 \xRightarrow{x_2} \cdots h q_n \xRightarrow{x_n} h q_{n+1}.$$

Then the label of lp' is w , lp' is valid for N , $h q_1 = h s_M = s_N$ and $h q_{n+1} \in A_N$, showing that $w \in L(N)$. \square

A consequence of the two preceding propositions is that isomorphic FAs are equivalent. Of course, the converse is not true, in general, since there are many FAs that accept the same language and yet don't have the same structure.

Proposition 3.5.3

Suppose M and N are isomorphic FAs. Then $M \approx N$.

Proof. Since $M \mathbf{iso} N$, we have that $N \mathbf{iso} M$, by Proposition 3.5.1. Thus, by Proposition 3.5.2, we have that $L(M) \subseteq L(N) \subseteq L(M)$. Hence $L(M) = L(N)$, i.e., $M \approx N$. \square

The function **renameStates** takes in a pair (M, f) , where $M \in \mathbf{FA}$ and f is a bijection from Q_M to some set of symbols, and returns the **FA** produced from M by renaming M 's states using the bijection f .

Proposition 3.5.4

Suppose M is an FA and f is a bijection from Q_M to some set of symbols. Then **renameStates** $(M, f) \mathbf{iso} M$.

The following function is a special case of **renameStates**. The function **renameStatesCanonically** $\in \mathbf{FA} \rightarrow \mathbf{FA}$ renames the states of an FA M to:

- A, B, etc., when the automaton has no more than 26 states (the smallest state of M will be renamed to A, the next smallest one to B, etc.); or
- $\langle 1 \rangle$, $\langle 2 \rangle$, etc., otherwise.

Of course, the resulting automaton will always be isomorphic to the original one.

Next, we consider an algorithm that finds an isomorphism from an FA M to an FA N , if one exists, and that indicates that no such isomorphism exists, otherwise.

Our algorithm is based on the following lemma.

Lemma 3.5.5

Suppose that h is a bijection from Q_M to Q_N . Then

$$\{ h q \xrightarrow{x} h r \mid q \xrightarrow{x} r \in T_M \} = T_N$$

iff, for all $(q, r) \in h$ and $x \in \mathbf{Str}$, there is a subset of h that is a bijection from

$$\{ p \in Q_M \mid q \xrightarrow{x} p \in T_M \}$$

to

$$\{ p \in Q_N \mid r \xrightarrow{x} p \in T_N \}.$$

If any of the following conditions are true, then the algorithm reports that there is no isomorphism from M to N :

- $|Q_M| \neq |Q_N|$;
- $|A_M| \neq |A_N|$;
- $|T_M| \neq |T_N|$;
- $s_M \in A_M$, but $s_N \notin A_N$; and
- $s_N \in A_N$, but $s_M \notin A_M$.

Otherwise, it calls its main function, **findIso**, which is defined by well-founded recursion.

The function **findIso** is called with an argument $(f, [C_1, \dots, C_n])$, where:

- f is a bijection from a subset of Q_M to a subset of Q_N , and

- the C_i are *constraints* of the form (X, Y) , where $X \subseteq Q_M$, $Y \subseteq Q_N$ and $|X| = |Y|$.

It returns an element of **Option** X , where X is the set of bijections from Q_M to Q_N . **none** is returned to indicate failure, and **some** h is returned when it has produced the bijection h . We say that a bijection *satisfies* a constraint (X, Y) iff it has a subset that is a bijection from X to Y .

We say that the *weight* of a constraint (X, Y) is $3^{|X|}$. Thus, we have the following facts:

- If (X, Y) is a constraint, then its weight is at least $3^0 = 1$.
- If $(\{p\} \cup X, \{q\} \cup Y)$ is a constraint, $p \notin X$, $q \notin Y$ and $|X| \geq 1$, then the weight of $(\{p\} \cup X, \{q\} \cup Y)$ is $3^{1+|X|} = 3 \cdot 3^{|X|}$, the weight of $(\{p\}, \{q\})$ is $3^1 = 3$, and the weight of (X, Y) is $3^{|X|}$. Because $|X| \geq 1$, it follows that the sum of the weights of $(\{p\}, \{q\})$ and (X, Y) ($3 + 3^{|X|}$) is strictly less than the weight of $(\{p\} \cup X, \{q\} \cup Y)$.

Each argument to a recursive call of **findIso** will be strictly smaller than the argument to the original call in the well-founded *termination relation* in which argument $(f, [C_1, \dots, C_n])$ is less than argument $(f', [C'_1, \dots, C'_m])$ iff either:

- $|f| > |f'|$ (remember that $|f| \leq |Q_M| = |Q_N|$); or
- $|f| = |f'|$ but the sum of the weights of the constraints C_1, \dots, C_n is strictly less than the sum of the weights of the constraints C'_1, \dots, C'_m .

When **findIso** is called with argument $(f, [C_1, \dots, C_n])$, the following property, which we call (\dagger) , will hold: for all bijections h from a subset of Q_M to a subset of Q_N , if $h \supseteq f$ and h satisfies all of the C_i 's, then:

- h is a bijection from Q_M to Q_N ;
- $h s_M = s_N$;
- $\{h q \mid q \in A_M\} = A_N$; and
- for all $(q, r) \in f$ and $x \in \mathbf{Str}$, there is a subset of h that is a bijection from $\{p \in Q_M \mid q, x \rightarrow p \in T_M\}$ to $\{p \in Q_N \mid r, x \rightarrow p \in T_N\}$.

Thus, if **findIso** is called with a bijection f and an empty list of constraints, it will follow, by Lemma 3.5.5, that f is an isomorphism from M to N .

Initially, the algorithm calls **findIso** with the *initial argument*:

$$(\emptyset, [(\{s_M\}, \{s_N\}), (U, V), (X, Y)]),$$

where $U = A_M - \{s_M\}$, $V = A_N - \{s_N\}$, $X = (Q_M - A_M) - \{s_M\}$ and $Y = (Q_N - A_N) - \{s_N\}$. Clearly the initial argument satisfies (\dagger) .

If **findIso** is called with argument $(f, [])$, then it returns **some** f .

Otherwise, if **findIso** is called with argument $(f, [(\emptyset, \emptyset), C_2, \dots, C_n])$, then it calls itself recursively with argument $(f, [C_2, \dots, C_n])$. (The size of the bijection has been preserved, but the sum of the weights of the constraints has gone down by one.)

Otherwise, if **findIso** is called with argument $(f, [(\{q\}, \{r\}), C_2, \dots, C_n])$, then it proceeds as follows:

- If $(q, r) \in f$, then it calls itself recursively with argument $(f, [C_2, \dots, C_n])$ and returns what the recursive call returns. (The size of the bijection has been preserved, but the sum of the weights of the constraints has gone down by three.)
- Otherwise, if $q \in \text{domain } f$ or $r \in \text{range } f$, then **findIso** returns **none**.
- Otherwise, it works its way through the strings appearing in the transitions of M and N , forming a list of new constraints, C'_1, \dots, C'_m . Given such a string, x , it lets $A_x = \{p \in Q_M \mid q, x \rightarrow p \in T_M\}$ and $B_x = \{p \in Q_N \mid r, x \rightarrow p \in T_N\}$. If $|A_x| \neq |B_x|$, then it returns **none**. Otherwise, it adds the constraint (A_x, B_x) to our list of new constraints. When all such strings have been exhausted, it calls itself recursively with argument $(f \cup \{(q, r)\}, [C'_1, \dots, C'_m, C_2, \dots, C_n])$ and returns what this recursive call returns. (The size of the bijection has been increased by one.)

Otherwise, **findIso** has been called with argument $(f, [(A, A'), C_2, \dots, C_n])$, where $|A| > 1$, and it proceeds as follows. It picks the smallest symbol $q \in A$, and lets $B = A - \{q\}$. Then, it works its way through the elements of A' . Given $r \in A'$, it lets $B' = A' - \{r\}$. Then, it tries calling itself recursively with argument $(f, [(\{q\}, \{r\}), (B, B'), C_2, \dots, C_n])$. (The size of the bijection has been preserved, but the sum of the sizes of the weights of the constraints has gone down by $2 \cdot 3^{|B|} - 3 \geq 3$.) If this call returns a result of form **some** h , then it returns this to its caller. But if the call returns **none**, it tries the next element of A' . If it exhausts the elements of A' , then it returns **none**.

Lemma 3.5.6

If **findIso** is called with an argument $(f, [C_1, \dots, C_n])$ satisfying property (\dagger) , then it returns **none**, if there is no isomorphism from M to N that is a superset of f and satisfies the constraints C_1, \dots, C_n , and returns **some** h where h is such an isomorphism, otherwise.

Proof. By well-founded induction on our termination relation. I.e., when proving the result for $(f, [C_1, \dots, C_n])$, we may assume that the result holds for all arguments $(f', [C'_1, \dots, C'_m])$ that are strictly smaller in our termination ordering. \square

Theorem 3.5.7

If **findIso** is called with its initial argument, then it returns **none**, if there is no isomorphism from M to N , and returns **some** h where h is an isomorphism from M to N , otherwise.

Proof. Follows easily from Lemma 3.5.6. \square

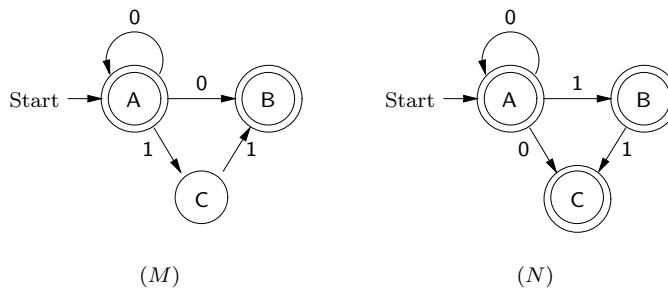
3.5.2 Isomorphism Finding/Checking in Forlan

The Forlan module FA also defines the functions

```
val isomorphism          : fa * fa * sym_rel -> bool
val findIsomorphism     : fa * fa -> sym_rel
val isomorphic          : fa * fa -> bool
val renameStates        : fa * sym_rel -> fa
val renameStatesCanonically : fa -> fa
```

The function **isomorphism** checks whether a relation on symbols is an isomorphism from one FA to another. The function **findIsomorphism** tries to find an isomorphism from one FA to another; it issues an error message if there isn't one. The function **isomorphic** checks whether two FAs are isomorphic. The function **renameStates** issues an error message if the supplied relation isn't a bijection from the set of states of the supplied FA to some set; otherwise, it returns the result of **renameStates**. And the function **renameStatesCanonically** acts like **renameStatesCanonically**.

Suppose **fa1** and **fa2** have been bound to our example our example finite automata M and N :



Then, here are some example uses of the above functions:

```
- val rel = FA.findIsomorphism(fa1, fa2);
val rel = - : sym_rel
- SymRel.output("", rel);
(A, A), (B, C), (C, B)
val it = () : unit
- FA.isomorphism(fa1, fa2, rel);
val it = true : bool
- FA.isomorphic(fa1, fa2);
```

```

val it = true : bool
- val rel' = FA.findIsomorphism(fa1, fa1);
val rel' = - : sym_rel
- SymRel.output("", rel');
(A, A), (B, B), (C, C)
val it = () : unit
- FA.isomorphism(fa1, fa1, rel');
val it = true : bool
- FA.isomorphism(fa1, fa2, rel');
val it = false : bool
- val rel'' = SymRel.input "";
@ (A, 2), (B, 1), (C, 0)
@ .
val rel'' = - : sym_rel
- val fa3 = FA.renameStates(fa1, rel'');
val fa3 = - : fa
- FA.output("", fa3);
{states} 0, 1, 2 {start state} 2 {accepting states} 0, 1, 2
{transitions} 0, 1 -> 1; 2, 0 -> 1 | 2; 2, 1 -> 0
val it = () : unit
- val fa4 = FA.renameStatesCanonically fa3;
val fa4 = - : fa
- FA.output("", fa4);
{states} A, B, C {start state} C {accepting states} A, B, C
{transitions} A, 1 -> B; C, 0 -> B | C; C, 1 -> A
val it = () : unit
- FA.equal(fa4, fa1);
val it = false : bool
- FA.isomorphic(fa4, fa1);
val it = true : bool

```

3.5.3 Notes

Books on formal language theory rarely formalize the isomorphism of finite automata, although most note or prove that the minimization of deterministic finite automata yields a result that is unique up to the renaming of states. Our algorithm for trying to find an isomorphism between finite automata will be unsurprising to those familiar with graph algorithms.

3.6 Checking Acceptance and Finding Accepting Paths

In this section we study algorithms for checking whether a string is accepted by a finite automaton, and for finding a labeled path that explains why a string is accepted by a finite automaton.


3.6.1 Processing a String from a Set of States

Suppose M is a finite automaton. We define a function $\Delta_M \in \mathcal{P} Q_M \times \mathbf{Str} \rightarrow \mathcal{P} Q_M$ by: $\Delta_M(P, w)$ is the set of all $r \in Q_M$ such that there is an $lp \in \mathbf{LP}$ such that

- w is the label of lp ;
- lp is valid for M ;
- the start state of lp is in P ; and
- r is the end state of lp .

In other words, $\Delta_M(P, w)$ consists of all of the states that can be reached from elements of P by labeled paths that are labeled by w and valid for M . When the FA M is clear from the context, we sometimes abbreviate Δ_M to Δ .

Suppose M is the finite automaton



Then, $\Delta_M(\{A\}, 12111111) = \{B, C\}$, since

$$A \xRightarrow{1} A \xRightarrow{2} B \xRightarrow{11} B \xRightarrow{11} B \xRightarrow{11} B \quad \text{and} \quad A \xRightarrow{1} A \xRightarrow{2} C \xRightarrow{111} C \xRightarrow{111} C$$

are all of the labeled paths that are labeled by 12111111, valid in M and whose start states are A. Furthermore, $\Delta_M(\{A, B, C\}, 11) = \{A, B\}$, since

$$A \xRightarrow{1} A \xRightarrow{1} A \quad \text{and} \quad B \xRightarrow{11} B$$

are all of the labeled paths that are labeled by 11 and valid in M .

Suppose M is a finite automaton, $P \subseteq Q_M$ and $w \in \mathbf{Str}$. We can calculate $\Delta_M(P, w)$ as follows.

Let S be the set of all suffixes of w . Given $y \in S$, we write $\mathbf{pre} y$ for the unique x such that $w = xy$.

First, we generate the least subset X of $Q_M \times S$ such that:

- (1) for all $p \in P$, $(p, w) \in X$; and
- (2) for all $q, r \in Q_M$ and $x, y \in \mathbf{Str}$, if $(q, xy) \in X$ and $q, x \rightarrow r \in T_M$, then $(r, y) \in X$.

We start by using rule (1), adding (p, w) to X , whenever $p \in P$. Then X (and any superset of X) will satisfy property (1). Then, rule (2) is used repeatedly to add more pairs to X . Since $Q_M \times S$ is a finite set, eventually X will satisfy property (2).

If M is our example finite automaton, then here are the elements of X , when $P = \{A\}$ and $w = 2111$:

- $(A, 2111)$;
- $(B, 111)$, because of $(A, 2111)$ and the transition $A, 2 \rightarrow B$;
- $(C, 111)$, because of $(A, 2111)$ and the transition $A, 2 \rightarrow C$ (now, we're done with $(A, 2111)$);
- $(B, 1)$, because of $(B, 111)$ and the transition $B, 11 \rightarrow B$ (now, we're done with $(B, 111)$);
- $(C, \%)$, because of $(C, 111)$ and the transition $C, 111 \rightarrow C$ (now, we're done with $(C, 111)$); and
- nothing can be added using $(B, 1)$ and $(C, \%)$, and so we've found all the elements of X .

The following lemma explains when pairs show up in X .

Lemma 3.6.1

For all $q \in Q_M$ and $y \in S$,

$$(q, y) \in X \quad \text{iff} \quad q \in \Delta_M(P, \mathbf{pre} y).$$

Proof. The “only if” (left-to-right) direction is by induction on X : we show that, for all $(q, y) \in X$, $q \in \Delta_M(P, \mathbf{pre} y)$.

- Suppose $p \in P$. Then $p \in \Delta_M(P, \%)$. But $\mathbf{pre} w = \%$, so that $p \in \Delta_M(P, \mathbf{pre} w)$.
- Suppose $q, r \in Q_M$, $x, y \in \mathbf{Str}$, $(q, xy) \in X$ and $(q, x, r) \in T_M$. Assume the inductive hypothesis: $q \in \Delta_M(P, \mathbf{pre}(xy))$. Thus there is an $lp \in \mathbf{LP}$ such that $\mathbf{pre}(xy)$ is the label of lp , lp is valid for M , the start state of lp is in P , and q is the end state of lp . Let $lp' \in \mathbf{LP}$ be the result of adding the step $q, x \Rightarrow r$ at the end of lp . Thus $\mathbf{pre} y$ is the label of lp' , lp' is valid for M , the start state of lp' is in P , and r is the end state of lp' , showing that $r \in \Delta_M(P, \mathbf{pre} y)$.

For the ‘if’ (right-to-left) direction, we have that there is a labeled path

$$q_1 \xRightarrow{x_1} q_2 \xRightarrow{x_2} \cdots q_{n-1} \xRightarrow{x_{n-1}} q_n,$$

that is valid for M and where $\mathbf{pre} y = x_1x_2 \cdots x_{n-1}$, $q_1 \in P$ and $q_n = q$. Since $q_1 \in P$ and $w = (\mathbf{pre} y)y = x_1x_2 \cdots x_{n-1}y$, we have that $(q_1, x_1x_2 \cdots x_{n-1}y) = (q_1, w) \in X$, by (1). But $(q_1, x_1, q_2) \in T_M$, and thus $(q_2, x_2 \cdots x_{n-1}y) \in X$, by (2). Continuing on in this way (we could do this by mathematical induction), we finally get that $(q, y) = (q_n, y) \in X$. \square

Lemma 3.6.2

For all $q \in Q_M$, $(q, \%) \in X$ iff $q \in \Delta_M(P, w)$.

Proof. Suppose $(q, \%) \in X$. Lemma 3.6.1 tells us that $q \in \Delta_M(P, \mathbf{pre} \%)$. But $\mathbf{pre} \% = w$, and thus $q \in \Delta_M(P, w)$.

Suppose $q \in \Delta_M(P, w)$. Since $w = \mathbf{pre} \%$, we have that $q \in \Delta_M(P, \mathbf{pre} \%)$. Lemma 3.6.1 tells us that $(q, \%) \in X$. \square

By Lemma 3.6.2, we have that

$$\Delta_M(P, w) = \{q \in Q_M \mid (q, \%) \in X\}.$$

Thus, we return the set of all states q that are paired with $\%$ in X .

3.6.2 Checking String Acceptance and Finding Accepting Paths

Proposition 3.6.3

Suppose M is a finite automaton. Then

$$L(M) = \{w \in \mathbf{Str} \mid \Delta_M(\{s_M\}, w) \cap A_M \neq \emptyset\}.$$

Proof. Suppose $w \in L(M)$. Then w is the label of a labeled path lp such that lp is valid in M , the start state of lp is s_M and the end state of lp is in A_M . Let q be the end state of lp . Thus $q \in \Delta_M(\{s_M\}, w)$ and $q \in A_M$, showing that $\Delta_M(\{s_M\}, w) \cap A_M \neq \emptyset$.

Suppose $\Delta_M(\{s_M\}, w) \cap A_M \neq \emptyset$, so that there is a q such that $q \in \Delta_M(\{s_M\}, w)$ and $q \in A_M$. Thus w is the label of a labeled path lp such that lp is valid in M , the start state of lp is s_M , and the end state of lp is $q \in A_M$. Thus $w \in L(M)$. \square

According to Proposition 3.6.3, to check if a string w is accepted by a finite automaton M , we simply use our algorithm to generate $\Delta_M(\{s_M\}, w)$, and then check if this set contains at least one accepting state.

Given a finite automaton M , subsets P, R of Q_M and a string w , how do we search for a labeled path that is labeled by w , valid in M , starts from an element of P , and ends with an element of R ? What we need to do is associate with each pair

$$(q, y)$$

of the set X that we generate when computing $\Delta_M(P, w)$ a labeled path lp such that lp is labeled by $\mathbf{pre}(y)$, lp is valid in M , the start state of lp is an element of P , and the end state of lp is q . If we process the elements of X in a breadth-first (rather than depth-first) manner, this will ensure that these labeled paths are as short as possible. As we generate the elements of X , we look for a pair of the form $(q, \%)$, where $q \in R$. Our answer will then be the labeled path associated with this pair.

The Forlan module **FA** also contains the following functions for processing strings and checking string acceptance:

```
val processStr      : fa -> sym set * str -> sym set
val accepted       : fa -> str -> bool
```


The function **processStr** takes in a finite automaton M , and returns a function that takes in a pair (P, w) and returns $\Delta_M(P, w)$. And, the function **accepted** takes in a finite automaton M , and returns a function that checks whether a string x is accepted by M .

The Forlan module **FA** also contains the following functions for finding labeled paths:

```
val findLP          : fa -> sym set * str * sym set -> lp
val findAcceptingLP : fa -> str -> lp
```

The function **findLP** takes in a finite automaton M , and returns a function that takes in a triple (P, w, R) and tries to find a labeled path lp that is labeled by w , valid for M , starts out with an element of P , and ends up at an element of R . It issues an error message when there is no such labeled path. The function **findAcceptingLP** takes in a finite automaton M , and returns a function that looks for a labeled path lp that explains why a string w is accepted by M . It issues an error message when there is no such labeled path. The labeled paths returned by these functions are always of minimal length.

Suppose **fa** is the finite automaton



We begin by applying our five functions to **fa**, and giving names to the resulting functions:

```
- val processStr = FA.processStr fa;
val processStr = fn : sym set * str -> sym set
- val accepted = FA.accepted fa;
val accepted = fn : str -> bool
```

```

- val findLP = FA.findLP fa;
val findLP = fn : sym set * str * sym set -> lp
- val findAcceptingLP = FA.findAcceptingLP fa;
val findAcceptingLP = fn : str -> lp

```

Next, we'll define a set of states and a string to use later:

```

- val bs = SymSet.input "";
@ A, B, C
@ .
val bs = - : sym set
- val x = Str.input "";
@ 11
@ .
val x = [-,-] : str

```

Here are some example uses of our functions:

```

- SymSet.output("", processStr(bs, x));
A, B
val it = () : unit
- accepted(Str.input "");
@ 12111111
@ .
val it = true : bool
- accepted(Str.input "");
@ 1211
@ .
val it = false : bool
- LP.output("", findLP(bs, x, bs));
B, 11 => B
val it = () : unit
- LP.output("", findAcceptingLP(Str.input ""));
@ 12111111
@ .
A, 1 => A, 2 => C, 111 => C, 111 => C
val it = () : unit
- LP.output("", findAcceptingLP(Str.input ""));
@ 222
@ .
no such labeled path exists

uncaught exception Error

```


3.6.3 Notes

The material in this section is original. Our definition of the meaning of FAs via labeled paths allows us not simply to test *whether* an FA accepts a string w , but to ask for evidence—in the form of a labeled path—for *why* FA accepts w .

3.7 Simplification of Finite Automata

In this section, we: say what it means for a finite automaton to be simplified; study an algorithm for simplifying finite automata; and see how finite automata can be simplified in Forlan.

Suppose M is the finite automaton



M is odd for two distinct reasons. First, there are no valid labeled paths from the start state to D and E, and so these states are redundant. Second, there are no valid labeled paths from C to an accepting state, and so it is also redundant. We will say that C is not “live” (C is “dead”), and that D and E are not “reachable”.

Suppose M is a finite automaton. We say that a state $q \in Q_M$ is:

- *reachable in M* iff there is a labeled path lp such that lp is valid for M , the start state of lp is s_M , and the end state of lp is q ;
- *live in M* iff there is a labeled path lp such that lp is valid for M , the start state of lp is q , and the end state of lp is in A_M ;
- *dead in M* iff q is not live in M ; and
- *useful in M* iff q is both reachable and live in M .

Let M be our example finite automaton. The reachable states of M are: A, B and C. The live states of M are: A, B, D and E. And, the useful states of M are: A and B.

There is a simple algorithm for generating the set of reachable states of a finite automaton M . We generate the least subset X of Q_M such that:

- $s_M \in X$; and
- for all $q, r \in Q_M$ and $x \in \mathbf{Str}$, if $q \in X$ and $(q, x, r) \in T_M$, then $r \in X$.

The start state of M is added to X , since s_M is always reachable, by the zero-length labeled path s_M . Then, if q is reachable, and (q, x, r) is a transition of M , then r is clearly reachable. Thus all of the elements of X are indeed reachable. And, it's not hard to show that every reachable state will be added to X .


Similarly, there is a simple algorithm for generating the set of live states of a finite automaton M . We generate the least subset Y of Q_M such that:

- $A_M \subseteq Y$; and
- for all $q, r \in Q_M$ and $x \in \mathbf{Str}$, if $r \in Y$ and $(q, x, r) \in T_M$, then $q \in Y$.

This time it's the accepting states of M that are initially added to our set, since each accepting state is trivially live. Then, if r is live, and (q, x, r) is a transition of M , then q is clearly live.

Thus, we can generate the set of useful states of an FA by generating the set of reachable states, generating the set of live states, and intersecting those sets of states.

Now, suppose N is the FA



Here, the transitions $(A, 0, B)$ and $(A, 1, B)$ are redundant, in the sense that if N' is the result of removing these transitions from N , we still have that $B \in \Delta_{N'}(\{A\}, 0)$ and $B \in \Delta_{N'}(\{A\}, 1)$.


Given an FA M and a finite subset U of $\{(q, x, r) \mid q, r \in Q_M \text{ and } x \in \mathbf{Str}\}$, we write M/U for the FA that is identical to M except that its set of transitions is U . If M is an FA and $(p, x, q) \in T_M$, we say that:

- (p, x, q) is *redundant* in M iff $q \in \Delta_N(\{p\}, x)$, where $N = M/(T_M - \{(p, x, q)\})$; and
- (p, x, q) is *irredundant* in M iff (p, x, q) is not redundant in M .

We say that a finite automaton M is *simplified* iff either

- every state of M is useful, and every transition of M is irredundant; or
- $|Q_M| = 1$ and $A_M = T_M = \emptyset$.

Thus the FA



is simplified, even though its start state is not live, and is thus not useful.

Proposition 3.7.1

If M is a simplified finite automaton, then $\mathbf{alphabet} M = \mathbf{alphabet}(L(M))$.

Proof. We always have that $\mathbf{alphabet}(L(M)) \subseteq \mathbf{alphabet} M$. But, because M is simplified, we also have that $\mathbf{alphabet} M \subseteq \mathbf{alphabet}(L(M))$, i.e., that every symbol appearing in a string of one of M 's transitions also appears in one of the strings accepted by M . This is because given any transition (p, x, q) of a simplified finite automaton, p is reachable and q is live. \square

To give our simplification algorithm for finite automata, we need an auxiliary function for removing redundant transitions from an FA. Given an FA M , $p, q \in Q_M$ and $x \in \mathbf{Str}$, we say that (p, x, q) is *implicit in M* iff $q \in \Delta_M(\{p\}, x)$.

Given an FA M , we define a function $\mathbf{remRedun}_M \in \mathcal{P} T_M \times \mathcal{P} T_M \rightarrow \mathcal{P} T_M$ (we often drop the M when it's clear from the context) by well-founded recursion on the size of its second argument. For $U, V \subseteq T_M$, $\mathbf{remRedun}(U, V)$ proceeds as follows:

- If $V = \emptyset$, then it returns U .
- Otherwise, let v be the greatest element of V , and $V' = V - \{v\}$. If v is implicit in $M/(U \cup V')$, then $\mathbf{remRedun}$ returns the result of evaluating $\mathbf{remRedun}(U, V')$. Otherwise, it returns the result of evaluating $\mathbf{remRedun}(U \cup \{v\}, V')$.

In general, there are multiple—incompatible—ways of removing redundant transitions from an FA. $\mathbf{remRedun}$ is defined so as to favor removing transitions that are larger in our total ordering on transitions.

Proposition 3.7.2

Suppose M is a finite automaton. For all $U, V \subseteq T_M$, if all the elements of U are irredundant in $M/(U \cup V)$, and, for all $p, q \in Q_M$ and $x \in \mathbf{Str}$, (p, x, q) is implicit in M iff (p, x, q) is implicit in $M/(U \cup V)$, then all the elements of $\mathbf{remRedun}(U, V)$ are irredundant in $M/\mathbf{remRedun}(U, V)$, and, for all $p, q \in Q_M$ and $x \in \mathbf{Str}$, (p, x, q) is implicit in M iff (p, x, q) is implicit in $M/\mathbf{remRedun}(U, V)$.

Proof. By well-founded induction on the size of the second argument to $\mathbf{remRedun}$. \square

Now we can give an algorithm for simplifying finite automata. We define a function $\mathbf{simplify} \in \mathbf{FA} \rightarrow \mathbf{FA}$ by: $\mathbf{simplify} M$ is the finite automaton N produced by the following process.

- First, the useful states of M are determined.
- If s_M is not useful in M , the N is defined by:
 - $Q_N = \{s_M\}$;
 - $s_N = s_M$;
 - $A_N = \emptyset$; and
 - $T_N = \emptyset$.
- And, if s_M is useful in M , then N is defined by:
 - $Q_N = \{q \in Q_M \mid q \text{ is useful in } M\}$;

- $s_N = s_M$;
- $A_N = A_M \cap Q_N = \{q \in A_M \mid q \in Q_N\}$; and
- $T_N = \mathbf{remRedun}(\emptyset, \{(q, x, r) \in T_M \mid q, r \in Q_N\})$.


Proposition 3.7.3

Suppose M is a finite automaton. Then:


- (1) **simplify** M is simplified;
- (2) **simplify** $M \approx M$; and
- (3) $\mathbf{alphabet}(\mathbf{simplify} M) = \mathbf{alphabet}(L(M)) \subseteq \mathbf{alphabet} M$.

Proof. Follows easily using Propositions 3.7.1 and 3.7.2. \square


If M is the finite automaton




then **simplify** M is the finite automaton



And if N is the finite automaton



then **simplify** N is the finite automaton



The Forlan module FA includes the following functions relating to the simplification of finite automata:


```

val simplify   : fa -> fa
val simplified : fa -> bool


```

The function `simplify` corresponds to **simplify**, and `simplified` tests whether an FA is simplified.


In the following, suppose `fa1` is the finite automaton



`fa2` is the finite automaton



and `fa3` is the finite automaton



Here are some example uses of `simplify` and `simplified`:

```
- FA.simplified fa1;
val it = false : bool
- val fa1' = FA.simplify fa1;
val fa1' = - : fa
- FA.output("", fa1');
{states} A, B {start state} A {accepting states} B
{transitions} A, % -> B; A, 0 -> A; B, 1 -> B
val it = () : unit
- FA.simplified fa1';
val it = true : bool
- val fa2' = FA.simplify fa2;
val fa2' = - : fa
- FA.output("", fa2');
{states} A, B {start state} A {accepting states} B
{transitions} A, % -> B; A, 0 -> A; B, 1 -> B
val it = () : unit
- val fa3' = FA.simplify fa3;
val fa3' = - : fa
- FA.output("", fa3');
{states} A, B, C {start state} A {accepting states} A
{transitions} A, % -> B | C; B, % -> A; C, % -> A
val it = () : unit
```

Thus the simplification of **fa3** resulted in the removal of the %-transitions between B and C.

Exercise 3.7.4

In the simplification of **fa3**, if transitions had been considered for removal due to being redundant in other orders, what FAs could have resulted.

3.7.1 Notes

The removal of useless states is analogous to the standard approach to ridding grammars of useless variables. The idea of removing redundant transitions, though, seems to be novel.

3.8 Proving the Correctness of Finite Automata

In this section, we consider techniques for proving the correctness of finite automata, i.e., for proving that finite automata accept the languages we want them to.

We begin by defining an indexed family of languages, Λ .

3.8.1 Definition of Λ

Proposition 3.8.1

Suppose M is a finite automaton.


- (1) For all $q \in Q_M$, $q \in \Delta_M(\{q\}, \%)$.
- (2) For all $q, r \in Q_M$ and $w \in \mathbf{Str}$, if $q, w \rightarrow r \in T_M$, then $r \in \Delta_M(\{q\}, w)$.
- (3) For all $p, q, r \in Q_M$ and $x, y \in \mathbf{Str}$, if $q \in \Delta_M(\{p\}, x)$ and $r \in \Delta_M(\{q\}, y)$, then $r \in \Delta_M(\{p\}, xy)$.

Suppose M is a finite automaton and $q \in Q_M$. Then we define

$$\Lambda_{M,q} = \{ w \in \mathbf{Str} \mid q \in \Delta_M(\{s_M\}, w) \}.$$

In other words, $\Lambda_{M,q}$ is the labels of all of the valid labeled paths for M that start at state s_M and end at q , i.e., it's all strings that can take us to state q when processed by M . Clearly, $\Lambda_{M,q} \subseteq (\mathbf{alphabet } M)^*$, for all FAs M and $q \in Q_M$. If it's clear which FA we are talking about, we sometimes abbreviate $\Lambda_{M,q}$ to Λ_q .

Let our example FA, M , be



Then:

- $01101 \in \Lambda_A$, because of the labeled path

$$A \xRightarrow{0} B \xRightarrow{1} B \xRightarrow{1} B \xRightarrow{0} A \xRightarrow{1} A,$$

- $01100 \in \Lambda_B$, because of the labeled path

$$A \xRightarrow{0} B \xRightarrow{1} B \xRightarrow{1} B \xRightarrow{0} A \xRightarrow{0} B.$$

Proposition 3.8.2

Suppose M is an FA. Then $L(M) = \bigcup \{ \Lambda_{M,q} \mid q \in A_M \}$, i.e., for all w , $w \in L(M)$ iff $w \in \Lambda_{M,q}$ for some $q \in A_M$.

Proof.

(only if) Suppose $w \in L(M)$. By Proposition 3.5.3, we have that $\Delta_M(s_M, w) \cap A_M \neq \emptyset$, so that there is a $q \in A_M$ such that $q \in \Delta_M(s_M, w)$. Thus $w \in \Lambda_{M,q}$.

(if) Suppose $w \in \Lambda_{M,q}$ for some $q \in A_M$. Thus $q \in \Delta_M(s_M, w)$ and $q \in A_M$, so that $\Delta_M(s_M, w) \cap A_M \neq \emptyset$. Hence $w \in L(M)$, by Proposition 3.5.3.

□

Proposition 3.8.3

Suppose M is a finite automaton.

(1) $\% \in \Lambda_{M,s_M}$.

(2) For all $q, r \in Q_M$ and $w, x \in \mathbf{Str}$. If $w \in \Lambda_{M,q}$ and $q, x \rightarrow r \in T_M$, then $wx \in \Lambda_{M,r}$.


Proof.

(1) By Proposition 3.8.1(1), we have that $s_M \in \Delta(\{s_M\}, \%)$, so that $\% \in \Lambda_{M,s_M}$.

(2) Suppose $q, r \in Q_M$, $w, x \in \mathbf{Str}$, $w \in \Lambda_{M,q}$ and $q, x \rightarrow r \in T_M$. Thus $q \in \Delta(\{s_M\}, w)$. Because $q, x \rightarrow r \in T_M$, Proposition 3.8.1(2) tells us that $r \in \Delta(\{q\}, x)$. Hence by Proposition 3.8.1(3), we have that $r \in \Delta(\{s_M\}, wx)$, so that $wx \in \Lambda_{M,r}$.

□

Our main example will be the FA, M :



Let

$$X = \{w \in \{0,1\}^* \mid w \text{ has an even number of 0's}\}, \text{ and}$$

$$Y = \{w \in \{0,1\}^* \mid w \text{ has an odd number of 0's}\}.$$

We want to prove that $L(M) = X$. Because $A_M = \{A\}$, Proposition 3.8.2 tells us that $L(M) = \Lambda_{M,A}$. Thus it will suffice to show that $\Lambda_{M,A} = X$. But our approach will also involve showing $\Lambda_{M,B} = Y$. We would cope with more states analogously, having one language per state.

3.8.2 Proving that Enough is Accepted

First, we study techniques for showing that everything we want an automaton to accept is really accepted.

Since $X, Y \subseteq \{0,1\}^*$, to prove that $X \subseteq \Lambda_{M,A}$ and $Y \subseteq \Lambda_{M,B}$, it will suffice to use strong string induction to show that, for all $w \in \{0,1\}^*$:

- (A) if $w \in X$, then $w \in \Lambda_{M,A}$; and
- (B) if $w \in Y$, then $w \in \Lambda_{M,B}$.

We proceed by strong string induction. Suppose $w \in \{0,1\}^*$, and assume the inductive hypothesis: for all $x \in \{0,1\}^*$, if x is a proper substring of w , then:

- (A) if $x \in X$, then $x \in \Lambda_A$; and
- (B) if $x \in Y$, then $x \in \Lambda_B$.

We must prove that:


- (A) if $w \in X$, then $w \in \Lambda_A$; and
- (B) if $w \in Y$, then $w \in \Lambda_B$.

There are two parts to show.

- (A) Suppose $w \in X$, so that w has an even number of 0's. We must show that $w \in \Lambda_A$. There are three cases to consider.
 - Suppose $w = \epsilon$. By Proposition 3.8.3(1), we have that $w = \epsilon \in \Lambda_A$.
 - Suppose $w = x0$, for some $x \in \{0,1\}^*$. Thus x has an odd number of 0's, so that $x \in Y$. Because x is a proper substring of w , part (B) of the inductive hypothesis tells us that $x \in \Lambda_B$. Furthermore, $B, 0 \rightarrow A \in T$, so that $w = x0 \in \Lambda_A$, by Proposition 3.8.3(2).

- Suppose $w = x1$, for some $x \in \{0, 1\}^*$. Thus x has an even number of 0's, so that $x \in X$. Because x is a proper substring of w , part (A) of the inductive hypothesis tells us that $x \in \Lambda_A$. Furthermore, $A, 1 \rightarrow A \in T$, so that $w = x1 \in \Lambda_A$, by Proposition 3.8.3(2).
- (B) Suppose $w \in Y$, so that w has an odd number of 0's. We must show that $w \in \Lambda_B$. There are three cases to consider.
- Suppose $w = \%$. But the number of 0's in $\%$ is 0, which is even—contradiction. Thus $w \in \Lambda_B$.
 - Suppose $w = x0$, for some $x \in \{0, 1\}^*$. Thus x has an even number of 0's, so that $x \in X$. Because x is a proper substring of w , part (A) of the inductive hypothesis tells us that $x \in \Lambda_A$. Furthermore, $A, 0 \rightarrow B \in T$, so that $w = x0 \in \Lambda_B$, by Proposition 3.8.3(2).
 - Suppose $w = x1$, for some $x \in \{0, 1\}^*$. Thus x has an odd number of 0's, so that $x \in Y$. Because x is a proper substring of w , part (B) of the inductive hypothesis tells us that $x \in \Lambda_B$. Furthermore, $B, 1 \rightarrow B \in T$, so that $w = x1 \in \Lambda_B$, by Proposition 3.8.3(2).

Let N be the finite automaton



Here we hope that $\Lambda_{N,A} = \{0\}^*$ and $L(N) = \Lambda_{N,B} = \{0\}^*\{11\}^*$, but if we try to prove that

$$\begin{aligned} \{0\}^* &\subseteq \Lambda_{N,A}, \text{ and} \\ \{0\}^*\{11\}^* &\subseteq \Lambda_{N,B} \end{aligned}$$

using our standard technique, there is a complication related to the $\%$ -transition.

We use strong string induction to show that, for all $w \in \{0, 1\}^*$:

- (A) if $w \in \{0\}^*$, then $w \in \Lambda_A$; and
- (B) if $w \in \{0\}^*\{11\}^*$, then $w \in \Lambda_B$.

In part (B), we assume that $w \in \{0\}^*\{11\}^*$, so that $w = 0^n(11)^m$ for some $n, m \in \mathbb{N}$. We must show that $w \in \Lambda_B$. We consider two cases: $m = 0$ and $m \geq 1$. The second of these is straightforward, so let's focus on the first. Then $w = 0^n \in \{0\}^*$. We want to use part (A) of the inductive hypothesis to conclude that $0^n \in \Lambda_A$, but there is a problem: 0^n is not a proper substring of $0^n = w$.

So, we must consider two subcases, when $n = 0$ and $n \geq 1$. In the first subcase, because $\% \in \Lambda_A$ and $A, \% \rightarrow B \in T$, we have that $w = \% = \% \% \in \Lambda_B$.

In the second subcase, we have that $w = 0^{n-1}0$. By part (A) of the inductive hypothesis, we have that $0^{n-1} \in \Lambda_A$. Thus, because $A, 0 \rightarrow A \in T$ and $A, \% \rightarrow B \in T$, we can conclude $w = 0^n = 0^{n-1}0\% \in \Lambda_B$.

Because there are no transitions from B back to A , we could first prove that, for all $w \in \{0, 1\}^*$,

(A) if $w \in \{0\}^*$, then $w \in \Lambda_A$,

and then use (A) to prove that for all $w \in \{0, 1\}^*$,

(B) if $w \in \{0\}^*\{11\}^*$, then $w \in \Lambda_B$.

This works whenever one part of a machine has transitions to another part, but there are no transitions from that second part back to the first part, i.e., when the two parts are not mutually recursive.

In the case of N , we could use mathematical induction instead of strong string induction:

(A) for all $n \in \mathbb{N}$, $0^n \in \Lambda_A$, and

(B) for all $n, m \in \mathbb{N}$, $0^n(11)^m \in \Lambda_B$ (do induction on m , fixing n).

3.8.3 Proving that Everything Accepted is Wanted

It's tempting to try to prove that everything accepted by a finite automaton is wanted using strong string induction, with implications like

(A) if $w \in \Lambda_A$, then $w \in X$.

Unfortunately, this doesn't work when a finite automaton contains $\%$ -transitions. Instead, we do such proofs using a new induction principle that we call induction on Λ .

Theorem 3.8.4 (Principle of Induction on Λ)

Suppose M is a finite automaton, and $P_q(w)$ is a property of a $w \in \Lambda_{M,q}$, for all $q \in Q_M$. If

- $P_{s_M}(\%)$ and
- for all $q, r \in Q_M$, $x \in \mathbf{Str}$ and $w \in \Lambda_{M,q}$,
if $q, x \rightarrow r \in T_M$ and $(\dagger) P_q(w)$, then $P_r(wx)$,

then

$$\text{for all } q \in Q_M, \text{ for all } w \in \Lambda_{M,q}, P_q(w).$$

We refer to (\dagger) as the inductive hypothesis.

Proof. It suffices to show that, for all $lp \in \mathbf{LP}$, for all $q \in Q_M$, if lp is valid for M , **startState** $lp = s_M$ and **endState** $lp = q$, then $P_q(\text{label } lp)$. We prove this by well-founded induction on the length of lp . \square

In the case of our example FA, M , we can let $P_A(w)$ and $P_B(w)$ be $w \in X$ and $w \in Y$, respectively, where, as before,

$$X = \{w \in \{0,1\}^* \mid w \text{ has an even number of 0's}\}, \text{ and}$$

$$Y = \{w \in \{0,1\}^* \mid w \text{ has an odd number of 0's}\}.$$

Then the principle of induction on Λ tells us that

(A) for all $w \in \Lambda_A$, $w \in X$, and

(B) for all $w \in \Lambda_B$, $w \in Y$,

follows from showing

(empty string) $\% \in X$,

(A, 0 \rightarrow B) for all $w \in \Lambda_A$, if $(\dagger) w \in X$, then $w0 \in Y$,

(A, 1 \rightarrow A) for all $w \in \Lambda_A$, if $(\dagger) w \in X$, then $w1 \in X$,

(B, 0 \rightarrow A) for all $w \in \Lambda_B$, if $(\dagger) w \in Y$, then $w0 \in X$, and

(B, 1 \rightarrow B) for all $w \in \Lambda_B$, if $(\dagger) w \in Y$, then $w1 \in Y$.

We refer to (\dagger) as the inductive hypothesis.

In fact, when setting this proof up, instead of explicitly mentioning P_A and P_B , we can simply say that we are proving

(A) for all $w \in \Lambda_A$, $w \in X$, and

(B) for all $w \in \Lambda_B$, $w \in Y$,

by induction on Λ .

There are five steps to show.

(empty string) Because $\% \in \{0,1\}^*$ and $\%$ has no 0's, we have that $\% \in X$.

(A, 0 \rightarrow B) Suppose $w \in \Lambda_A$, and assume the inductive hypothesis: $w \in X$. Hence $w \in \{0,1\}^*$ and w has an even number of 0's. Thus $w0 \in \{0,1\}^*$ and $w0$ has an odd number of 0's, so that $w0 \in Y$.

(A, 1 \rightarrow A) Suppose $w \in \Lambda_A$, and assume the inductive hypothesis: $w \in X$. Then $w1 \in X$.

(B, 0 \rightarrow A) Suppose $w \in \Lambda_B$, and assume the inductive hypothesis: $w \in Y$. Then $w0 \in X$.

(B, $1 \rightarrow B$) Suppose $w \in \Lambda_B$, and assume the inductive hypothesis: $w \in Y$.
Then $w1 \in Y$.

Because of


(A) for all $w \in \Lambda_A$, $w \in X$, and

(B) for all $w \in \Lambda_B$, $w \in Y$,

we have that $\Lambda_A \subseteq X$ and $\Lambda_B \subseteq Y$.

Because $X \subseteq \Lambda_A$ and $Y \subseteq \Lambda_B$, we can conclude that $L(M) = \Lambda_A = X$ and $\Lambda_B = Y$.

Consider our second example, N , again:



We can use induction on Λ to prove that

(A) for all $w \in \Lambda_A$, $w \in \{0\}^*$, and

(B) for all $w \in \Lambda_B$, $w \in \{0\}^*\{11\}^*$.

Thus $\Lambda_A \subseteq \{0\}^*$ and $\Lambda_B \subseteq \{0\}^*\{11\}^*$. Because $\{0\}^* \subseteq \Lambda_A$ and $\{0\}^*\{11\}^* \subseteq \Lambda_B$, we can conclude that $\Lambda_A = \{0\}^*$ and $L(N) = \Lambda_B = \{0\}^*\{11\}^*$.

3.8.4 Notes

Books on formal language theory typically give short shrift to the proof of correctness of finite automata, carrying out one or two correctness proofs using induction on the length of strings. In contrast, we have introduced and applied elegant techniques for proving the correctness of FAs. Of particular note is our principle of induction on Λ .

3.9 Empty-string Finite Automata

In this and the following two sections, we will study three progressively more restricted kinds of finite automata:

- empty-string finite automata (EFAs);
- nondeterministic finite automata (NFAs); and
- deterministic finite automata (DFAs).

Every DFA will be an NFA; every NFA will be an EFA; and every EFA will be an FA. Thus, $L(M)$ will be well-defined, if M is a DFA, NFA or EFA.

The more restricted kinds of automata will be easier to process on the computer than the more general kinds; they will also have nicer reasoning principles than the more general kinds. We will give algorithms for converting the more general kinds of automata into the more restricted kinds. Thus even the deterministic finite automata will accept the same set of languages as the finite automata. On the other hand, it will sometimes be easier to find one of the more general kinds of automata that accepts a given language rather than one of the more restricted kinds accepting the language. And, there are languages where the smallest DFA accepting the language is exponentially bigger than the smallest FA accepting the language.

3.9.1 Definition of EFAs

An *empty-string finite automaton* (EFA) M is a finite automaton such that

$$T_M \subseteq \{ q, x \rightarrow r \mid q, r \in \mathbf{Sym} \text{ and } x \in \mathbf{Str} \text{ and } |x| \leq 1 \}.$$

In other words, an FA is an EFA iff every string of every transition of the FA is either ϵ or has a single symbol.

For example, $A, \epsilon \rightarrow B$ and $A, 1 \rightarrow B$ are legal EFA transitions, but $A, 11 \rightarrow B$ is not legal. We write **EFA** for the set of all empty-string finite automata. Thus $\mathbf{EFA} \subsetneq \mathbf{FA}$.

The following proposition obviously holds.

Proposition 3.9.1

Suppose M is an EFA.

- For all $N \in \mathbf{FA}$, if M **iso** N , then N is an EFA.
- For all bijections f from Q_M to some set of symbols, **renameStates**(M, f) is an EFA.
- **renameStatesCanonically** M is an EFA.
- **simplify** M is an EFA.

3.9.2 Converting FAs to EFAs

If we want to convert an FA into an equivalent EFA, we can proceed as follows. Every state of the FA will be a state of the EFA, the start and accepting states are unchanged, and every transition of the FA that is a legal EFA transition will be a transition of the EFA. If our FA has a transition


$$p, b_1 b_2 \cdots b_n \rightarrow r,$$

where $n \geq 2$ and the b_i are symbols, then we replace this transition with the n transitions


$$p \xrightarrow{b_1} q_1, q_1 \xrightarrow{b_2} q_2, \dots, q_{n-1} \xrightarrow{b_n} r,$$

where q_1, \dots, q_{n-1} are $n - 1$ new, non-accepting, states.

For example, we can convert the FA



into the EFA



We have to be careful how we choose our new states. The symbols we choose can't be states of the original machine, and we can't choose the same symbol twice. Instead of making a series of random choices, we will use structured symbols in such a way that one will be able to look at a resulting EFA and tell what the original FA was.

First, the algorithm renames each old state q to $\langle 1, q \rangle$. Then it can replace a transition

$$p \xrightarrow{b_1 b_2 \dots b_n} r,$$

where $n \geq 2$ and the b_i are symbols, with the transitions

$$\begin{aligned} \langle 1, p \rangle &\xrightarrow{b_1} \langle 2, \langle p, b_1, b_2 \dots b_n, r \rangle \rangle, \\ \langle 2, \langle p, b_1, b_2 \dots b_n, r \rangle \rangle &\xrightarrow{b_2} \langle 2, \langle p, b_1 b_2, b_3 \dots b_n, r \rangle \rangle, \\ &\dots, \\ \langle 2, \langle p, b_1 b_2 \dots b_{n-1}, b_n, r \rangle \rangle &\xrightarrow{b_n} \langle 1, r \rangle. \end{aligned}$$

We define a function $\mathbf{faToEFA} \in \mathbf{FA} \rightarrow \mathbf{EFA}$ that converts FAs into EFAs by saying that $\mathbf{faToEFA} M$ is the result of running the above algorithm on input M .

Theorem 3.9.2

For all $M \in \mathbf{FA}$:

- $\mathbf{faToEFA} M \approx M$; and
- $\mathbf{alphabet}(\mathbf{faToEFA} M) = \mathbf{alphabet} M$.

3.9.3 Processing EFAs in Forlan

The Forlan module `EFA` defines an abstract type `efa` (in the top-level environment) of empty-string finite automata, along with various functions for processing EFAs. Values of type `efa` are implemented as values of type `fa`, and the module `EFA` provides functions

```
val injToFA    : efa -> fa
val projFromFA : fa -> efa
```

for making a value of type `efa` have type `fa`, i.e., “injecting” an `efa` into type `fa`, and for making a value of type `fa` that is an EFA have type `efa`, i.e., “projecting” an `fa` that is an EFA to type `efa`. If one tries to project an `fa` that is not an EFA to type `efa`, an error is signaled. The functions `injToFA` and `projFromFA` are available in the top-level environment as `injEFAToFA` and `projFAToEFA`, respectively.

The module `EFA` also defines the functions:

```
val input  : string -> efa
val fromFA : fa -> efa
```

The function `input` is used to input an EFA, i.e., to input a value of type `fa` using `FA.input`, and then attempt to project it to type `efa`. The function `fromFA` corresponds to our conversion function `faToEFA`, and is available in the top-level environment with that name:

```
val faToEFA : fa -> efa
```

Finally, most of the functions for processing FAs that were introduced in previous sections are inherited by `EFA`:


```
val output          : string * efa -> unit
val numStates       : efa -> int
val numTransitions  : efa -> int
val equal           : efa * efa -> bool
val alphabet        : efa -> sym set
val checkLP         : efa -> lp -> unit
val validLP         : efa -> lp -> bool
val isomorphism     : efa * efa * sym_rel -> bool
val findIsomorphism : efa * efa -> sym_rel
val isomorphic      : efa * efa -> bool
val renameStates    : efa * sym_rel -> efa
```

```

val renameStatesCanonically : efa -> efa
val processStr                : efa -> sym set * str -> sym set
val accepted                  : efa -> str -> bool
val findLP                    : efa -> sym set * str * sym set -> lp
val findAcceptingLP           : efa -> str -> lp
val simplified                  : efa -> bool
val simplify                   : efa -> efa

```

Suppose that *fa* is the finite automaton



Here are some example uses of a few of the above functions:

```

- projFAToEFA fa;
invalid label in transition: "12"

uncaught exception Error
- val efa = faToEFA fa;
val efa = - : efa
- EFA.output("", efa);
{states}
<1,A>, <1,B>, <2,<A,1,2,B>>, <2,<B,3,45,B>>, <2,<B,34,5,B>>
{start state} <1,A> {accepting states} <1,B>
{transitions}
<1,A>, 0 -> <1,A>; <1,A>, 1 -> <2,<A,1,2,B>>;
<1,B>, 3 -> <2,<B,3,45,B>>; <2,<A,1,2,B>>, 2 -> <1,B>;
<2,<B,3,45,B>>, 4 -> <2,<B,34,5,B>>; <2,<B,34,5,B>>, 5 -> <1,B>
val it = () : unit
- val efa' = EFA.renameStatesCanonically efa;
val efa' = - : efa
- EFA.output("", efa');
{states} A, B, C, D, E {start state} A {accepting states} B
{transitions}
A, 0 -> A; A, 1 -> C; B, 3 -> D; C, 2 -> B; D, 4 -> E; E, 5 -> B
val it = () : unit
- val rel = EFA.findIsomorphism(efa, efa');
val rel = - : sym_rel
- SymRel.output("", rel);
(<1,A>, A), (<1,B>, B), (<2,<A,1,2,B>>, C), (<2,<B,3,45,B>>, D),
(<2,<B,34,5,B>>, E)
val it = () : unit
- LP.output("", FA.findAcceptingLP fa (Str.input ""));
@ 012345
@ .
A, 0 => A, 12 => B, 345 => B

```

```

val it = () : unit
- LP.output("", EFA.findAcceptingLP efa' (Str.input ""));
@ 012345
@ .
A, 0 => A, 1 => C, 2 => B, 3 => D, 4 => E, 5 => B
val it = () : unit

```

3.9.4 Notes

The algorithm for converting FAs to EFAs is obvious, but our use of structured state names so as to make the resulting EFAs self-documenting is novel.

3.10 Nondeterministic Finite Automata

In this section, we study the second of our more restricted kinds of finite automata: nondeterministic finite automata.

3.10.1 Definition of NFAs

A *nondeterministic finite automaton* (NFA) M is a finite automaton such that

$$T_M \subseteq \{ q, x \rightarrow r \mid q, r \in \mathbf{Sym} \text{ and } x \in \mathbf{Str} \text{ and } |x| = 1 \}.$$

In other words, an FA is an NFA iff every string of every transition of the FA has a single symbol. For example, $A, 1 \rightarrow B$ is a legal NFA transition, but $A, \% \rightarrow B$ and $A, 11 \rightarrow B$ are not legal. We write **NFA** for the set of all nondeterministic finite automata. Thus $\mathbf{NFA} \subsetneq \mathbf{EFA} \subsetneq \mathbf{FA}$.

The following proposition obviously holds.


Proposition 3.10.1

Suppose M is an NFA.

- For all $N \in \mathbf{FA}$, if $M \text{ iso } N$, then N is an NFA.
- For all bijections f from Q_M to some set of symbols, $\text{renameStates}(M, f)$ is an NFA.
- $\text{renameStatesCanonically } M$ is an NFA.
- $\text{simplify } M$ is an NFA.

3.10.2 Converting EFAs to NFAs

Suppose M is the EFA



To convert M into an equivalent NFA, we will have to:

- replace the transitions $A, \% \rightarrow B$ and $B, \% \rightarrow C$ with legal transitions (for example, because of the valid labeled path

$$A \xRightarrow{\%} B \xRightarrow{1} B \xRightarrow{\%} C,$$

we will add the transition $A, 1 \rightarrow C$;

- make (at least) A be an accepting state (so that $\%$ is accepted by the NFA).

Before defining a general procedure for converting EFAs to NFAs, we first say what we mean by the empty-closure of a set of states. Suppose M is a finite automaton and $P \subseteq Q_M$. The *empty-closure* of P (**emptyClose** $_M P$) is the least subset X of Q_M such that

- $P \subseteq X$; and
- for all $q, r \in Q_M$, if $q \in X$ and $q, \% \rightarrow r \in T_M$, then $r \in X$.

We sometimes abbreviate **emptyClose** $_M P$ to **emptyClose** P , when M is clear from the context. For example, if M is our example EFA and $P = \{A\}$, then:

- $A \in X$;
- $B \in X$, since $A \in X$ and $A, \% \rightarrow B \in T_M$;
- $C \in X$, since $B \in X$ and $B, \% \rightarrow C \in T_M$.

Thus **emptyClose** $P = \{A, B, C\}$.

Suppose M is a finite automaton and $P \subseteq Q_M$. The *backwards empty-closure* of P (**emptyCloseBackwards** $_M P$) is the least subset X of Q_M such that

- $P \subseteq X$; and
- for all $q, r \in Q_M$, if $r \in X$ and $q, \% \rightarrow r \in T_M$, then $q \in X$.

We sometimes drop the M from **emptyCloseBackwards** $_M$, when it's clear from the context. For example, if M is our example EFA and $P = \{C\}$, then:

- $C \in X$;
- $B \in X$, since $C \in X$ and $B, \% \rightarrow C \in T_M$;

- $A \in X$, since $B \in X$ and $A, \% \rightarrow B \in T_M$.

Thus $\text{emptyCloseBackwards } P = \{A, B, C\}$.

Proposition 3.10.2

Suppose M is a finite automaton. For all $P \subseteq Q_M$,

$$\text{emptyClose}_M P = \Delta_M(P, \%).$$

In other words, $\text{emptyClose}_M P$ is all of the states that can be reached from elements of P by sequences of $\%$ -transitions.

Proposition 3.10.3

Suppose M is a finite automaton. For all $P \subseteq Q_M$,

$$\text{emptyCloseBackwards}_M P = \{q \in Q_M \mid \Delta_M(\{q\}, \%) \cap P \neq \emptyset\}.$$

In other words, $\text{emptyCloseBackwards}_M P$ is all of the states from which it is possible to reach elements of P by sequences of $\%$ -transitions.

We define a function/algorithm $\text{efaToNFA} \in \mathbf{EFA} \rightarrow \mathbf{NFA}$ that converts EFAs into NFAs by saying that $\text{efaToNFA } M$ is the NFA N such that:

- $Q_N = Q_M$;
- $s_N = s_M$;
- $A_N = \text{emptyCloseBackwards } A_M$; and
- T_N is the set of all transitions $q', a \rightarrow r'$ such that $q', r' \in Q_M$, $a \in \mathbf{Sym}$, and there are $q, r \in Q_M$ such that:
 - $q, a \rightarrow r \in T_M$;
 - $q' \in \text{emptyCloseBackwards } \{q\}$; and
 - $r' \in \text{emptyClose } \{r\}$.

To compute the set T_N , we process each transition $q, x \rightarrow r$ of M as follows. If $x = \%$, then we generate no transitions. Otherwise, our transition is $q, a \rightarrow r$ for some symbol a . We then compute the backwards empty-closure of $\{q\}$, and call the result X , and compute the (forwards) empty-closure of $\{r\}$, and call the result Y . We then add all of the elements of

$$\{q', a \rightarrow r' \mid q' \in X \text{ and } r' \in Y\}$$


to T_N .

Because the algorithm defines A_N to be $\text{emptyCloseBackwards } A_M$, it could let T_N be the set of all transitions $q', a \rightarrow r$ such that $q', r \in Q_M$, $a \in \mathbf{Sym}$, and there is a $q \in Q_M$ such that:

- $q, a \rightarrow r \in T_M$; and
- $q' \in \mathbf{emptyCloseBackwards} \{q\}$.

This would mean that N would have fewer transitions. However, for esthetic reasons, we'll stick with the symmetric definition of T_N .

Let M be our example EFA




and let $N = \mathbf{efaToNFA} \ M$. Then

- $Q_N = Q_M = \{A, B, C\}$;
- $s_N = s_M = A$;
- $A_N = \mathbf{emptyCloseBackwards} \ A_M = \mathbf{emptyCloseBackwards} \ \{C\} = \{A, B, C\}$.

Now, let's work out what T_N is, by processing each of M 's transitions.

- From the transitions $A, \% \rightarrow B$ and $B, \% \rightarrow C$, we get no elements of T_N .
- Consider the transition $A, 0 \rightarrow A$. Since $\mathbf{emptyCloseBackwards} \ \{A\} = \{A\}$ and $\mathbf{emptyClose} \ \{A\} = \{A, B, C\}$, we add $A, 0 \rightarrow A$, $A, 0 \rightarrow B$ and $A, 0 \rightarrow C$ to T_N .
- Consider the transition $B, 1 \rightarrow B$. Since $\mathbf{emptyCloseBackwards} \ \{B\} = \{A, B\}$ and $\mathbf{emptyClose} \ \{B\} = \{B, C\}$, we add $A, 1 \rightarrow B$, $A, 1 \rightarrow C$, $B, 1 \rightarrow B$ and $B, 1 \rightarrow C$ to T_N .
- Consider the transition $C, 2 \rightarrow C$. Since $\mathbf{emptyCloseBackwards} \ \{C\} = \{A, B, C\}$ and $\mathbf{emptyClose} \ \{C\} = \{C\}$, we add $A, 2 \rightarrow C$, $B, 2 \rightarrow C$ and $C, 2 \rightarrow C$ to T_N .

Thus our NFA N is



Theorem 3.10.4

For all $M \in \mathbf{EFA}$:


- $\mathbf{efaToNFA} \ M \approx M$; and
- $\mathbf{alphabet}(\mathbf{efaToNFA} \ M) = \mathbf{alphabet} \ M$.

3.10.3 Converting EFAs to NFAs, and Processing NFAs in Forlan

The Forlan module `FA` defines the following functions for computing forwards and backwards empty-closures:

```
val emptyClose      : fa -> sym set -> sym set
val emptyCloseBackwards : fa -> sym set -> sym set
```

For example, if `fa` is bound to the finite automaton



then we can compute the empty-closure of $\{A\}$ as follows:

```
- SymSet.output("", FA.emptyClose fa (SymSet.input ""));
@ A
@ .
A, B, C
val it = () : unit
```

The Forlan module `NFA` defines an abstract type `nfa` (in the top-level environment) of nondeterministic finite automata, along with various functions for processing NFAs. Values of type `nfa` are implemented as values of type `fa`, and the module `NFA` provides the following injection and projection functions:

```
val injToFA      : nfa -> fa
val injToEFA     : nfa -> efa
val projFromFA   : fa -> nfa
val projFromEFA  : efa -> nfa
```

The functions `injToFA`, `injToEFA`, `projFromFA` and `textttprojFromEFA` are available in the top-level environment as `textttinjNFAToFA`, `injNFAToEFA`, `projFAToNFA` and `textttprojEFAToNFA`, respectively.

The module `NFA` also defines the functions:

```
val input      : string -> nfa
val fromEFA    : efa -> nfa
```

The function `input` is used to input an NFA, and the function `fromEFA` corresponds to our conversion function `efaToNFA`, and is available in the top-level environment with that name:

```
val efaToNFA : efa -> nfa
```

Most of the functions for processing FAs that were introduced in previous sections are inherited by `NFA`:

```

val output          : string * nfa -> unit
val numStates       : nfa -> int
val numTransitions  : nfa -> int
val alphabet        : nfa -> sym set
val equal           : nfa * nfa -> bool
val checkLP         : nfa -> lp -> unit
val validLP         : nfa -> lp -> bool
val isomorphism      : nfa * nfa * sym_rel -> bool
val findIsomorphism  : nfa * nfa -> sym_rel
val isomorphic       : nfa * nfa -> bool
val renameStates     : nfa * sym_rel -> nfa
val renameStatesCanonically : nfa -> nfa
val processStr       : nfa -> sym set * str -> sym set
val accepted        : nfa -> str -> bool
val findLP          : nfa -> sym set * str * sym set -> lp
val findAcceptingLP : nfa -> str -> lp
val simplified       : nfa -> bool
val simplify        : nfa -> nfa

```


Finally, the functions for computing forwards and backwards empty-closures are inherited by the EFA module

```

val emptyClose      : efa -> sym set -> sym set
val emptyCloseBackwards : efa -> sym set -> sym set

```

Suppose that `efa` is the `efa`



Here are some example uses of a few of the above functions:

```

- projEFAToNFA efa;
invalid label in transition: "%"

uncaught exception Error
- val nfa = efaToNFA efa;
val nfa = - : nfa
- NFA.output("", nfa);
{states} A, B, C {start state} A {accepting states} A, B, C
{transitions}
A, 0 -> A | B | C; A, 1 -> B | C; A, 2 -> C; B, 1 -> B | C;
B, 2 -> C; C, 2 -> C
val it = () : unit
- LP.output("", EFA.findAcceptingLP efa (Str.input ""));
@ 012
@ .
A, 0 => A, % => B, 1 => B, % => C, 2 => C

```

```

val it = () : unit
- LP.output("", NFA.findAcceptingLP nfa (Str.input ""));
@ 012
@ .
A, 0 => A, 1 => B, 2 => C
val it = () : unit

```

3.10.4 Notes

Because we have defined the meaning of finite automata via labeled paths instead of transition functions, our EFA to NFA conversion algorithm is easy to understand and prove correct.

3.11 Deterministic Finite Automata

In this section, we study the third of our more restricted kinds of finite automata: deterministic finite automata.


3.11.1 Definition of DFAs

A *deterministic finite automaton* (DFA) M is a finite automaton such that:


- $T_M \subseteq \{ q, x \rightarrow r \mid q, r \in \mathbf{Sym} \text{ and } x \in \mathbf{Str} \text{ and } |x| = 1 \}$; and
- for all $q \in Q_M$ and $a \in \mathbf{alphabet } M$, there is a unique $r \in Q_M$ such that $q, a \rightarrow r \in T_M$.

In other words, an FA is a DFA iff it is an NFA and, for every state q of the automaton and every symbol a of the automaton's alphabet, there is exactly one state that can be entered from state q by reading a from the automaton's input. We write **DFA** for the set of all deterministic finite automata. Thus **DFA** \subsetneq **NFA** \subsetneq **EFA** \subsetneq **FA**.

Let M be the finite automaton



It turns out that $L(M) = \{ w \in \{0, 1\}^* \mid 000 \text{ is not a substring of } w \}$. Although M is an NFA, it's not a DFA, since $0 \in \mathbf{alphabet } M$ but there is no transition of the form $C, 0 \rightarrow r$. However, we can make M into a DFA by adding a dead state D:



We will never need more than one dead state in a DFA.

The following proposition obviously holds.

Proposition 3.11.1

Suppose M is a DFA.

- For all $N \in \mathbf{FA}$, if $M \text{ iso } N$, then N is a DFA.
- For all bijections f from Q_M to some set of symbols, $\text{renameStates}(M, f)$ is a DFA.
- $\text{renameStatesCanonically } M$ is a DFA.

Now, we prove a proposition that doesn't hold for arbitrary NFAs.

Proposition 3.11.2

Suppose M is a DFA. For all $q \in Q_M$ and $w \in (\text{alphabet } M)^*$, $|\Delta_M(\{q\}, w)| = 1$.

Proof. An easy left string induction on w . \square

Suppose M is a DFA. Because of Proposition 3.11.2, we can define *the transition function* δ_M for M , $\delta_M \in Q_M \times (\text{alphabet } M)^* \rightarrow Q_M$, by:


$$\delta_M(q, w) = \text{the unique } r \in Q_M \text{ such that } r \in \Delta_M(\{q\}, w).$$

In other words, $\delta_M(q, w)$ is the unique state r of M that is the end of a valid labeled path for M that starts at q and is labeled by w . Thus, for all $q, r \in Q_M$ and $w \in (\text{alphabet } M)^*$,

$$\delta_M(q, w) = r \quad \text{iff} \quad r \in \Delta_M(\{q\}, w).$$

We sometimes abbreviate $\delta_M(q, w)$ to $\delta(q, w)$.

For example, if M is the DFA



then $\delta(A, \%) = A$, $\delta(A, 0100) = C$ and $\delta(B, 000100) = D$.

Having defined the δ function, we can study its properties.

Proposition 3.11.3

Suppose M is a DFA.

- (1) For all $q \in Q_M$, $\delta_M(q, \%) = q$.
- (2) For all $q \in Q_M$ and $a \in \mathbf{alphabet} M$, $\delta_M(q, a) =$ the unique $r \in Q_M$ such that $q, a \rightarrow r \in T_M$.
- (3) For all $q \in Q_M$ and $x, y \in (\mathbf{alphabet} M)^*$, $\delta_M(q, xy) = \delta_M(\delta_M(q, x), y)$.

Suppose M is a DFA. By part (2) of the proposition, we have that, for all $q, r \in Q_M$ and $a \in \mathbf{alphabet} M$,

$$\delta_M(q, a) = r \quad \text{iff} \quad q, a \rightarrow r \in T_M.$$

Now we can use the δ function to explain when a string is accepted by an FA.


Proposition 3.11.4

Suppose M is a DFA. $L(M) = \{ w \in (\mathbf{alphabet} M)^* \mid \delta_M(s_M, w) \in A_M \}$.

Proof. To see that the left-hand side is a subset of the right-hand side, suppose $w \in L(M)$. Then $w \in (\mathbf{alphabet} M)^*$ and there is a $q \in A_M$ such that $q \in \Delta_M(\{s_M\}, w)$. Thus $\delta_M(s_M, w) = q \in A_M$.

To see that the right-hand side is a subset of the left-hand side, suppose $w \in (\mathbf{alphabet} M)^*$ and $\delta_M(s_M, w) \in A_M$. Then $\delta_M(s_M, w) \in \Delta_M(\{s_M\}, w)$, and thus $w \in L(M)$. \square

The preceding propositions give us an efficient algorithm for checking whether a string is accepted by a DFA. For example, suppose M is the DFA



To check whether 0100 is accepted by M , we need to determine whether $\delta(A, 0100) \in \{A, B, C\}$.

For instance, we have that:

$$\begin{aligned}
 \delta(A, 0100) &= \delta(\delta(A, 0), 100) \\
 &= \delta(B, 100) \\
 &= \delta(\delta(B, 1), 00) \\
 &= \delta(A, 00) \\
 &= \delta(\delta(A, 0), 0) \\
 &= \delta(B, 0) \\
 &= C \\
 &\in \{A, B, C\}.
 \end{aligned}$$

Thus 0100 is accepted by M .

3.11.2 Proving the Correctness of DFAs

Since every DFA is an FA, we could prove the correctness of DFAs using the techniques that we have already studied. But it turns out that giving a separate proof that enough is accepted by a DFA is unnecessary—it will follow from the proof that everything accepted is wanted.

Proposition 3.11.5

Suppose M is a DFA. Then, for all $w \in (\text{alphabet } M)^*$ and $q \in Q_M$,

$$w \in \Lambda_{M,q} \quad \text{iff} \quad \delta_M(s_M, w) = q.$$

Proof. Suppose $w \in (\text{alphabet } M)^*$ and $q \in Q_M$.

(only if) Suppose $w \in \Lambda_q$. Then $q \in \Delta(\{s\}, w)$. Thus $\delta(s, w) = q$.

(if) Suppose $\delta(s, w) = q$. Then $q \in \Delta(\{s\}, w)$, so that $w \in \Lambda_q$.

□

We already know that, if M is an FA, then $L(M) = \bigcup \{ \Lambda_q \mid q \in A_M \}$.


Proposition 3.11.6

Suppose M is a DFA.

(1) $(\text{alphabet } M)^* = \bigcup \{ \Lambda_q \mid q \in Q_M \}$.

(2) For all $q, r \in Q_M$, if $q \neq r$, then $\Lambda_q \cap \Lambda_r = \emptyset$.

Suppose M is the DFA



and let $X = \{w \in \{0,1\}^* \mid 000 \text{ is not a substring of } w\}$. We will show that $L(M) = X$. Note that, for all $w \in \{0,1\}^*$:

- $w \in X$ iff 000 is not a substring of w .
- $w \notin X$ iff 000 is a substring of w .

First, we use induction on Λ , to prove that:

- (A) for all $w \in \Lambda_A$, $w \in X$ and 0 is not a suffix of w ;
- (B) for all $w \in \Lambda_B$, $w \in X$ and 0, but not 00, is a suffix of w ;
- (C) for all $w \in \Lambda_C$, $w \in X$ and 00 is a suffix of w ; and
- (D) for all $w \in \Lambda_D$, $w \notin X$.

There are nine steps (1 + the number of transitions) to show.

(empty string) We must show that $\epsilon \in X$ and 0 is not a suffix of ϵ . This follows since ϵ has no 0's.

(A, 0 \rightarrow B) Suppose $w \in \Lambda_A$, and assume the inductive hypothesis: $w \in X$ and 0 is not a suffix of w . We must show that $w0 \in X$ and 0, but not 00, is a suffix of $w0$. Because $w \in X$ and 0 is not a suffix of w , we have that $w0 \in X$. Clearly, 0 is a suffix of $w0$. And, since 0 is not a suffix of w , we have that 00 is not a suffix of $w0$.

(A, 1 \rightarrow A) Suppose $w \in \Lambda_A$, and assume the inductive hypothesis: $w \in X$ and 0 is not a suffix of w . We must show that $w1 \in X$ and 0 is not a suffix of $w1$. Since $w \in X$, we have that $w1 \in X$. And, 0 is not a suffix of $w1$.

(B, 0 \rightarrow C) Suppose $w \in \Lambda_B$, and assume the inductive hypothesis: $w \in X$ and 0, but not 00, is a suffix of w . We must show that $w0 \in X$ and 00 is a suffix of $w0$. Because $w \in X$ and 00 is not suffix of w , we have that $w0 \in X$. And since 0 is a suffix of w , it follows that 00 is a suffix of $w0$.

(B, 1 \rightarrow A) Suppose $w \in \Lambda_B$, and assume the inductive hypothesis: $w \in X$ and 0, but not 00, is a suffix of w . We must show that $w1 \in X$ and 0 is not a suffix of $w1$. Because $w \in X$, we have that $w1 \in X$. And, 0 is not a suffix of $w1$.

(C, $0 \rightarrow D$) Suppose $w \in \Lambda_C$, and assume the inductive hypothesis: $w \in X$ and 00 is a suffix of w . We must show that $w0 \notin X$. Because 00 is a suffix of w , we have that 000 is a suffix of $w0$. Thus $w0 \notin X$.

(C, $1 \rightarrow A$) Suppose $w \in \Lambda_C$, and assume the inductive hypothesis: $w \in X$ and 00 is a suffix of w . We must show that $w1 \in X$ and 0 is not a suffix of $w1$. Because $w \in X$, we have that $w1 \in X$. And, 0 is not a suffix of $w1$.

(D, $0 \rightarrow D$) Suppose $w \in \Lambda_D$, and assume the inductive hypothesis: $w \notin X$. We must show that $w0 \notin X$. Because $w \notin X$, we have that $w0 \notin X$.

(D, $1 \rightarrow D$) Suppose $w \in \Lambda_D$, and assume the inductive hypothesis: $w \notin X$. We must show that $w1 \notin X$. Because $w \notin X$, we have that $w1 \notin X$.

Now, we use the result of our induction on Λ to show that $L(M) = X$.

($L(M) \subseteq X$) Suppose $w \in L(M)$. Because $A_M = \{A, B, C\}$, we have that $w \in L(M) = \Lambda_A \cup \Lambda_B \cup \Lambda_C$. Thus, by parts (A)–(C), we have that $w \in X$.


($X \subseteq L(M)$) Suppose $w \in X$. Since $X \subseteq \{0, 1\}^*$, we have that $w \in \{0, 1\}^*$. Suppose, toward a contradiction, that $w \notin L(M)$. Because $w \notin L(M) = \Lambda_A \cup \Lambda_B \cup \Lambda_C$ and $w \in \{0, 1\}^* = (\text{alphabet } M)^* = \Lambda_A \cup \Lambda_B \cup \Lambda_C \cup \Lambda_D$, we must have that $w \in \Lambda_D$. But then part (D) tells us that $w \notin X$ —contradiction. Thus $w \in L(M)$.

For the above approach to work, when proving $L(M) = X$ for a DFA M and language X , we simply need that:

- the property associated with each accepting state implies being in X ; and
- the property associated with each non-accepting state implies not being in X .

3.11.3 Simplification of DFAs

Let M be our example DFA



Then M is not simplified, since the state D is dead. But if we get rid of D, then we won't have a DFA anymore. Thus, we will need:


- a notion of when a DFA is simplified that is more liberal than our standard notion; and

- a corresponding simplification procedure for DFAs.

We say that a DFA M is *deterministically simplified* iff

- every element of Q_M is reachable; and
- at most one element of Q_M is dead.

For example, both of the following DFAs, which accept \emptyset , are deterministically simplified:



We define a simplification algorithm for DFAs that takes in

- a DFA M and
- an alphabet Σ

and returns a DFA N such that

- N is deterministically simplified,
- $N \approx M$, and
- **alphabet** $N = \mathbf{alphabet}(L(M)) \cup \Sigma$.

Thus, the alphabet of N will consist of all symbols that either appear in strings that are accepted by M or are in Σ .

The algorithm begins by letting the FA M' be **simplify** M , i.e., the result of running our simplification algorithm for FAs on M . M' will have the following properties.

- M' is simplified.
- $M' \approx M$.
- **alphabet** $M' = \mathbf{alphabet}(L(M')) = \mathbf{alphabet}(L(M))$.
- For all $q \in Q_{M'}$ and $a \in \mathbf{alphabet} M'$, there is at most one $r \in Q_{M'}$ such that $q, a \rightarrow r \in T_{M'}$. This property holds since M is a DFA and M' was formed by removing states and transitions from M .

Let $\Sigma' = \mathbf{alphabet} M' \cup \Sigma = \mathbf{alphabet}(L(M)) \cup \Sigma$. If M' is a DFA and **alphabet** $M' = \Sigma'$, then the algorithm returns M' . Otherwise, it must turn M' into a DFA whose alphabet is Σ' . We have that

- **alphabet** $M' \subseteq \Sigma'$; and
- for all $q \in Q_{M'}$ and $a \in \Sigma'$, there is at most one $r \in Q_{M'}$ such that $q, a \rightarrow r \in T_{M'}$.

Since M' is simplified, there are two cases to consider. If M' has no accepting states, then $s_{M'}$ is the only state of M' and M' has no transitions. Thus the algorithm can return the DFA N defined by:

- $Q_N = Q_{M'} = \{s_{M'}\}$;
- $s_N = s_{M'}$;
- $A_N = A_{M'} = \emptyset$; and
- $T_N = \{s_{M'}, a \rightarrow s_{M'} \mid a \in \Sigma'\}$.

Alternatively, M' has at least one accepting state. Thus, M' has no dead states. So it can return the DFA N defined by:

- $Q_N = Q_{M'} \cup \{\langle \text{dead} \rangle\}$ (we put enough brackets around $\langle \text{dead} \rangle$ so that it's not in $Q_{M'}$);
- $s_N = s_{M'}$;
- $A_N = A_{M'}$; and
- $T_N = T_{M'} \cup T'$, where T' is the set of all transitions $q, a \rightarrow \langle \text{dead} \rangle$ such that either
 - $q \in Q_{M'}$ and $a \in \Sigma'$, but there is no $r \in Q_{M'}$ such that $q, a \rightarrow r \in T_{M'}$;
 - or
 - $q = \langle \text{dead} \rangle$ and $a \in \Sigma'$.


We define a function **determSimplify** $\in \mathbf{DFA} \times \mathbf{Alp} \rightarrow \mathbf{DFA}$ by: **determSimplify**(M, Σ) is the result of running the above algorithm on M and Σ .

Theorem 3.11.7


For all $M \in \mathbf{DFA}$ and $\Sigma \in \mathbf{Alp}$:

- **determSimplify**(M, Σ) is deterministically simplified;
- **determSimplify**(M, Σ) $\approx M$; and
- **alphabet**(**determSimplify**(M, Σ)) = **alphabet**($L(M)$) $\cup \Sigma$.

For example, suppose M is the DFA




Then $\text{determinize}(\text{determinize}(M, \{2\}))$ is the DFA



3.11.4 Converting NFAs to DFAs

Suppose M is the NFA



Our algorithm for converting NFAs to DFAs will convert M into a DFA N whose states represent the elements of the set

$$\{ \Delta_M(\{A\}, w) \mid w \in \{0, 1\}^* \}.$$

For example, one of the states of N will be $\langle A, B \rangle$, which represents $\{A, B\} = \Delta_M(\{A\}, 1)$. This is the state that our DFA will be in after processing 1 from the start state.

Before describing our conversion algorithm, we first state a proposition concerning the Δ function for NFAs, and say how we will represent finite sets of symbols as symbols.

Proposition 3.11.8

Suppose M is an NFA.

- (1) For all $P \subseteq Q_M$, $\Delta_M(P, \epsilon) = P$.
- (2) For all $P \subseteq Q_M$ and $a \in \text{alphabet } M$,

$$\Delta_M(P, a) = \{ r \in Q_M \mid p, a \rightarrow r \in T_M, \text{ for some } p \in P \}.$$

(3) For all $P \subseteq Q_M$ and $x, y \in (\mathbf{alphabet} M)^*$,

$$\Delta_M(P, xy) = \Delta_M(\Delta_M(P, x), y).$$

Given a finite set of symbols P , we write \overline{P} for the symbol

$$\langle a_1, \dots, a_n \rangle,$$

where a_1, \dots, a_n are all of the elements of P , in order according to our total ordering on **Sym**, and without repetition. For example, $\overline{\{B, A\}} = \langle A, B \rangle$ and $\overline{\emptyset} = \langle \rangle$. It is easy to see that, if P and R are finite sets of symbols, then $\overline{P} = \overline{R}$ iff $P = R$.

We convert an NFA M into a DFA N as follows. First, we generate the least subset X of $\mathcal{P}Q_M$ such that:


- $\{s_M\} \in X$; and
- for all $P \in X$ and $a \in \mathbf{alphabet} M$, $\Delta_M(P, a) \in X$.

Thus $|X| \leq 2^{|Q_M|}$. Then we define the DFA N as follows:

- $Q_N = \{\overline{P} \mid P \in X\}$;
- $s_N = \overline{\{s_M\}} = \langle s_M \rangle$;
- $A_N = \{\overline{P} \mid P \in X \text{ and } P \cap A_M \neq \emptyset\}$; and
- $T_N = \{(\overline{P}, a, \overline{\Delta_M(P, a)}) \mid P \in X \text{ and } a \in \mathbf{alphabet} M\}$.

Then N is a DFA with alphabet **alphabet** M and, for all $P \in X$ and $a \in \mathbf{alphabet} M$, $\delta_N(\overline{P}, a) = \overline{\Delta_M(P, a)}$.

Suppose M is the NFA




Let's work out what the DFA N is.

- To begin with, $\{A\} \in X$, so that $\langle A \rangle \in Q_N$. And $\langle A \rangle$ is the start state of N . It is not an accepting state, since $A \notin A_M$.
- Since $\{A\} \in X$, and $\Delta(\{A\}, 0) = \emptyset$, we add \emptyset to X , $\langle \rangle$ to Q_N and $\langle A \rangle, 0 \rightarrow \langle \rangle$ to T_N .

Since $\{A\} \in X$, and $\Delta(\{A\}, 1) = \{A, B\}$, we add $\{A, B\}$ to X , $\langle A, B \rangle$ to Q_N and $\langle A \rangle, 1 \rightarrow \langle A, B \rangle$ to T_N .

- Since $\emptyset \in X$, $\Delta(\emptyset, 0) = \emptyset$ and $\emptyset \in X$, we don't have to add anything to X or Q_N , but we add $\langle \rangle, 0 \rightarrow \langle \rangle$ to T_N .
Since $\emptyset \in X$, $\Delta(\emptyset, 1) = \emptyset$ and $\emptyset \in X$, we don't have to add anything to X or Q_N , but we add $\langle \rangle, 1 \rightarrow \langle \rangle$ to T_N .
- Since $\{A, B\} \in X$, $\Delta(\{A, B\}, 0) = \emptyset$ and $\emptyset \in X$, we don't have to add anything to X or Q_N , but we add $\langle A, B \rangle, 0 \rightarrow \langle \rangle$ to T_N .
Since $\{A, B\} \in X$, $\Delta(\{A, B\}, 1) = \{A, B\} \cup \{C\} = \{A, B, C\}$, we add $\{A, B, C\}$ to X , $\langle A, B, C \rangle$ to Q_N , and $\langle A, B \rangle, 1 \rightarrow \langle A, B, C \rangle$ to T_N . Since $\{A, B, C\}$ contains (the only) one of M 's accepting states, we add $\langle A, B, C \rangle$ to A_N .
- Since $\{A, B, C\} \in X$ and $\Delta(\{A, B, C\}, 0) = \emptyset \cup \emptyset \cup \{C\} = \{C\}$, we add $\{C\}$ to X , $\langle C \rangle$ to Q_N and $\langle A, B, C \rangle, 0 \rightarrow \langle C \rangle$ to T_N . Since $\{C\}$ contains one of M 's accepting states, we add $\langle C \rangle$ to A_N .
Since $\{A, B, C\} \in X$, $\Delta(\{A, B, C\}, 1) = \{A, B\} \cup \{C\} \cup \emptyset = \{A, B, C\}$ and $\{A, B, C\} \in X$, we don't have to add anything to X or Q_N , but we add $\langle A, B, C \rangle, 1 \rightarrow \langle A, B, C \rangle$ to T_N .
- Since $\{C\} \in X$, $\Delta(\{C\}, 0) = \{C\}$ and $\{C\} \in X$, we don't have to add anything to X or Q_N , but we add $\langle C \rangle, 0 \rightarrow \langle C \rangle$ to T_N .
Since $\{C\} \in X$, $\Delta(\{C\}, 1) = \emptyset$ and $\emptyset \in X$, we don't have to add anything to X or Q_N , but we add $\langle C \rangle, 1 \rightarrow \langle \rangle$ to T_N .

Since there are no more elements to add to X , we are done. Thus, the DFA N is



The following two lemmas show why our conversion process is correct.

Lemma 3.11.9

For all $w \in (\text{alphabet } M)^*$:

- $\Delta_M(\{s_M\}, w) \in X$; and
- $\delta_N(s_N, w) = \overline{\Delta_M(\{s_M\}, w)}$.

Proof. By left string induction.

(Basis Step) We have that $\Delta_M(\{s_M\}, \%) = \{s_M\} \in X$ and $\delta_N(s_N, \%) = s_N = \overline{\{s_M\}} = \overline{\Delta_M(\{s_M\}, \%)}$.

(Inductive Step) Suppose $a \in \mathbf{alphabet} M$ and $w \in (\mathbf{alphabet} M)^*$. Assume the inductive hypothesis: $\Delta_M(\{s_M\}, w) \in X$ and $\delta_N(s_N, w) = \overline{\Delta_M(\{s_M\}, w)}$. Since $\Delta_M(\{s_M\}, w) \in X$ and $a \in \mathbf{alphabet} M$, we have that $\Delta_M(\{s_M\}, wa) = \Delta_M(\Delta_M(\{s_M\}, w), a) \in X$. Thus

$$\begin{aligned} \delta_N(s_N, wa) &= \delta_N(\delta_N(s_N, w), a) \\ &= \delta_N(\overline{\Delta_M(\{s_M\}, w)}, a) \quad (\text{ind. hyp.}) \\ &= \overline{\Delta_M(\Delta_M(\{s_M\}, w), a)} \\ &= \overline{\Delta_M(\{s_M\}, wa)}. \end{aligned}$$

□

Lemma 3.11.10

$L(N) = L(M)$.

Proof.

$(L(M) \subseteq L(N))$ Suppose $w \in L(M)$, so that $w \in (\mathbf{alphabet} M)^* = (\mathbf{alphabet} N)^*$ and $\Delta_M(\{s_M\}, w) \cap A_M \neq \emptyset$. By Lemma 3.11.9, we have that $\Delta_M(\{s_M\}, w) \in X$ and $\delta_N(s_N, w) = \overline{\Delta_M(\{s_M\}, w)}$. Since $\Delta_M(\{s_M\}, w) \in X$ and $\Delta_M(\{s_M\}, w) \cap A_M \neq \emptyset$, it follows that $\delta_N(s_N, w) = \overline{\Delta_M(\{s_M\}, w)} \in A_N$. Thus $w \in L(N)$.

$(L(N) \subseteq L(M))$ Suppose $w \in L(N)$, so that $w \in (\mathbf{alphabet} N)^* = (\mathbf{alphabet} M)^*$ and $\delta_N(s_N, w) \in A_N$. By Lemma 3.11.9, we have that $\delta_N(s_N, w) = \overline{\Delta_M(\{s_M\}, w)}$. Thus $\overline{\Delta_M(\{s_M\}, w)} \in A_N$, so that $\Delta_M(\{s_M\}, w) \cap A_M \neq \emptyset$. Thus $w \in L(M)$.

□

We define a function $\mathbf{nfaToDFA} \in \mathbf{NFA} \rightarrow \mathbf{DFA}$ by: $\mathbf{nfaToDFA} M$ is the result of running the preceding algorithm with input M .

Theorem 3.11.11

For all $M \in \mathbf{NFA}$:

- $\mathbf{nfaToDFA} M \approx M$; and
- $\mathbf{alphabet}(\mathbf{nfaToDFA} M) = \mathbf{alphabet} M$.

3.11.5 Processing DFAs in Forlan

The Forlan module `DFA` defines an abstract type `dfa` (in the top-level environment) of deterministic finite automata, along with various functions for processing DFAs. Values of type `dfa` are implemented as values of type `fa`, and the module `DFA` provides the following injection and projection functions

```
val injToFA      : dfa -> fa
val injToEFA     : dfa -> efa
val injToNFA     : dfa -> nfa
val projFromFA   : fa -> dfa
val projFromEFA  : efa -> dfa
val projFromNFA  : nfa -> dfa
```

These functions are available in the top-level environment with the names `injDFAToFA`, `injDFAToEFA`, `injDFAToNFA`, `projFAToDFA`, `projEFAToDFA` and `projNFAToDFA`.

The module `DFA` also defines the functions:

```
val input        : string -> dfa
val determProcStr : dfa -> sym * str -> sym
val determAccepted : dfa -> str -> bool
val determSimplified : dfa -> bool
val determSimplify : dfa * sym set -> dfa
val fromNFA       : nfa -> dfa
```

The function `input` is used to input a DFA. The function `determProcStr` is used to compute $\delta_M(q, w)$ for a DFA M , using the properties of δ_M . The function `determAccepted` uses `determProcStr` to check whether a string is accepted by a DFA. The function `determSimplified` tests whether a DFA is deterministically simplified, and the function `determSimplify` corresponds to `determSimplify`. The function `fromNFA` corresponds to our conversion function `nfaToDFA`, and is available in the top-level environment with that name:

```
val nfaToDFA : nfa -> dfa
```

Most of the functions for processing FAs that were introduced in previous sections are inherited by `DFA`:


```
val output          : string * dfa -> unit
val numStates       : dfa -> int
val numTransitions  : dfa -> int
val alphabet        : dfa -> sym set
val equal           : dfa * dfa -> bool
val checkLP         : dfa -> lp -> unit
val validLP         : dfa -> lp -> bool
val isomorphism     : dfa * dfa * sym_rel -> bool
val findIsomorphism : dfa * dfa -> sym_rel
val isomorphic      : dfa * dfa -> bool
```

```

val renameStates      : dfa * sym_rel -> dfa
val renameStatesCanonically : dfa -> dfa
val processStr        : dfa -> sym set * str -> sym set
val accepted         : dfa -> str -> bool
val findLP           : dfa -> sym set * str * sym set -> lp
val findAcceptingLP  : dfa -> str -> lp

```

Suppose `dfa` is the DFA




We can turn `dfa` into an equivalent deterministically simplified DFA whose alphabet is the union of the alphabet of the language of `dfa` and $\{2\}$, i.e., whose alphabet is $\{0, 1, 2\}$, as follows:

```


- val dfa' = DFA.determSimplify(dfa, SymSet.input "");
@ 2
@ .
val dfa' = - : dfa
- DFA.output("", dfa');
{states} A, B, C, <dead> {start state} A
{accepting states} A, B, C
{transitions}
A, 0 -> B; A, 1 -> A; A, 2 -> <dead>; B, 0 -> C; B, 1 -> A;
B, 2 -> <dead>; C, 0 -> <dead>; C, 1 -> A; C, 2 -> <dead>;
<dead>, 0 -> <dead>; <dead>, 1 -> <dead>; <dead>, 2 -> <dead>
val it = () : unit

```

Thus `dfa'` is




Suppose that `nfa` is the NFA



We can convert `nfa` to a DFA as follows:

```
- val dfa = nfaToDFA nfa;
val dfa = - : dfa
- DFA.output("", dfa);
{states} <>, <A>, <C>, <A,B>, <A,B,C> {start state} <A>
{accepting states} <C>, <A,B,C>
{transitions}
<>, 0 -> <>; <>, 1 -> <>; <A>, 0 -> <>; <A>, 1 -> <A,B>;
<C>, 0 -> <C>; <C>, 1 -> <>; <A,B>, 0 -> <>; <A,B>, 1 -> <A,B,C>;
<A,B,C>, 0 -> <C>; <A,B,C>, 1 -> <A,B,C>
val it = () : unit
```


Thus `dfa` is



And we can see why `nfa` and `dfa` accept 111100, as follows:

```
- LP.output
= ("",
= NFA.findAcceptingLP nfa (Str.fromString "111100"));
A, 1 => A, 1 => A, 1 => B, 1 => C, 0 => C, 0 => C
val it = () : unit
- LP.output
= ("",
= DFA.findAcceptingLP dfa (Str.fromString "111100"));
<A>, 1 => <A,B>, 1 => <A,B,C>, 1 => <A,B,C>, 1 => <A,B,C>, 0 =>
<C>, 0 => <C>
val it = () : unit
```

Finally, we see an example in which an NFA with 4 states is converted to a DFA with 16 states; there is a state of the DFA corresponding to every element of the power set of the set of states of the NFA. Suppose `nfa'` is the NFA



Then we can convert `nfa'` into a DFA, as follows:

```

- val dfa' = nfaToDFA nfa';
val dfa' = - : dfa
- DFA.numStates dfa';
val it = 16 : int
- DFA.output("", dfa');
{states}
<>, <A>, <B>, <C>, <D>, <A,B>, <A,C>, <A,D>, <B,C>, <B,D>, <C,D>,
<A,B,C>, <A,B,D>, <A,C,D>, <B,C,D>, <A,B,C,D>
{start state} <A>
{accepting states}
<D>, <A,D>, <B,D>, <C,D>, <A,B,D>, <A,C,D>, <B,C,D>, <A,B,C,D>
{transitions}
<>, 0 -> <>; <>, 1 -> <>; <>, 2 -> <>; <A>, 0 -> <A>;
<A>, 1 -> <A,B>; <A>, 2 -> <>; <B>, 0 -> <C>; <B>, 1 -> <C>;
<B>, 2 -> <B>; <C>, 0 -> <D>; <C>, 1 -> <D>; <C>, 2 -> <C>;
<D>, 0 -> <>; <D>, 1 -> <>; <D>, 2 -> <D>; <A,B>, 0 -> <A,C>;
<A,B>, 1 -> <A,B,C>; <A,B>, 2 -> <B>; <A,C>, 0 -> <A,D>;
<A,C>, 1 -> <A,B,D>; <A,C>, 2 -> <C>; <A,D>, 0 -> <A>;
<A,D>, 1 -> <A,B>; <A,D>, 2 -> <D>; <B,C>, 0 -> <C,D>;
<B,C>, 1 -> <C,D>; <B,C>, 2 -> <B,C>; <B,D>, 0 -> <C>;
<B,D>, 1 -> <C>; <B,D>, 2 -> <B,D>; <C,D>, 0 -> <D>;
<C,D>, 1 -> <D>; <C,D>, 2 -> <C,D>; <A,B,C>, 0 -> <A,C,D>;
<A,B,C>, 1 -> <A,B,C,D>; <A,B,C>, 2 -> <B,C>; <A,B,D>, 0 -> <A,C>;
<A,B,D>, 1 -> <A,B,C>; <A,B,D>, 2 -> <B,D>; <A,C,D>, 0 -> <A,D>;
<A,C,D>, 1 -> <A,B,D>; <A,C,D>, 2 -> <C,D>; <B,C,D>, 0 -> <C,D>;
<B,C,D>, 1 -> <C,D>; <B,C,D>, 2 -> <B,C,D>;
<A,B,C,D>, 0 -> <A,C,D>; <A,B,C,D>, 1 -> <A,B,C,D>;
<A,B,C,D>, 2 -> <B,C,D>
val it = () : unit

```

In Section ??, we will use Forlan to show that there is no DFA with fewer than 16 states that accepts the language accepted by `nfa'` and `dfa'`.

3.11.6 Notes

In contrast to the standard approach, the transition function δ for a DFA M is not part of the definition of M , but is derived from the definition. Our approach to proving the correctness of DFAs, using induction on Λ plus proof by contradiction, is novel, simple and elegant. The material on deterministic simplification is original, but straightforward. And the algorithm for converting NFAs to DFAs is standard.

Bibliography

- [AM91] A. W. Appel and D. B. MacQueen. Standard ML of New Jersey. In *Programming Language Implementation and Logic Programming*, volume 528 of *Lecture Notes in Computer Science*, pages 1–26. Springer-Verlag, 1991.
- [BE93] J. Barwise and J. Etchemendy. *Turing’s World 3.0 for Mac: An Introduction to Computability Theory*. Cambridge University Press, 1993.
- [BLP⁺97] A. O. Bilska, K. H. Leider, M. Procopiuc, O. Procopiuc, S. H. Rodger, J. R. Salemme, and E. Tsang. A collection of tools for making automata theory and formal languages come alive. In *Twenty-eighth ACM SIGCSE Technical Symposium on Computer Science Education*, pages 15–19. ACM Press, 1997.
- [End77] H. B. Enderton. *Elements of Set Theory*. Academic Press, 1977.
- [HMU01] J. E. Hopcroft, R. Motwani, and J. D. Ullman. *Introduction to Automata Theory, Languages and Computation*. Addison-Wesley, second edition, 2001.
- [HR00] T. Hung and S. H. Rodger. Increasing visualization and interaction in the automata theory course. In *Thirty-first ACM SIGCSE Technical Symposium on Computer Science Education*, pages 6–10. ACM Press, 2000.
- [Jon97] N. J. Jones. *Computability and Complexity: From a Programming Perspective*. The MIT Press, 1997.
- [Koz97] D. C. Kozen. *Automata and Computability*. Springer-Verlag, 1997.
- [Lei10] H. Leiß. The Automata Library. <http://www.cis.uni-muenchen.de/~leiss/sml-automata.html>, 2010.
- [Lin01] P. Linz. *An Introduction to Formal Languages and Automata*. Jones and Bartlett Publishers, 2001.

- [LP98] H. R. Lewis and C. H. Papadimitriou. *Elements of the Theory of Computation*. Prentice Hall, second edition, 1998.
- [LRGS04] S. Lombardy, Y. Régis-Gianas, and J. Sakarovitch. Introducing vaucanson. *Theoretical Computer Science*, 328:77–96, 2004.
- [Mar91] J. C. Martin. *Introduction to Languages and the Theory of Computation*. McGraw Hill, second edition, 1991.
- [MTHM97] R. Milner, M. Tofte, R. Harper, and D. MacQueen. *The Definition of Standard ML—Revised 1997*. The MIT Press, 1997.
- [Pau96] L. C. Paulson. *ML for the Working Programmer*. Cambridge University Press, second edition, 1996.
- [RHND99] M. B. Robinson, J. A. Hamshar, J. E. Novillo, and A. T. Duchowski. A Java-based tool for reasoning about models of computation through simulating finite automata and turing machines. In *Thirtieth ACM SIGCSE Technical Symposium on Computer Science Education*, pages 105–109. ACM Press, 1999.
- [Rod06] S. H. Rodger. *JFLAP: An Interactive Formal Languages and Automata Package*. Jones and Bartlett Publishers, 2006.
- [RW94] D. Raymond and D. Wood. Grail: A C++ library for automata and expressions. *Journal of Symbolic Computation*, 17:341–350, 1994.
- [RWYC17] D. Raymond, D. Wood, S. Yu, and C. Câmpeanu. Grail+: A symbolic computation environment for finite-state machines, regular expressions, and finite languages. <http://www.csit.upei.ca/~ccampeanu/Grail/>, 2017.
- [Sar02] J. Saraiva. HaLeX: A Haskell library to model, manipulate and animate regular languages. In *ACM Workshop on Functional and Declarative Programming in Education (FDPE/PLI’02)*, Pittsburgh, October 2002.
- [Sto05] A. Stoughton. Experimenting with formal languages. In *Thirty-sixth ACM SIGCSE Technical Symposium on Computer Science Education*, page 566. ACM Press, 2005.
- [Sto08] A. Stoughton. Experimenting with formal languages using Forlan. In *FDPE ’08: Proceedings of the 2008 International Workshop on Functional and Declarative Programming in Education*, pages 41–50, New York, NY, USA, 2008. ACM.

-
- [Sut92] K. Sutner. Implementing finite state machines. In N. Dean and G. E. Shannon, editors, *Computational Support for Discrete Mathematics, DIMACS Series in Discrete Mathematics and Theoretical Computer Science*, volume 15, pages 347–363. American Mathematical Society, 1992.
- [Ull98] J. D. Ullman. *Elements of ML Programming: ML97 Edition*. Prentice Hall, 1998.

Index

- @, 15, 70
- \circ , 6, 8
- $-$, 4
- \in , 2
- \emptyset , 1
- $()$, 51
- $=$, 2
- \approx , 83
- $\cdot\cdot$, 8
- \cap , 5
- \cup , 5
- \sharp , 10
- $\cdot(\cdot)$, 9
- \cap , 4
- $[\cdot : \cdot]$, 3
- $\cdot^{-1}(\cdot)$, 22
- $\cdot^{-1}(\cdot)$, 9
- $[\cdot, \dots, \cdot]$, 15
- \preceq , 13
- (\cdot, \cdot) , 4, 65
- (\cdot, \cdot, \cdot) , 4
- $\%$, 36
- \cdot , 37, 70
- \times , 4
- \subsetneq , 2
- \supsetneq , 2
- $\cdot|$, 9
- \cdot_R , 43
- \cong , 11
- $\{\dots\}$, 1
- $\{\dots | \dots\}$, 3, 6
- \rightarrow , 8
- $|\cdot|$, 11, 37
- \cdot^* , 39, 40
- \subseteq , 2
- \supseteq , 2
- \cup , 4
- $\cdot[\cdot \mapsto \cdot]$, 9
- a, b, c , 36
- α, β, γ , 73
- alphabet, 39
 - \cdot^* , 39
 - language, 40, 72
 - Σ , 39
- alphabet**, 39, 40, 44, 72, 76
- and**, 14
- antisymmetric, 7
- application, *see* function, application
- associative
 - function composition, 9
 - intersection, 4
 - language concatenation, 70
 - list concatenation, 15
 - relation composition, 6
 - string concatenation, 37
 - union, 4
- Axiom of Choice, 14
- bijection from set to set, 10
- Bool**, 14
- `bool`, 52
- boolean, 14
 - and**, 14
 - Bool**, 14
 - conjunction, 14
 - disjunction, 14
 - false**, 14
 - negation, 14
 - not**, 14
 - or**, 14
 - true**, 14
- bound variable, 3
- cardinality, 10–14
- closure
 - language, 71
 - regular expression, 73

- closure complexity, 96
- commutative
 - intersection, 4
 - union, 4
- composition
 - function, *see* function, composition
 - relation, *see* relation, composition
- concatenation
 - language, 69
 - associative, 70
 - identity, 70
 - power, 70
 - zero, 70
 - list, 15
 - associative, 15
 - identity, 15
 - regular expression, 73
 - string, 37
 - associative, 37
 - identity, 37
 - power, 37
- conservative approximation to subset
 - testing, 112
- contradiction, 18
 - proof by, 12
- countable, 11, 40
- countably infinite, 11, 36, 37, 39, 40
- curried function, 66
- data structure, 14–16
- dead state, 147
- Δ ., 142
- deterministic finite automaton
 - exponential blowup, 184
- diagonalization
 - cardinality, 12
- diff**, 46
- difference
 - set, 4
- difference function, 46
- disjoint sets, 4
- distributivity, 4
- domain, 6
- domain**, 6
- element of, 2
- empty set, 1
- equal
 - set, 2
- existentially quantified, 3
- FA**, 140, 145, 150
 - accepted**, 145
 - findAcceptingLP**, 145
 - findIsomorphism**, 140
 - findLP**, 145
 - isomorphic**, 140
 - isomorphism**, 140
 - processStr**, 145
 - renameStates**, 140
 - renameStatesCanonically**, 140
 - simplified**, 151
 - simplify**, 151
- false**, 14
- finite, 11
- finite automaton
 - calculating $\Delta(\cdot, \cdot)$, 142
 - characterizing $L(\cdot)$, 144
 - checking for string acceptance, 141
 - dead state, 147
 - Δ ., 142
 - iso**, 135
 - isomorphic, 135
 - isomorphism, 135–141
 - checking whether FAs are isomorphic, 137
 - isomorphism from FA to FA, 135
 - $L(\cdot)$, 144
 - live state, 147
 - reachable state, 147
 - renameStates**, 136
 - renameStatesCanonically**, 137
 - renaming states, 136
 - searching for labeled paths, 141, 144
 - simplification, 147–152
 - simplification algorithm, 149
 - simplified, 148
 - simplify**, 149
 - useful state, 147
- fn**, 55
- Forlan, 50–68
 - exiting, 51
 - input prompt, 59
 - installing, 50
 - interrupting, 51
 - primary prompt, 51
 - regular expression syntax, 79

- running, 50
 - secondary prompt, 54
 - string syntax, 62
 - formal language, *see* language
 - forming sets, 3, 6
 - function, 7, 54
 - $\cdot \cdot$, 8
 - \circ , 8
 - $\cdot(\cdot)$, 9
 - $\cdot^{-1}(\cdot)$, 22
 - $\cdot^{-1}(\cdot)$, 9
 - $\cdot \cdot$, 9
 - $\cdot[\cdot \mapsto \cdot]$, 9
 - application, 8
 - bijection from set to set, 10
 - composition, 8
 - associative, 9
 - identity, 9
 - equality, 8
 - from set to set, 8
 - id**, 8
 - identity, 8
 - image under, 9
 - injection, 10
 - injection from set to set, 10
 - injective, 10
 - inverse image of relation under, 22
 - inverse image under, 9
 - restriction, 9
 - surjection from set to set, 13
 - updating, 9
- generalized intersection, 5
- generalized union, 5
- hasEmp**, 111
- hasSym**, 111
- id**, 6, 8
- idempotent
 - intersection, 4
 - union, 4
- identity
 - function composition, 9
 - language concatenation, 70
 - list concatenation, 15
 - relation composition, 6
 - string concatenation, 37
 - union, 4
- identity function, 8
- identity relation, 6, 8
- iff, 2
- image under
 - function, *see* function, image under
- inclusion, 83
- induction, 16–23
 - mathematical, *see* mathematical induction
 - strong, *see* strong induction
 - well-founded, *see* well-founded induction
- inductive definition, 44, 47
 - induction principle, 45, 47
- inductive hypothesis
 - mathematical induction, 16
 - strong induction, 18
 - well-founded induction, 21
- infinite, 11
 - countably, *see* countably infinite
- injection, 10
- injection from set to set, 10
- injective, 10
- int**, 52
- integer, 1
- integers
 - $[\cdot : \cdot]$, 3
 - interval, 3
- interactive input, 59
- intersection
 - language, 69
 - set, 4
 - associative, 4
 - commutative, 4
 - generalized, 5
 - idempotent, 4
 - zero, 4
- inverse image under
 - function, *see* function, inverse image under
- iso**, 135
 - reflexive, 136
 - symmetric, 136
 - transitive, 136
- isomorphic
 - finite automaton, 135
- isomorphism
 - finite automaton, 135–141

- checking whether FAs are isomorphic, 137
- iso**, 135
- isomorphic, 135
- isomorphism from FA to FA, 135
- JForlan, x, xiii, 82
- Kleene closure, *see* closure
- $L(\cdot)$, 75, 144
- Lan**, 40
- language, 40, 64
 - @, 70
 - \cdot , 70
 - alphabet, 40, 72
 - alphabet**, 40, 72
 - closure, 86
 - concatenation, 69, 85
 - associative, 70
 - identity, 70
 - power, 70
 - zero, 70
 - Lan**, 40
 - operation precedence, 72
 - regular, *see* regular language
 - Σ -language, 40
- left string induction, 42
- length
 - string, 37
- List** \cdot , 16
- list, 15–16, 36
 - @, 15
 - concatenation, 15
 - associative, 15
 - identity, 15
 - List** \cdot , 16
 - list, 16
- list, 16
- lists
 - $[\cdot, \dots, \cdot]$, 15
- live state, 147
- logical contradiction, 6
- mathematical induction, 16, 37
 - inductive hypothesis, 16
- \mathbb{N} , 1
- natural number, 1
- natural numbers
 - $[\cdot : \cdot]$, 3
 - interval, 3
 - no bigger, 13
 - none**, 14
 - not**, 14
 - numConcats**, 100
 - numSyms**, 100
 - obviousSubset**, 112
 - one-to-one correspondence, 10
 - Option**, 14
 - option, 14
 - none**, 14
 - Option**, 14
 - some** \cdot , 14
 - or**, 14
 - ordered pair, 4, 65
 - ordered triple, 4
 - ordered n -tuple, 4
 - palindrome, 40, 44
 - powerset, 4
 - \mathcal{P} , 4
 - predecessor, 20
 - pred** $_{\mathbb{N}}$, 22
 - predessor relation, 22
 - prefix, 38
 - proper, 38
 - product, 4, 10
 - projection, 10
 - prompt
 - input, 59
 - primary, 51
 - secondary, 54
 - proof by contradiction, 12
 - proper
 - prefix, 38
 - subset, 2
 - substring, 38
 - suffix, 38
 - superset, 2
 - quantification
 - existential, 3
 - universal, 3
 - \mathbb{R} , 1
 - range, 6
 - range**, 6
 - reachable state, 147

- real number, 1
- recursion, 28–31
 - natural numbers, 37
 - string, 39, 43
 - left, 39
 - right, 39
- reflexive on set, 7
 - \approx , 84
 - iso**, 136
- Reg**, 73
- Reg, 79
 - allStr, 79
 - allSym, 79
 - alphabet, 79
 - closure, 73
 - compare, 79
 - concat, 79
 - concatToList, 79
 - emptySet, 79
 - emptyStr, 79
 - equal, 79
 - fromStr, 79
 - fromStrSet, 79
 - fromSym, 79
 - genConcat, 79
 - genUnion, 79
 - height, 79
 - input, 79
 - numLeaves, 79
 - output, 79
 - power, 79
 - reg, 79
 - rightConcat, 79
 - rightUnion, 79
 - size, 79
 - sortUnions, 79
 - union, 79
 - unionsToList, 79
- reg, 79
- RegLab**, 73
- RegLan**, 78
- regular expression, 73–125
 - \approx , 83
 - reflexive, 84
 - symmetric, 84
 - transitive, 84
 - abbreviated notation, 74
 - α , β , γ , 73
 - alphabet, 76
 - closure, 73
 - closure complexity, 96
 - concatenation, 73
 - conservative approximation to subset testing, 112
 - conservative subset test, 112
 - design(, 88
 - design), 96
 - equivalence, 83–88
 - Forlan syntax, 79
 - hasEmp**, 111
 - hasSym**, 111
 - $L(\cdot)$, 75
 - label, 73
 - language generated by, 75
 - meaning, 75
 - number of concatenations, 100
 - number of symbols, 100
 - numConcat**, 100
 - numSyms**, 100
 - obviousSubset**, 112
 - operator associativity, 74
 - operator precedence, 74
 - order, 74
 - power, 76
 - proof of correctness(, 88
 - proof of correctness), 96
 - simplification, 96–125
 - standardized, 100
 - structural rule, 115
 - testing for membership of empty string, 111
 - testing for membership of symbol, 111
 - union, 73
 - weakly simplified, 104
- regular expression), 96
- regular language, 78
- relation, 6, 65
 - \circ , 6, 8
 - antisymmetric, 7
 - composition, 6, 8
 - associative, 6
 - identity, 6
 - domain, 6
 - domain**, 6
 - function, *see* function
 - id**, 6
 - identity, 6, 8

- inverse, 7
- range, 6
- range**, 6
- reflexive on set, 7
- symmetric, 7
- total, 7
- transitive, 7
- well-founded, 20
 - $\cdot \triangleright \cdot$, 22
 - lexicographic order, 22
 - R -eqtxtminimal, 20
 - predecessor, 20
 - pred** _{\mathbb{N}} , 22
 - predecessor relation, 22
- relation from set to set, 6
- renameStates**, 136
- renameStatesCanonically**, 137
- restriction
 - function, *see* function, restriction
- reversal
 - string, 43
- right string induction, 42, 43
- Russell's Paradox, 3
- same size, 11
- Schröder-Bernstein Theorem, 14
- Set**, 60
 - empty**, 60
 - filter**, 60
 - 'a set**, 60
 - sing**, 60
 - size**, 60
 - toList**, 60
- set, 1–16
 - $-$, 4
 - \in , 2
 - \emptyset , 1
 - $=$, 2
 - \bigcap , 5
 - \bigcup , 5
 - \cap , 4
 - \subsetneq , 13
 - \times , 4
 - \subsetneq , 2
 - \supsetneq , 2
 - \cong , 11
 - $\{\dots\}$, 1
 - $\{\dots \mid \dots\}$, 3, 6
 - \rightarrow , 8
 - $|\cdot|$, 11
 - \subseteq , 2
 - \supseteq , 2
 - \cup , 4
 - cardinality, 10–14
 - countable, 11
 - difference, 4
 - disjoint, 4
 - element of, 2
 - empty, 1
 - equal, 2
 - finite, 11, 60
 - formation, 3, 6
 - inclusion, 83
 - infinite, 11
 - countably, *see* countably infinite
 - intersection, *see* intersection, set
 - no bigger, 13
 - powerset, 4
 - \mathcal{P} , 4
 - product, 4
 - same size, 11
 - singleton, 1
 - size, 10–14
 - strictly smaller, 13
 - subset, 2
 - proper, 2
 - superset, 2
 - proper, 2
 - uncountable, *see* uncountable
 - union, *see* union, set
- 'a set**, 60, 61, 64
- set difference
 - language, 69
- Σ , 39
- Σ -language, 40
- simplification
 - finite automaton, 147–152
 - algorithm, 149
 - simplified, 148
 - simplify**, 149
 - regular expression, 96–125
 - structural rule, 115
 - weakly simplified, 104
- simplified
 - finite automaton, 148
- simplify**, 149
- singleton set, 1
- size

- set, 10–14
- some** ·, 14
- Standard ML, 59
 - o, 55
 - associativity, 56
 - bool, 52
 - composition, 55
 - curried function, 66
 - declaration, 53
 - function, 54
 - curried, 66
 - recursive, 57
 - function type, 55
 - int, 52
 - NONE, 53
 - option type, 53
 - precedence, 55
 - product type, 52
 - recursive datatype, 58
 - ;, 52, 54
 - NONE, 53
 - string, 52
 - tail recursion, 58
 - tree, 58
 - type, 52
 - value, 52
- standardized, 100
- Str**, 36, 40
- Str**, 62
 - allButLast, 62
 - alphabet, 62
 - compare, 62
 - input, 62
 - last, 62
 - output, 62
 - power, 62
 - prefix, 62
 - str, 62
 - substr, 62
 - suffix, 62
- str**, 62
- strictly smaller, 13
- string, 36–39, 62
 - %, 36
 - , 37
 - \cdot^R , 43
 - $|\cdot|$, 37
 - alphabet, 39, 44
 - alphabet**, 39, 44
 - concatenation, 37
 - associative, 37
 - identity, 37
 - power, 37
 - diff**, 46
 - difference function, 46
 - empty, 36
 - Forlan syntax, 62
 - length, 37
 - ordering, 36
 - palindrome, 40, 44
 - power, 37
 - prefix, 38
 - proper, 38
 - reversal, 43
 - Str**, 36
 - substring, 38
 - proper, 38
 - suffix, 38
 - proper, 38
 - u, v, w, x, y, z , 36
- string**, 52
- string induction
 - left, *see* left string induction
 - right, *see* right string induction
 - strong, *see* strong string induction
- strong induction, 17
 - inductive hypothesis, 18
- strong string induction, 42, 48
- StrSet**, 64
 - alphabet, 64
 - concat, 72
 - equal, 64
 - fromList, 64
 - getInter, 64
 - genUnion, 64
 - input, 64
 - inter, 64
 - memb, 64
 - minus, 64
 - output, 64
 - power, 72
 - subset, 64
 - union, 64
- strToReg**, 80
- structural rule, 115
- subset, 2
 - proper, 2
- substring, 38

- proper, 38
- suffix, 38
 - proper, 38
- superset, 2
 - proper, 2
- surjection from set to set, 13
- Sym**, 36, 40
- Sym, 59
 - compare, 59
 - fromString, 59
 - input, 59
 - output, 59
 - sym, 59
 - toString, 59
- sym, 59
- sym_rel, 65
- symbol, 35–36, 59
 - a, b, c , 36
 - ordering, 36
 - Sym**, 36
- symmetric, 7
 - \approx , 84
- symmetry
 - iso**, 136
- SymRel, 65
 - antisymmetric, 65
 - applyFunction, 65
 - bijectionFromTo, 65
 - compose, 65
 - domain, 65
 - equal, 65
 - fromList, 65
 - function, 65
 - functionFromTo, 65
 - genInter, 65
 - genUnion, 65
 - injection, 65
 - input, 65
 - inter, 65
 - inverse, 65
 - memb, 65
 - minus, 65
 - output, 65
 - range, 65
 - reflexive, 65
 - relationFromTo, 65
 - subset, 65
 - sym_rel, 65
 - symmetric, 65
 - total, 65
 - transitive, 65
 - union, 65
- SymSet, 61
 - equal, 61
 - fromList, 61
 - genInter, 61
 - genUnion, 61
 - input, 61
 - inter, 61
 - memb, 61
 - minus, 61
 - output, 61
 - subset, 61
 - union, 61
- symToReg, 80
- tail recursion, 58
- total, 7
- total ordering, 8
- transitive, 7
 - \approx , 84
 - iso**, 136
- tree, 24–27, 73
 - Tree_X**, 73
- Tree_X**, 73
- true**, 14
- tuple
 - projection, 10
- u, v, w, x, y, z , 36
- uncountable, 11, 12, 40
- union
 - language, 69
 - regular expression, 73
 - set, 4
 - associative, 4
 - commutative, 4
 - generalized, 5
 - idempotent, 4
 - identity, 4
- unit, 51
- universally quantified, 3
- updating
 - function, *see* function, updating
- use, 57
- useful state, 147
- val, 53

- weakly simplified, 104
- well-founded induction, 20
 - inductive hypothesis, 21
- well-founded relation, 20
 - $\cdot \triangleright \cdot$, 22
 - lexicographic order, 22
 - R -eqtxtminimal, 20
 - predecessor, 20
 - pred** _{\mathbb{N}} , 22
 - predecessor relation, 22
- \mathbb{Z} , 1
- zero
 - intersection, 4
 - language concatenation, 70